

Modeling_Katrina

Katrina Truebebach

March 18, 2019

```
rm(list = ls())
```

Load cleaned data

```
load(file = '~/DS5110/data/proj_cleaned_dta.RData')

# need to drop years before 1980: too sparse
# most of those years have 1 or 0 observations. If including year in the model, we aren't getting any
# only necessary when including year in model, but hard to compare different models then b/c data differ
train <- train %>% filter(as.integer(as.character(year)) >= 1980)
valid <- valid %>% filter(as.integer(as.character(year)) >= 1980)

# calculate logs
train <- train %>% mutate(real_gross_log = log10(real_gross))
valid <- valid %>% mutate(real_gross_log = log10(real_gross))
train <- train %>%
  mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
    'imdb_score'), funs(log = log10(.)))
valid <- valid %>%
  mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
    'imdb_score'), funs(log = log10(.)))
```

Write Functions to Automate

Write function to automate stepwise

Note: not using the step() function because can't fit and find RMSE on different datasets (train, valid)

```
# function to automate each step of stepwise variable selection
# df_vars is the dataset with only the relevant variables
# var_lst is the list of variables that are in the base model
# formula is the formula with those variables besides the y variable
step_wise_step <- function(df_vars, var_lst = NULL, formula = NULL) {
  # if first step
  if (length(var_lst) == 0) {
    # rmse with each variable against real_gross
    rmse_vars <- sapply(names(df_vars), function(var) {
      # rmse of model
      rmse(lm(as.formula(str_c('real_gross_log ~', var)), data = train), data = valid)
    })
  } else {
    # if > first step: exclude variables from var_lst from data and include in model formula
    rmse_vars <- sapply(names(df_vars) %>% select(-var_lst)), function(var) {
      # rmse of model
      rmse(lm(as.formula(str_c('real_gross_log ~', formula, ' + ', var)), data = train), data = valid)
    })
  }
}
```

```

# return the name and value of the genre that resulted in the lowest RMSE
return(rmse_vars[which.min(rmse_vars)])
}

# function to loop through each step wise loop
# adding optional starting vars and formula in case want to build off of an existing formula
step_wise_loop <- function(df_vars, starting_vars = NULL, starting_formula = NULL) {
  # list to store min RMSE from each step in
  rmse_lst <- c()

  # first step: no genre_lst or formula (default values NULL)
  min_rmse_var <- step_wise_step(df = df_vars, var_lst = starting_vars, formula = starting_formula)
  print(min_rmse_var)

  # add to list of genres, formula, and min RMSE list
  var_lst <- c(starting_vars, names(min_rmse_var))
  formula <- str_c(starting_formula, '+', names(min_rmse_var))
  rmse_lst <- c(rmse_lst, min(min_rmse_var))

  # if have starting variables, take those out of the number we are iterating through
  if (!is.null(starting_vars)) {
    df_vars_seq <- df_vars %>% select(-starting_vars)
  } else {
    df_vars_seq <- df_vars
  }
  # loop through until have considered every variable
  for (i in seq(1:(ncol(df_vars_seq)-1))) {
    print(i)
    # step
    min_rmse_var <- step_wise_step(df = df_vars, var_lst = var_lst, formula = formula)
    print(min_rmse_var)

    # add to lists
    var_lst <- c(var_lst, names(min_rmse_var))
    formula <- str_c(formula, ' + ', names(min_rmse_var))
    rmse_lst <- c(rmse_lst, min(min_rmse_var))
  }
  return(rmse_lst)
}

```

Function to graph the residuals from a model against all potential variables (included and excluded)

```

gr_resid <- function(mod) {
  # graph residuals
  # get log versions of variables since residuals are log: same scale
  df_resid <- train %>%
    add_residuals(mod, 'lresid') %>%
    mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
      'imdb_score'), funs(log = log10(.)))

  # graph each against log residual: continuous
  lapply(c('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
    'imdb_score'), function(var) {
    print(df_resid %>%

```

```

    ggplot() +
      geom_point(aes_string(str_c(var, '_log'), y = 'lresid')))
  }

# categorical
# can't log categorical variables
lapply(c('content_rating', 'year', 'total_oscars_actor', 'total_oscars_director', all_genre_vars), fun
  print(df_resid %>%
    filter(!is.na(!!rlang::sym(var))) %>%
    ggplot() +
    geom_jitter(aes_string(var, 'lresid'), alpha = .3))
}

# qq plot of residuals
df_resid %>% ggplot() +
  geom_qq(aes(sample = lresid))
}

```

Fit Model with Genre Variables vs Real Revenue

Step Wise Selection

End model includes (in order of steps): ‘Adventure’, ‘Action’, ‘Family’, ‘Mystery’, ‘Documentary’, ‘Drama’, ‘History’, ‘Romance’

Dependent variable is $\log(\text{real_gross})$. Makes model look better *and* a lot of the relationships with other variables are more linear with log, so we will need to use this as y variable in the main model.

This model selection by and large makes sense. All included variables are significant at some level. However, according to Qiang’s graphs in EDA, some of the included genres do not make a real difference to real_gross . Especially History. Also, some genres that look like they would make a significant difference are not included. For example, Animation.

Thoughts:

- There are a few genres that define almost all of the movies (For example, almost 80% of the movies are either Adventure, Action, Romance, or Drama). Thus, the relationship between revenue and some genres can be explained by other genres. For example, 93 out of 99 Animation movies are also Family. So Animation’s effect on revenue may already be captured by Family, which is included in the model.
- On the flip side, History is included even though it seems to have a negligible effect on revenue based on the EDA bar graphs. I don’t have a great explanation for this other than it was close to the cutoff RMSE for being included. 53 out of 55 History movies are also Drama. So unclear why included.

Also, the residuals are debatably random vs included and excluded variables in model (not sure if these are not-random enough to matter – see graphs).

More concerning is the fact that the residuals themselves are not Normal. See QQ-Plot (close-ish...)

```

train %>% filter(Animation == 1, Family == 1) %>% count() # 93
train %>% filter(Animation == 1) %>% count() # 101

train %>% filter(History == 1) %>% count() # 52
train %>% filter(History == 1, Drama == 1) %>% count() # 51

```

Model Fit

Which genres should we be using?

```

# version of train set with just genre columns to loop through
all_genre_vars <- c('Action', 'Adventure', 'Animation', 'Biography', 'Comedy', 'Crime', 'Documentary',
                     'Family', 'Fantasy', 'History', 'Horror', 'Music', 'Musical', 'Mystery',
                     'Romance', 'SciFi', 'Sport', 'Thriller', 'War', 'Western')

train_genre_only <- train %>% select(all_genre_vars)

# step wise implement
# return list of all min RMSE from each step -> graph
rmse_lst <- step_wise_loop(df = train_genre_only)

## Adventure
## 0.9065803
## [1] 1
## Action
## 0.8939041
## [1] 2
## Family
## 0.8870258
## [1] 3
## Mystery
## 0.8818136
## [1] 4
## Romance
## 0.8785826
## [1] 5
## Drama
## 0.8762839
## [1] 6
## History
## 0.8743844
## [1] 7
## Documentary
## 0.8729182
## [1] 8
## Musical
## 0.8721386
## [1] 9
## War
## 0.8715984
## [1] 10
## SciFi
## 0.8715957
## [1] 11
## Crime
## 0.8716081
## [1] 12
## Fantasy
## 0.8715244
## [1] 13
## Sport
## 0.8714654
## [1] 14
## Music

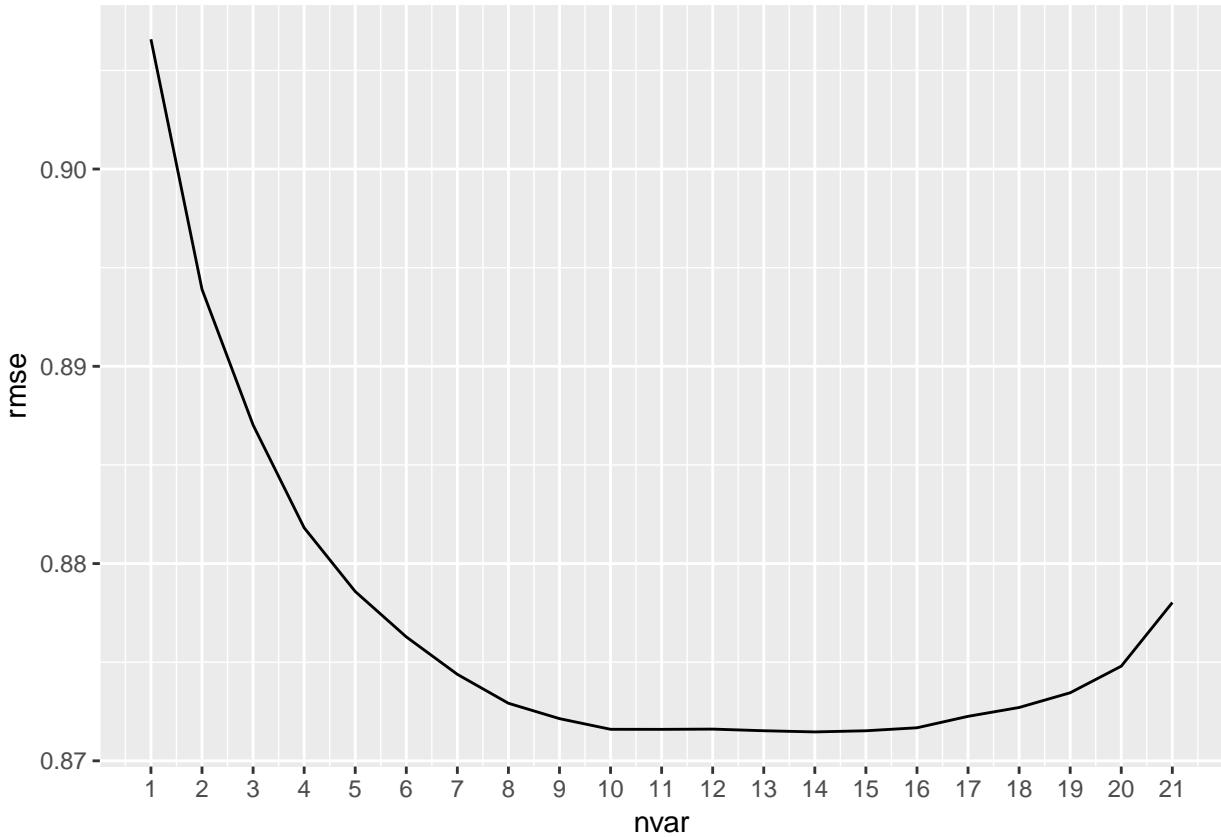
```

```
## 0.8715242
## [1] 15
## Biography
## 0.8716754
## [1] 16
## Comedy
## 0.8722553
## [1] 17
## Horror
## 0.8727044
## [1] 18
## Animation
## 0.8734501
## [1] 19
## Thriller
## 0.8747985
## [1] 20
## Western
## 0.8780216
```

Graph RMSE vs number of variables: how many to include?

Specify ‘final’ model

```
# graph RMSE at each step
fit_rmse <- tibble(nvar = 1:length(rmse_lst),
                    rmse = rmse_lst)
ggplot(fit_rmse) + geom_line(aes(x = nvar, y = rmse))+
  scale_x_continuous(breaks = seq(1, length(rmse_lst), by = 1))
```



```
# after var 8, decreases too small or increase (debatably 10?)

# model based off of step wise
# HOWEVER some of these variables are insignificant
# (see pvalues and graphs from Qiang's EDA where barely any difference in revenue from genre)
mod_genre <- lm(real_gross_log ~ Adventure + Action + Family + Mystery + Romance + Drama + History + Documentary, data = train)

summary(mod_genre)

##
## Call:
## lm(formula = real_gross_log ~ Adventure + Action + Family + Mystery +
##     Romance + Drama + History + Documentary, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -4.0859 -0.3211  0.1680  0.5636  1.7697 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 7.21040   0.04041 178.418 < 2e-16 ***
## Adventure1  0.25572   0.05959   4.291 1.87e-05 ***
## Action1     0.37493   0.05337   7.025 2.99e-12 ***
## Family1     0.46184   0.06733   6.859 9.43e-12 ***
## Mystery1    0.19739   0.07021   2.811 0.004985 **
```

```

## Romance1      0.10130   0.04886   2.073 0.038297 *
## Drama1       -0.23074   0.04398  -5.247 1.73e-07 ***
## History1      0.45334   0.12298   3.686 0.000234 ***
## Documentary1 -1.10142   0.13044  -8.444 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8582 on 1842 degrees of freedom
## Multiple R-squared:  0.1648, Adjusted R-squared:  0.1612
## F-statistic: 45.44 on 8 and 1842 DF,  p-value: < 2.2e-16
rmse(mod_genre, data = valid)

## [1] 0.8729182

# list of these variables for future use
genre_xvar <- c('Adventure', 'Action', 'Family', 'Mystery',
                 'Documentary', 'Drama', 'History', 'Romance')

```

Graph genres in and out of model against residuals. Most are fairly evenly distributed around residual. Worst is probably Western.

I like this geom jitter view better. Can see individual points. Most movies have some outliers where actual makes less money than predicted based on the included genres. But tricky because movies are multiple genres. So could be because that movie is also another genre that makes less money. Bulk of observations around zero.

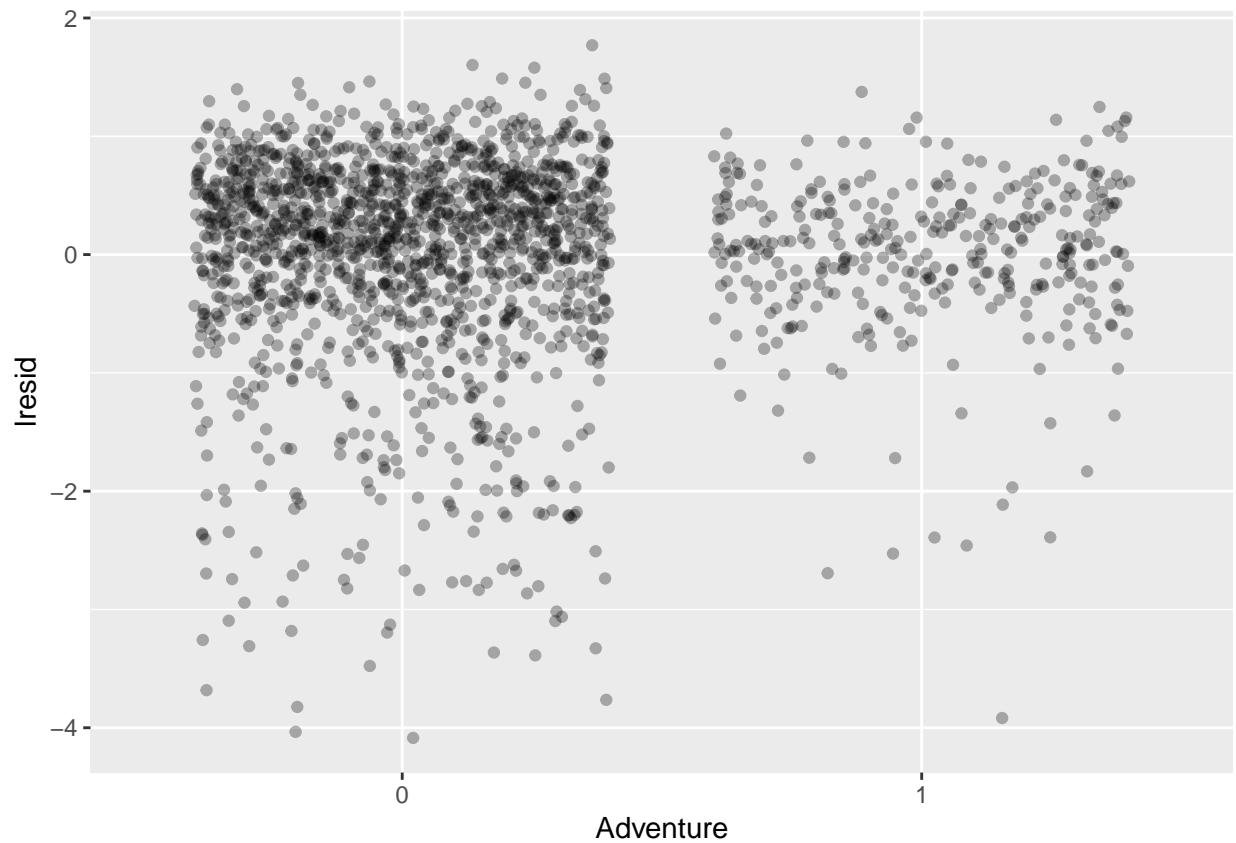
```

# graph residuals against each variable included in the model
# most look random except Adventure
train_resid <- train %>%
  add_residuals(mod_genre, 'lresid')

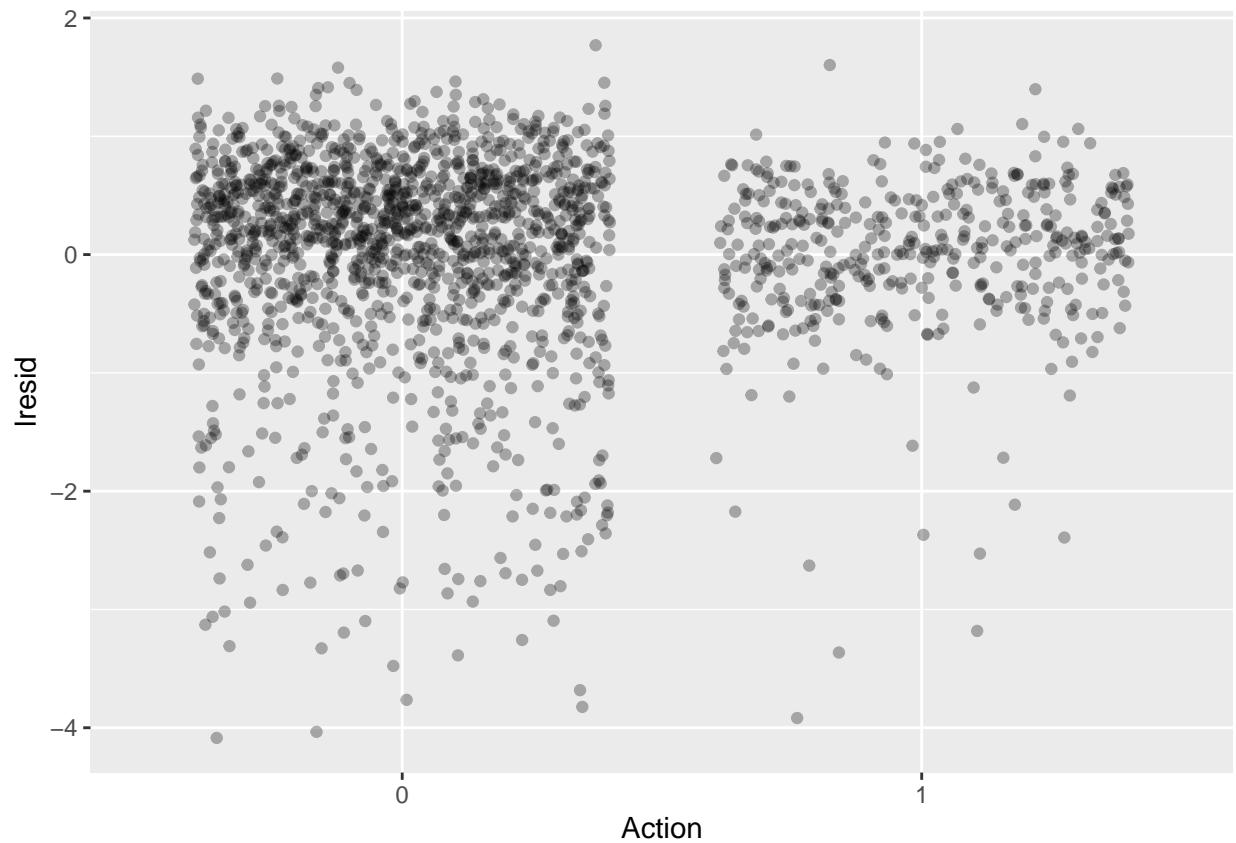
lapply(genre_xvar, function(var) {
  train_resid %>%
    ggplot() +
    geom_jitter(aes_string(var, y = 'lresid'), alpha = .3)
})

## [[1]]

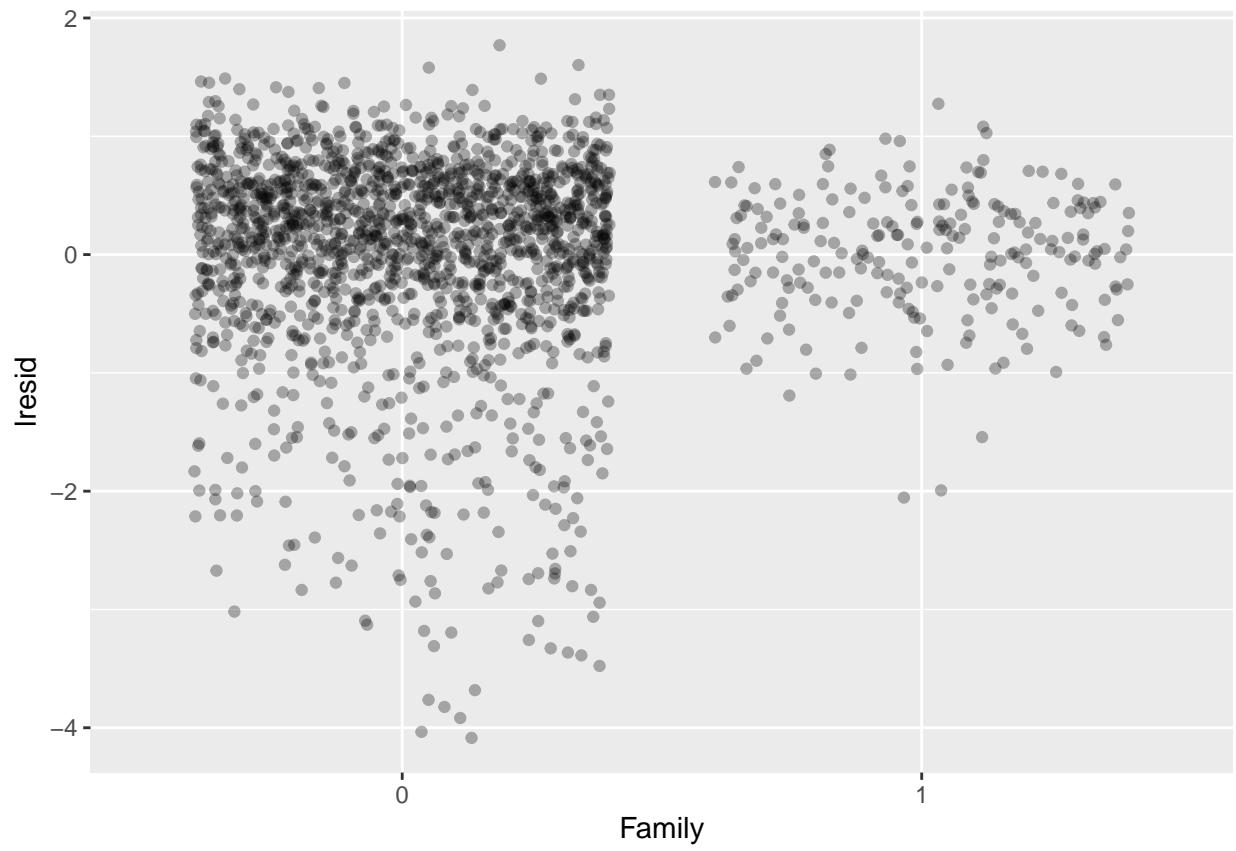
```



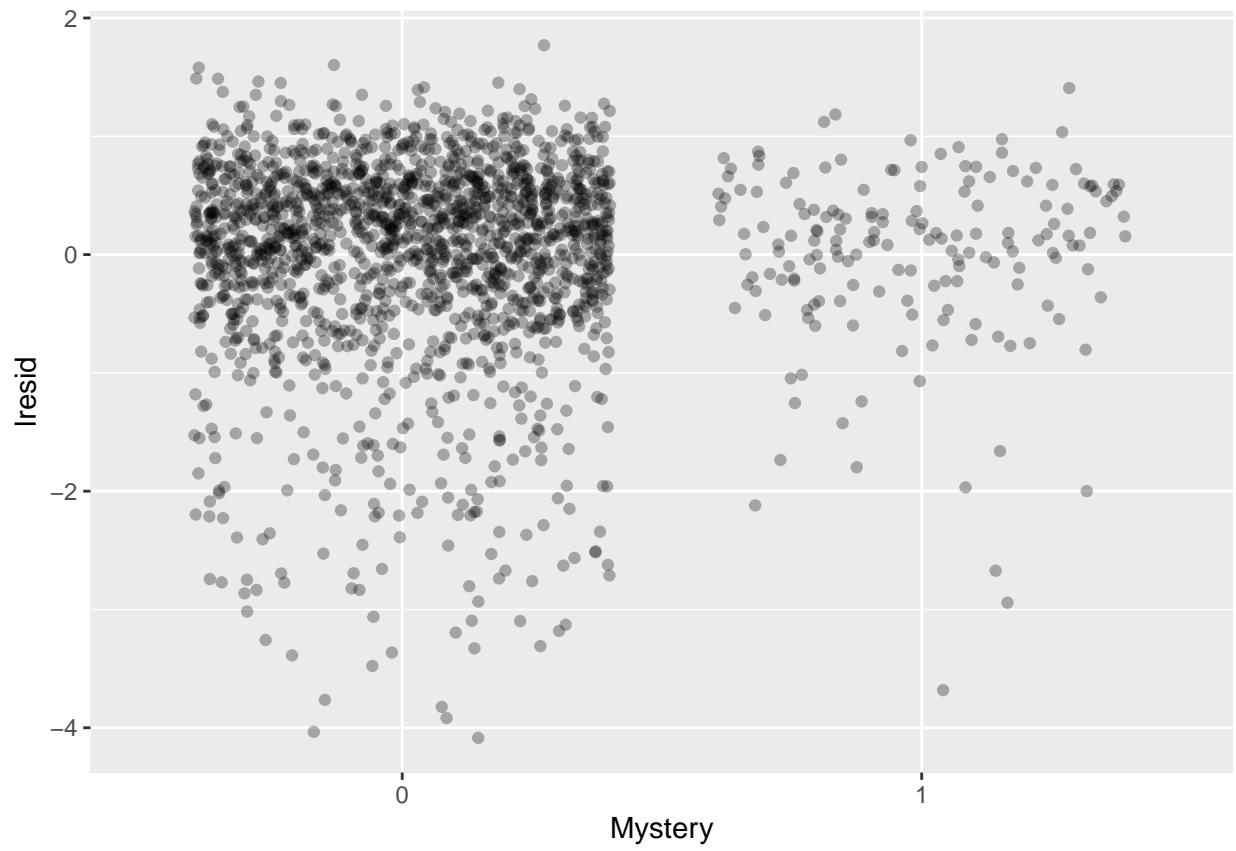
```
##  
## [[2]]
```



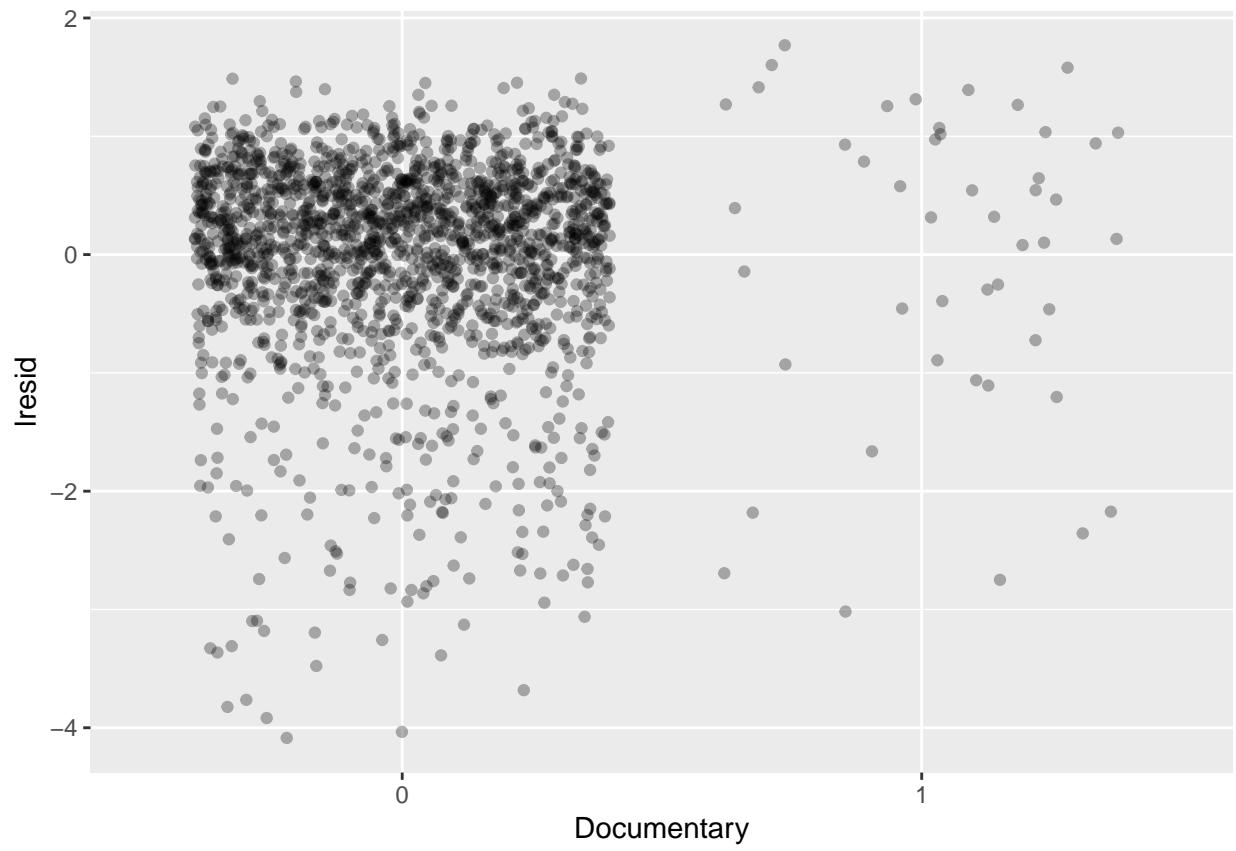
```
##  
## [[3]]
```



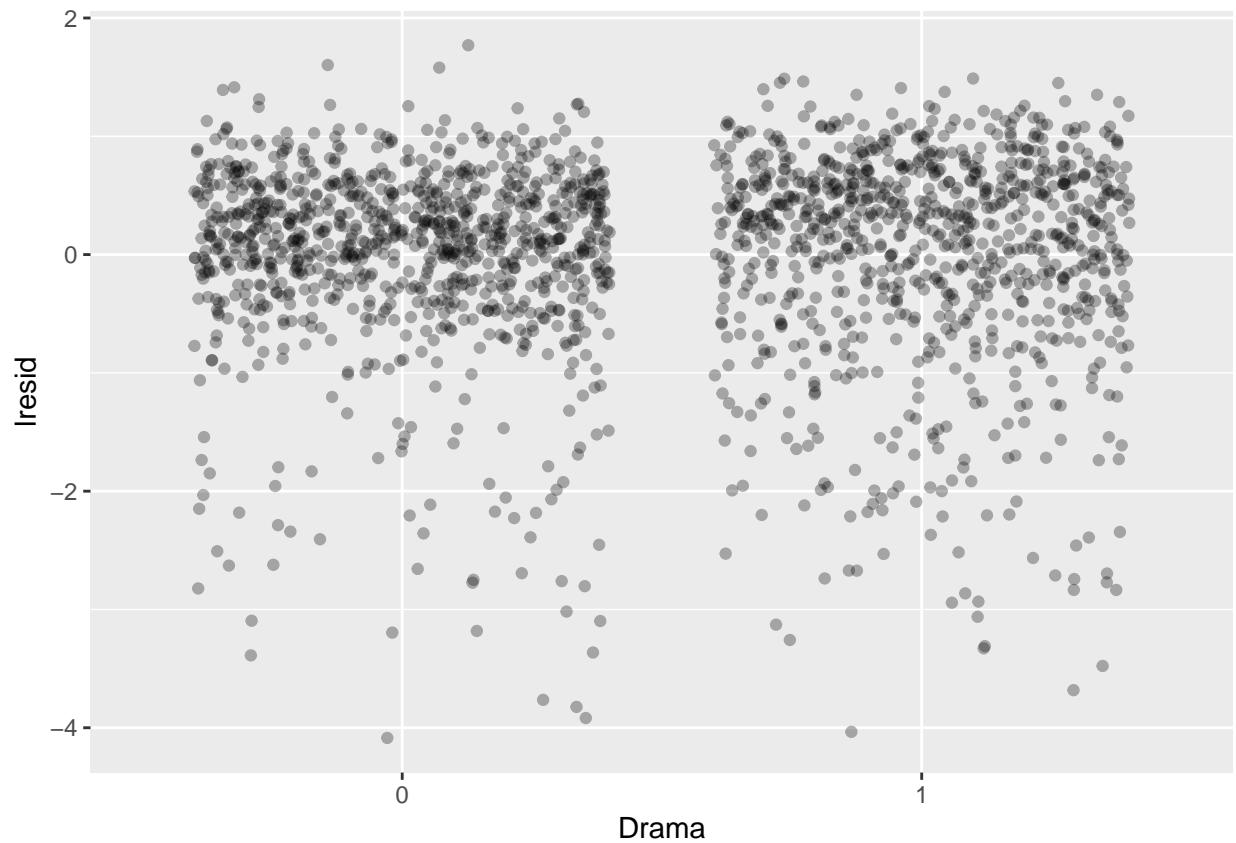
```
##  
## [[4]]
```



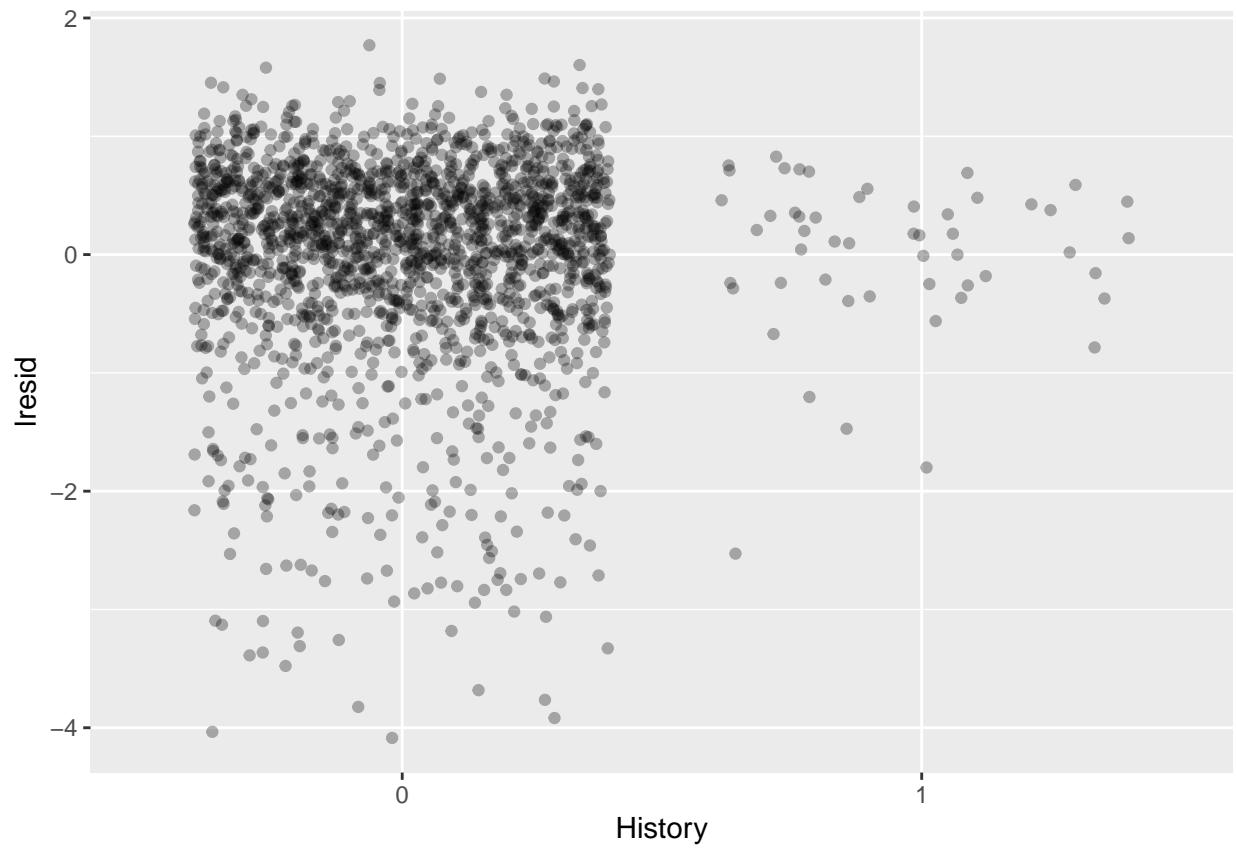
```
##  
## [[5]]
```



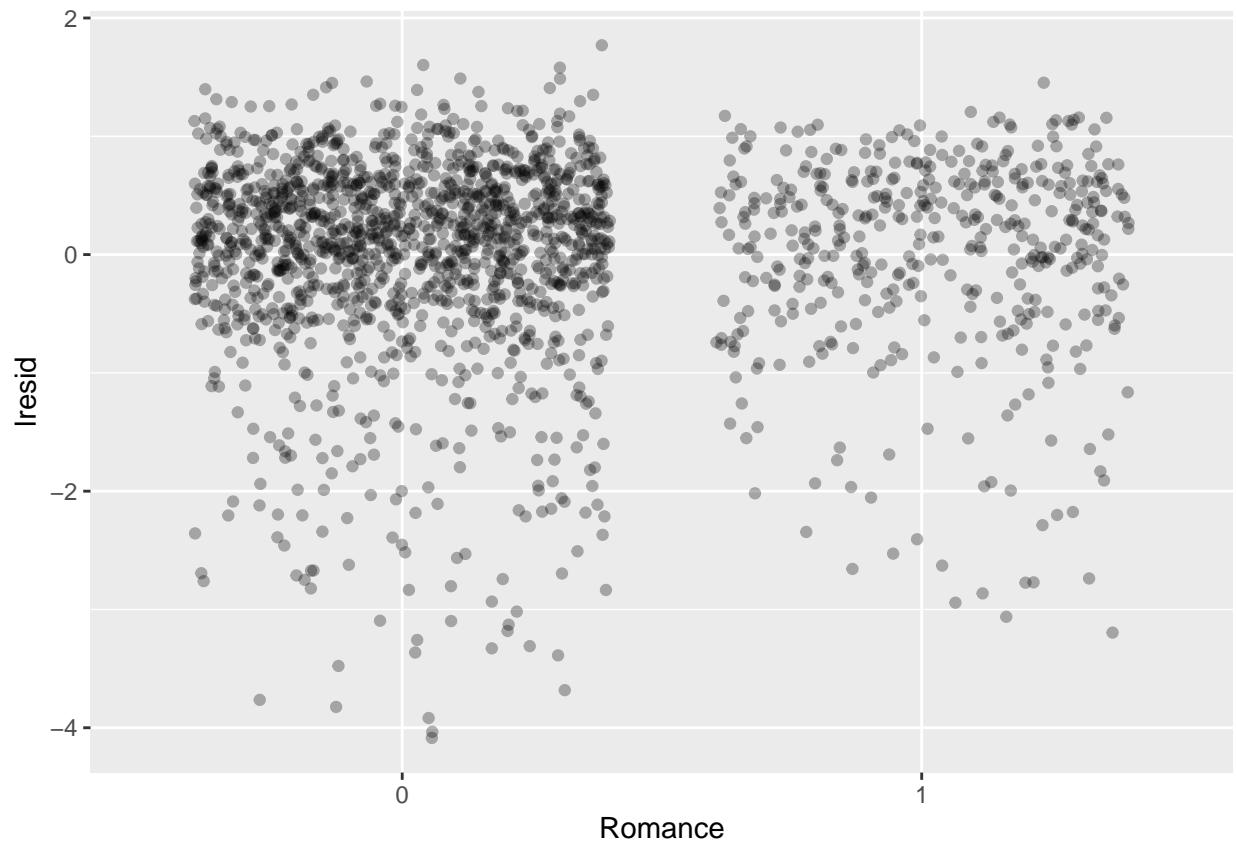
```
##  
## [[6]]
```



```
##  
## [[7]]
```

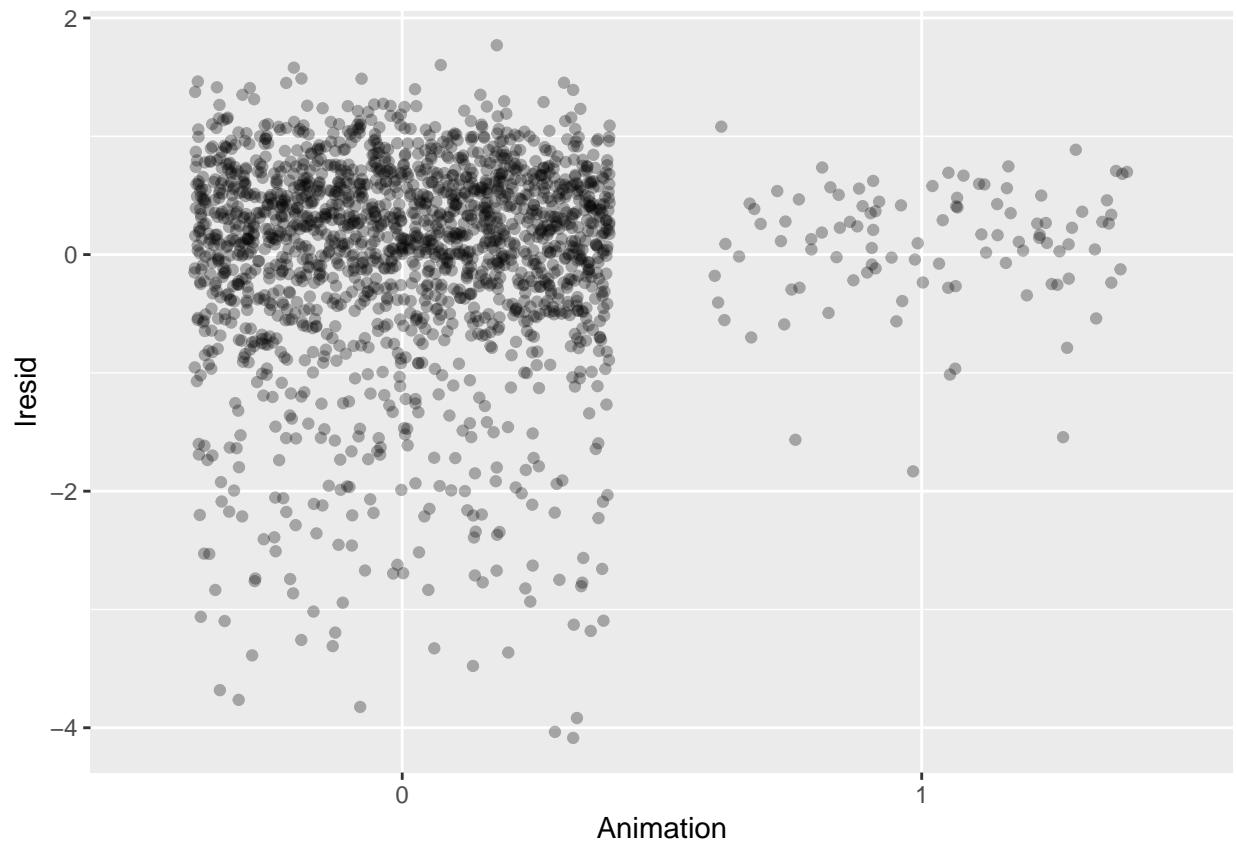


```
##  
## [[8]]
```

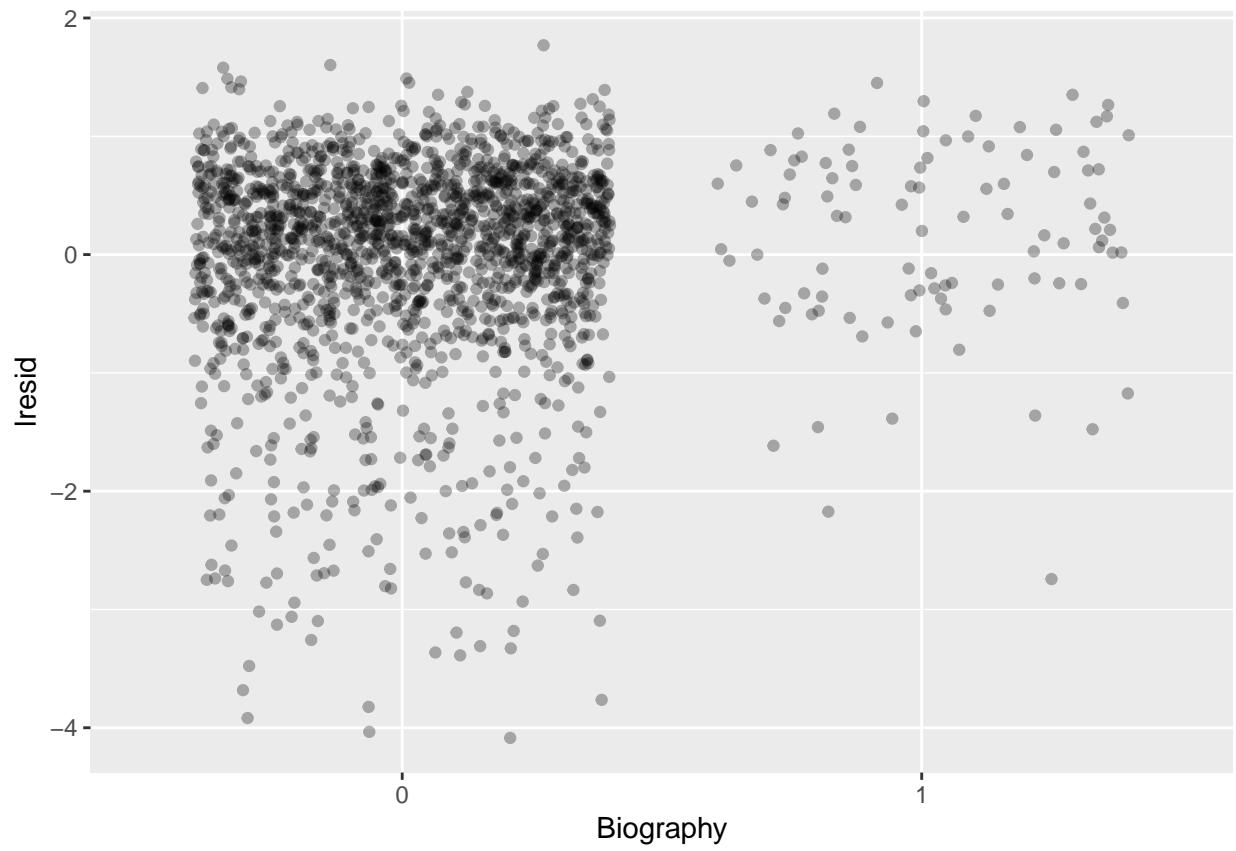


```
# graph residuals against each genre not included in the model
lapply(names(train_genre_only %>% select(-genre_xvar)), function(var) {
  train_resid %>%
    ggplot() +
    geom_jitter(aes_string(var, y = 'lresid'), alpha = .3)
})
```

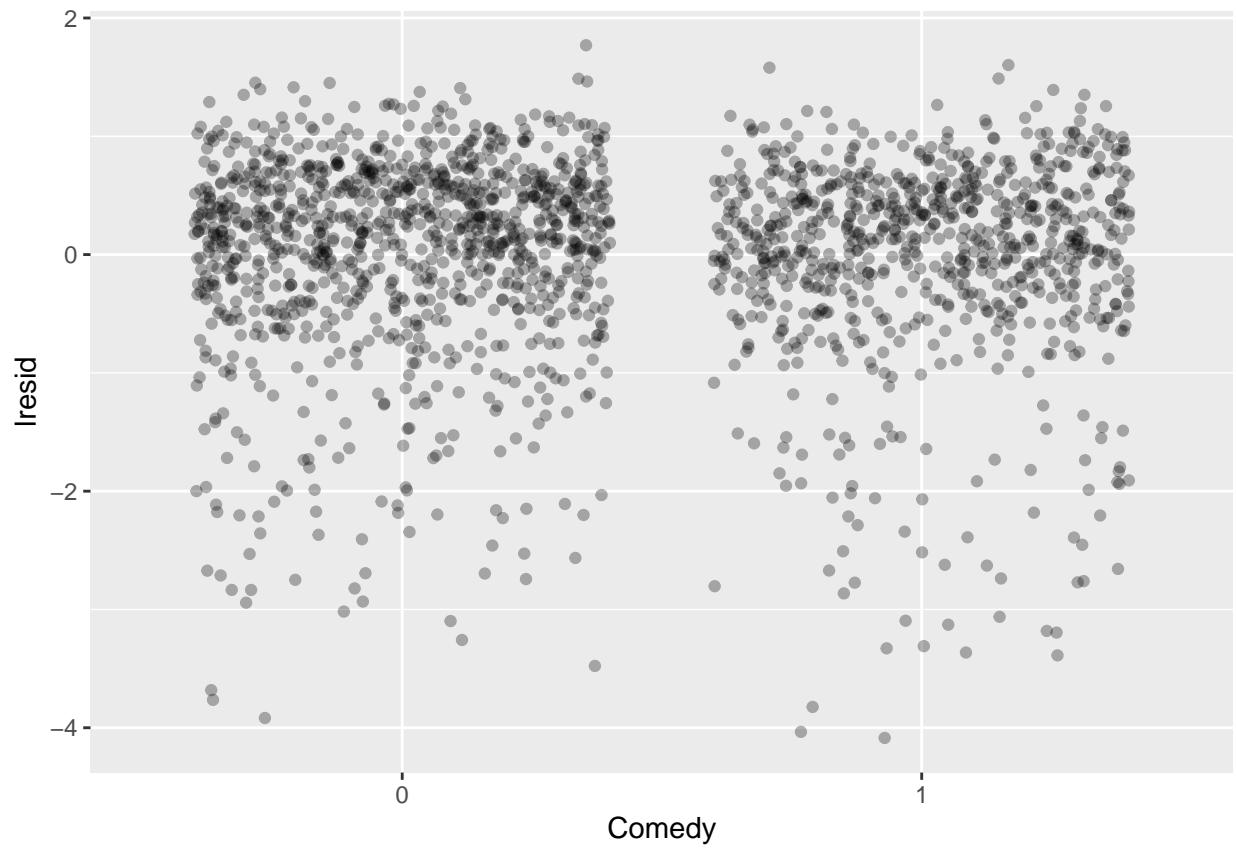
[[1]]



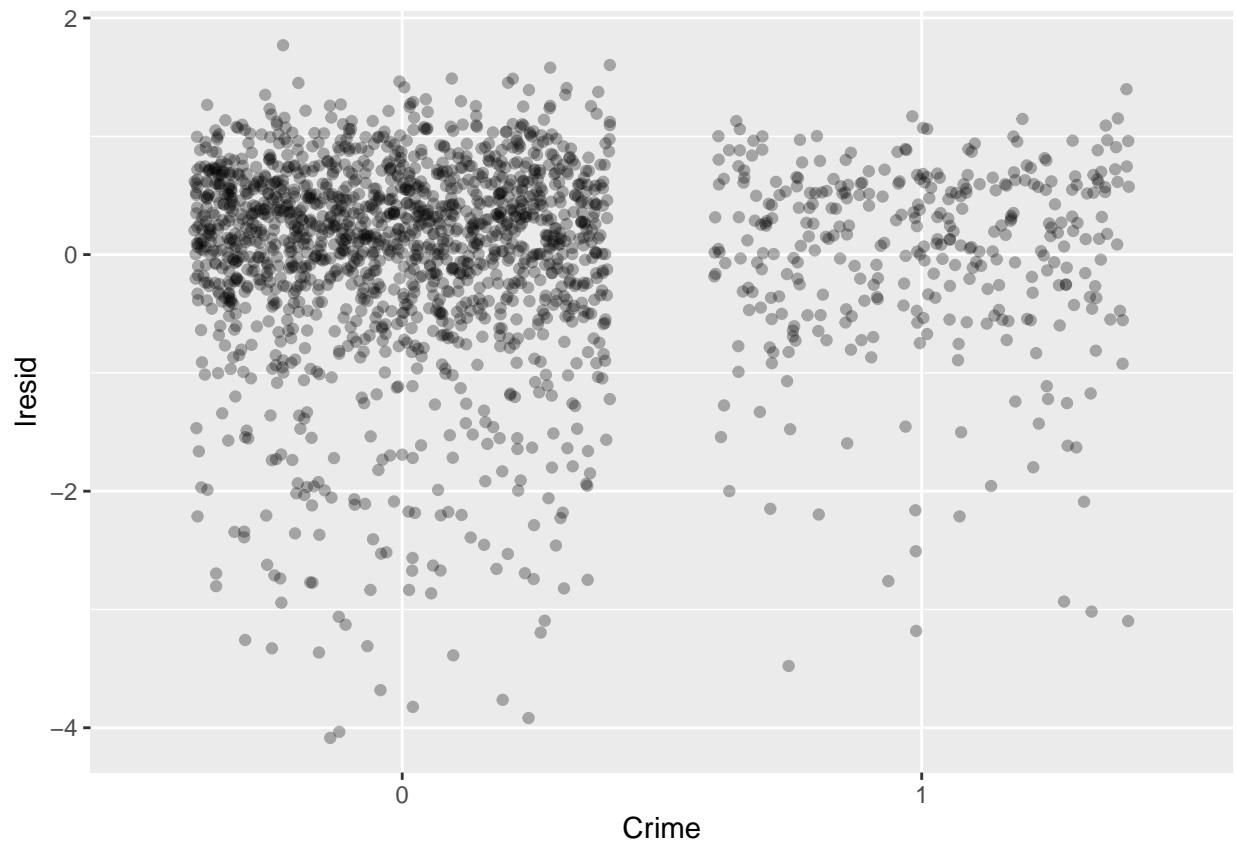
```
##  
## [[2]]
```



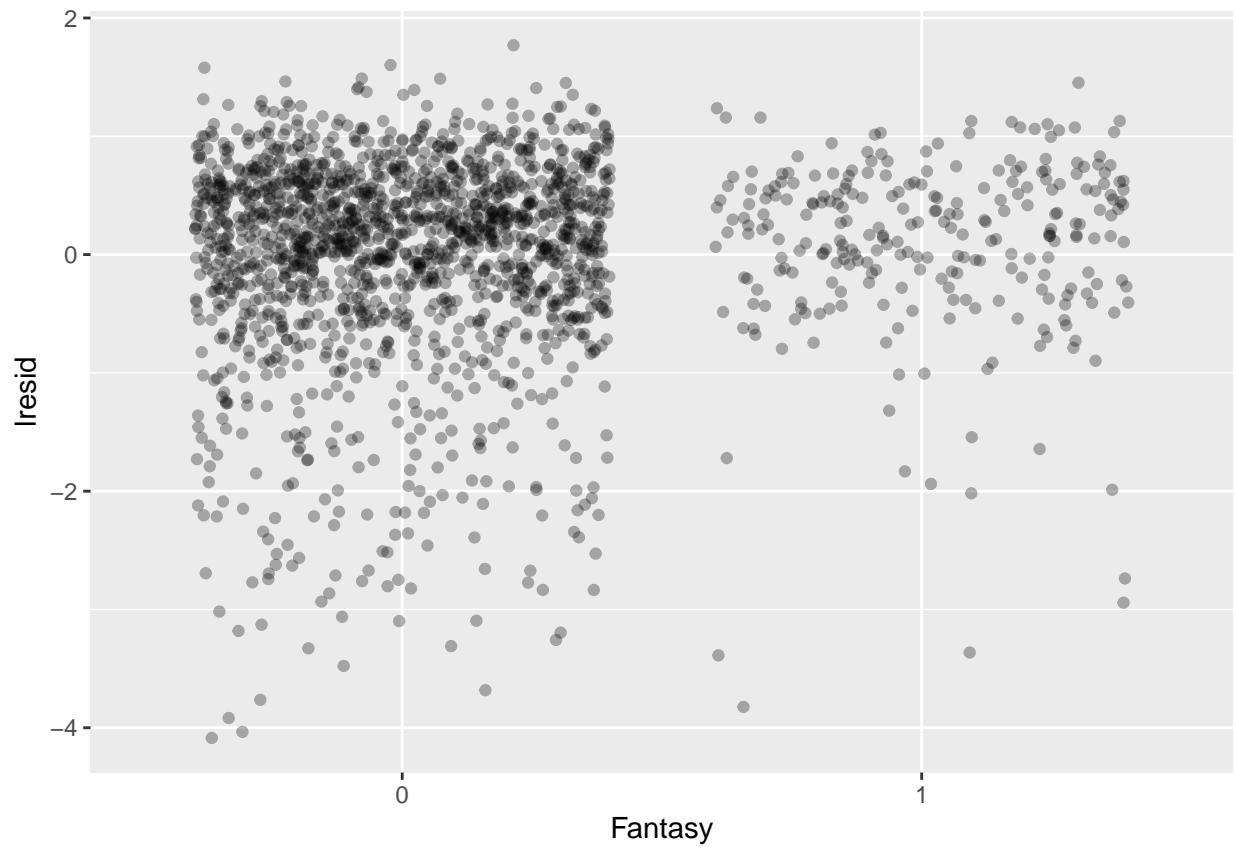
```
##  
## [3]
```



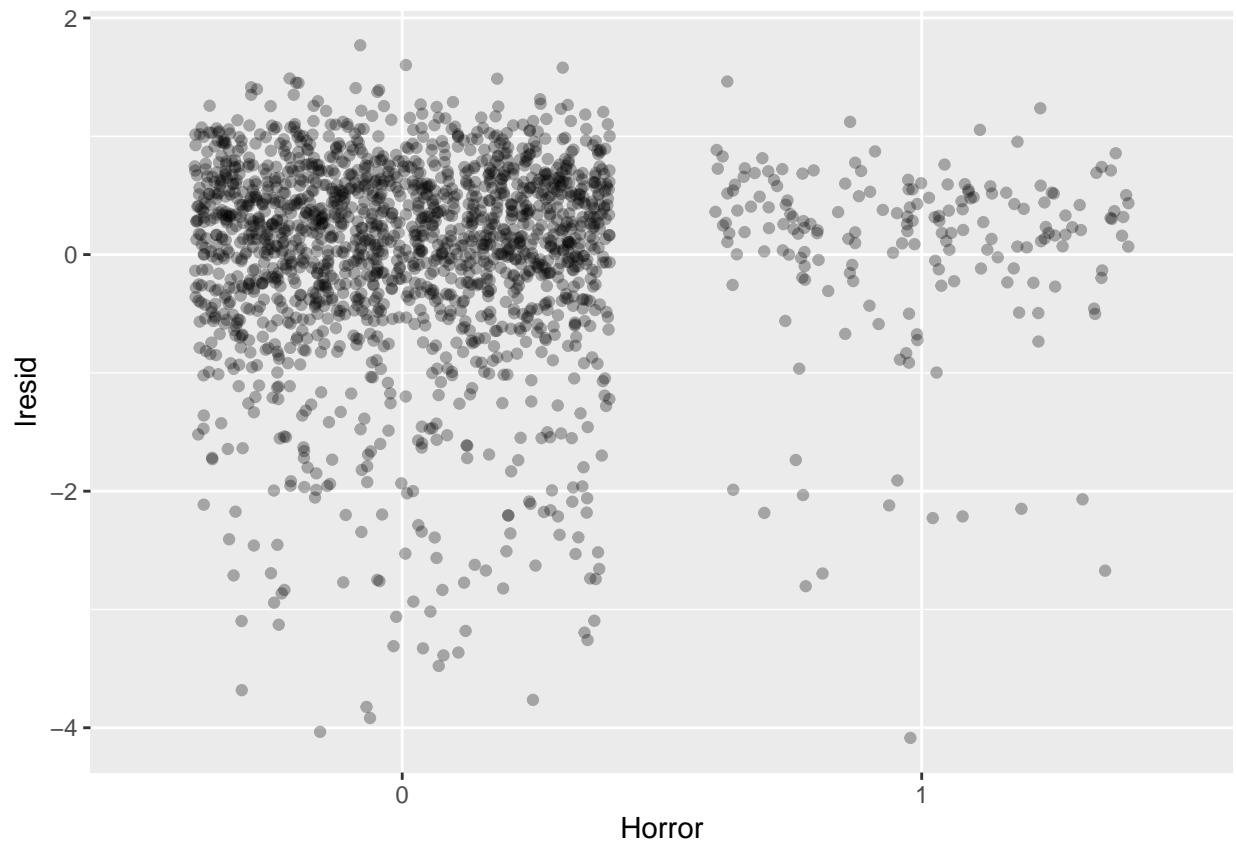
```
##  
## [[4]]
```



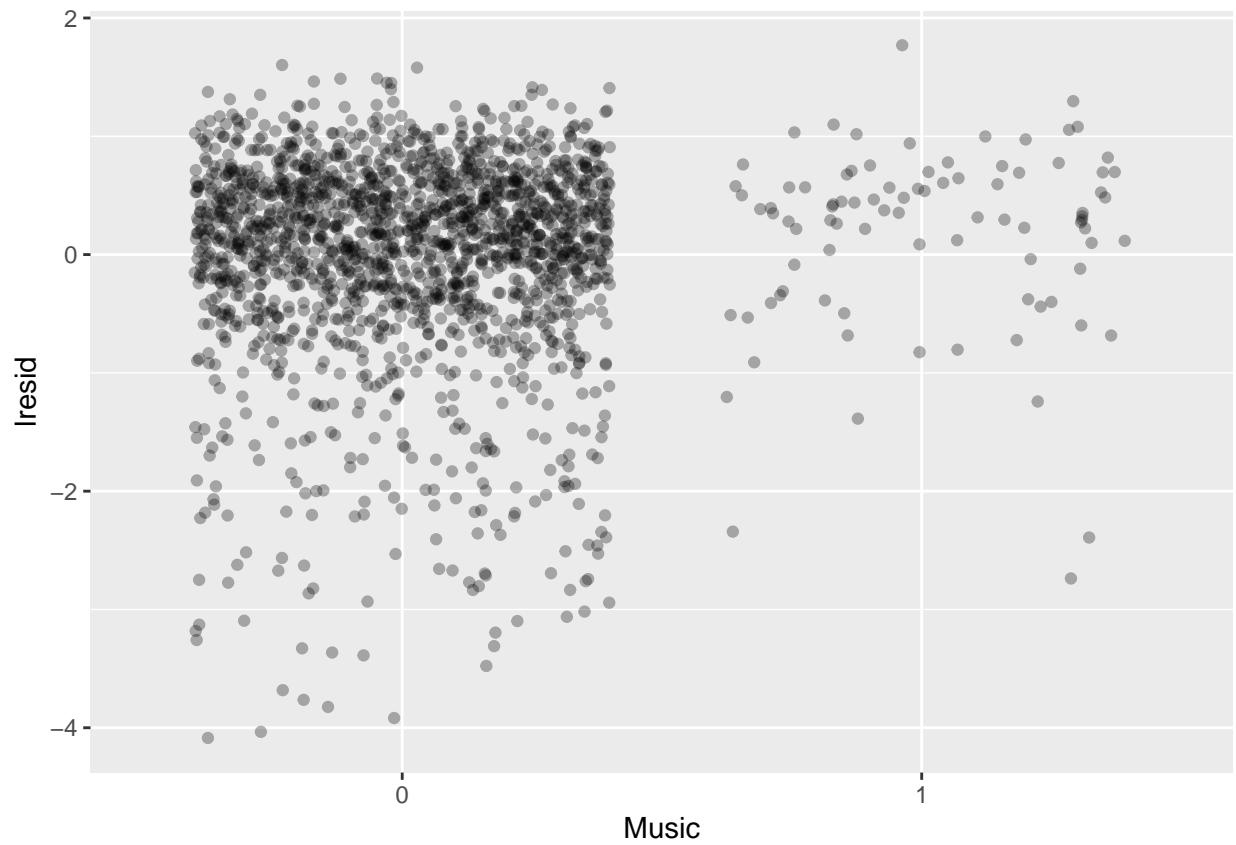
```
##  
## [[5]]
```



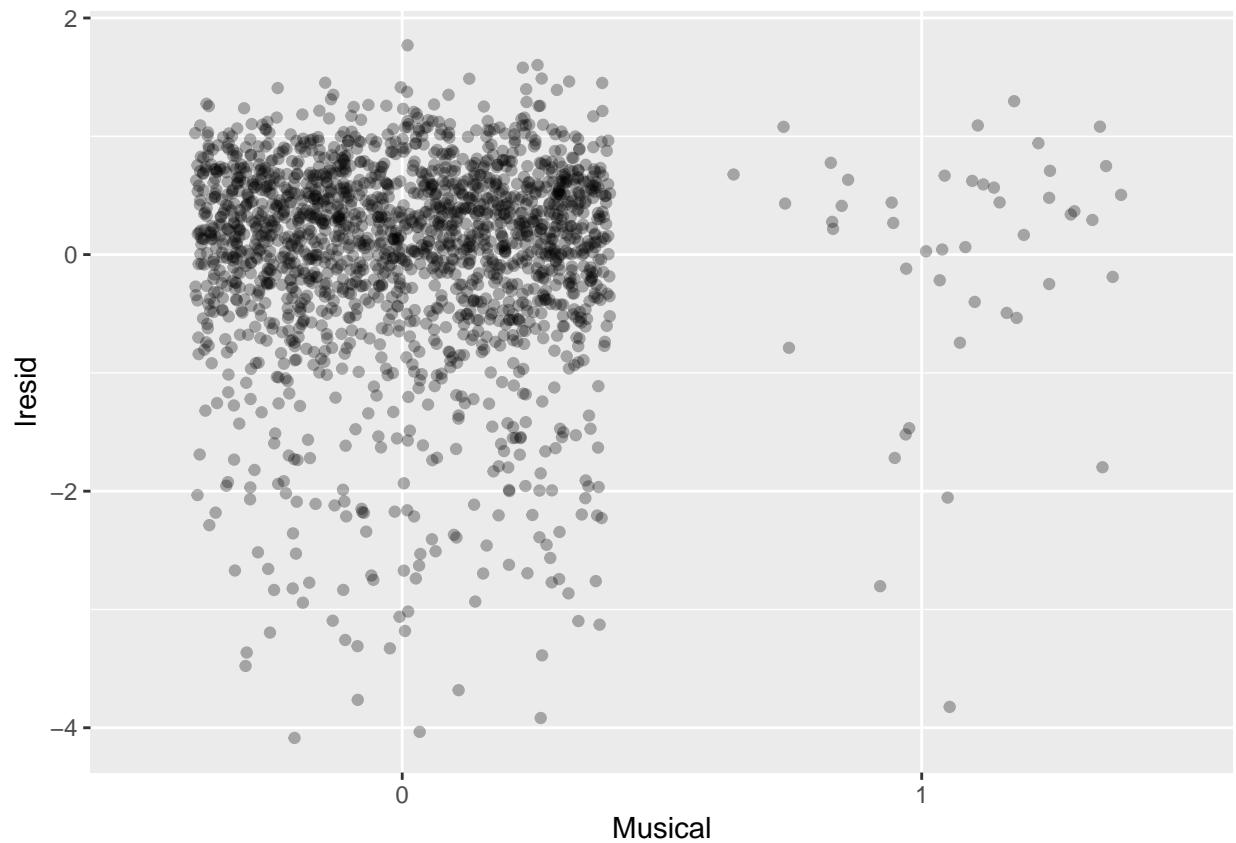
```
##  
## [[6]]
```



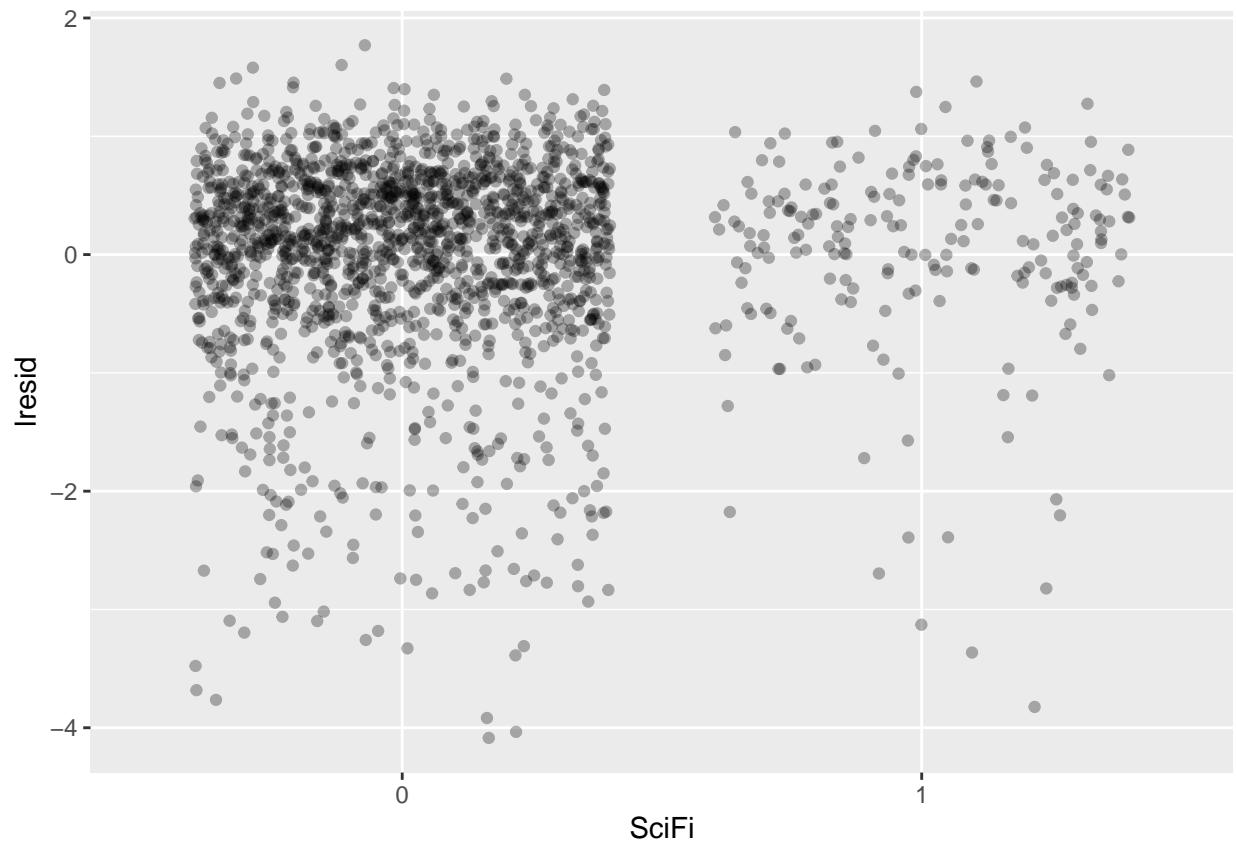
```
##  
## [[7]]
```



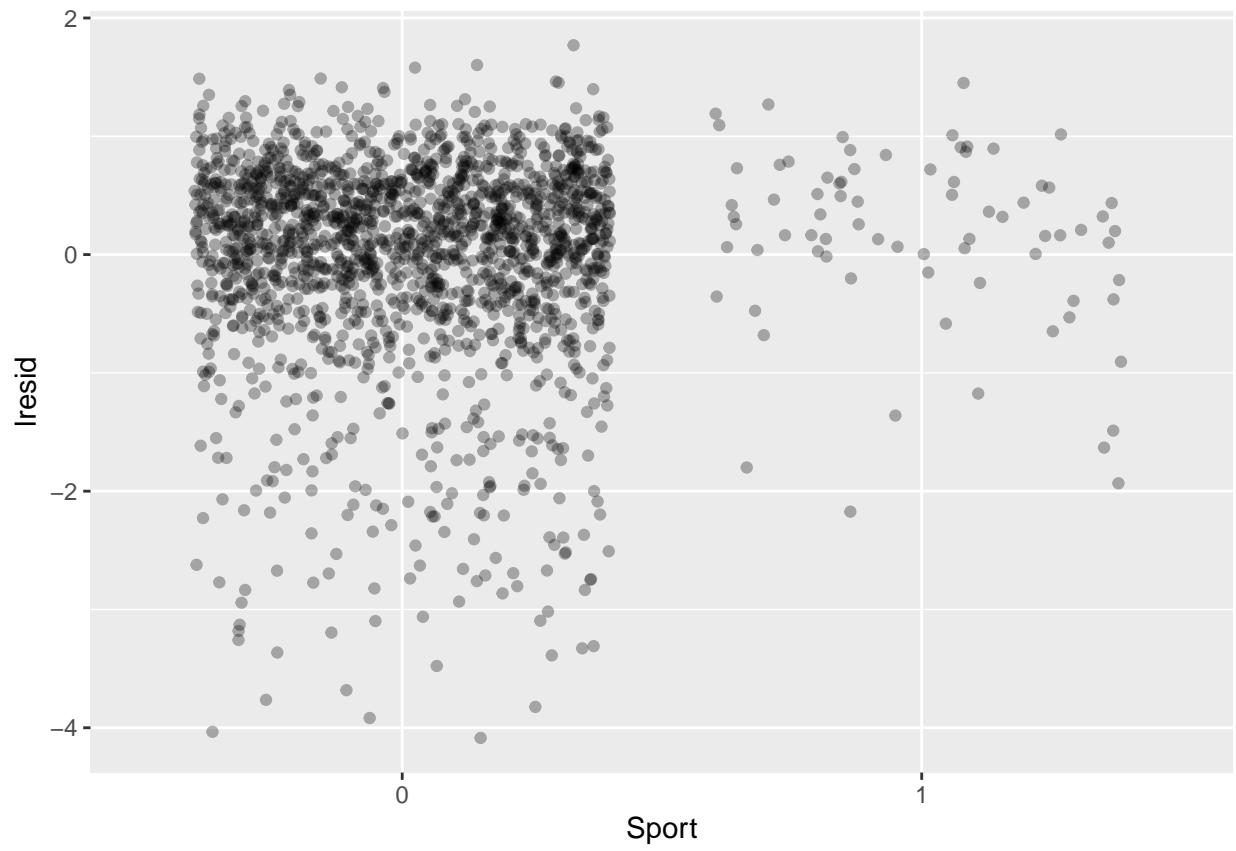
```
##  
## [8]
```



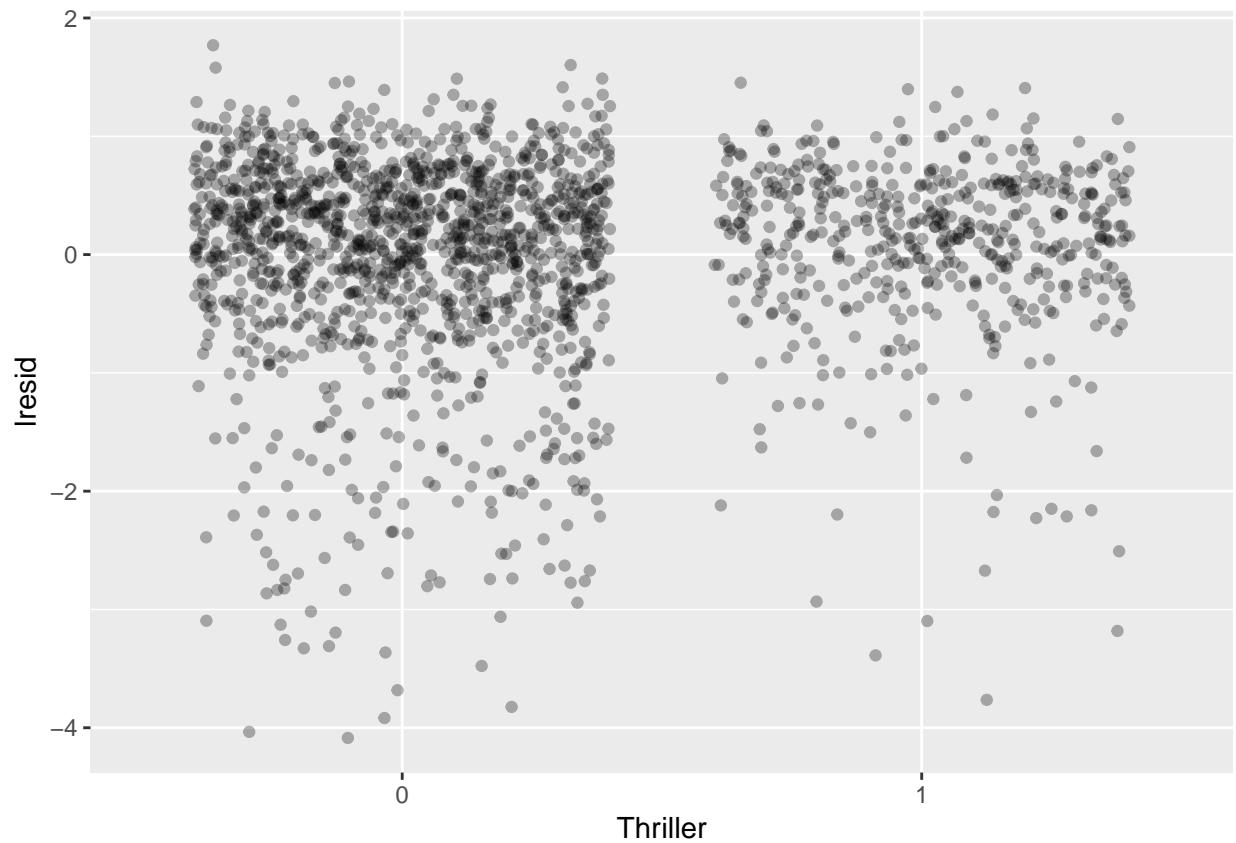
```
##  
## [[9]]
```



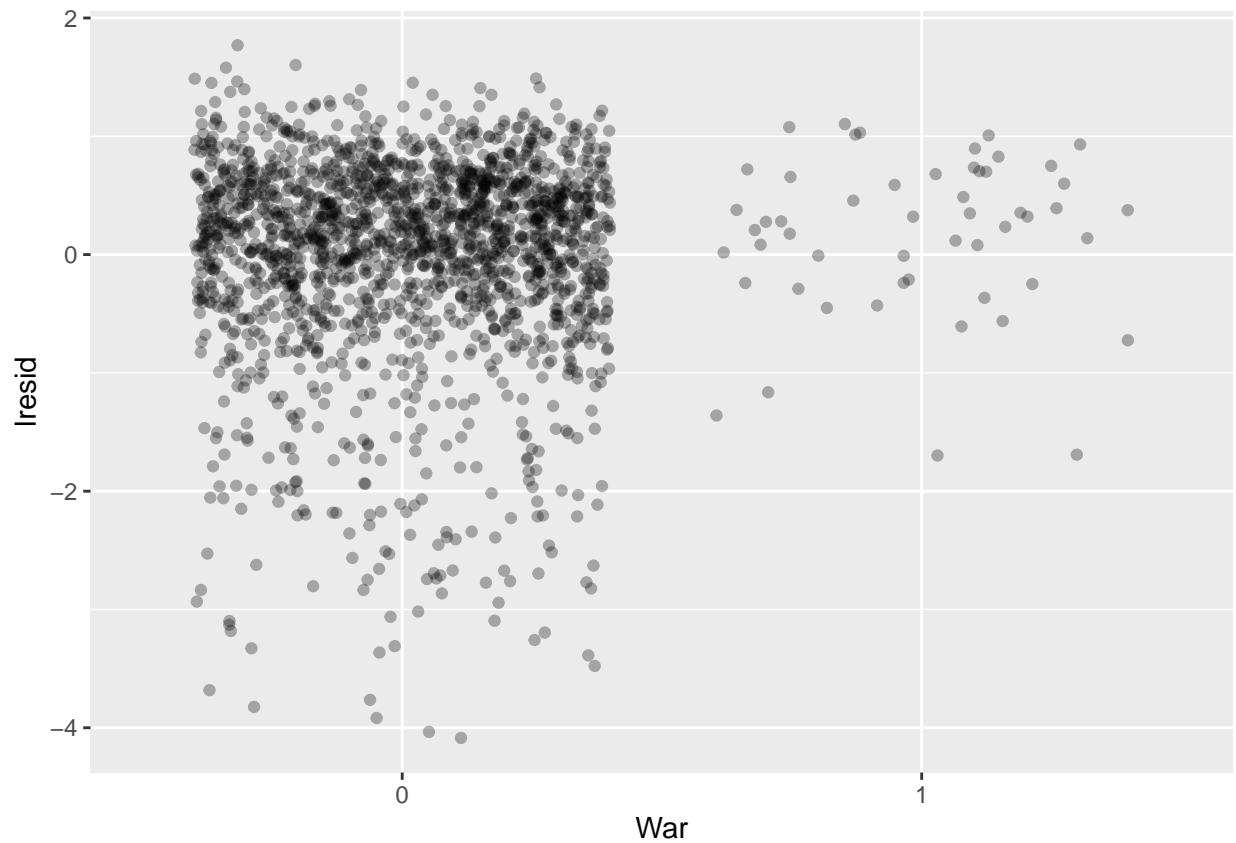
```
##  
## [[10]]
```



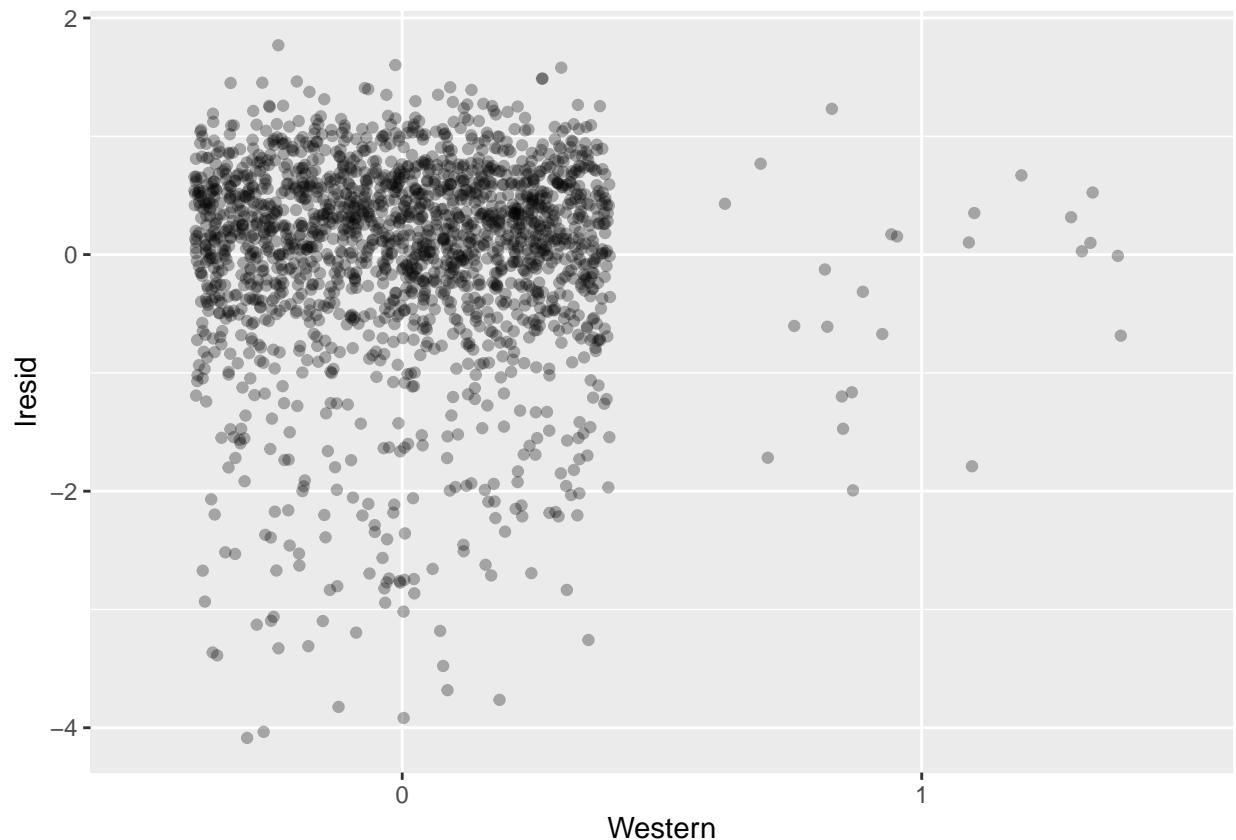
```
##  
## [[11]]
```



```
##  
## [[12]]
```

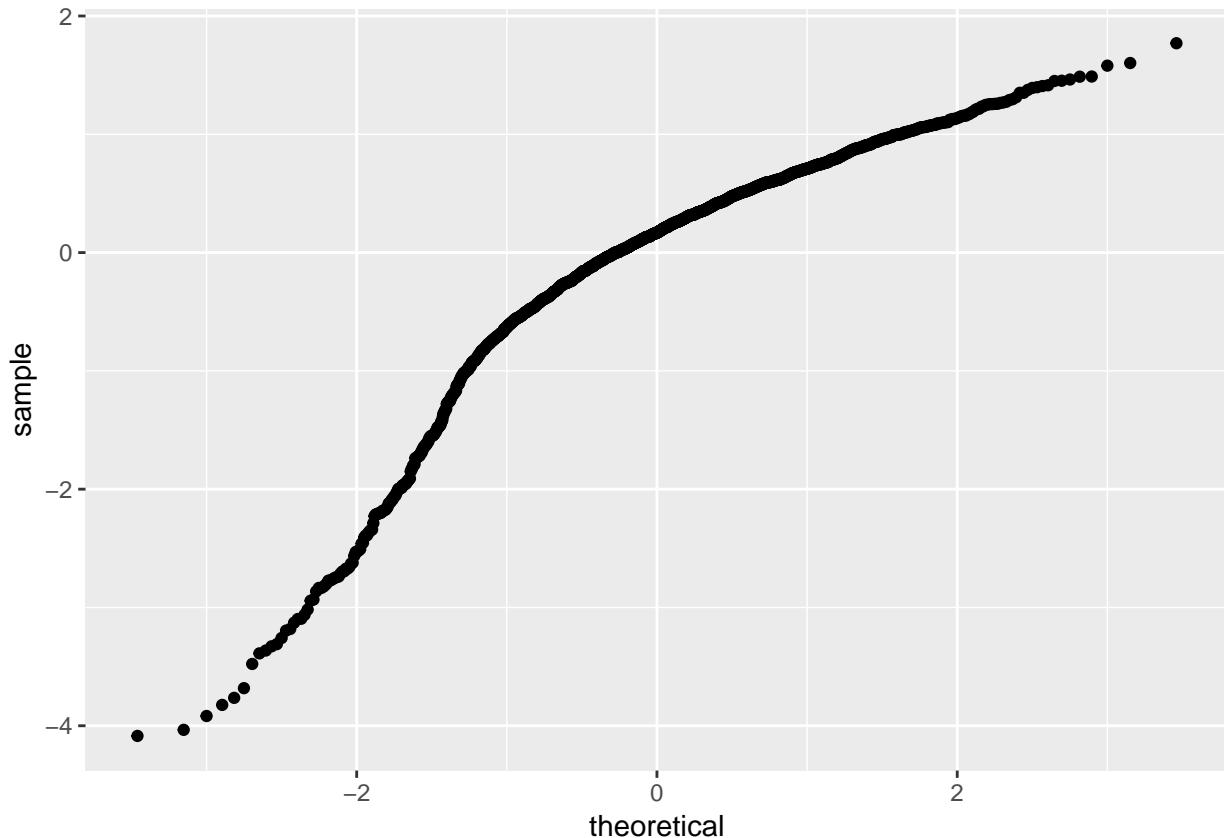


```
##  
## [13]
```



Plot QQ plot for residuals. Not normally distributed, but close-ish.

```
# residuals themselves are NOT normally distributed
# qq plot
train_resid %>% ggplot() +
  geom_qq(aes(sample = lresid))
```

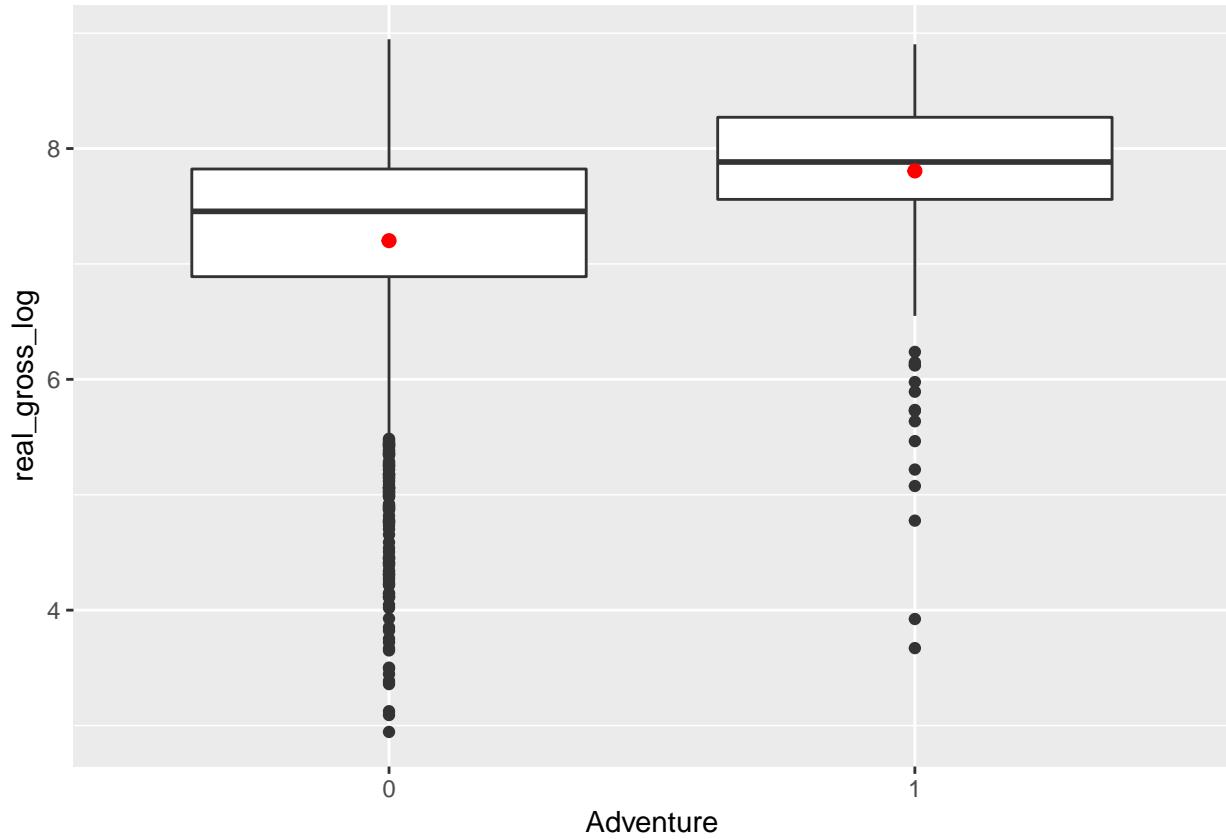


Plot Predictions

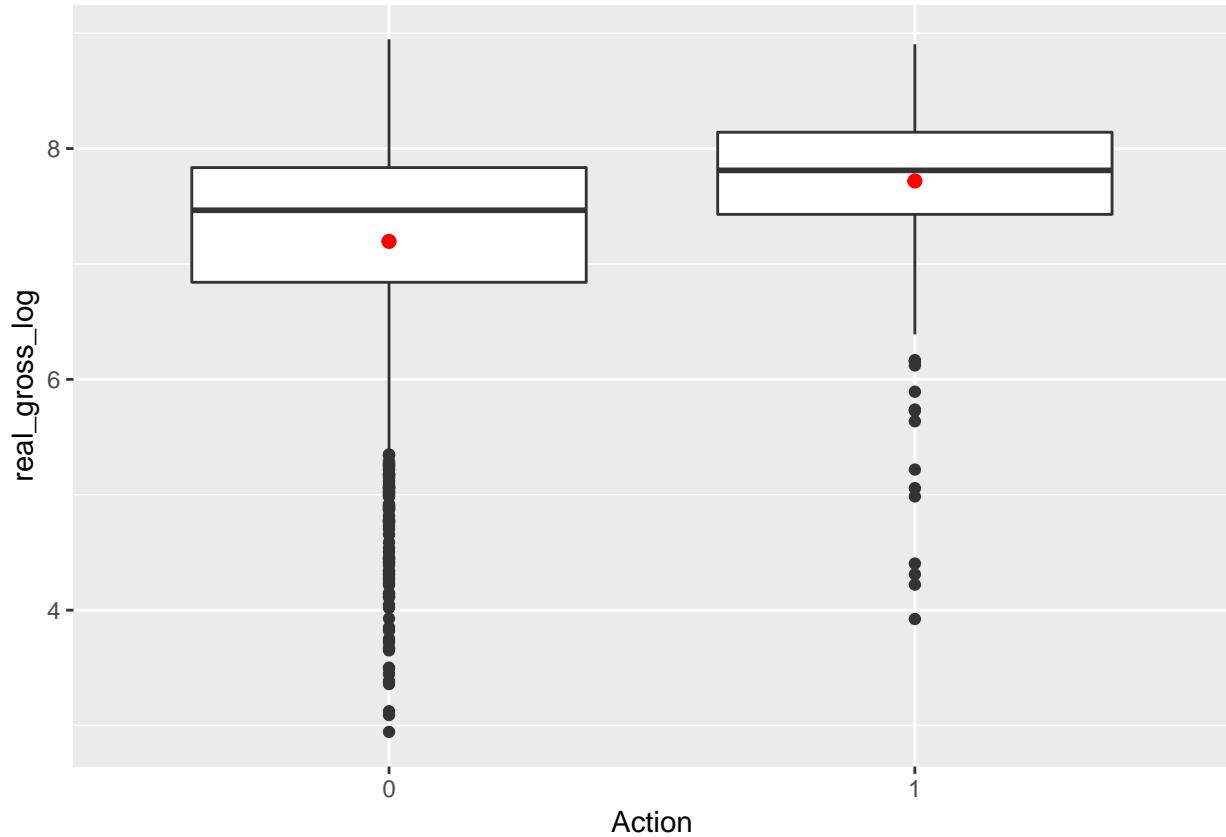
Plot prediction for mean real revenue against each genre included in the model.

```
train_pred <- train %>% add_predictions(mod_genre, 'lpred')

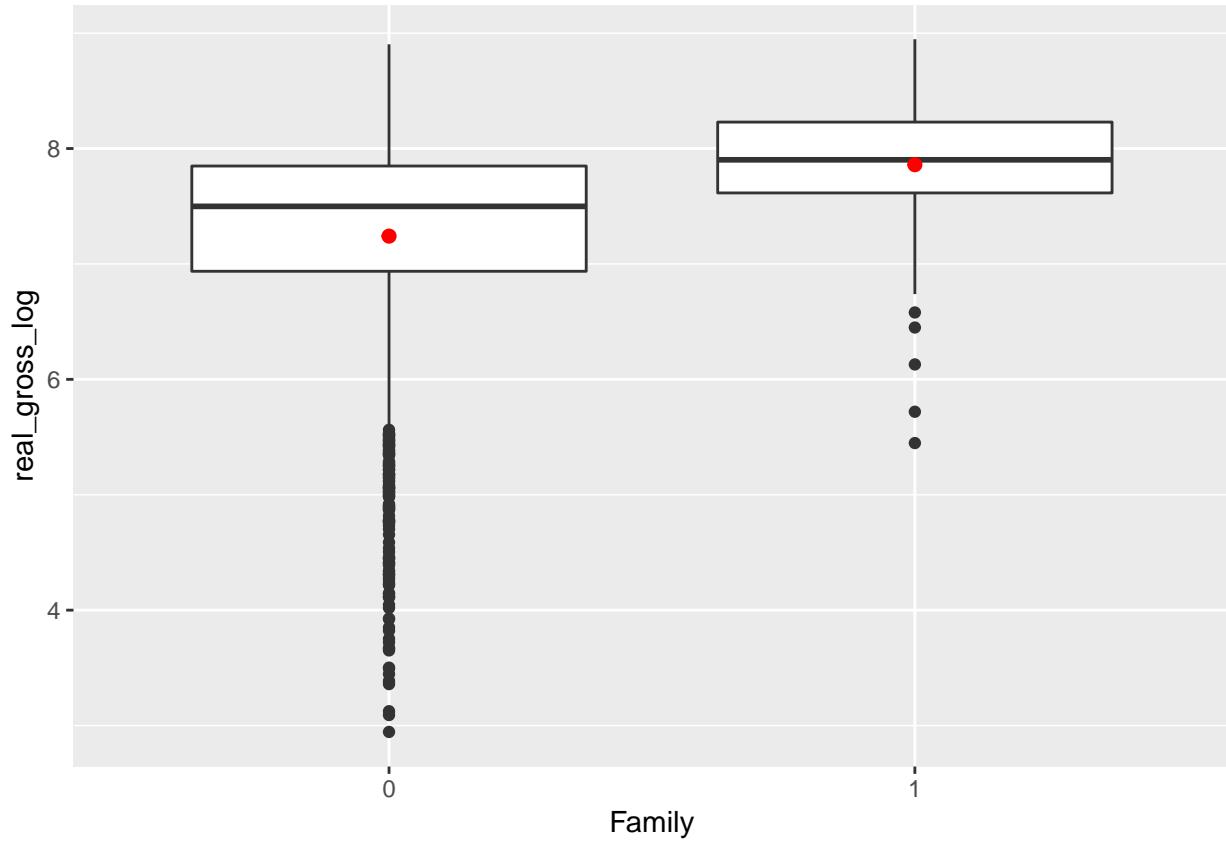
lapply(genre_xvar, function(var) {
  train_pred %>%
    ggplot(aes_string(x = var)) +
    geom_boxplot(aes(y = real_gross_log)) +
    geom_point(data = train_pred %>% group_by(!rlang::sym(var)) %>% summarize(mean = mean(lpred)),
               aes(y = mean), color = 'red', size = 2)
})  
## [[1]]
```



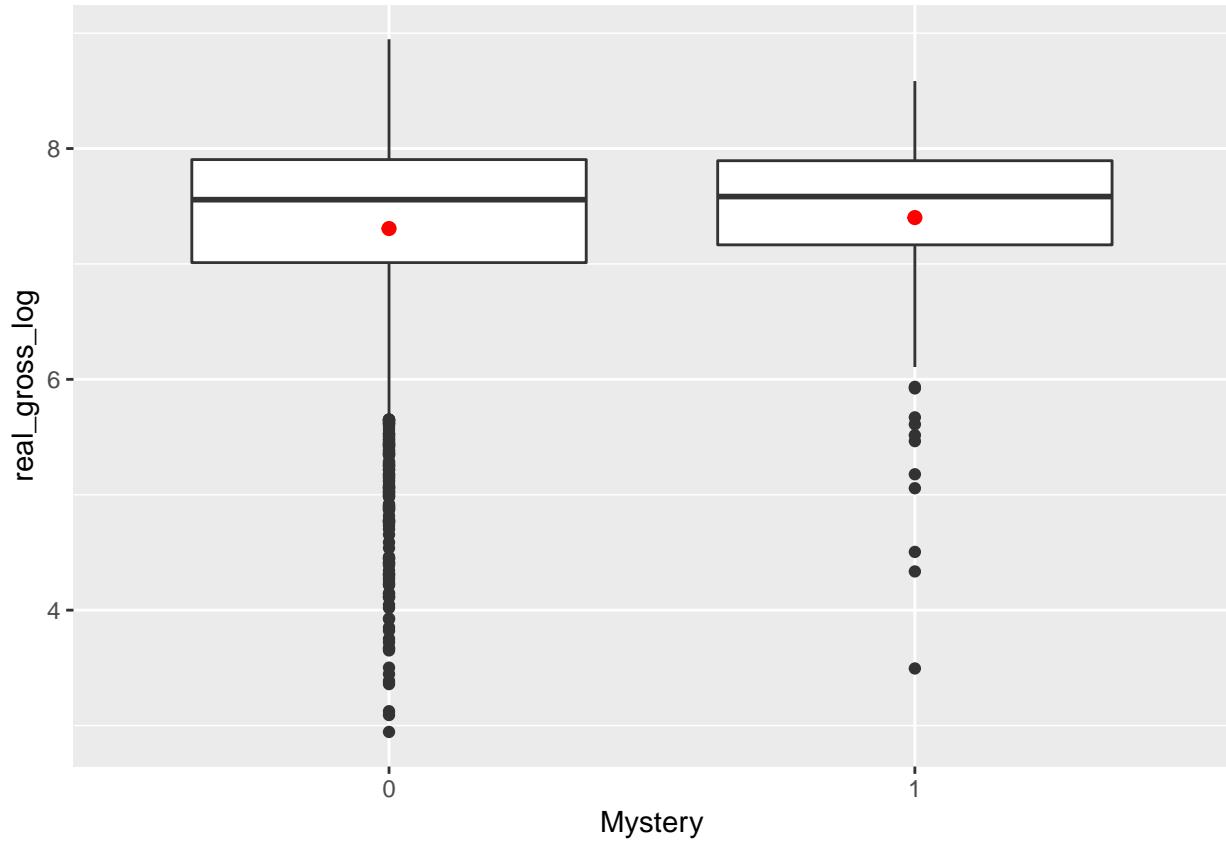
```
##  
## [[2]]
```



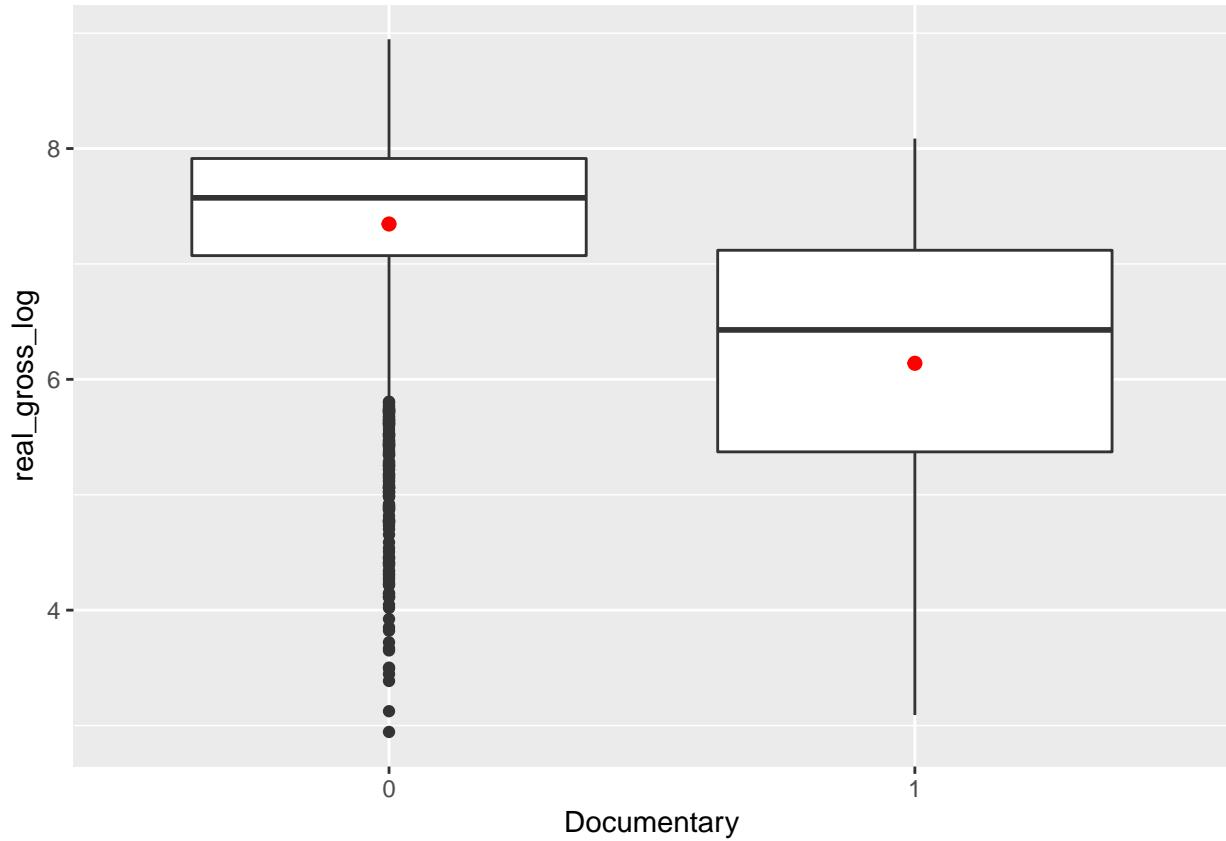
```
##  
## [[3]]
```



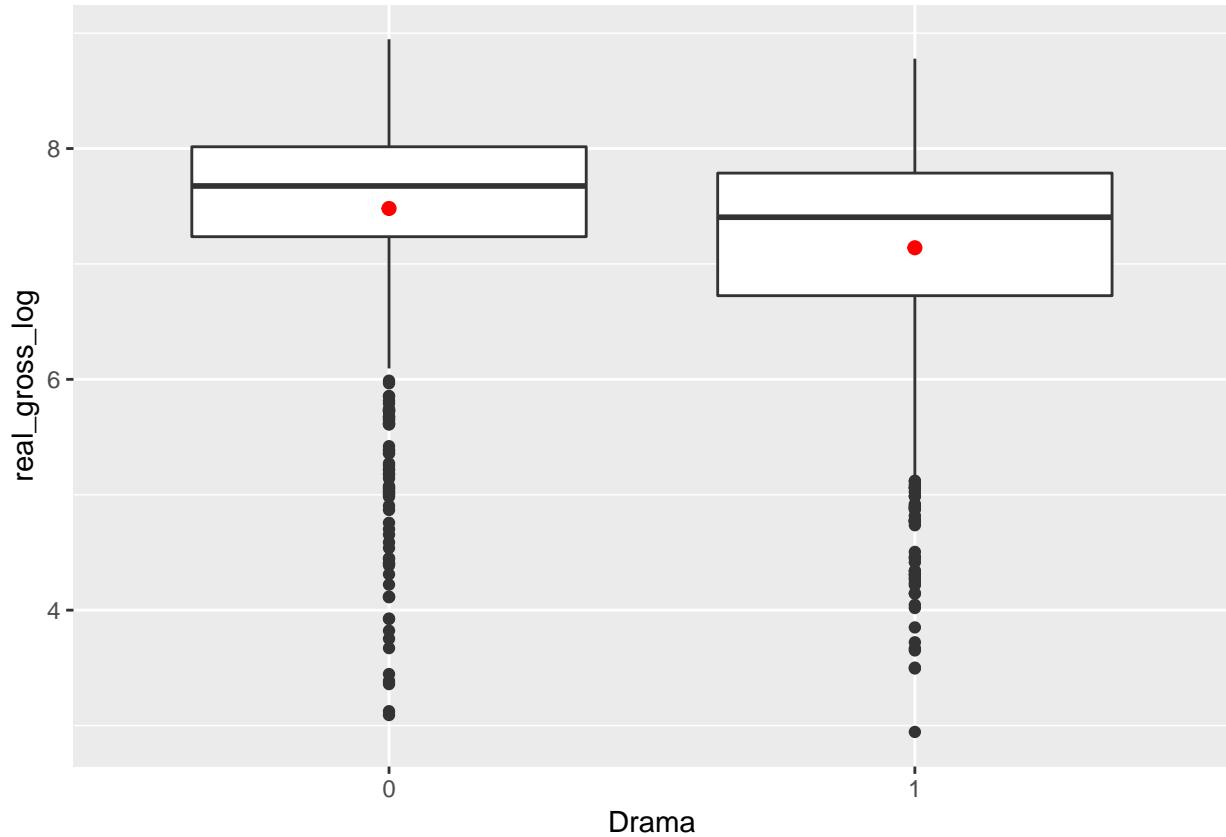
```
##  
## [[4]]
```



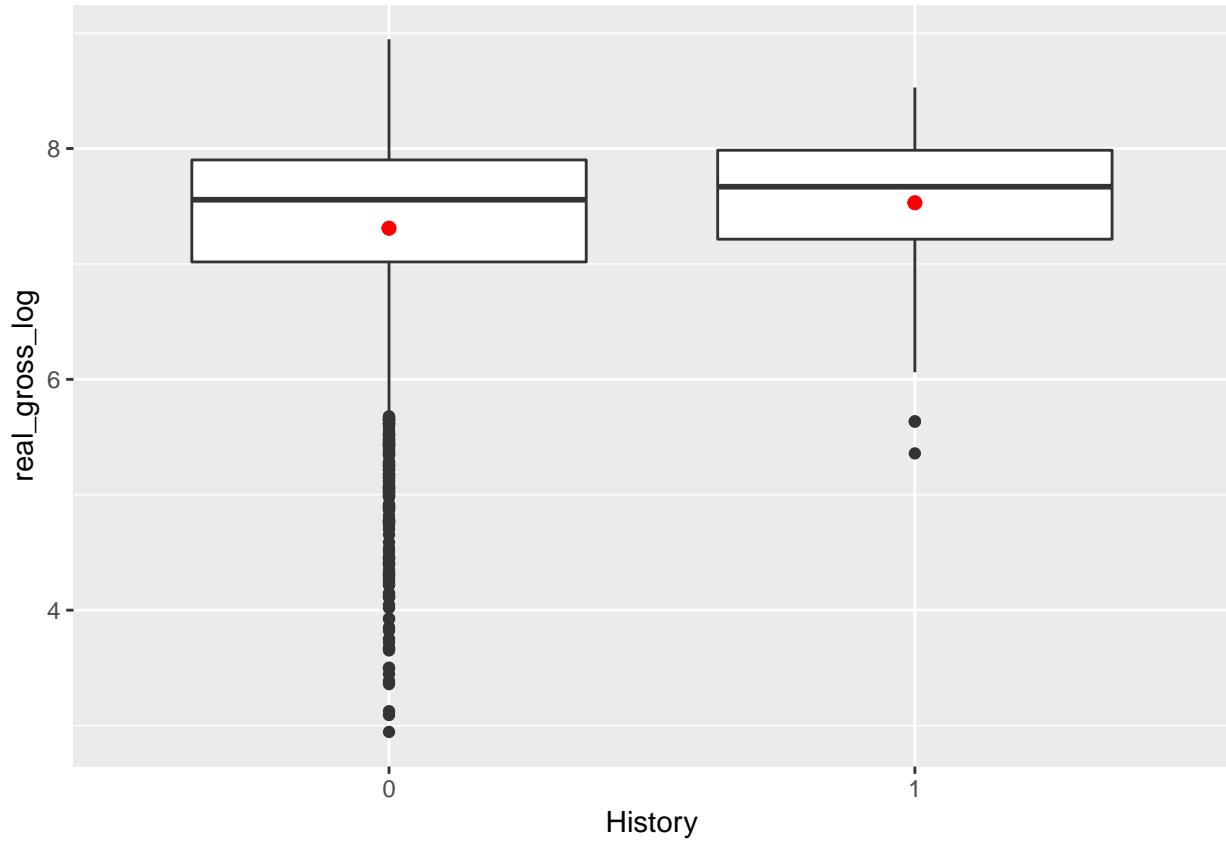
```
##  
## [[5]]
```



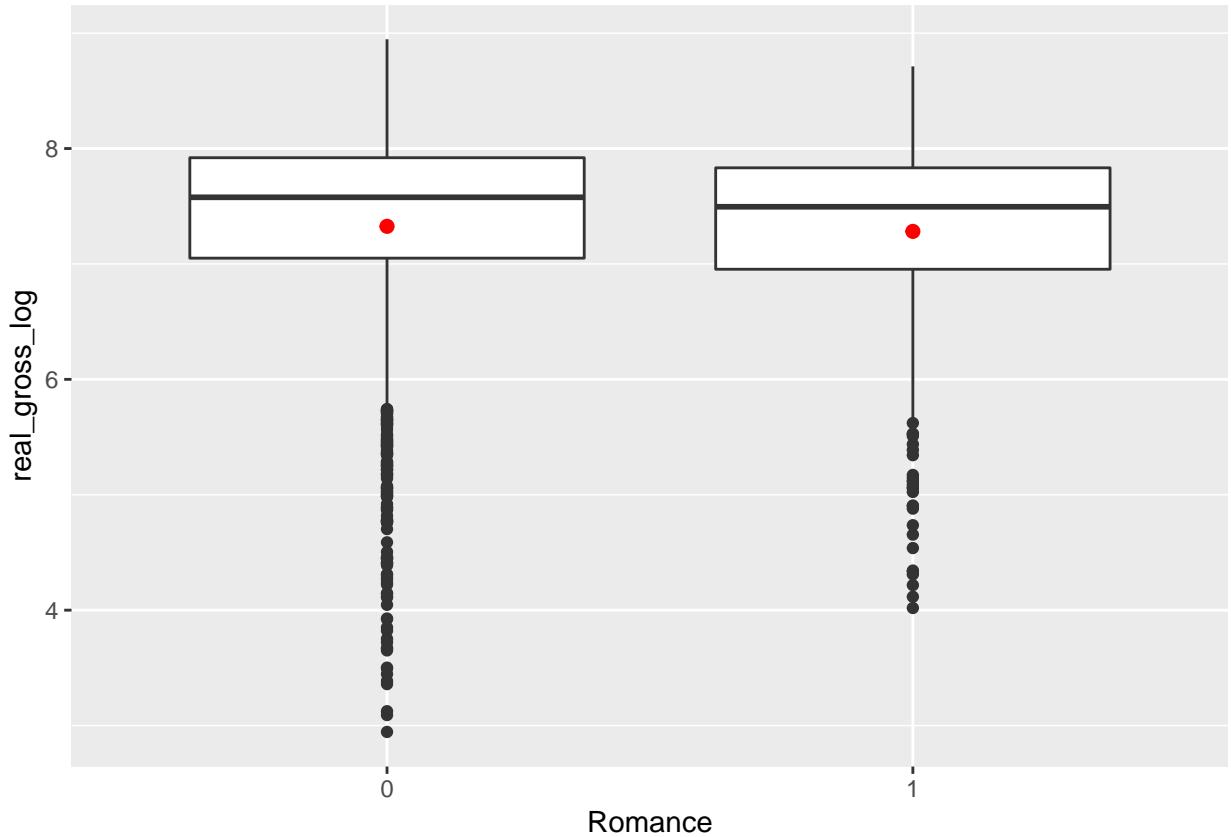
```
##  
## [[6]]
```



```
##  
## [[7]]
```



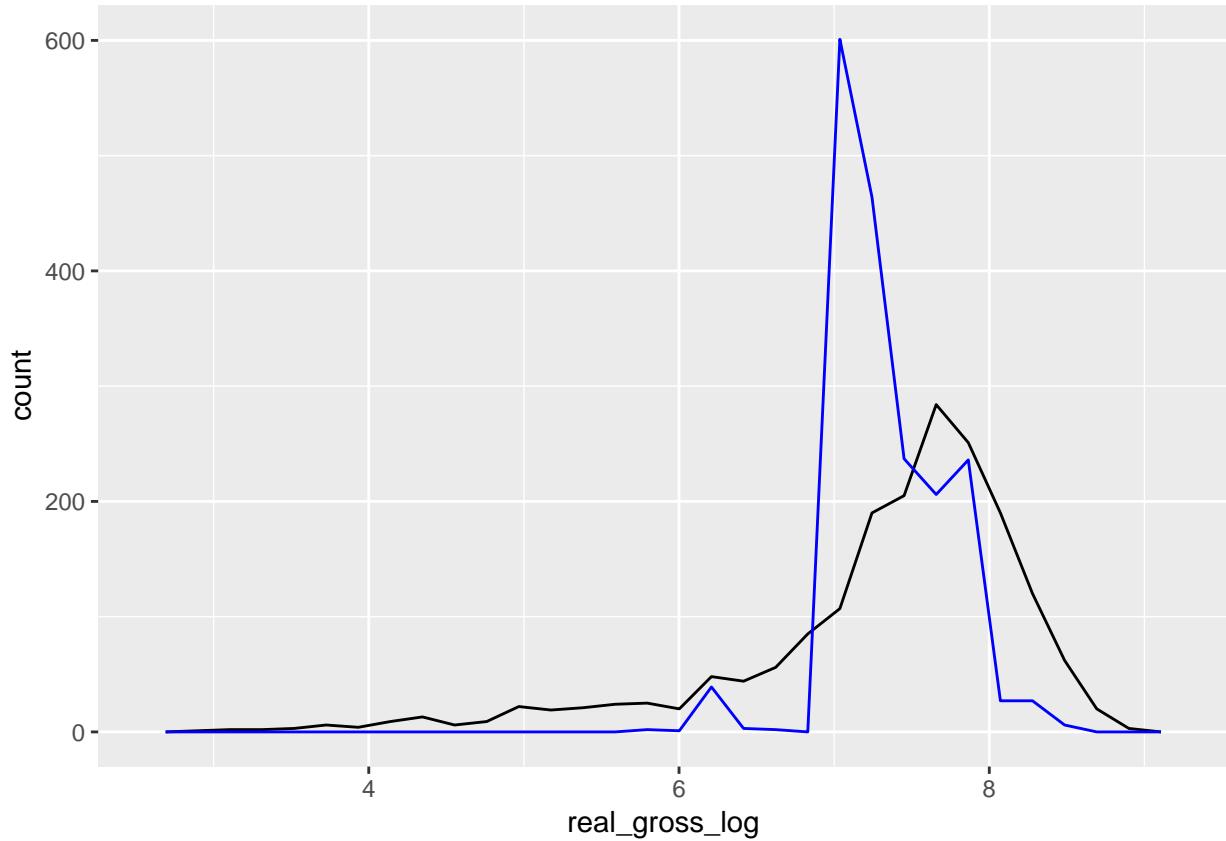
```
##  
## [[8]]
```



Overall predictions: clearly not enough to just specify genres

```
train %>%
  add_predictions(mod_genre, 'lpred') %>%
  ggplot() +
  geom_freqpoly(aes(x = real_gross_log)) +
  geom_freqpoly(aes(x = lpred), color = 'blue')

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Glmnet: sparse

Quickly try this new method from class instead of stepwise. The sparse version does give us a lot of the same variables as stepwise. Good sign!

Can't do statistical tests, so not useful for analysis, but can use to aid justification.

```
library(glmnet)

## Warning: package 'glmnet' was built under R version 3.5.3
## Loading required package: Matrix

##
## Attaching package: 'Matrix'
## The following object is masked from 'package:tidyverse':
##     expand
## Loading required package: foreach

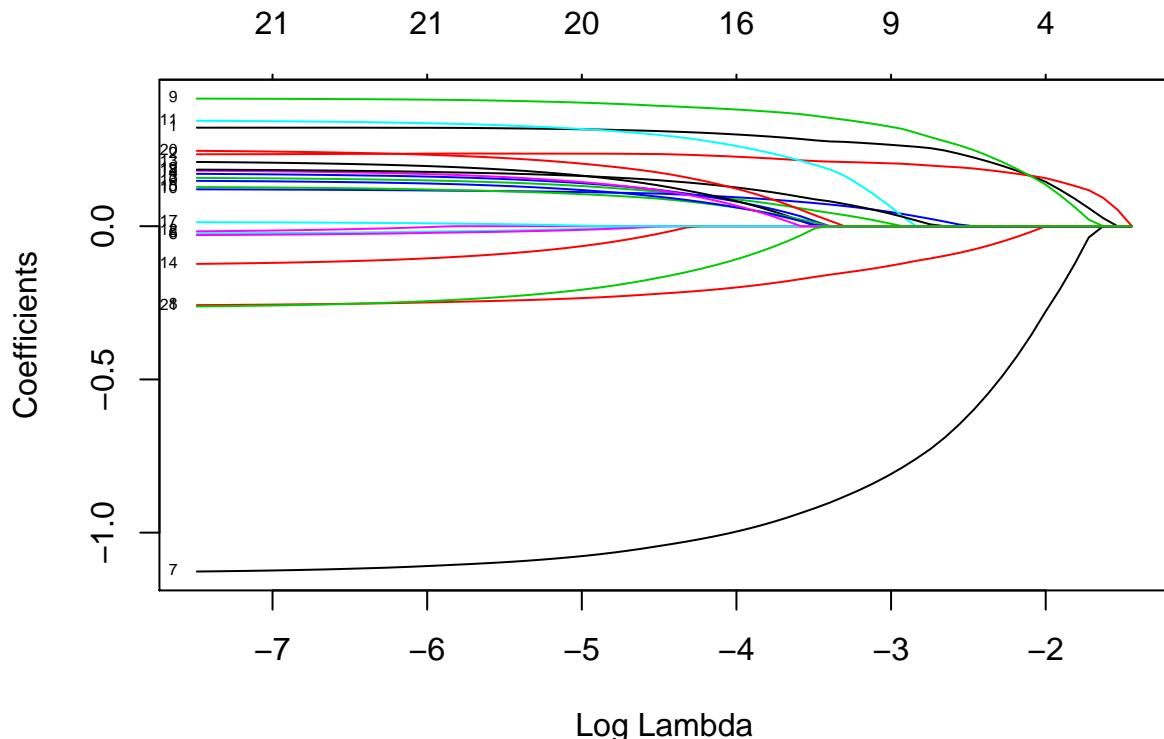
## Warning: package 'foreach' was built under R version 3.5.3
##
## Attaching package: 'foreach'
## The following objects are masked from 'package:purrr':
##     accumulate, when
## Loaded glmnet 2.0-16
```

```

# matrix of x and y variables
x <- as.matrix(train_genre_only %>% mutate_all(funs(as.numeric(as.character()))))
y <- as.matrix(train$real_gross_log)

# glmnet process form class
mod_sparse <- glmnet(x, y, family = 'gaussian')
plot(mod_sparse, xvar = 'lambda', label = TRUE)

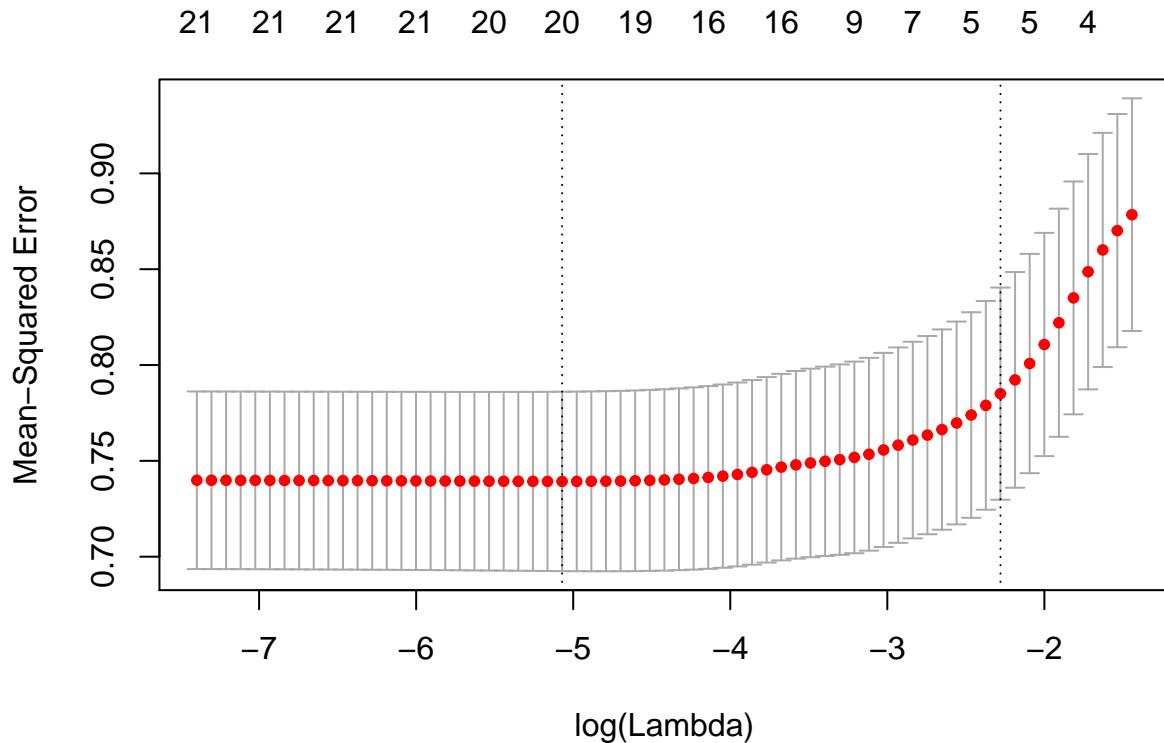
```



```

mod_sparse <- cv.glmnet(x, y)
plot(mod_sparse)

```



```
coef(mod_sparse, s = 'lambda.min') # use min lambda
```

```
## 22 x 1 sparse Matrix of class "dgCMatrix"
##           1
## (Intercept) 7.158319467
## Action      0.316866270
## Adventure   0.236784610
## Animation   0.133081838
## Biography   0.142672310
## Comedy      -0.009843859
## Crime       -0.012866133
## Documentary -1.080069143
## Drama       -0.236062771
## Family      0.404262630
## Fantasy     0.111824039
## History     0.319273704
## Horror      .
## Music       0.168382172
## Musical     -0.069374568
## Mystery     0.106229872
## Romance    0.120292330
## SciFi       0.002684535
## Sport       0.146848534
## Thriller    0.165198422
## War         0.207283391
## Western    -0.211255187
```

```

coef(mod_sparse, s = 'lambda.1se') # use most sparse

## 22 x 1 sparse Matrix of class "dgCMatrix"
##                1
## (Intercept) 7.24823459
## Action      0.19684629
## Adventure   0.17743337
## Animation   .
## Biography   .
## Comedy      .
## Crime       .
## Documentary -0.48840116
## Drama       -0.04977566
## Family      0.20946931
## Fantasy    .
## History    .
## Horror     .
## Music      .
## Musical    .
## Mystery    .
## Romance    .
## SciFi      .
## Sport      .
## Thriller   .
## War        .
## Western    .

```

Fit model with other variables

Plot residuals of other variables based on the genre model.

All of these plots indicate a relationship that is not fully represented in the model yet and thus all are valid candidates for including in the model (also given their relatively linear relationships)

For example, movies with lower budgets make less revenue than predicted by the genres in the model (negative residual) and movies with higher budgets make more revenue than predicted by genre (positive residual). Cast facebook likes, director facebook likes, and IMDB score follow a similar pattern. Many years have revenue either higher or lower than that predicted by genre.

Content rating has more of a random relationship with the residual. Perhaps this is because genres and content ratings have some correlation (i.e. Family movies tend to be G or PG while Horror tend to be R) and thus this relationship may have already been captured.

There is some indication that R movies may make less revenue than predicted and PG-13 movies make more revenue than predicted.

```

# get log versions of variables since residuals are log: same scale
train_resid <- train_resid %>%
  mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
  'imdb_score'), funs(log = log10(.)))

# graph each against log residual: continuous
lapply(c('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
  'imdb_score'), function(var) {
  print(var)
  train_resid %>%

```

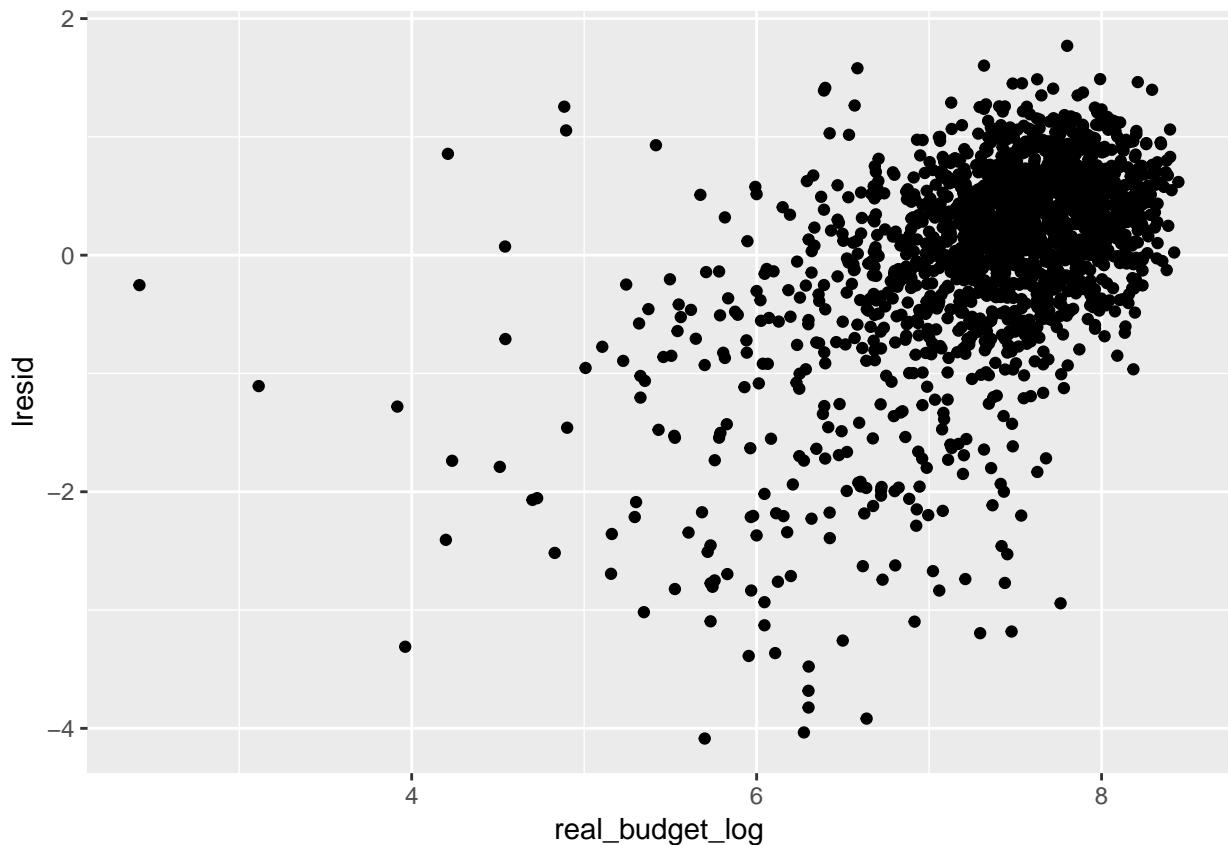
```

ggplot() +
  geom_point(aes_string(str_c(var, '_log'), y = 'lresid'))
}

## [1] "real_budget"
## [1] "director_facebook_likes"
## [1] "cast_total_facebook_likes"
## [1] "imdb_score"
## [[1]]

## Warning: Removed 102 rows containing missing values (geom_point).

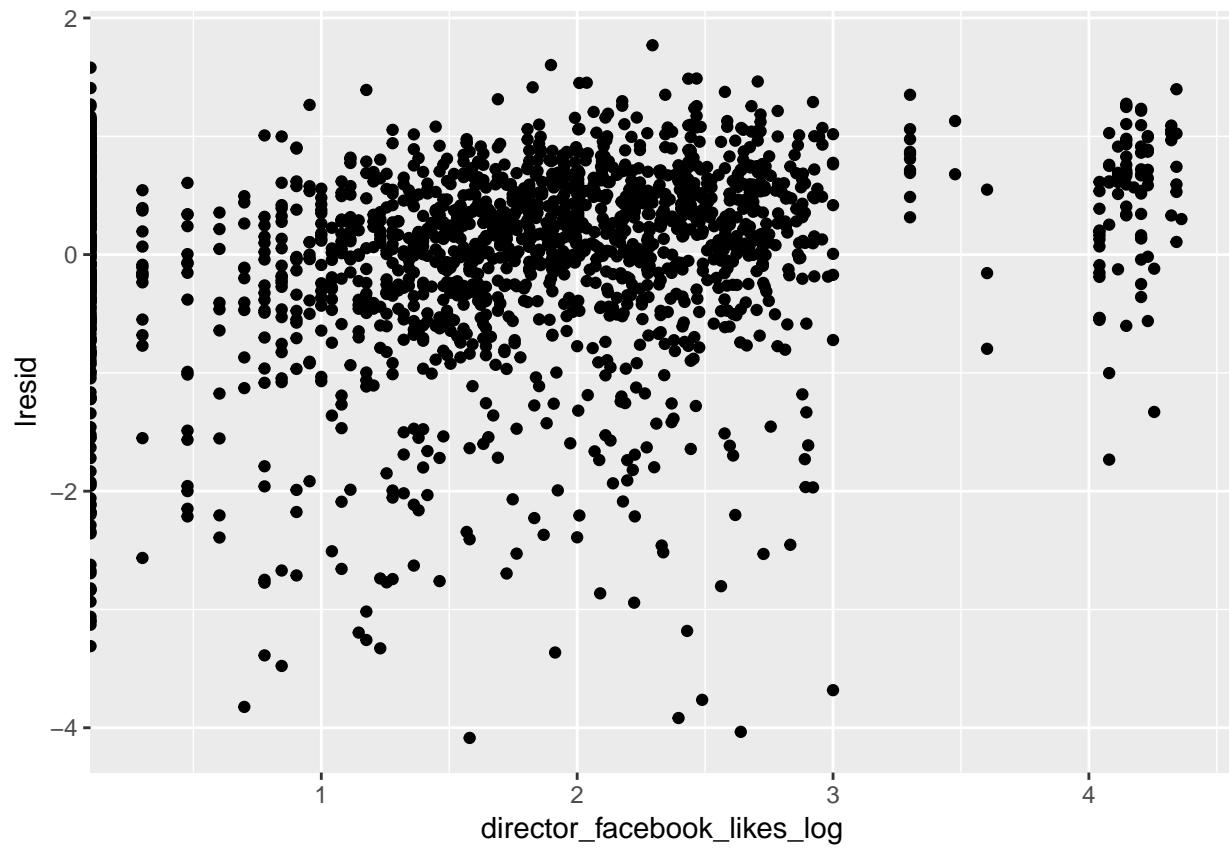
```



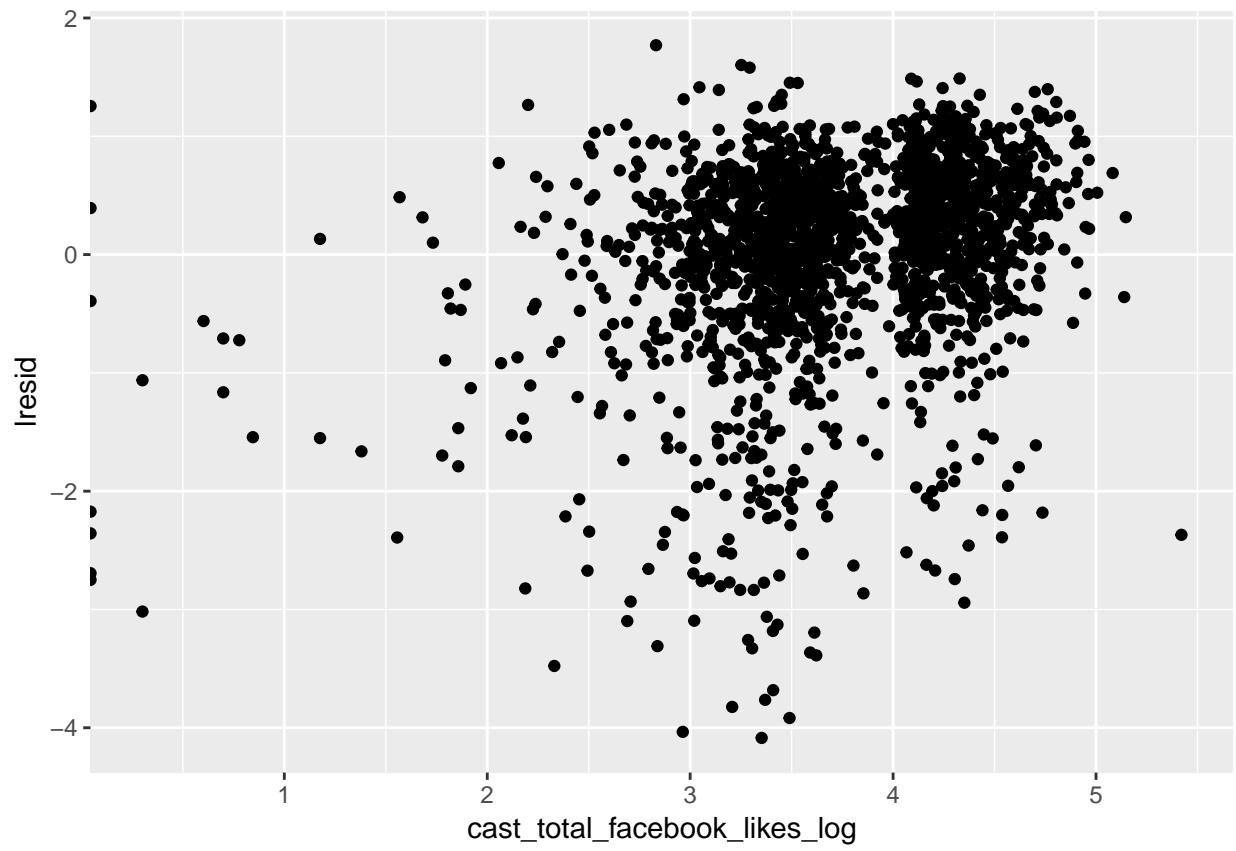
```

## 
## [[2]]

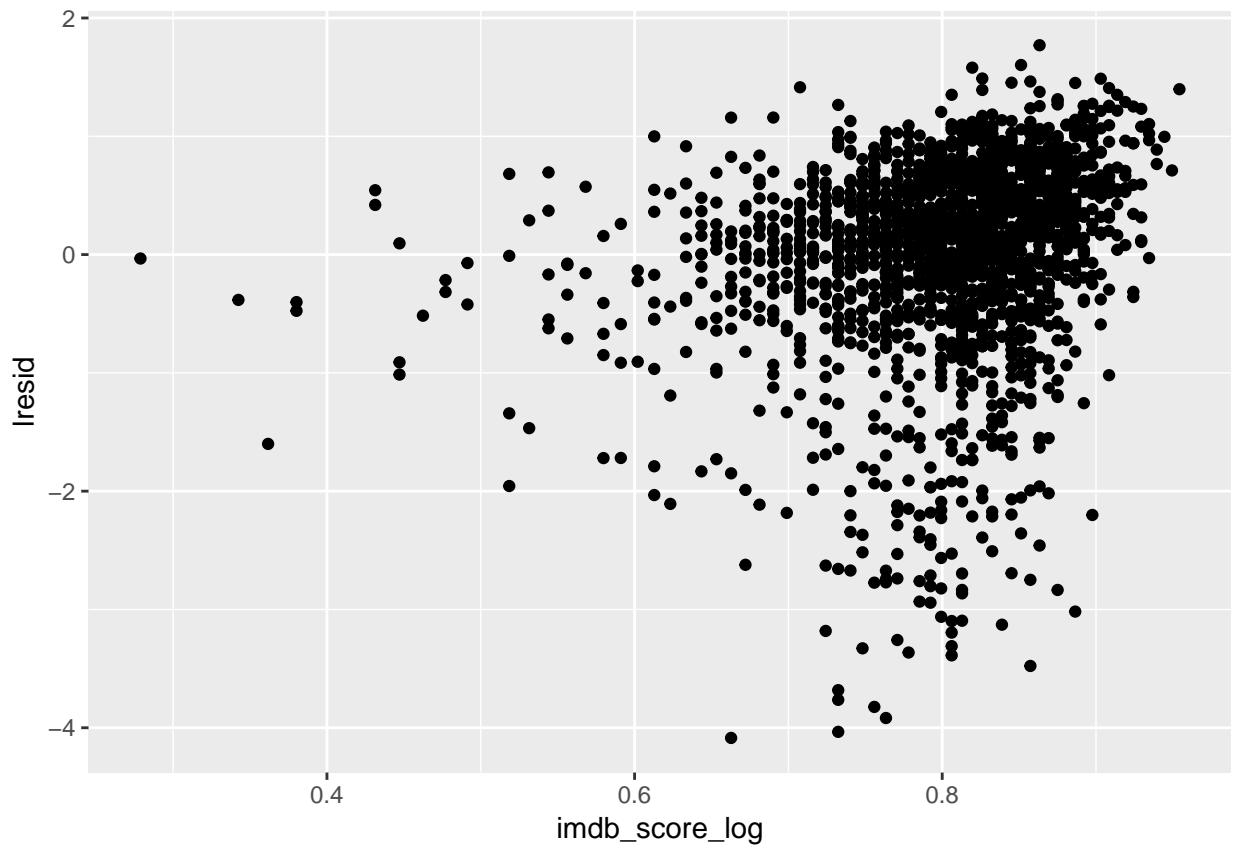
```



```
##  
## [[3]]
```



```
##  
## [[4]]
```

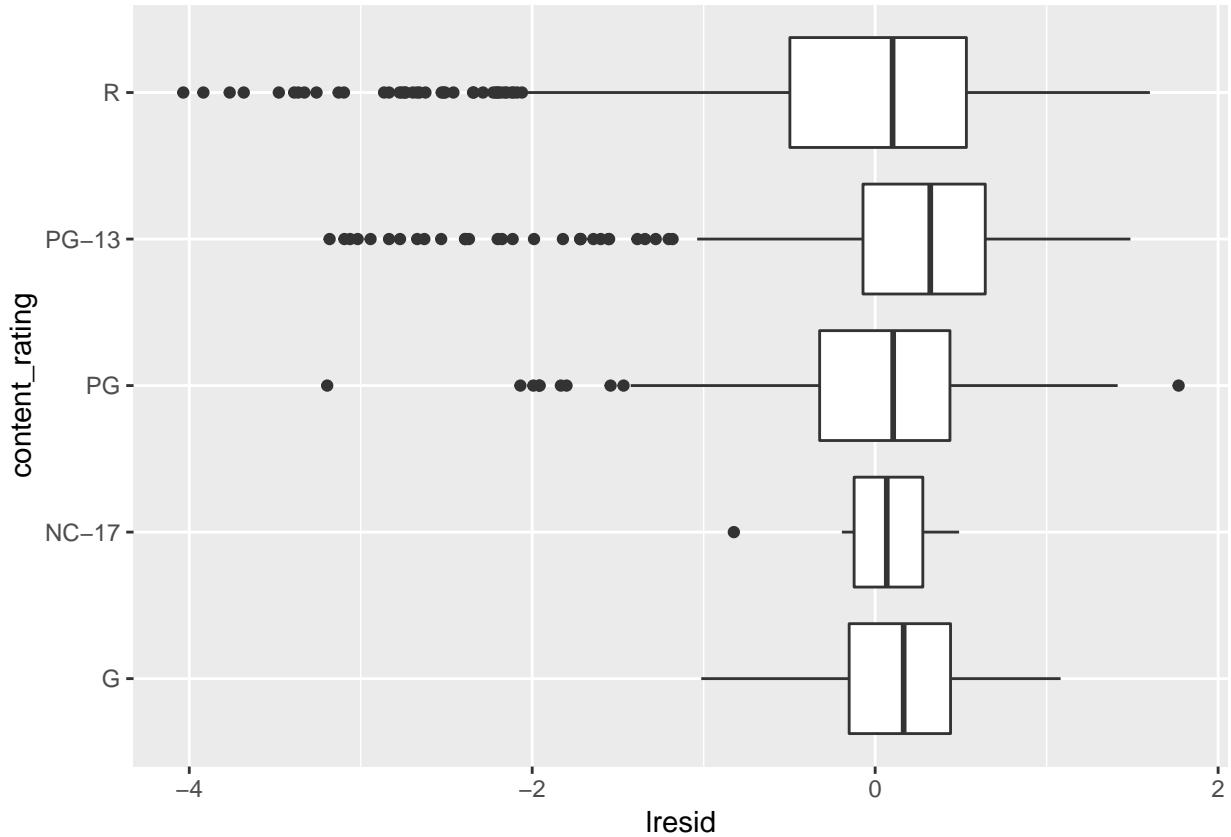


```

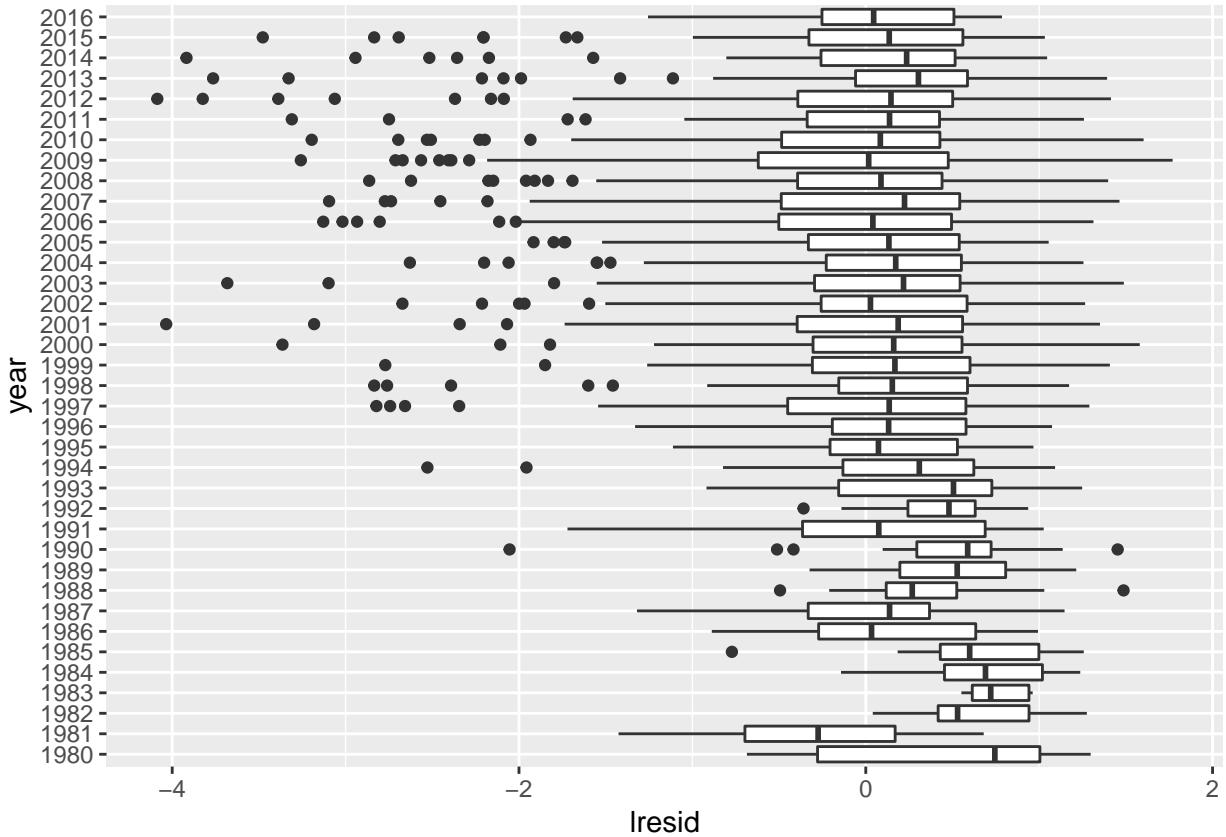
# categorical
# can't log categorical variables
lapply(c('content_rating', 'year'), function(var) {
  train_resid %>%
    filter(!is.na (!!rlang::sym(var))) %>%
    ggplot() +
    geom_boxplot(aes_string(var, 'lresid')) +
    coord_flip()
})

## [[1]]

```



```
##  
## [[2]]
```



Stepwise: Genre as Base

Try stepwise selection with these other variables given that none had fully random relationships with the residual from the genre model. Use the fitted genre model as a base.

For factor variables (content_rating, total_oscars) use normal versions of variables.

For facebook likes, use log versions as those were more linear with log(real_gross).

For budget and IMDB score, I think log versions are better, but try the non-log versions too. Both had some linearity.

For year, use normal version.

```
# create log versions of continuous variables
# also turn -Inf from log(0) to NA
train <- train %>%
  mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
    'imdb_score'), funs(log = log10(.))) %>%
  mutate_at(vars(contains('log')), funs(ifelse(is.infinite(.), NA, .)))
valid <- valid %>%
  mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
    'imdb_score'), funs(log = log10(.))) %>%
  mutate_at(vars(contains('log')), funs(ifelse(is.infinite(.), NA, .)))

# starting formula: genre
starting_formula = 'Adventure + Action + Family + Mystery + Documentary + Drama + History + Romance'

# stepwise starting with genre
rmse_lst <- step_wise_loop(df = train %>% select(genre_xvar, content_rating, real_budget, year,
```

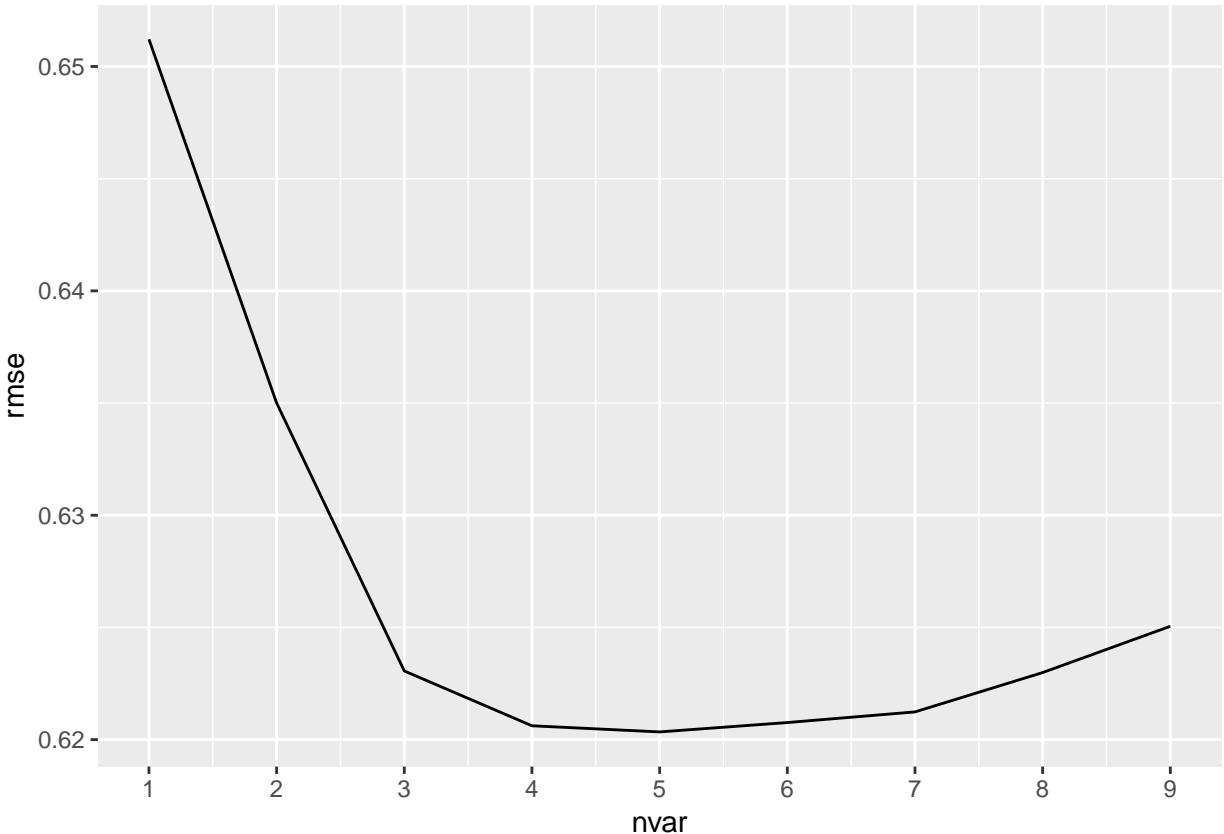
```

total_oscars_actor, total_oscars_director,
imdb_score_log, real_budget_log,
director_facebook_likes_log, cast_total_facebook_likes
starting_vars = genre_xvar,
starting_formula = starting_formula)

## real_budget_log
##          0.6512142
## [1] 1
## imdb_score_log
##          0.6350029
## [1] 2
##      year
## 0.6230605
## [1] 3
## content_rating
##          0.6206177
## [1] 4
## total_oscars_director
##          0.6203438
## [1] 5
## cast_total_facebook_likes_log
##          0.6207615
## [1] 6
## real_budget
##          0.6212351
## [1] 7
## total_oscars_actor
##          0.6229884
## [1] 8
## director_facebook_likes_log
##          0.6250515

# graph RMSE vs number of variables
fit_rmse <- tibble(nvar = 1:length(rmse_lst),
                    rmse = rmse_lst)
ggplot(fit_rmse) + geom_line(aes(x = nvar, y = rmse))+
  scale_x_continuous(breaks = seq(1, length(rmse_lst), by = 1))

```



```

# after var 4, decreases too small or increase

# model with extra 4 variables
mod_all <- lm(real_gross_log ~ Adventure + Action + Family + Mystery +
               Documentary + Drama + History + Romance +
               real_budget_log + imdb_score_log + year + content_rating,
               data = train)

summary(mod_all)

##
## Call:
## lm(formula = real_gross_log ~ Adventure + Action + Family + Mystery +
##     Documentary + Drama + History + Romance + real_budget_log +
##     imdb_score_log + year + content_rating, data = train)
##
## Residuals:
##      Min        1Q    Median        3Q       Max
## -3.4485 -0.2238  0.0878  0.3407  3.3377
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.016486   0.394298 -0.042  0.96665
## Adventure1  -0.142890   0.045332 -3.152  0.00165 **
## Action1     -0.011306   0.040640 -0.278  0.78090
## Family1      0.311938   0.079689  3.914 9.43e-05 ***

```

```

## Mystery1          0.042585  0.051695  0.824  0.41019
## Documentary1     0.163354  0.141606  1.154  0.24884
## Drama1           -0.257951  0.034509 -7.475 1.24e-13 ***
## History1          -0.017739  0.089809 -0.198  0.84345
## Romance1          0.015177  0.037044  0.410  0.68208
## real_budget_log   0.812038  0.029073 27.931 < 2e-16 ***
## imdb_score_log    2.160658  0.206270 10.475 < 2e-16 ***
## year1981          -0.209854  0.368926 -0.569  0.56955
## year1982          0.182927  0.339404  0.539  0.58998
## year1983          0.247796  0.384325  0.645  0.51917
## year1984          0.500267  0.329092  1.520  0.12866
## year1985          0.288020  0.339707  0.848  0.39665
## year1986          0.065028  0.324926  0.200  0.84140
## year1987          0.172547  0.316476  0.545  0.58568
## year1988          0.147460  0.310528  0.475  0.63494
## year1989          0.279267  0.305019  0.916  0.36002
## year1990          0.062553  0.306795  0.204  0.83846
## year1991          0.044401  0.310206  0.143  0.88620
## year1992          -0.007279  0.315891 -0.023  0.98162
## year1993          0.025560  0.307311  0.083  0.93372
## year1994          -0.222206  0.301314 -0.737  0.46095
## year1995          -0.050133  0.293633 -0.171  0.86445
## year1996          -0.217635  0.286206 -0.760  0.44712
## year1997          -0.171164  0.286195 -0.598  0.54988
## year1998          -0.272722  0.285619 -0.955  0.33980
## year1999          -0.270022  0.282260 -0.957  0.33889
## year2000          -0.173448  0.283525 -0.612  0.54078
## year2001          -0.287595  0.281570 -1.021  0.30721
## year2002          -0.262480  0.281560 -0.932  0.35135
## year2003          -0.250251  0.282916 -0.885  0.37653
## year2004          -0.266357  0.282475 -0.943  0.34585
## year2005          -0.255937  0.281463 -0.909  0.36332
## year2006          -0.346753  0.281963 -1.230  0.21895
## year2007          -0.347511  0.283316 -1.227  0.22015
## year2008          -0.374356  0.281500 -1.330  0.18375
## year2009          -0.371088  0.280547 -1.323  0.18611
## year2010          -0.402398  0.281138 -1.431  0.15253
## year2011          -0.279867  0.283485 -0.987  0.32367
## year2012          -0.184045  0.281991 -0.653  0.51406
## year2013          -0.120662  0.281346 -0.429  0.66807
## year2014          -0.154492  0.283282 -0.545  0.58558
## year2015          -0.267933  0.285375 -0.939  0.34793
## year2016          -0.176387  0.301510 -0.585  0.55862
## content_ratingNC-17 -0.036289  0.275191 -0.132  0.89510
## content_ratingPG   -0.035081  0.119686 -0.293  0.76948
## content_ratingPG-13  0.139695  0.137054  1.019  0.30822
## content_ratingR    -0.052452  0.136983 -0.383  0.70184
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6071 on 1668 degrees of freedom
##   (132 observations deleted due to missingness)
## Multiple R-squared:  0.495, Adjusted R-squared:  0.4799
## F-statistic: 32.71 on 50 and 1668 DF, p-value: < 2.2e-16

```

```

rmse(mod_all, data = valid)

## [1] 0.6206177

# when consider the factors as one variable, they are significant
anova(mod_all)

## Analysis of Variance Table
##
## Response: real_gross_log
##                               Df Sum Sq Mean Sq F value    Pr(>F)
## Adventure                  1  75.64   75.64 205.2075 < 2.2e-16 ***
## Action                     1  24.10   24.10  65.3761 1.178e-15 ***
## Family                     1  32.07   32.07  86.9917 < 2.2e-16 ***
## Mystery                    1   2.40    2.40   6.5082  0.01083 *
## Documentary                1   7.70    7.70  20.8945 5.210e-06 ***
## Drama                      1  14.60   14.60  39.6072 3.960e-10 ***
## History                    1   5.95    5.95  16.1431 6.134e-05 ***
## Romance                    1   1.86    1.86   5.0542  0.02470 *
## real_budget_log            1 347.64  347.64 943.0834 < 2.2e-16 ***
## imdb_score_log              1  46.38   46.38 125.8221 < 2.2e-16 ***
## year                       36  33.63   0.93   2.5341 1.764e-06 ***
## content_rating              4  10.83   2.71   7.3440 7.330e-06 ***
## Residuals                  1668 614.85   0.37
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# number of observations
# lost about 150 observations to missings
nobs(mod_all)

## [1] 1719

```

New Residuals

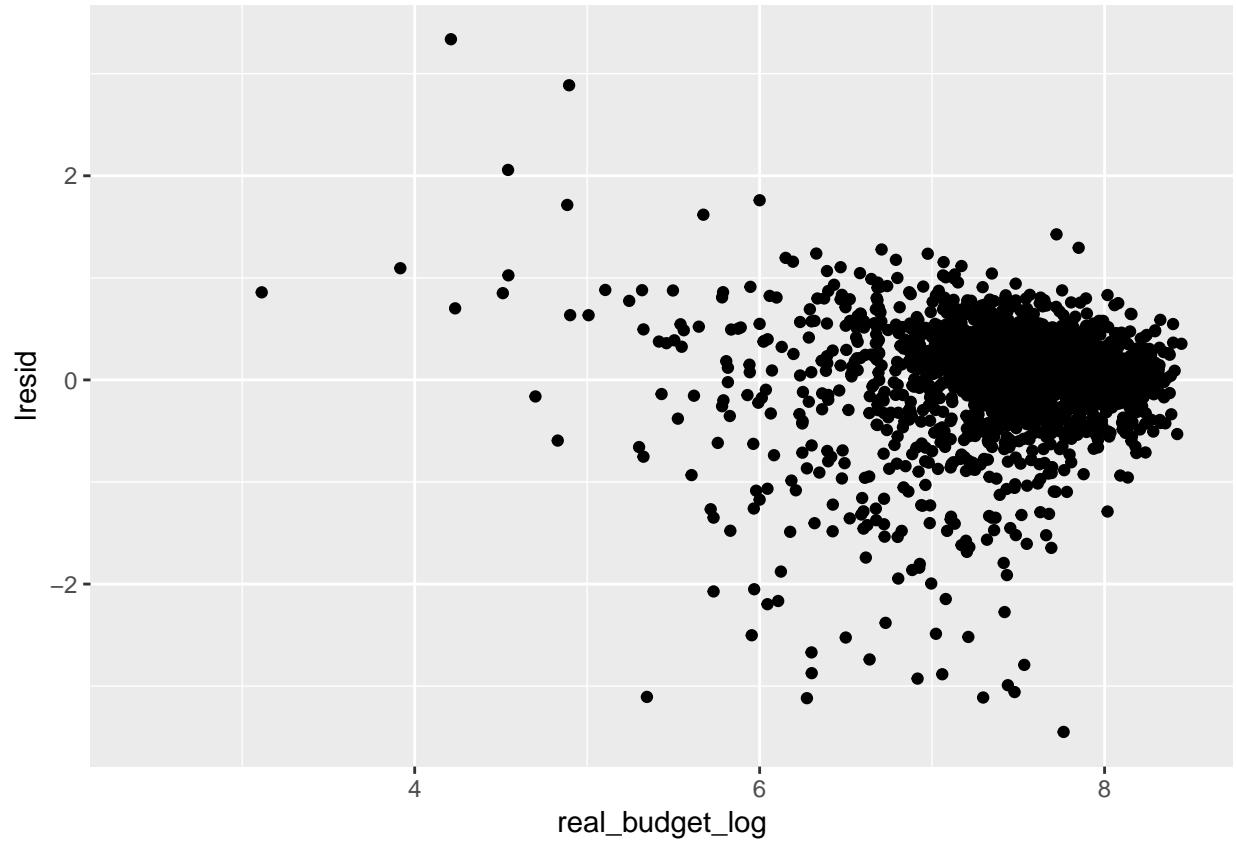
Graph residuals of included and excluded variables: have we captured all of the relationships?

```

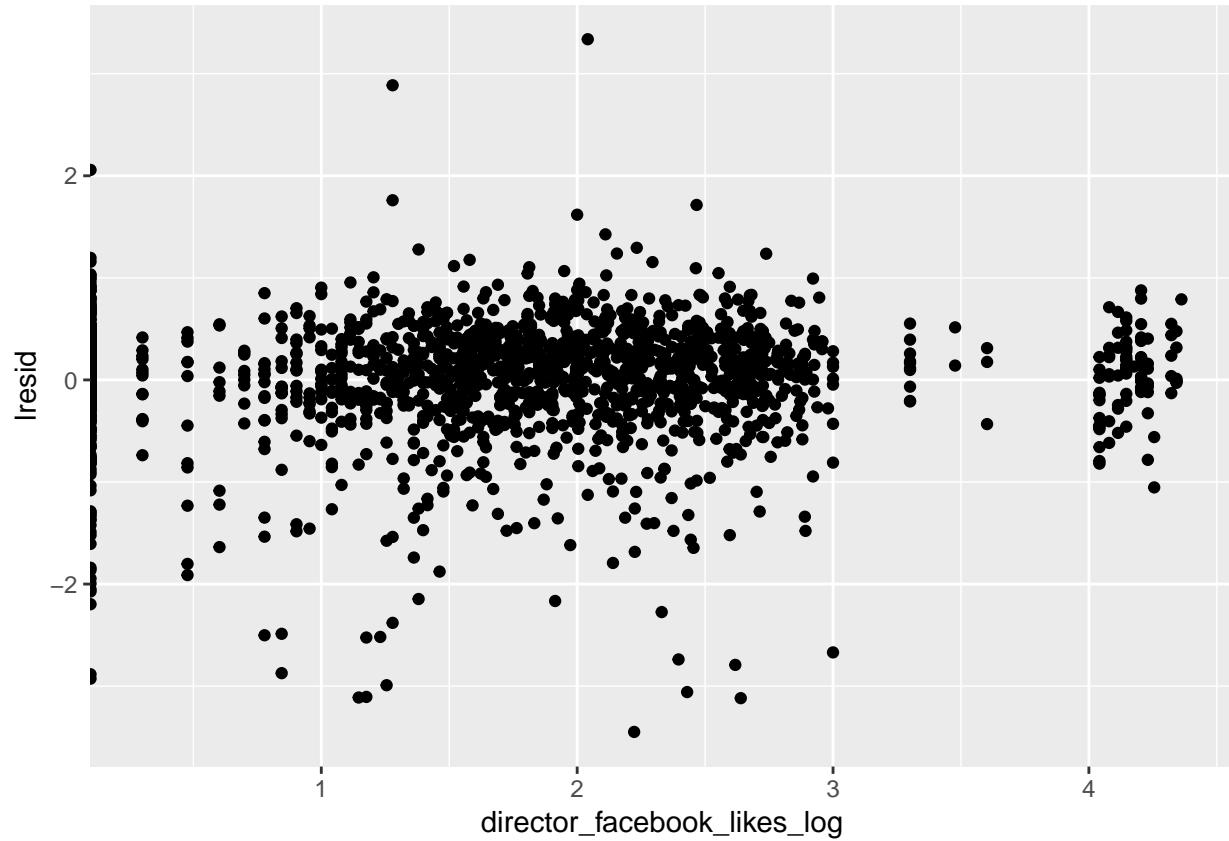
gr_resid(mod_all)

## Warning: Removed 132 rows containing missing values (geom_point).

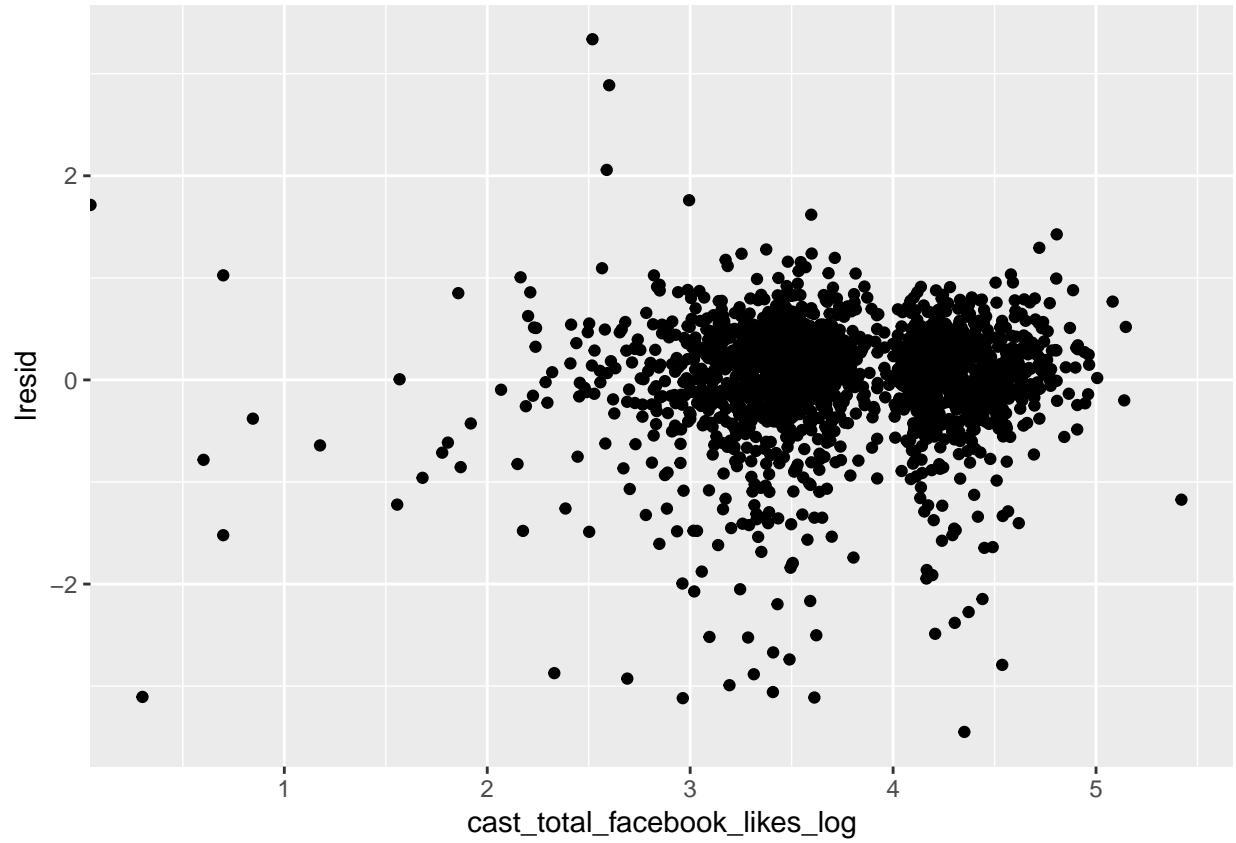
```

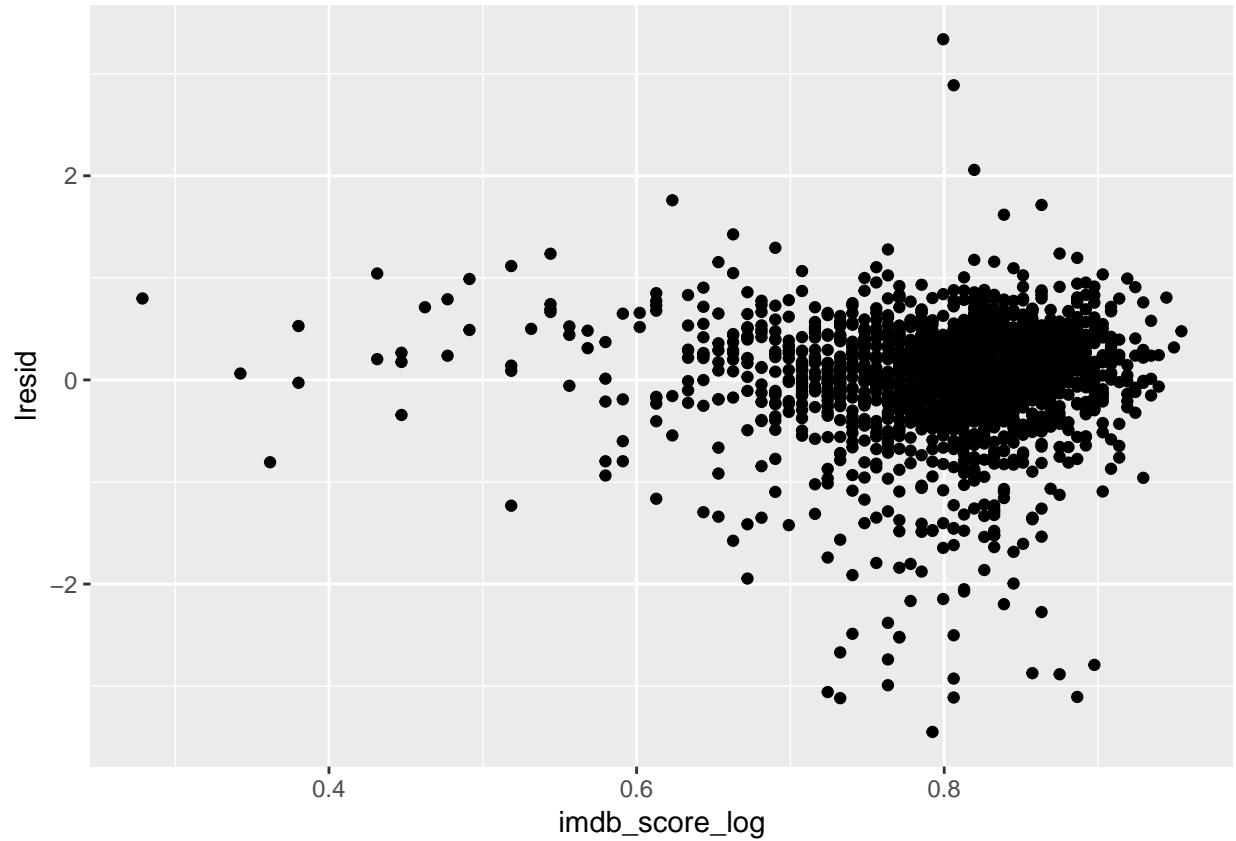


```
## Warning: Removed 132 rows containing missing values (geom_point).
```

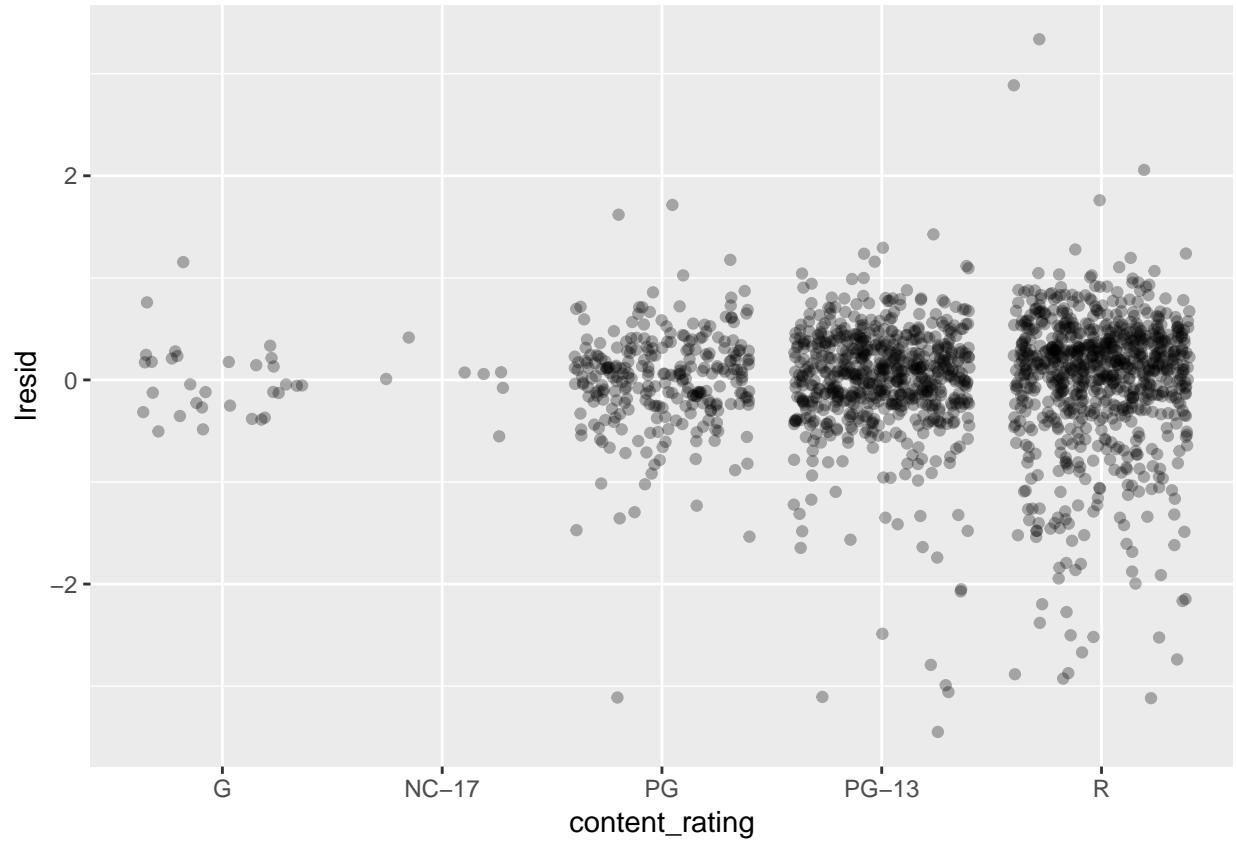


```
## Warning: Removed 132 rows containing missing values (geom_point).
```

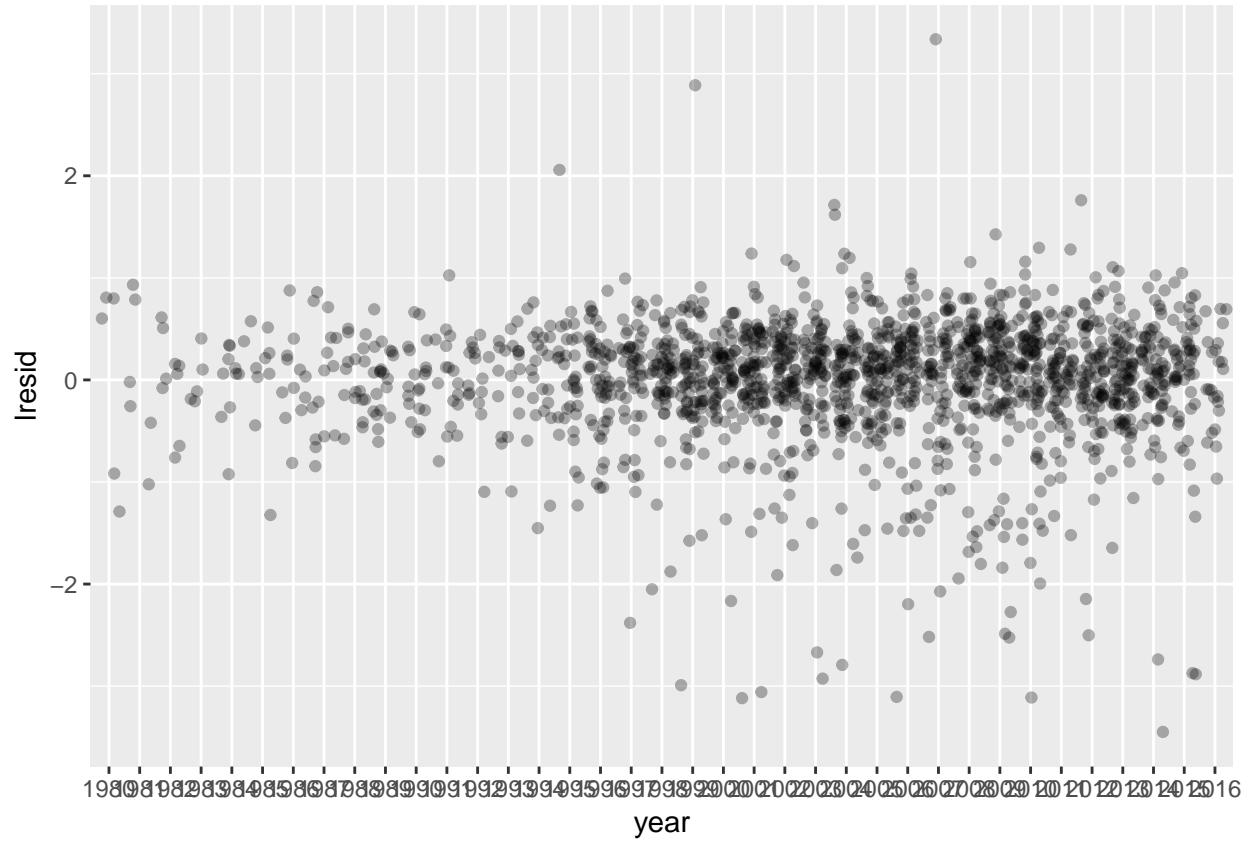




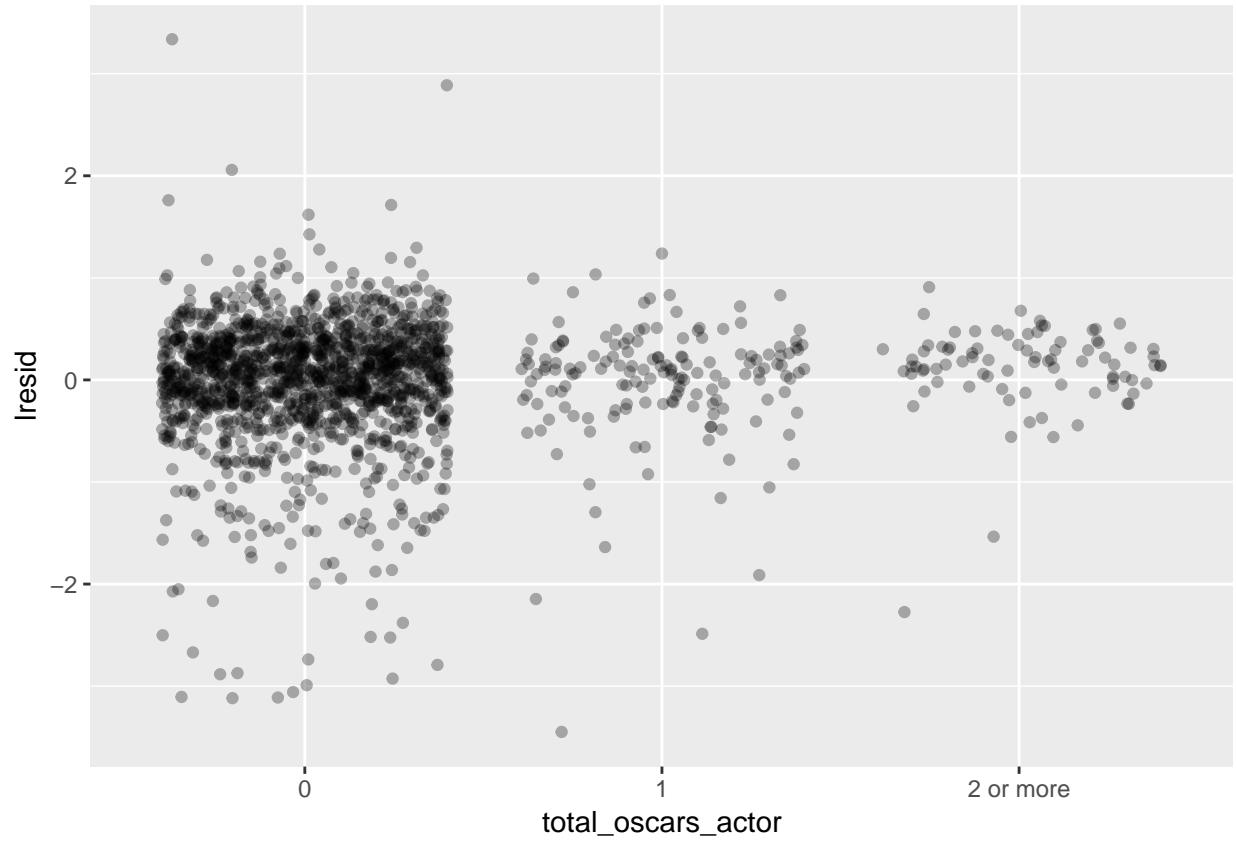
```
## Warning: Removed 94 rows containing missing values (geom_point).
```



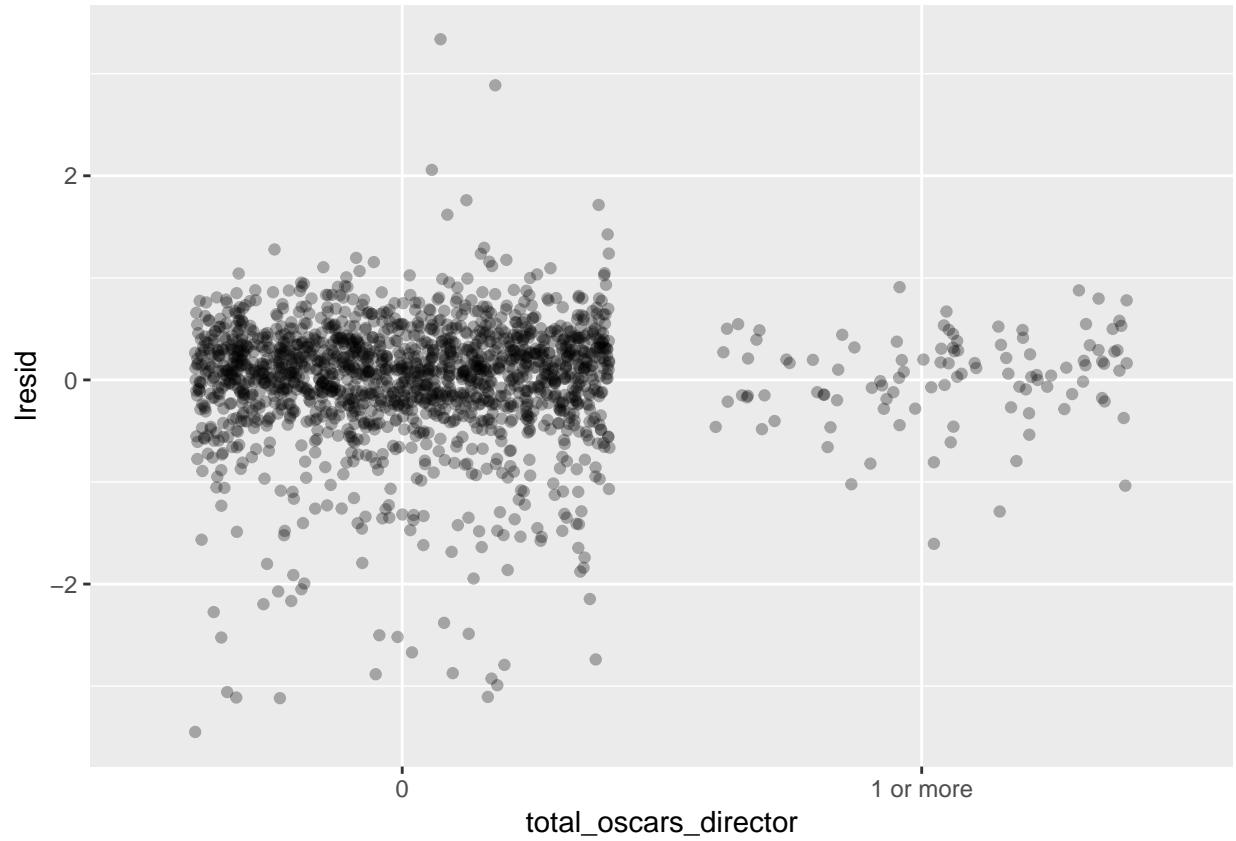
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



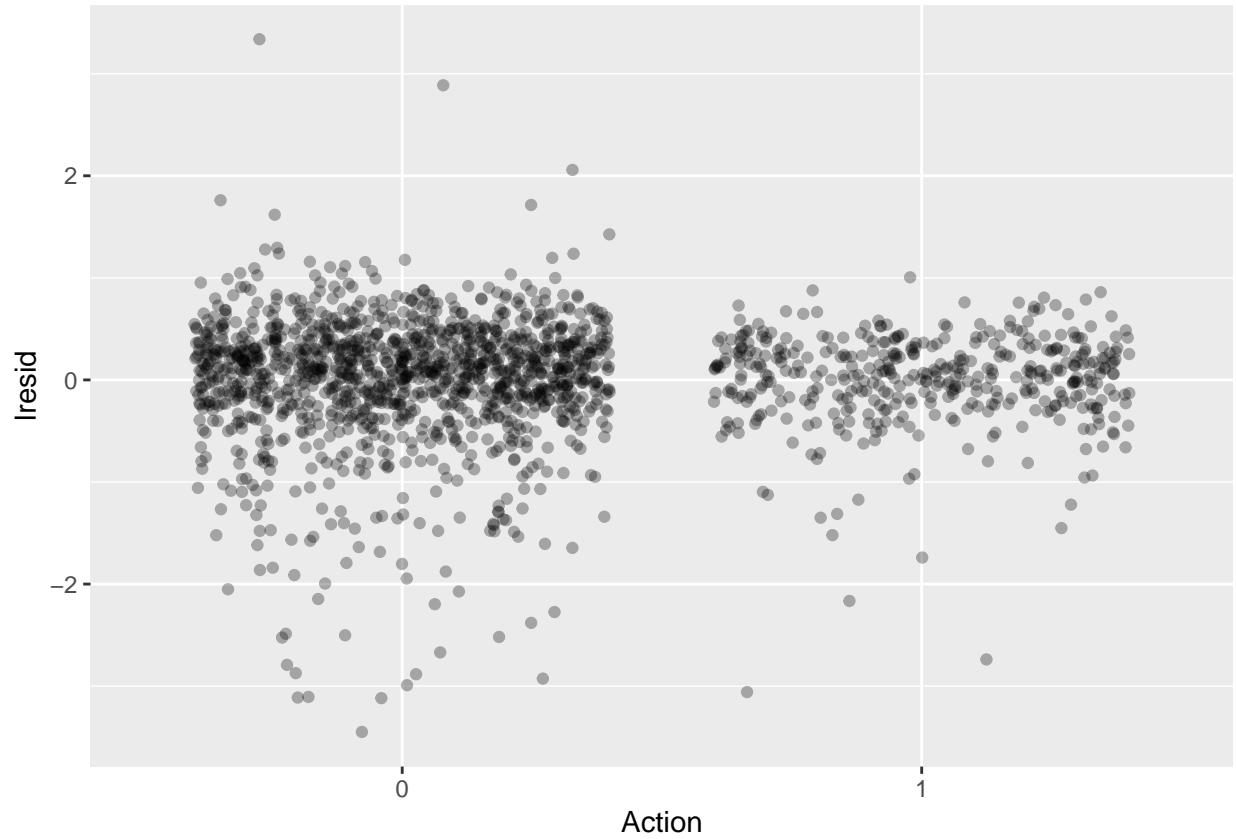
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



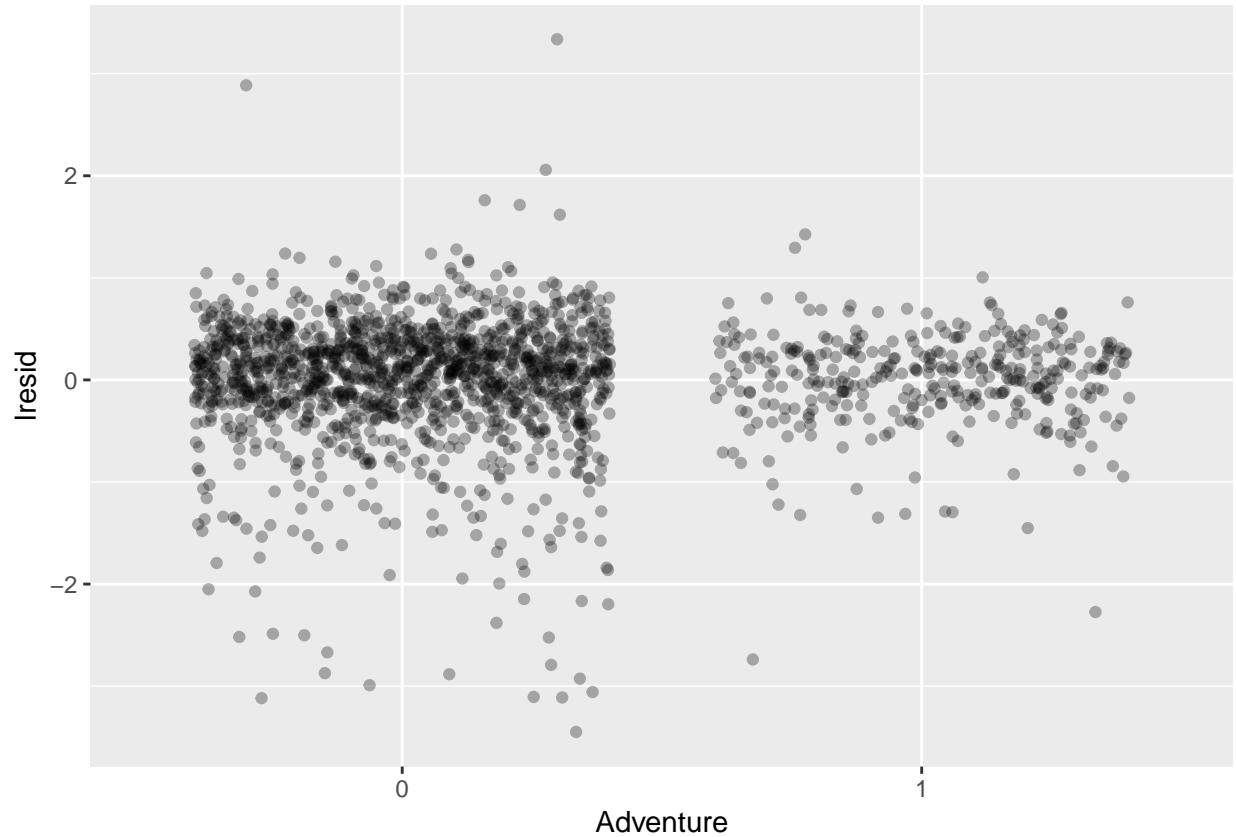
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



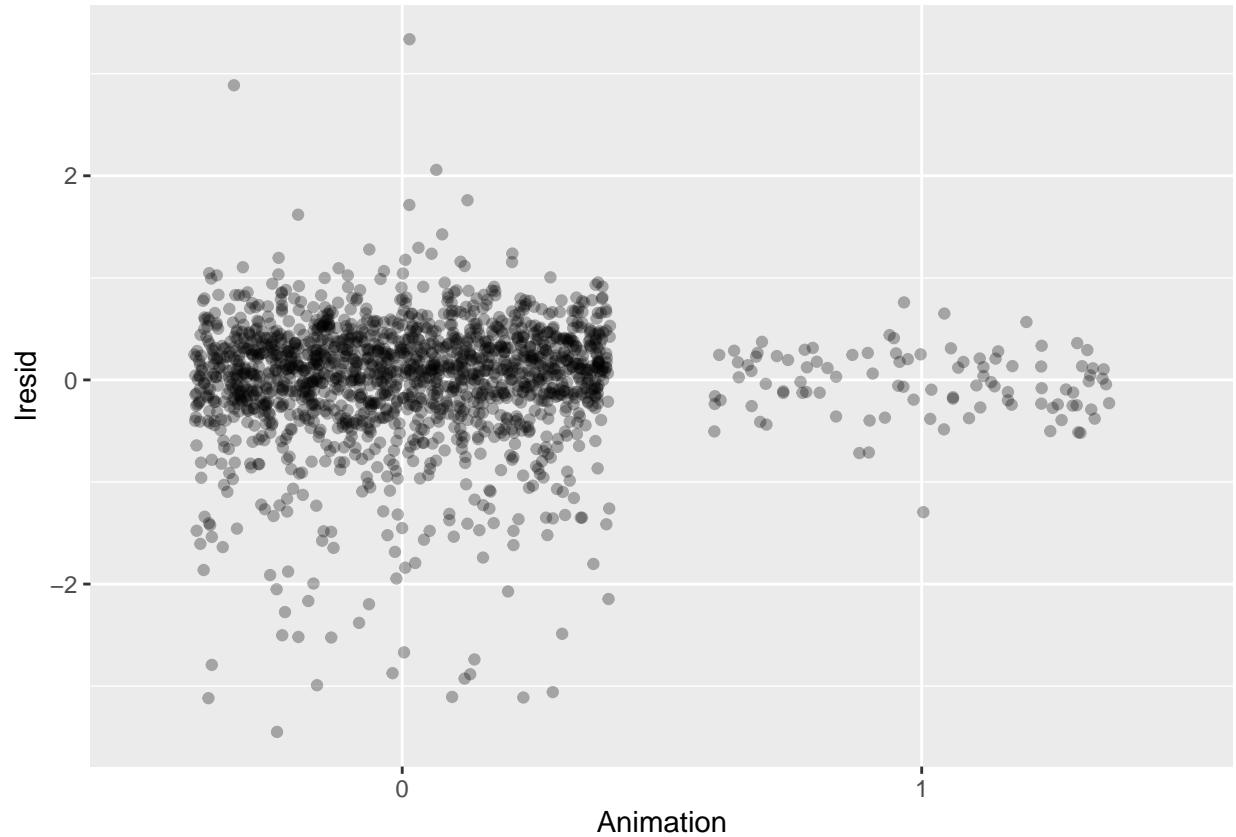
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



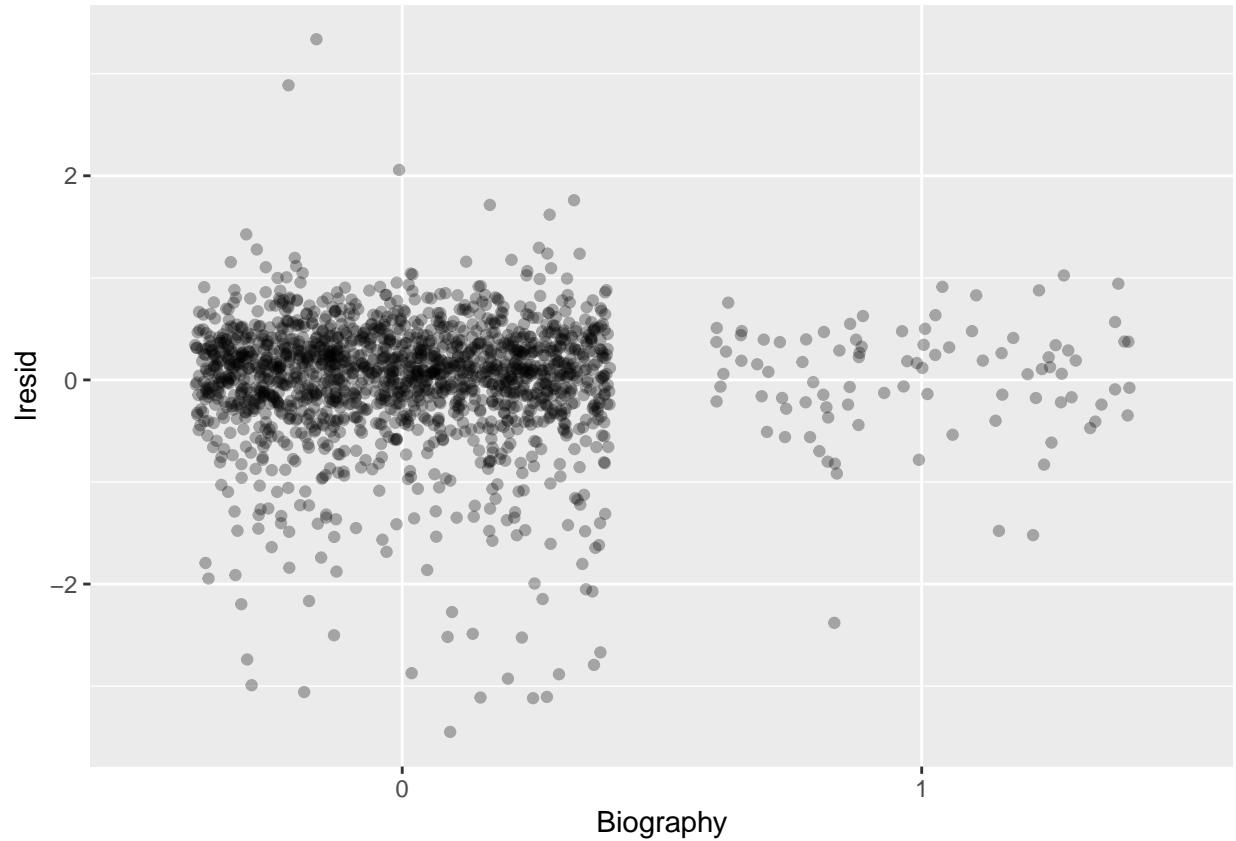
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



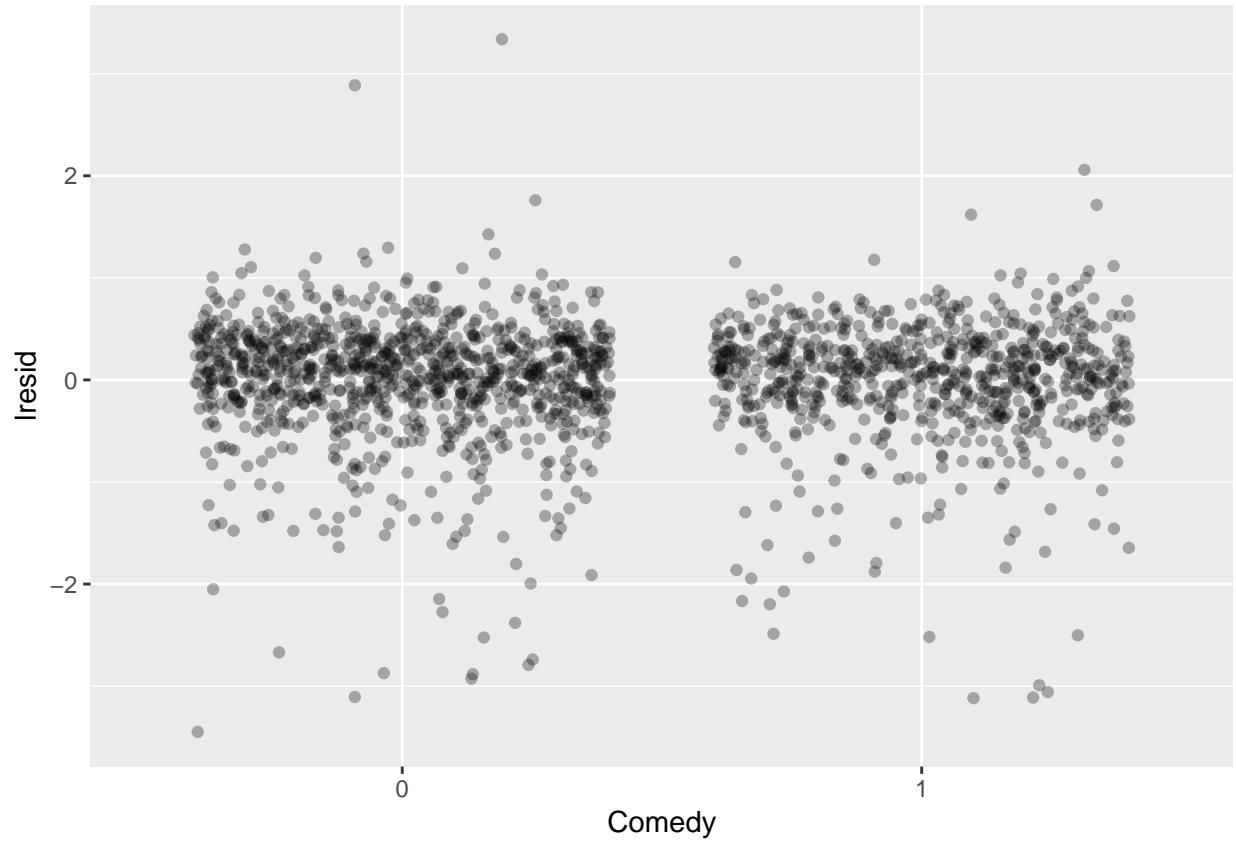
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



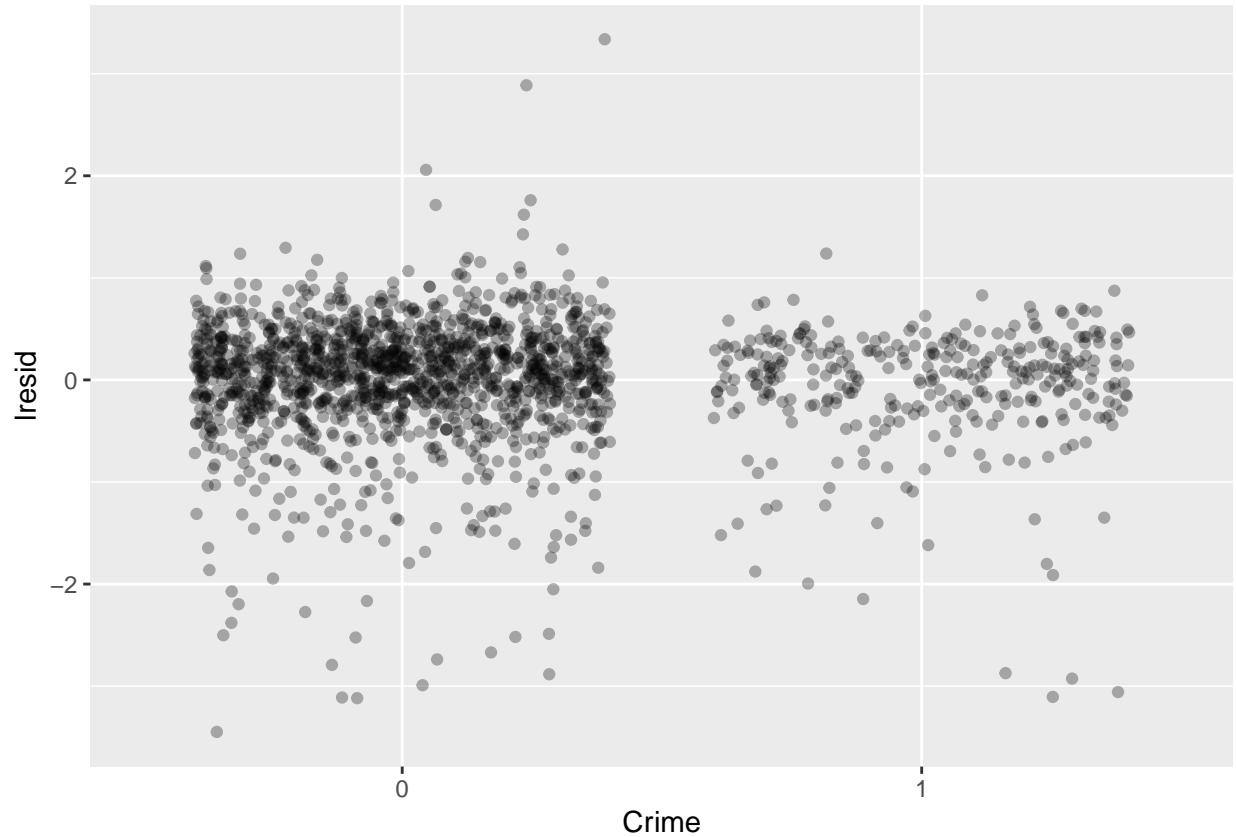
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



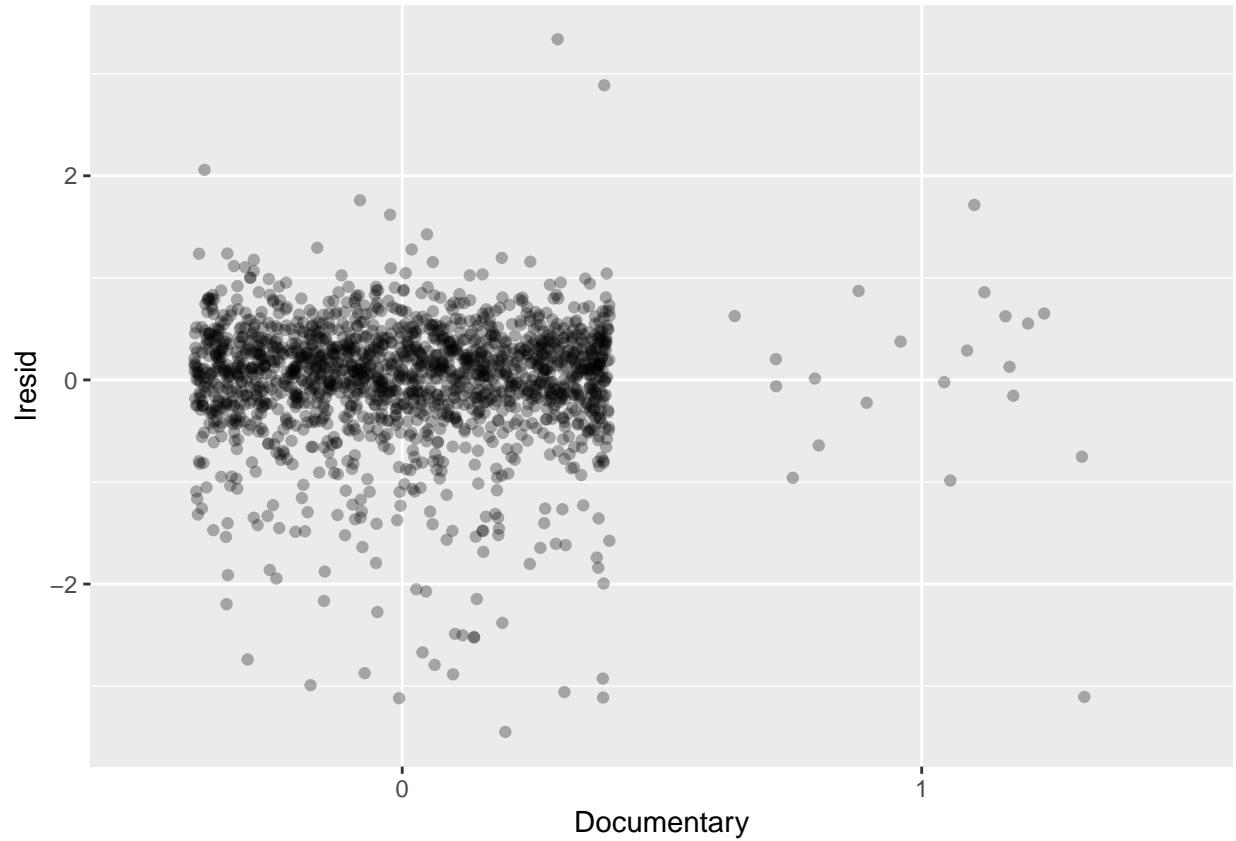
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



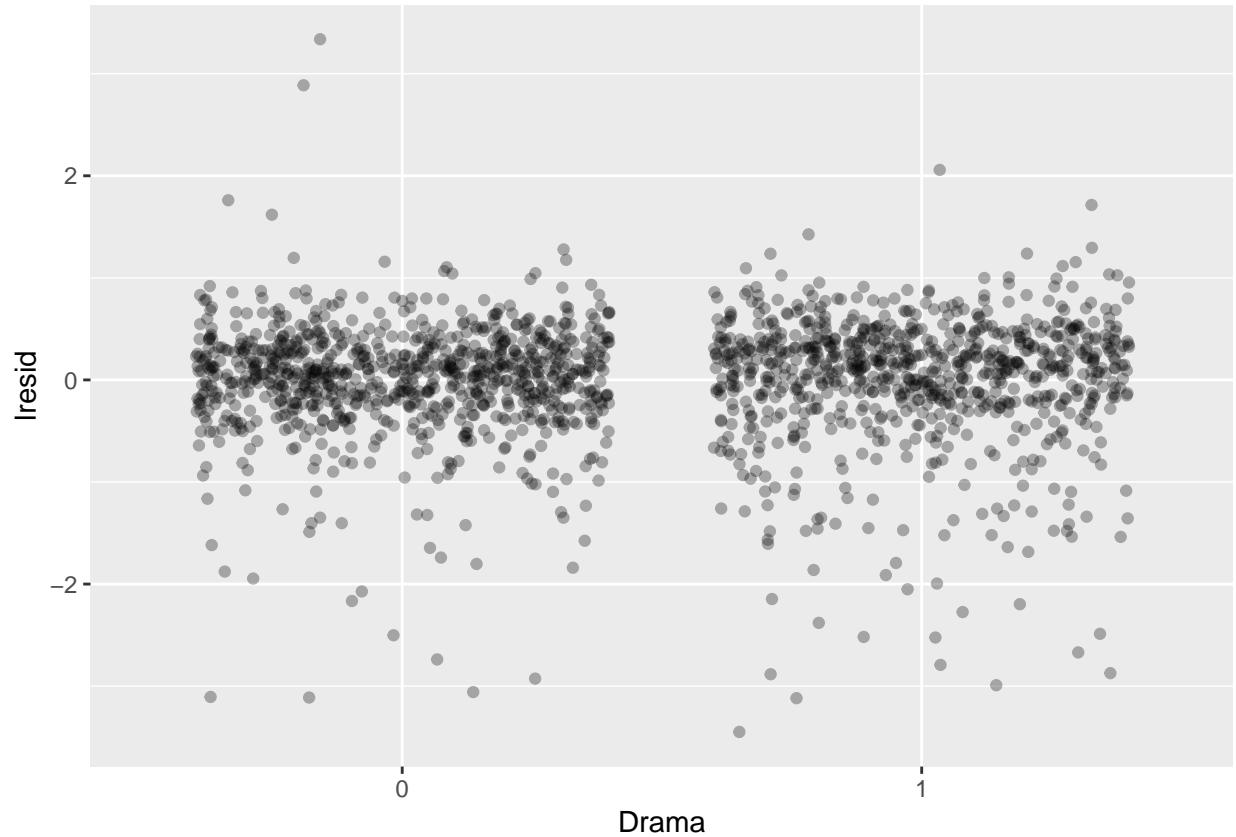
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



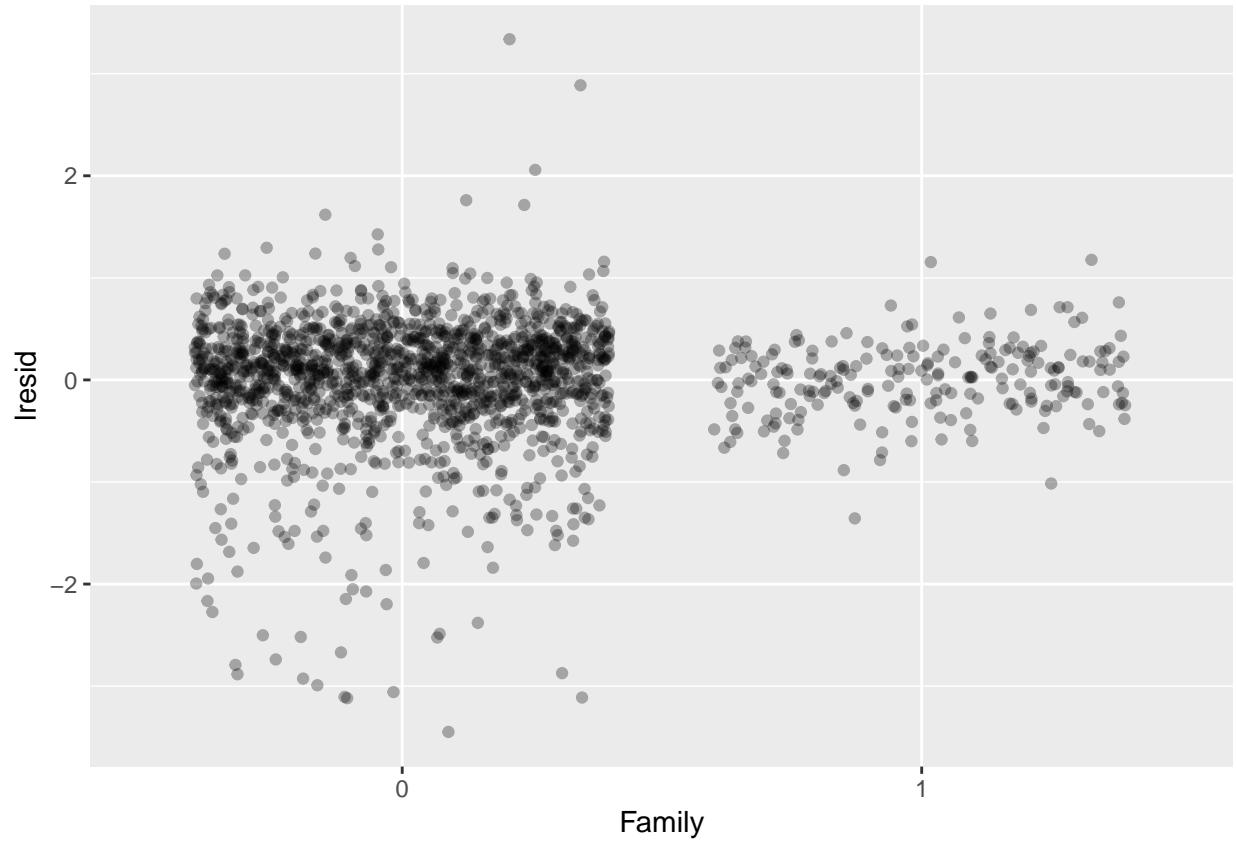
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



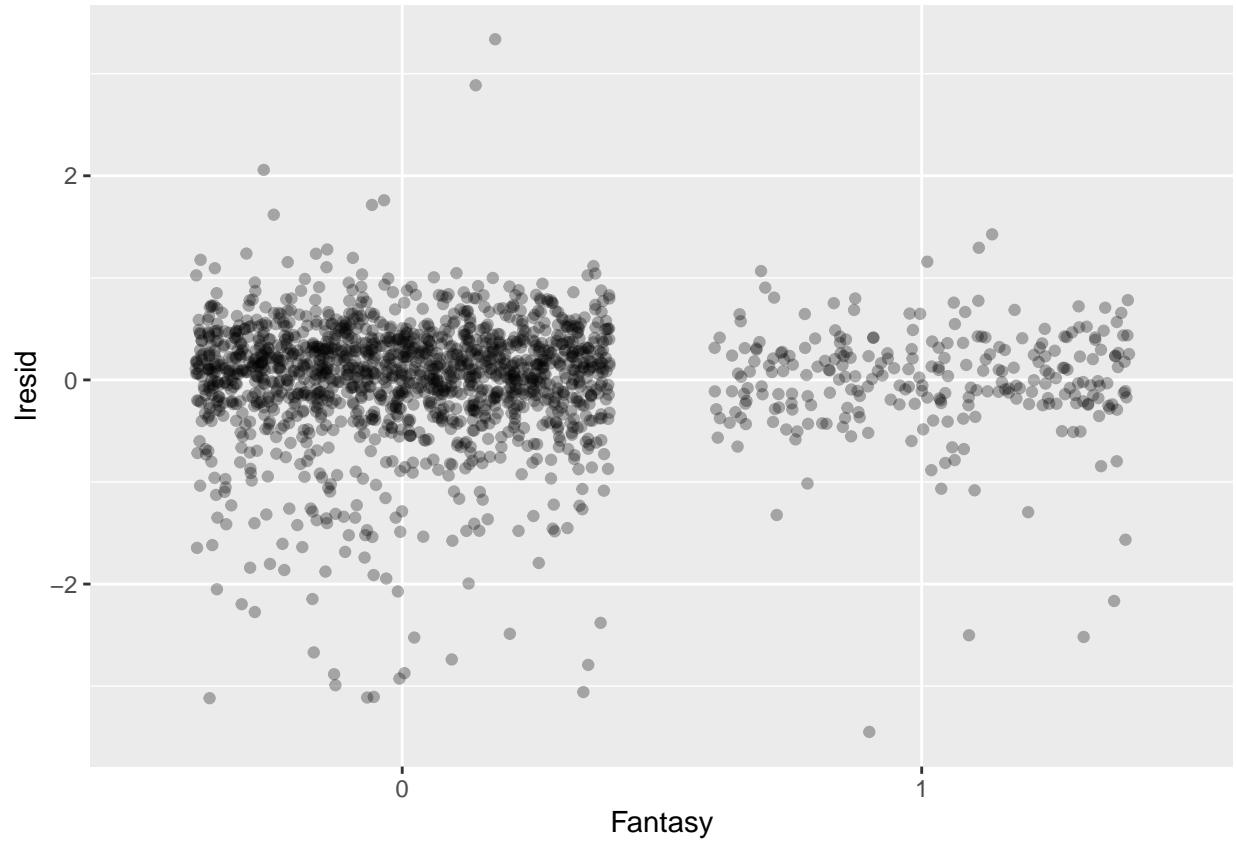
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



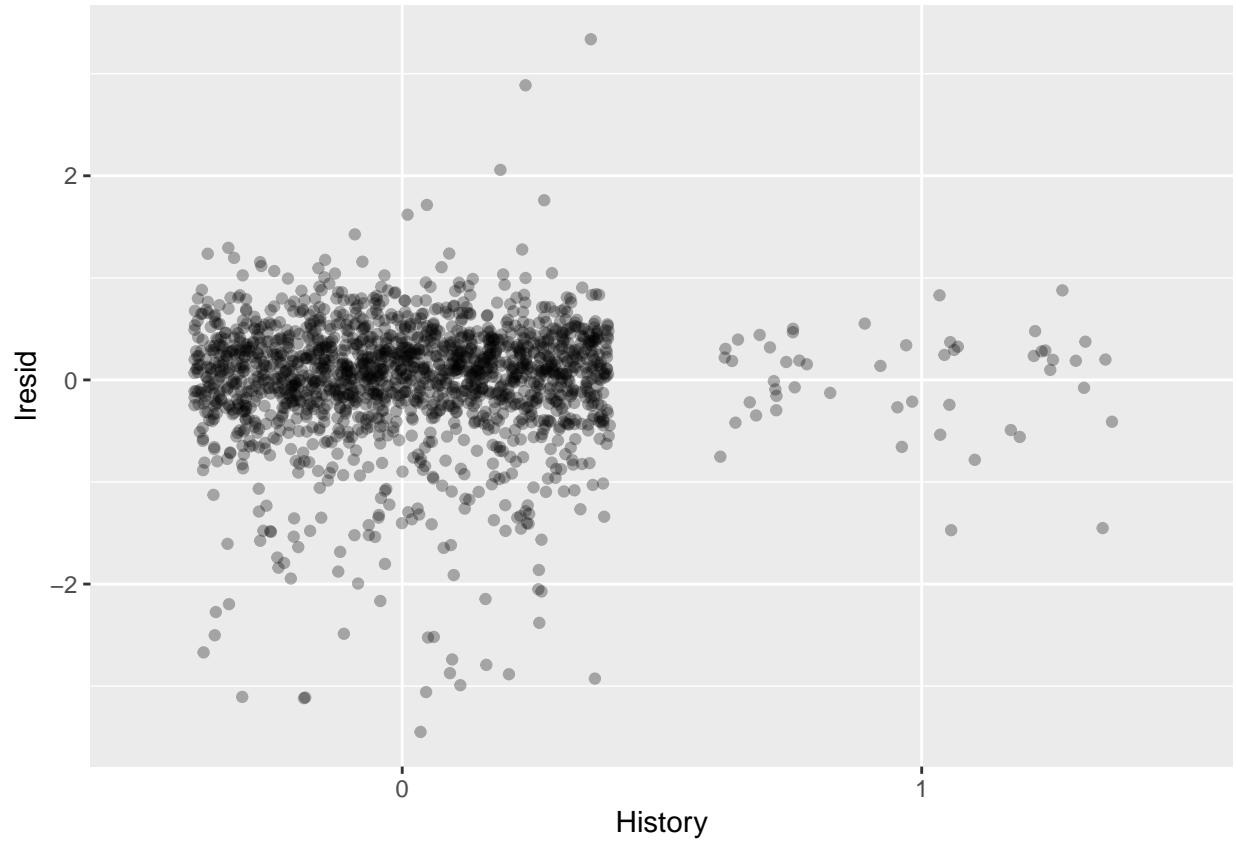
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



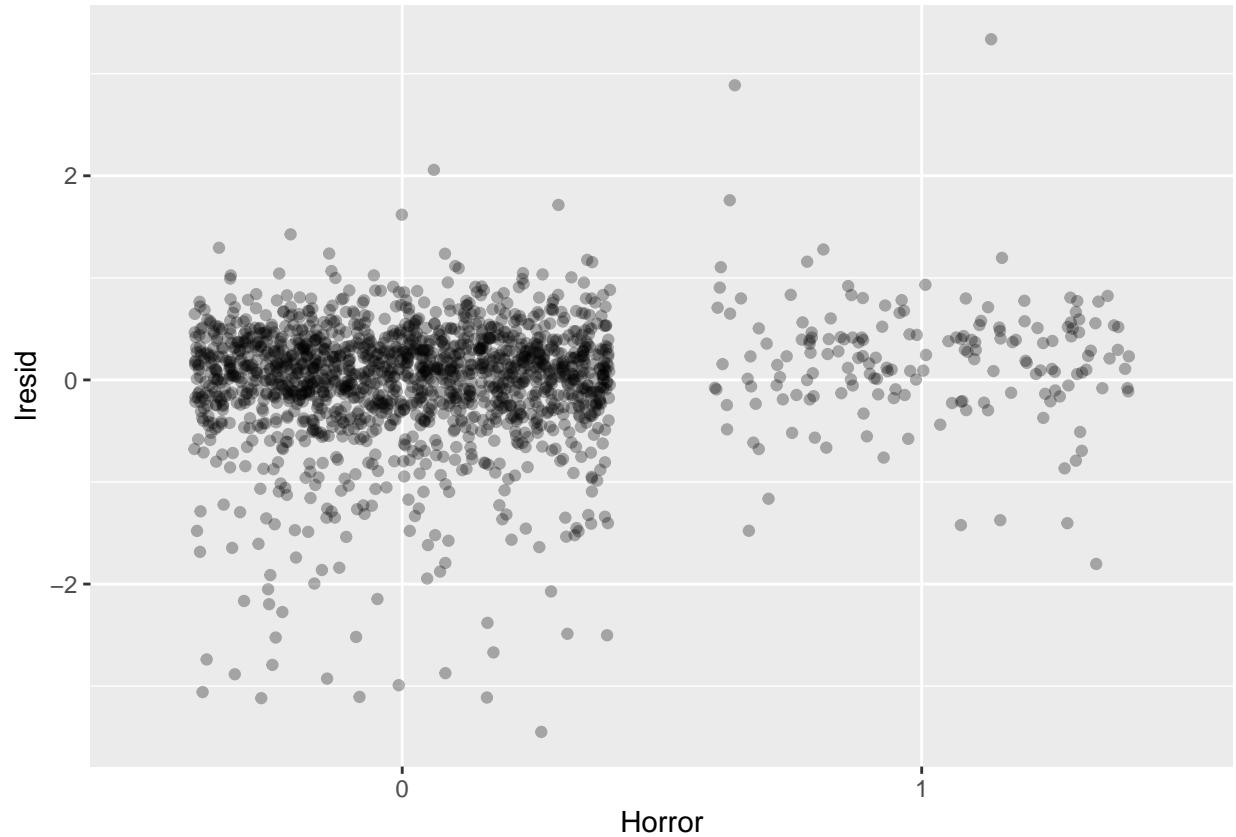
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



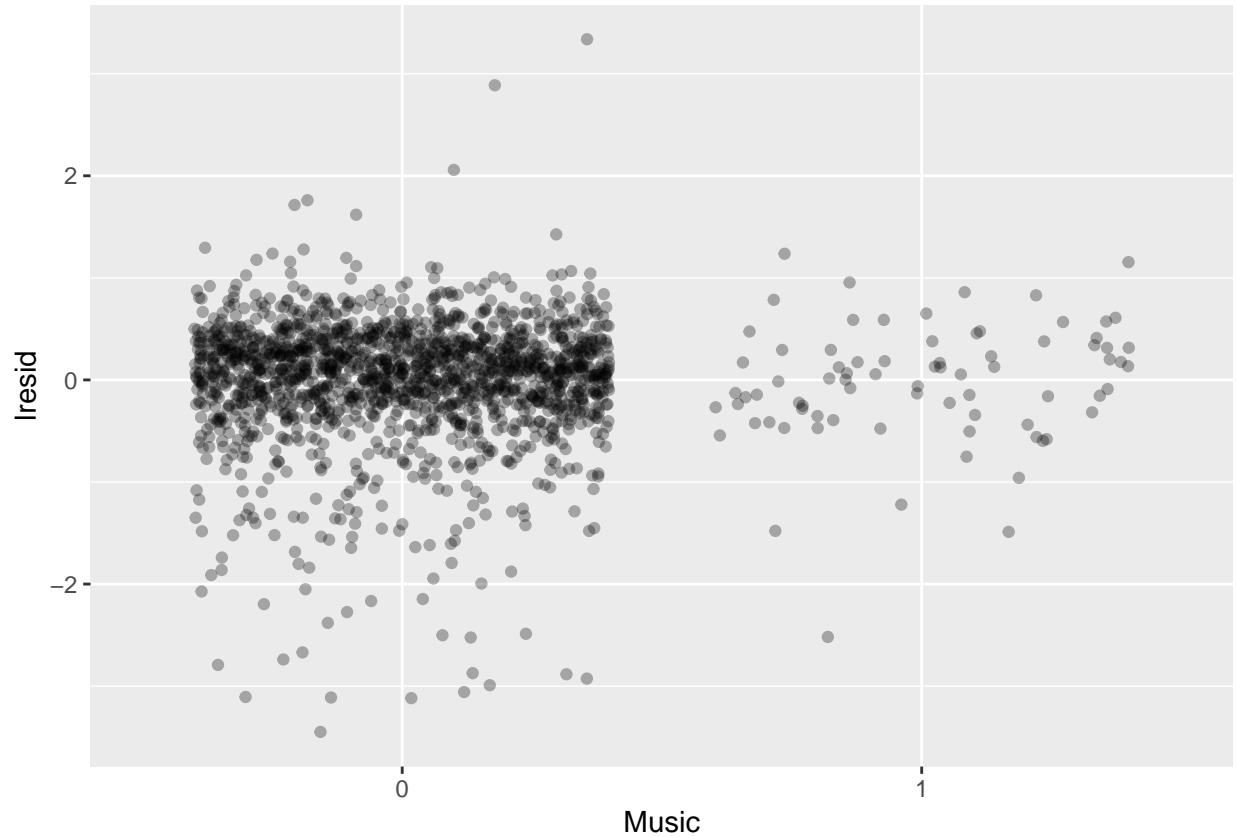
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



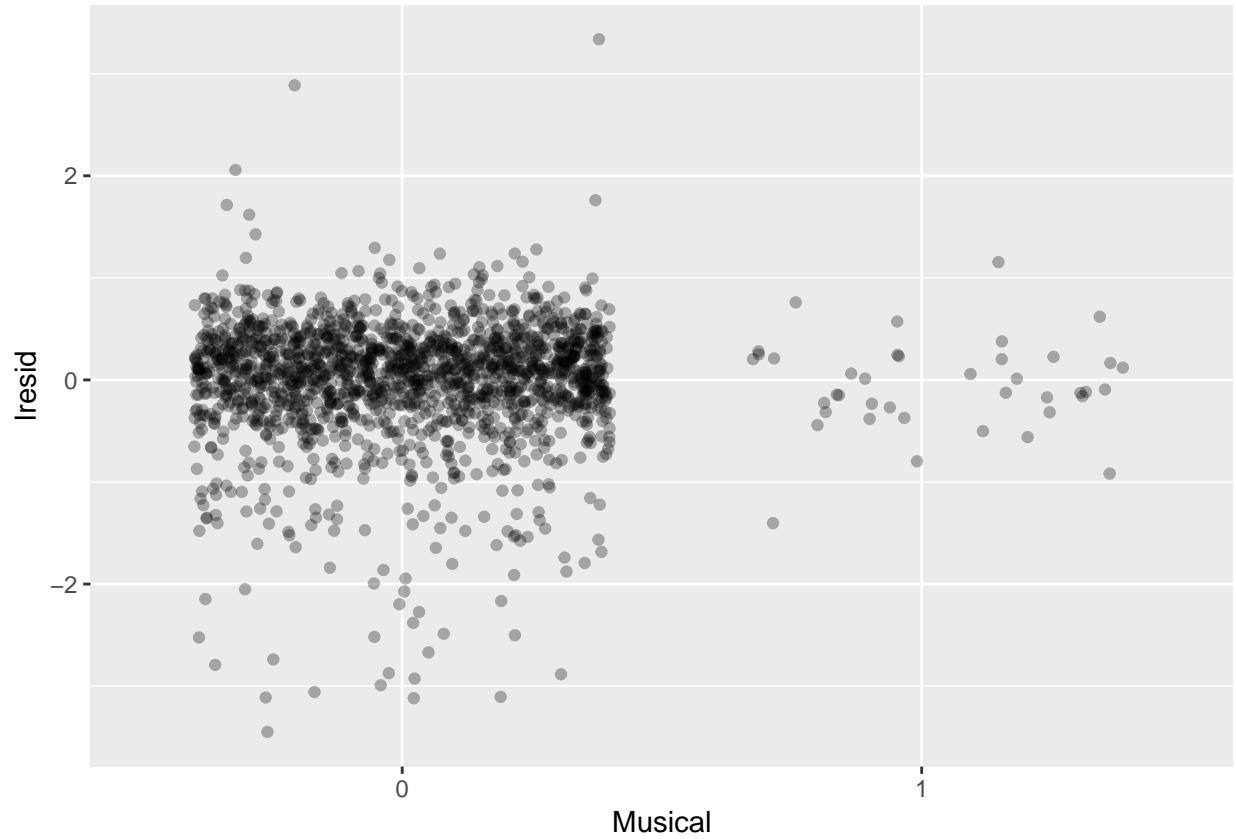
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



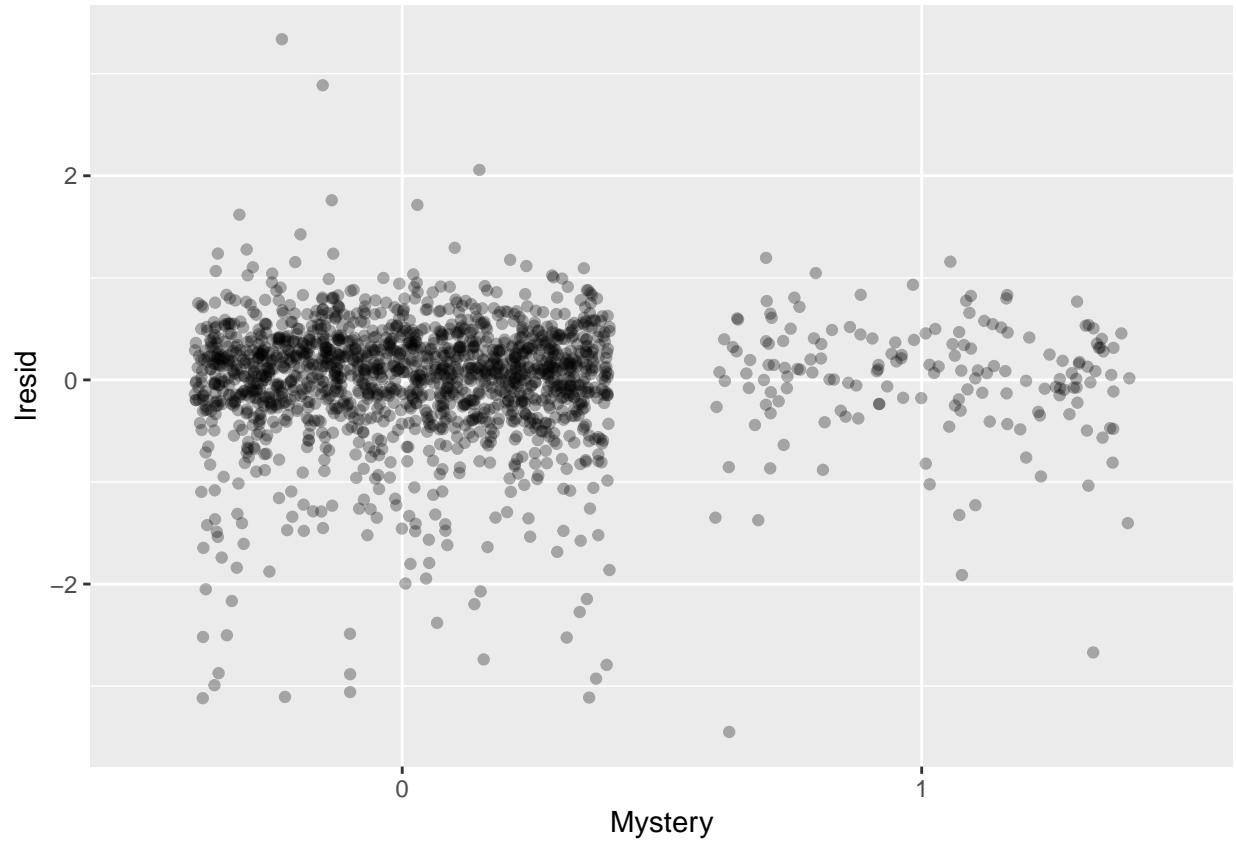
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



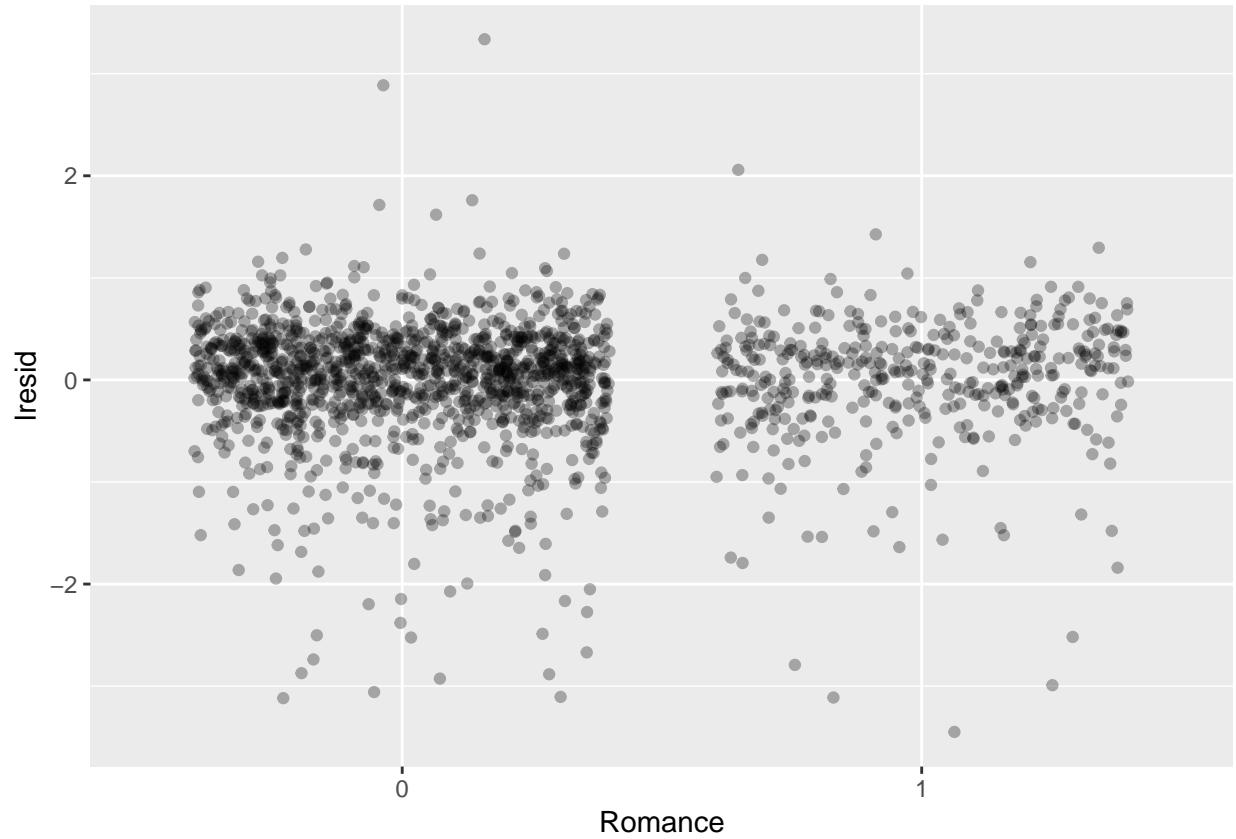
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



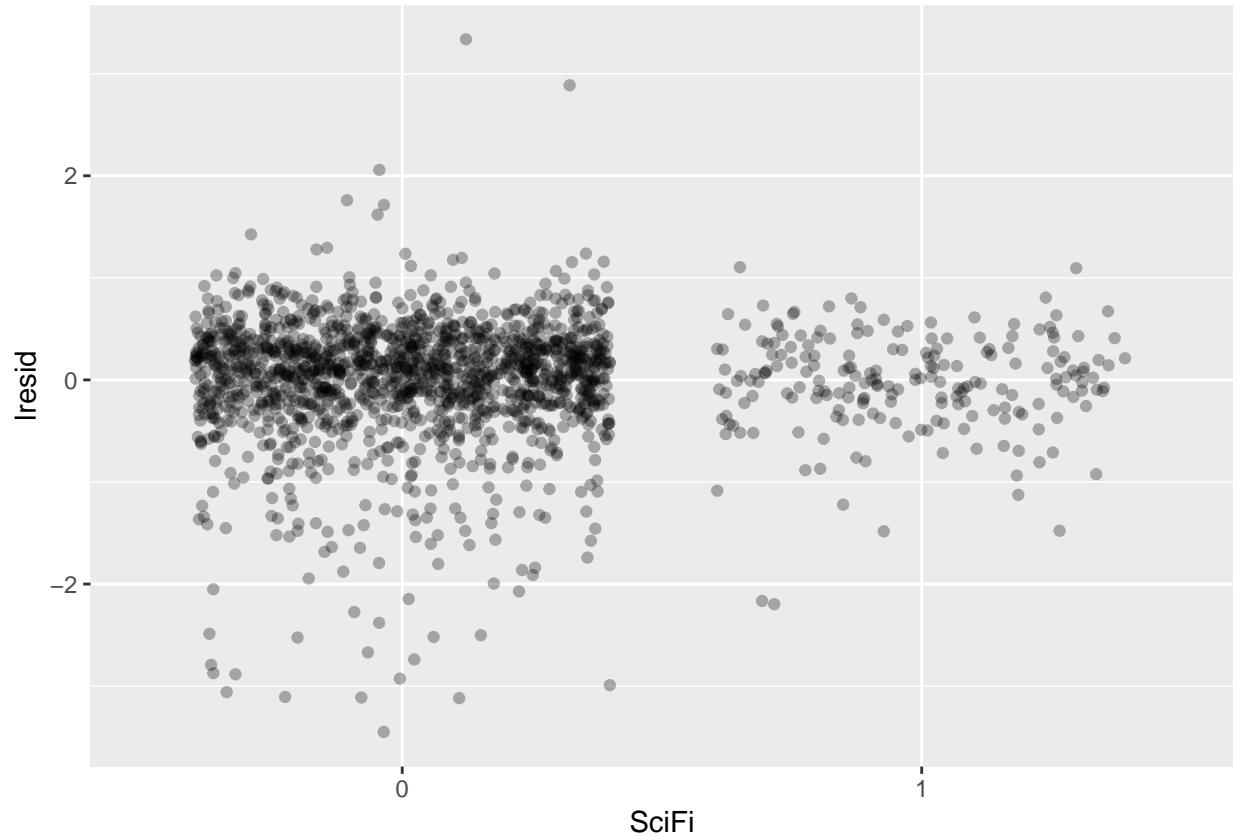
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



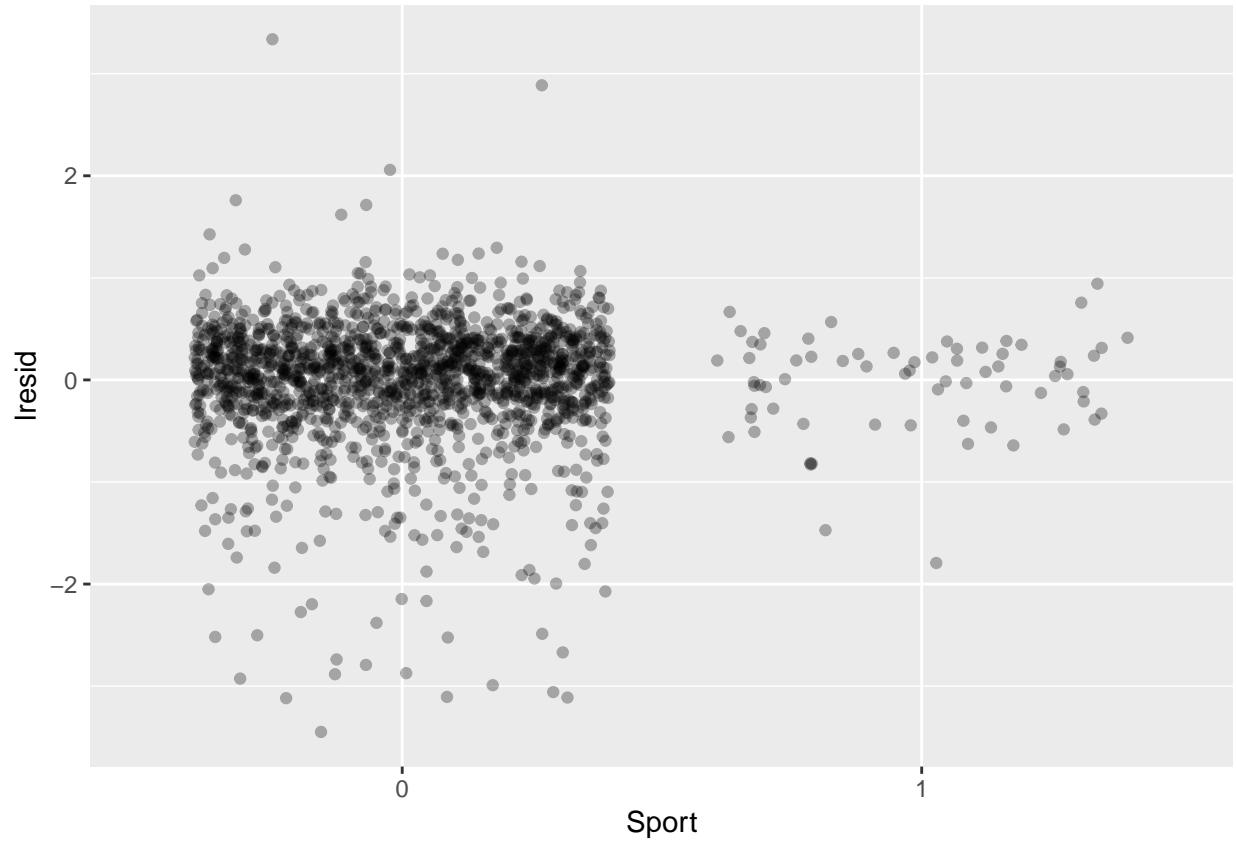
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



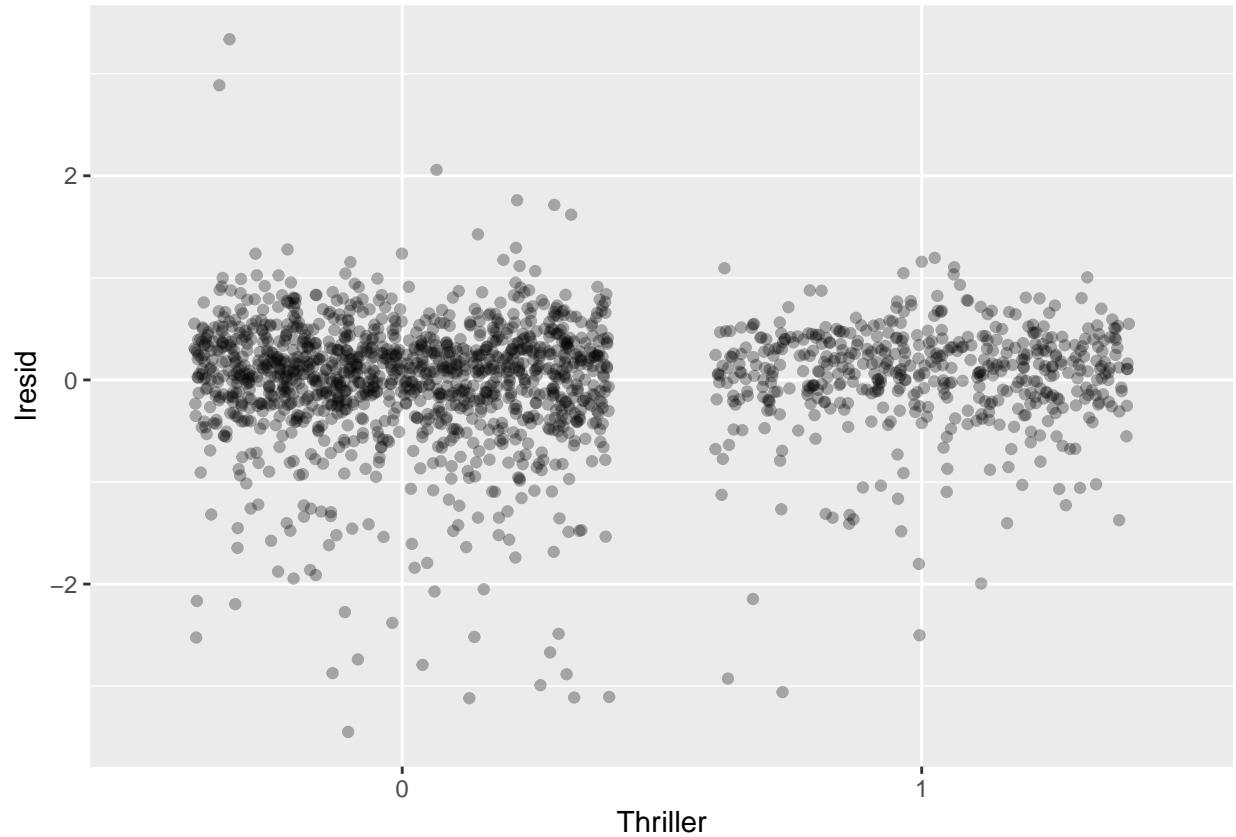
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



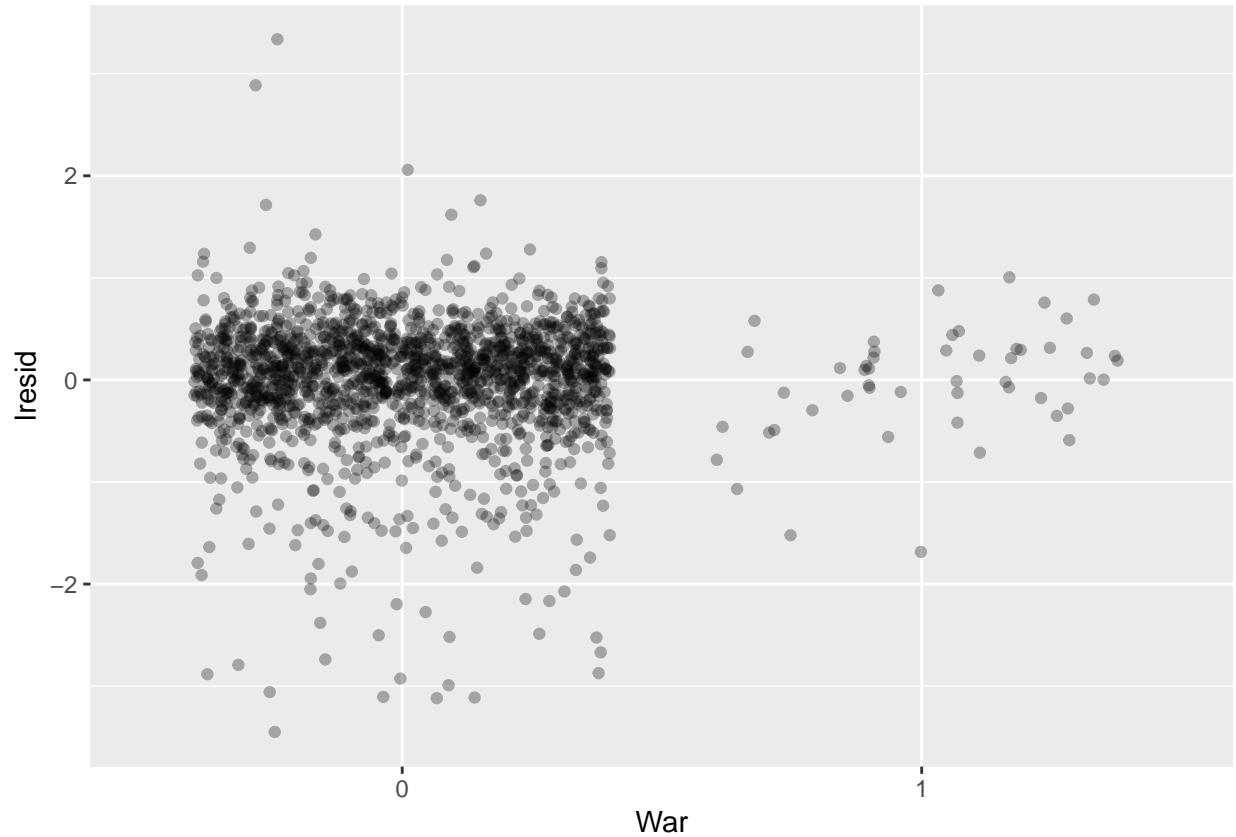
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



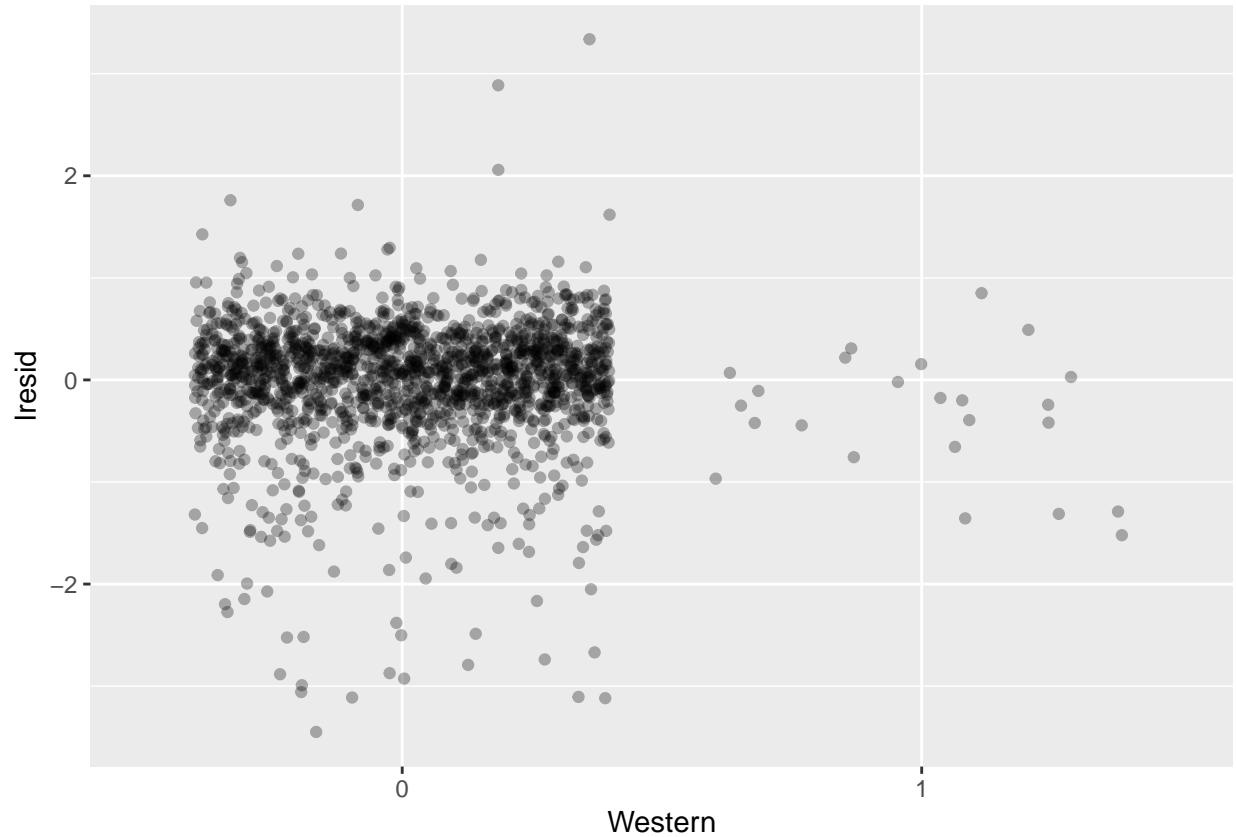
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



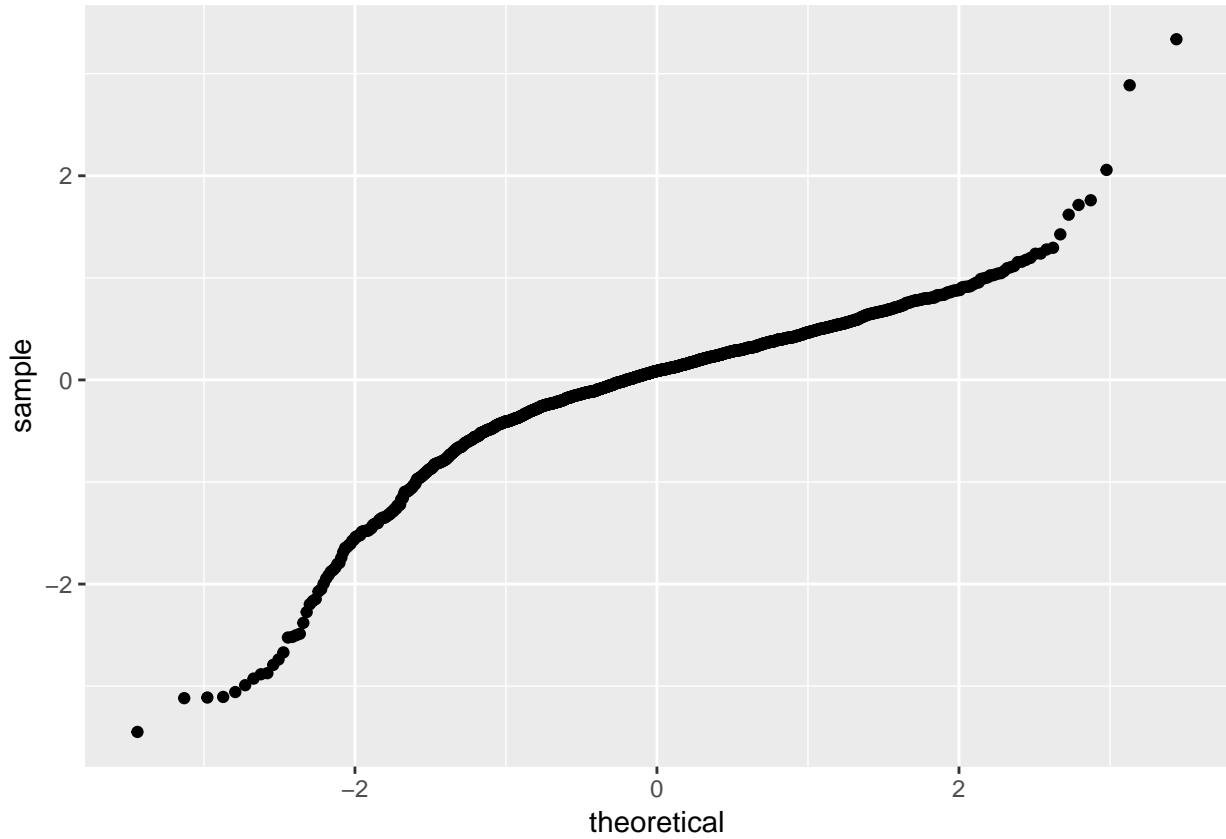
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



```
## Warning: Removed 132 rows containing missing values (geom_point).
```



```
## Warning: Removed 132 rows containing non-finite values (stat_qq).
```



```
# two points with consistently really high residuals (greater than 2) i.e. based on all of their factors
#### make a graph labeling these and pointing out
# often many points with large negative results less than -2: often get movies that are a flop. High budget
train %>%
  add_residuals(mod_all, 'lresid') %>%
  filter(lresid > 2.1)

## # A tibble: 2 x 58
##   director_name num_critic_for_~ director_facebook_likes actor_3_facebook_likes
##   <chr>           <dbl>                 <dbl>                  <dbl>
## 1 Daniel Myrick      360                  19                   39
## 2 Oren Peli          409                 110                   21
## # ... with 54 more variables: actor_2_name <chr>,
## #   actor_1_facebook_likes <dbl>, gross <dbl>, genres <chr>,
## #   actor_1_name <chr>, movie_title <chr>,
## #   cast_total_facebook_likes <dbl>, actor_3_name <chr>,
## #   num_user_for_reviews <dbl>, language <chr>, country <chr>,
## #   content_rating <fct>, budget <dbl>, year <fct>,
## #   actor_2_facebook_likes <dbl>, imdb_score <dbl>,
## #   movie_facebook_likes <dbl>, Action <fct>, Adventure <fct>,
## #   Animation <fct>, Biography <fct>, Comedy <fct>, Crime <fct>,
## #   Documentary <fct>, Drama <fct>, Family <fct>, Fantasy <fct>,
## #   History <fct>, Horror <fct>, Music <fct>, Musical <fct>,
## #   Mystery <fct>, Romance <fct>, SciFi <fct>, Sport <fct>,
## #   Thriller <fct>, War <fct>, Western <fct>, gdpd <dbl>,
## #   real_budget <dbl>, real_gross <dbl>, rev_budget <dbl>,
```

```

## #   profit_dum <dbl>, total_oscars_actor1 <dbl>,
## #   total_oscars_actor2 <dbl>, total_oscars_actor3 <dbl>,
## #   total_oscars_director <fct>, total_oscars_actor <fct>,
## #   real_gross_log <dbl>, real_budget_log <dbl>,
## #   director_facebook_likes_log <dbl>,
## #   cast_total_facebook_likes_log <dbl>, imdb_score_log <dbl>,
## #   lresid <dbl>

train %>%
  add_residuals(mod_all, 'lresid') %>%
  filter(lresid < -2)

## # A tibble: 23 x 58
##   director_name num_critic_for_~ director_facebo~ actor_3_faceboo~
##   <chr>           <dbl>           <dbl>           <dbl>
## 1 Akiva Goldsm~        189            167            778
## 2 Nick Cassave~       177            415            545
## 3 John Hillcoat      355            214           3000
## 4 George Gallo        45             269            384
## 5 John Duigan        16              18            186
## 6 Anand Tucker       134             14            356
## 7 Jake Paltrow        50              17             66
## 8 Timothy Hines        1               0             247
## 9 Michael Mere~       23              7             875
## 10 Damian Nieman      25              0             83
## # ... with 13 more rows, and 54 more variables: actor_2_name <chr>,
## #   actor_1_facebook_likes <dbl>, gross <dbl>, genres <chr>,
## #   actor_1_name <chr>, movie_title <chr>,
## #   cast_total_facebook_likes <dbl>, actor_3_name <chr>,
## #   num_user_for_reviews <dbl>, language <chr>, country <chr>,
## #   content_rating <fct>, budget <dbl>, year <fct>,
## #   actor_2_facebook_likes <dbl>, imdb_score <dbl>,
## #   movie_facebook_likes <dbl>, Action <fct>, Adventure <fct>,
## #   Animation <fct>, Biography <fct>, Comedy <fct>, Crime <fct>,
## #   Documentary <fct>, Drama <fct>, Family <fct>, Fantasy <fct>,
## #   History <fct>, Horror <fct>, Music <fct>, Musical <fct>,
## #   Mystery <fct>, Romance <fct>, SciFi <fct>, Sport <fct>,
## #   Thriller <fct>, War <fct>, Western <fct>, gdpp <dbl>,
## #   real_budget <dbl>, real_gross <dbl>, rev_budget <dbl>,
## #   profit_dum <dbl>, total_oscars_actor1 <dbl>,
## #   total_oscars_actor2 <dbl>, total_oscars_actor3 <dbl>,
## #   total_oscars_director <fct>, total_oscars_actor <fct>,
## #   real_gross_log <dbl>, real_budget_log <dbl>,
## #   director_facebook_likes_log <dbl>,
## #   cast_total_facebook_likes_log <dbl>, imdb_score_log <dbl>,
## #   lresid <dbl>

```

Prediction

```

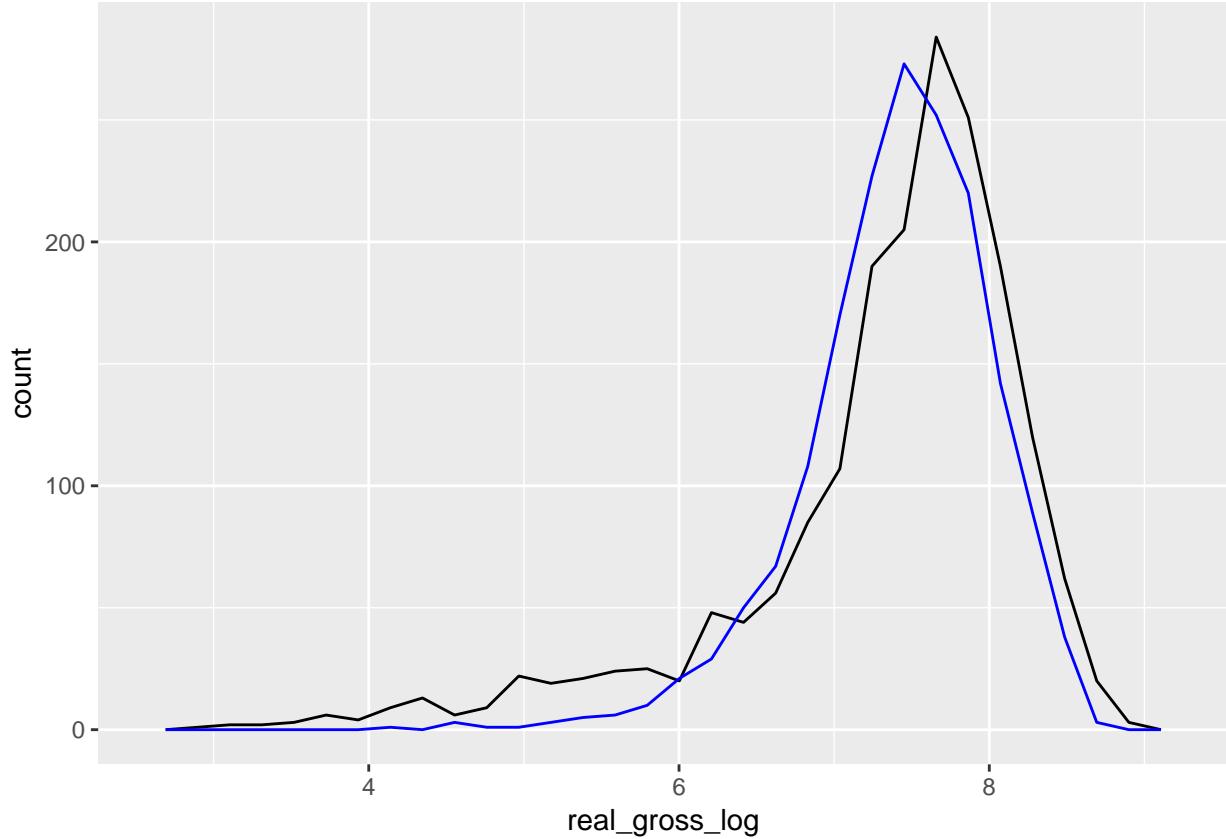
train %>%
  add_predictions(mod_all, 'lpred') %>%
  ggplot() +
  geom_freqpoly(aes(x = real_gross_log)) +
  geom_freqpoly(aes(x = lpred), color = 'blue')

```

```

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 132 rows containing non-finite values (stat_bin).

```



Fit Model W/O Assuming Genre

```

# stepwise
# ALL potentially relevant variables
rmse_lst <- step_wise_loop(df = train %>% select(all_genre_vars, content_rating, real_budget, year,
                                                 total_oscars_actor, total_oscars_director,
                                                 imdb_score_log, real_budget_log,
                                                 director_facebook_likes_log, cast_total_facebook_likes)

## real_budget_log
##      0.6497192
## [1] 1
## imdb_score_log
##      0.6382915
## [1] 2
##      year
## 0.6281035
## [1] 3
## Comedy
## 0.6221681
## [1] 4

```

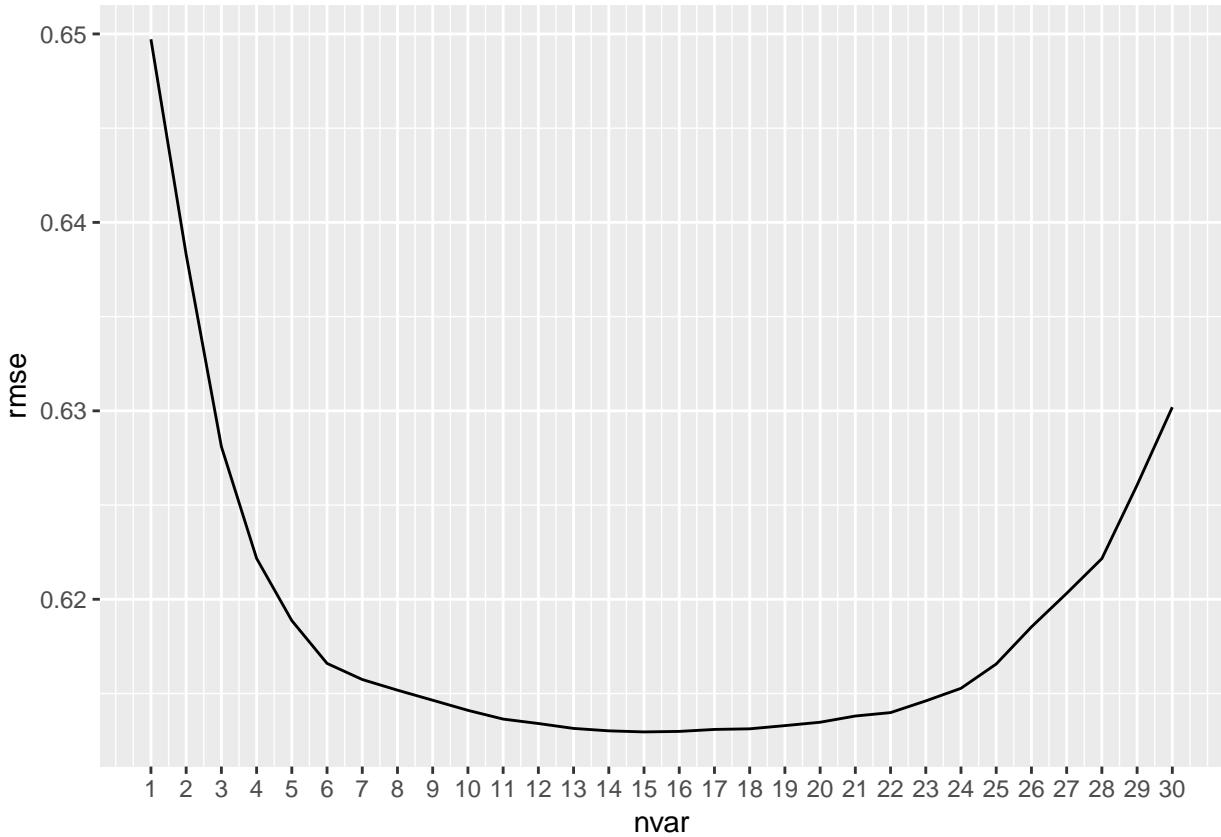
```
## content_rating
##      0.6188688
## [1] 5
##   Mystery
## 0.6165971
## [1] 6
##   Biography
## 0.6157421
## [1] 7
##   Musical
## 0.6151764
## [1] 8
##   Sport
## 0.6146438
## [1] 9
##   Crime
## 0.6141069
## [1] 10
## Documentary
## 0.6136432
## [1] 11
##   Music
## 0.6134072
## [1] 12
##   SciFi
## 0.6131476
## [1] 13
##   Action
## 0.6130183
## [1] 14
##   History
## 0.6129618
## [1] 15
## total_oscars_director
##      0.6129888
## [1] 16
##   Fantasy
## 0.613094
## [1] 17
## real_budget
## 0.6131265
## [1] 18
## Animation
## 0.6132966
## [1] 19
##   Romance
## 0.6134755
## [1] 20
##   War
## 0.6138056
## [1] 21
##   Drama
## 0.6139871
## [1] 22
```

```

## Thriller
## 0.6146081
## [1] 23
## cast_total_facebook_likes_log
## 0.6152761
## [1] 24
## total_oscars_actor
## 0.6165643
## [1] 25
## Family
## 0.6185332
## [1] 26
## director_facebook_likes_log
## 0.6203153
## [1] 27
## Horror
## 0.6221584
## [1] 28
## Western
## 0.6260611
## [1] 29
## Adventure
## 0.6301929

# graph RMSE vs number of variables
fit_rmse <- tibble(nvar = 1:length(rmse_lst),
                     rmse = rmse_lst)
ggplot(fit_rmse) + geom_line(aes(x = nvar, y = rmse))+
  scale_x_continuous(breaks = seq(1, length(rmse_lst), by = 1))

```



```
# after var 6, decreases too small or increase

mod_all2 <- lm(real_gross_log ~ real_budget_log + imdb_score_log + year + Comedy + content_rating + Mystery,
                 data = train)

summary(mod_all2) # still a lot of insignificant

## 
## Call:
## lm(formula = real_gross_log ~ real_budget_log + imdb_score_log +
##     year + Comedy + content_rating + Mystery, data = train)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -3.5488 -0.2234  0.0941  0.3462  3.5384 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 0.290225  0.396873  0.731   0.46471    
## real_budget_log 0.827045  0.026177 31.594  < 2e-16 *** 
## imdb_score_log 1.710134  0.201331  8.494  < 2e-16 *** 
## year1981    -0.254810  0.376404 -0.677   0.49852    
## year1982     0.296995  0.346191  0.858   0.39107    
## year1983     0.246156  0.392671  0.627   0.53083    
## year1984     0.506559  0.335606  1.509   0.13139    
## year1985     0.345134  0.346559  0.996   0.31945
```

```

## year1986          0.079291  0.331303  0.239  0.81088
## year1987          0.212157  0.322316  0.658  0.51048
## year1988          0.227782  0.315951  0.721  0.47105
## year1989          0.277787  0.311206  0.893  0.37219
## year1990          0.162527  0.312245  0.521  0.60278
## year1991          0.019237  0.315800  0.061  0.95143
## year1992          0.039691  0.321044  0.124  0.90162
## year1993          0.005509  0.312227  0.018  0.98593
## year1994          -0.199457  0.306368 -0.651  0.51511
## year1995          -0.051840  0.298095 -0.174  0.86196
## year1996          -0.220062  0.291028 -0.756  0.44966
## year1997          -0.138061  0.291017 -0.474  0.63527
## year1998          -0.266021  0.290209 -0.917  0.35946
## year1999          -0.248466  0.286897 -0.866  0.38659
## year2000          -0.155180  0.288267 -0.538  0.59043
## year2001          -0.261791  0.286269 -0.914  0.36059
## year2002          -0.237985  0.286021 -0.832  0.40550
## year2003          -0.193529  0.287432 -0.673  0.50085
## year2004          -0.213507  0.287003 -0.744  0.45703
## year2005          -0.226318  0.285994 -0.791  0.42886
## year2006          -0.316366  0.286338 -1.105  0.26937
## year2007          -0.284521  0.287894 -0.988  0.32316
## year2008          -0.335182  0.285977 -1.172  0.24134
## year2009          -0.342031  0.285352 -1.199  0.23084
## year2010          -0.350964  0.285611 -1.229  0.21931
## year2011          -0.232108  0.287990 -0.806  0.42038
## year2012          -0.124210  0.286558 -0.433  0.66474
## year2013          -0.046933  0.286103 -0.164  0.86972
## year2014          -0.112968  0.288063 -0.392  0.69499
## year2015          -0.223600  0.290358 -0.770  0.44136
## year2016          -0.086210  0.306680 -0.281  0.77866
## Comedy1           0.083397  0.032751  2.546  0.01097 *
## content_ratingNC-17 -0.160263  0.272686 -0.588  0.55680
## content_ratingPG    -0.118755  0.119473 -0.994  0.32037
## content_ratingPG-13 -0.156513  0.116101 -1.348  0.17782
## content_ratingR     -0.325984  0.116336 -2.802  0.00514 **
## Mystery1           0.088559  0.053305  1.661  0.09683 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6203 on 1674 degrees of freedom
##   (132 observations deleted due to missingness)
## Multiple R-squared:  0.471, Adjusted R-squared:  0.4571
## F-statistic: 33.87 on 44 and 1674 DF, p-value: < 2.2e-16
rmse(mod_all2, data = valid) # fit is somewhat better (and fewer variables)

## [1] 0.6165971
# when consider factors as one variable, they are significant
anova(mod_all2)

## Analysis of Variance Table
##
## Response: real_gross_log

```

```

##          Df Sum Sq Mean Sq F value    Pr(>F)
## real_budget_log     1 491.67 491.67 1277.7221 < 2.2e-16 ***
## imdb_score_log     1   28.18   28.18   73.2444 < 2.2e-16 ***
## year              36   36.51    1.01    2.6359 5.733e-07 ***
## Comedy             1    4.01    4.01   10.4110  0.001277 **
## content_rating      4   12.05    3.01    7.8296 2.994e-06 ***
## Mystery            1    1.06    1.06    2.7601  0.096829 .
## Residuals         1674 644.16    0.38
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
# number of observations
nobs(mod_all2)

```

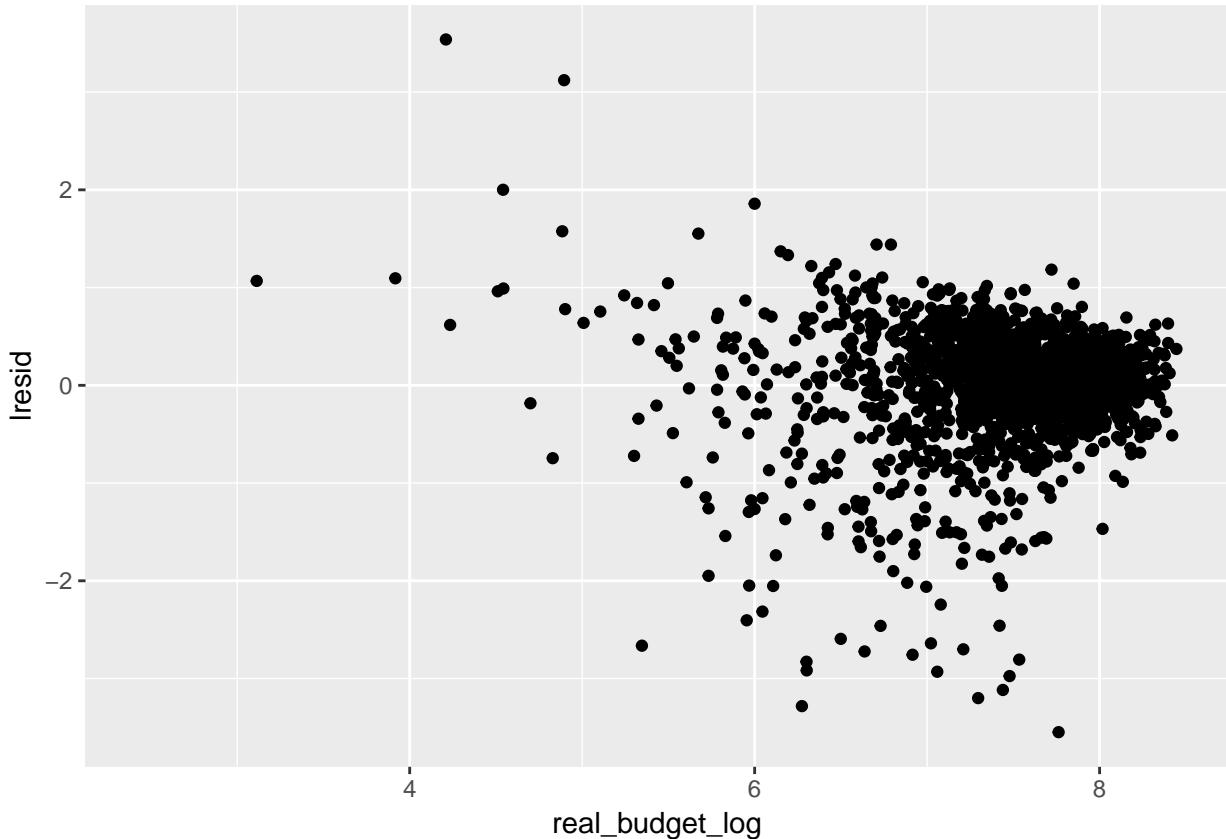
```
## [1] 1719
```

New Residuals

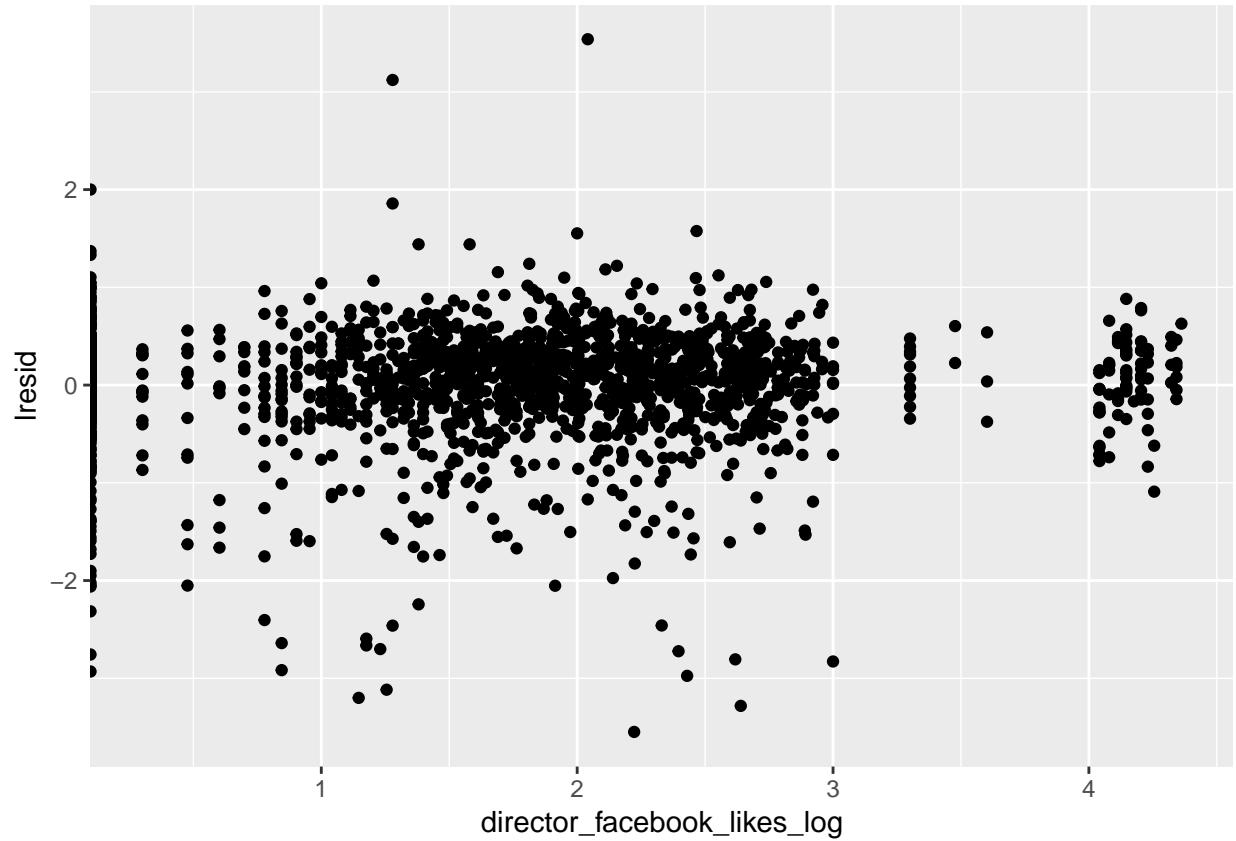
Graph residuals of included and excluded variables: have we captured all of the relationships?

```
gr_resid(mod_all2)
```

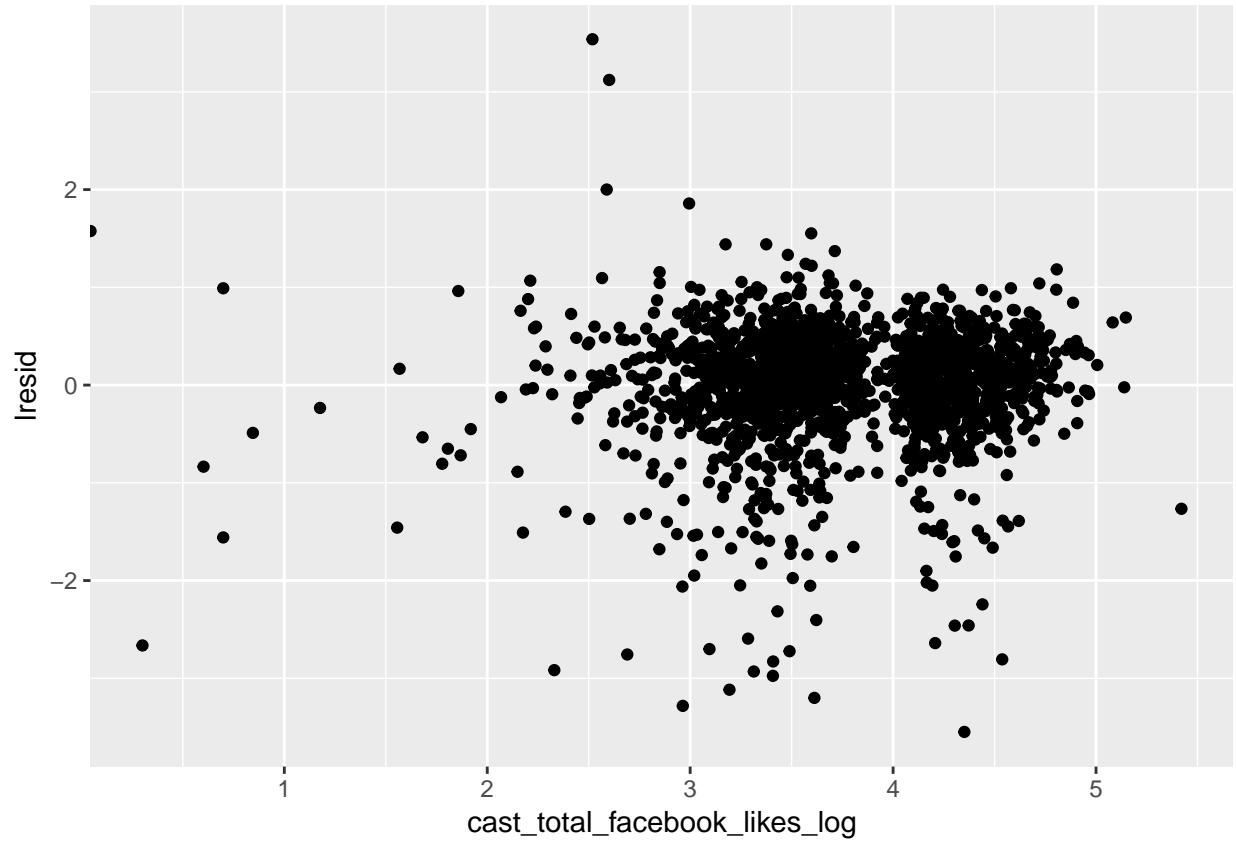
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



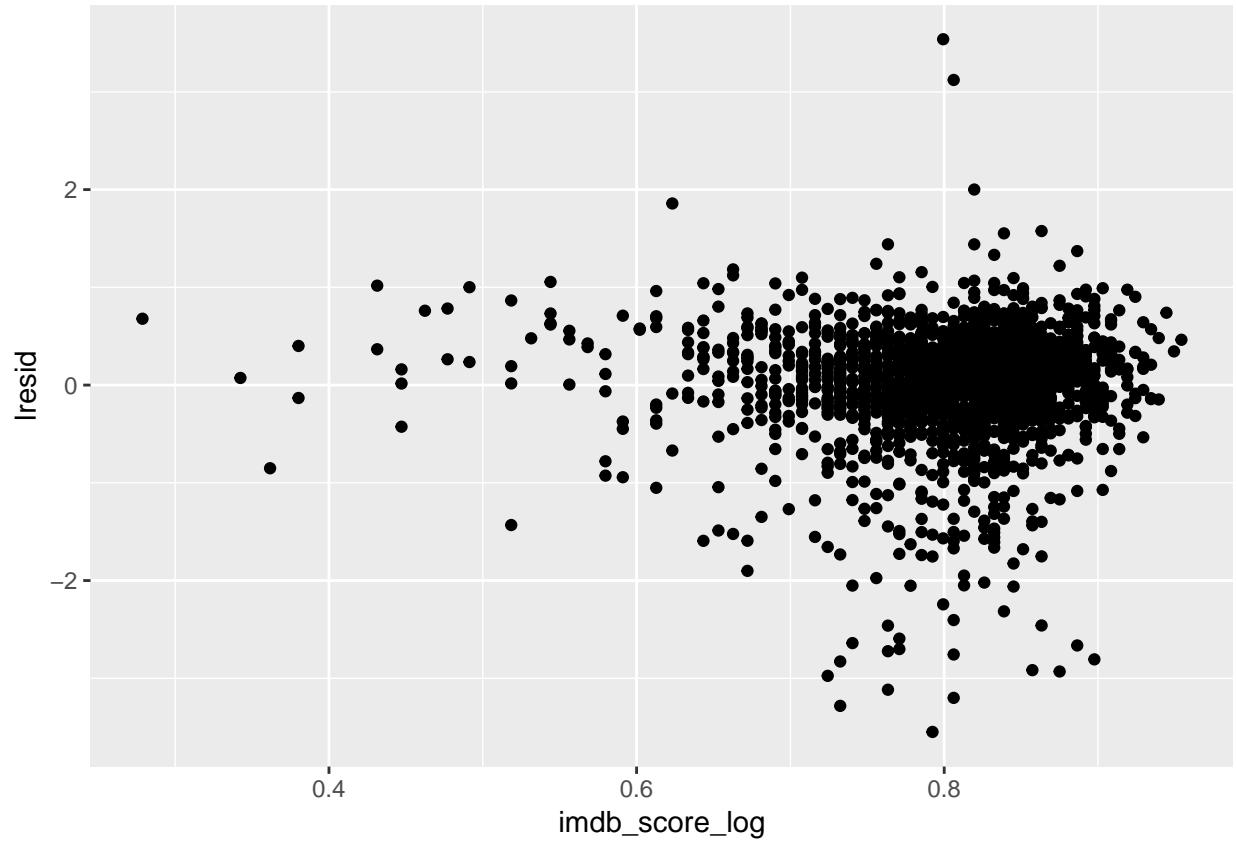
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



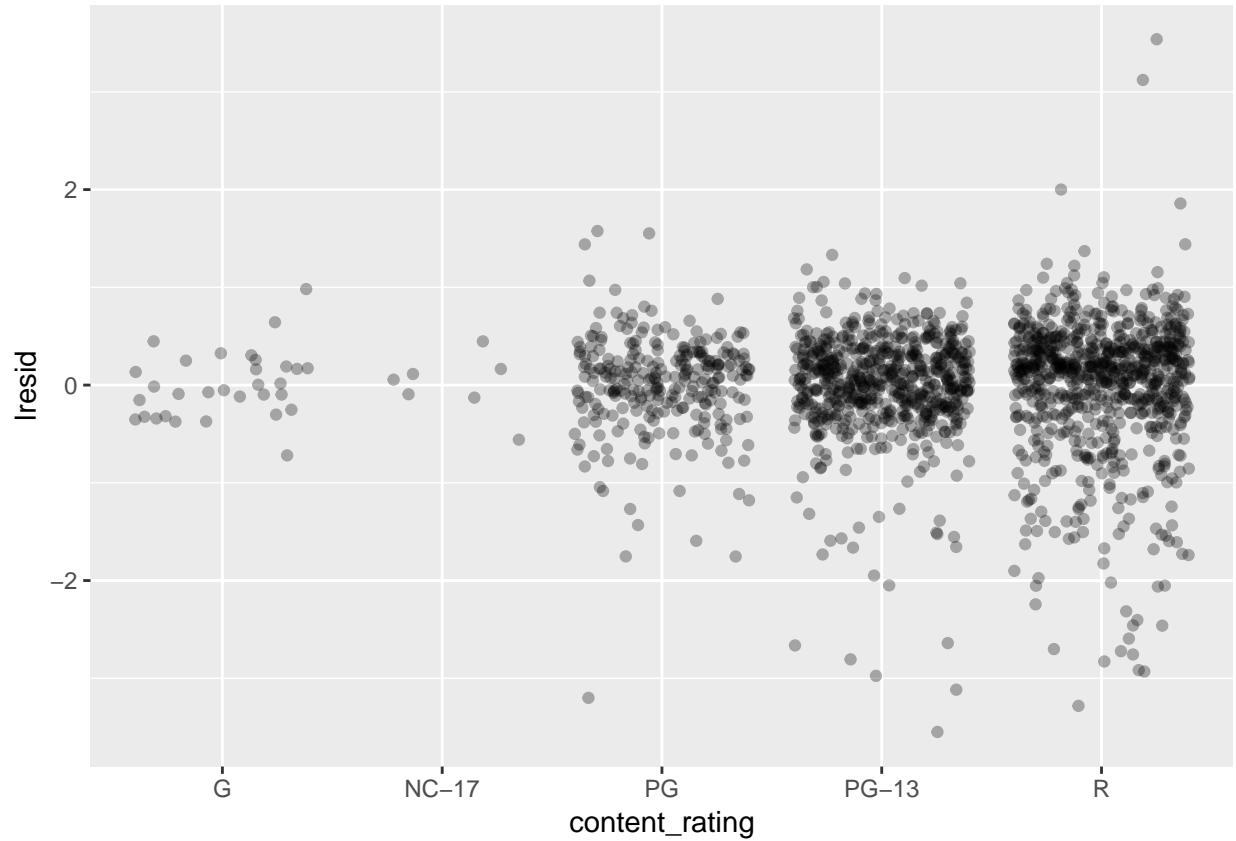
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



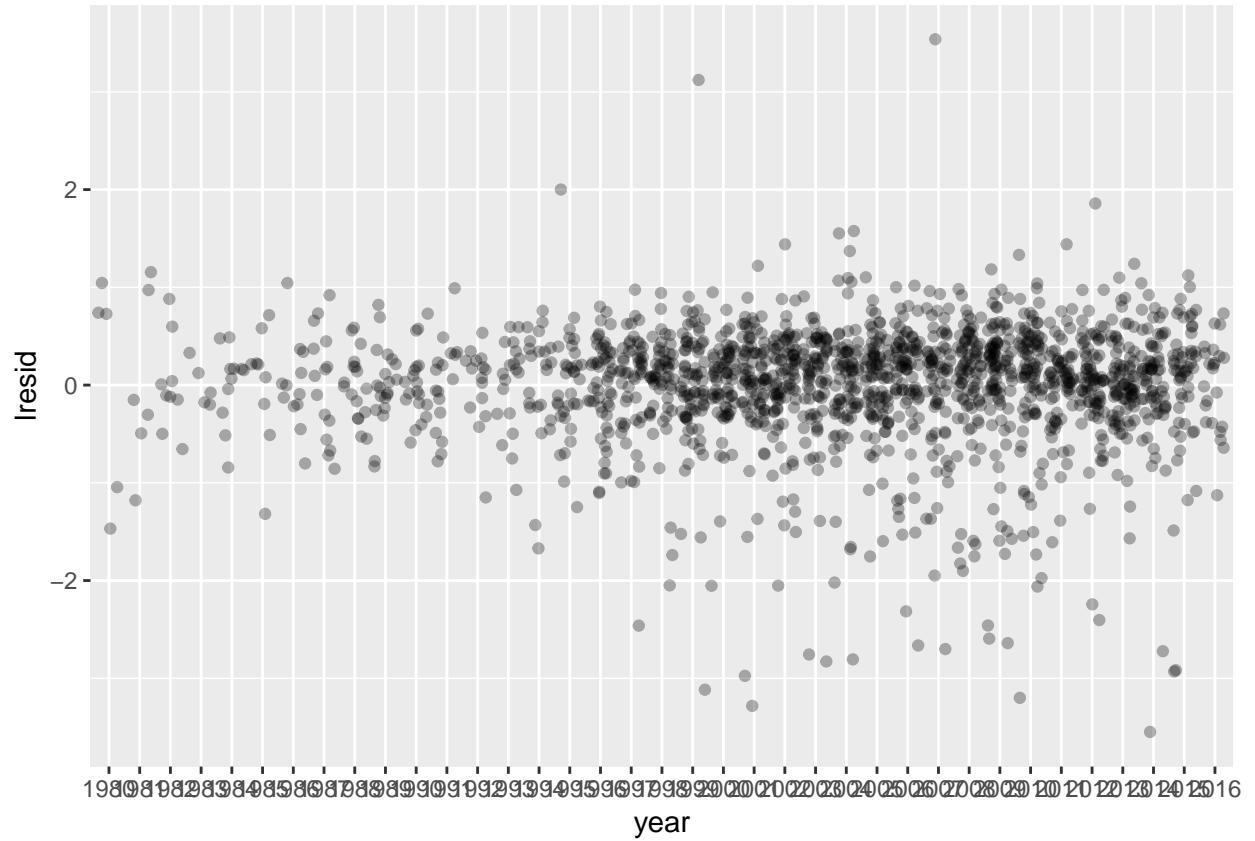
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



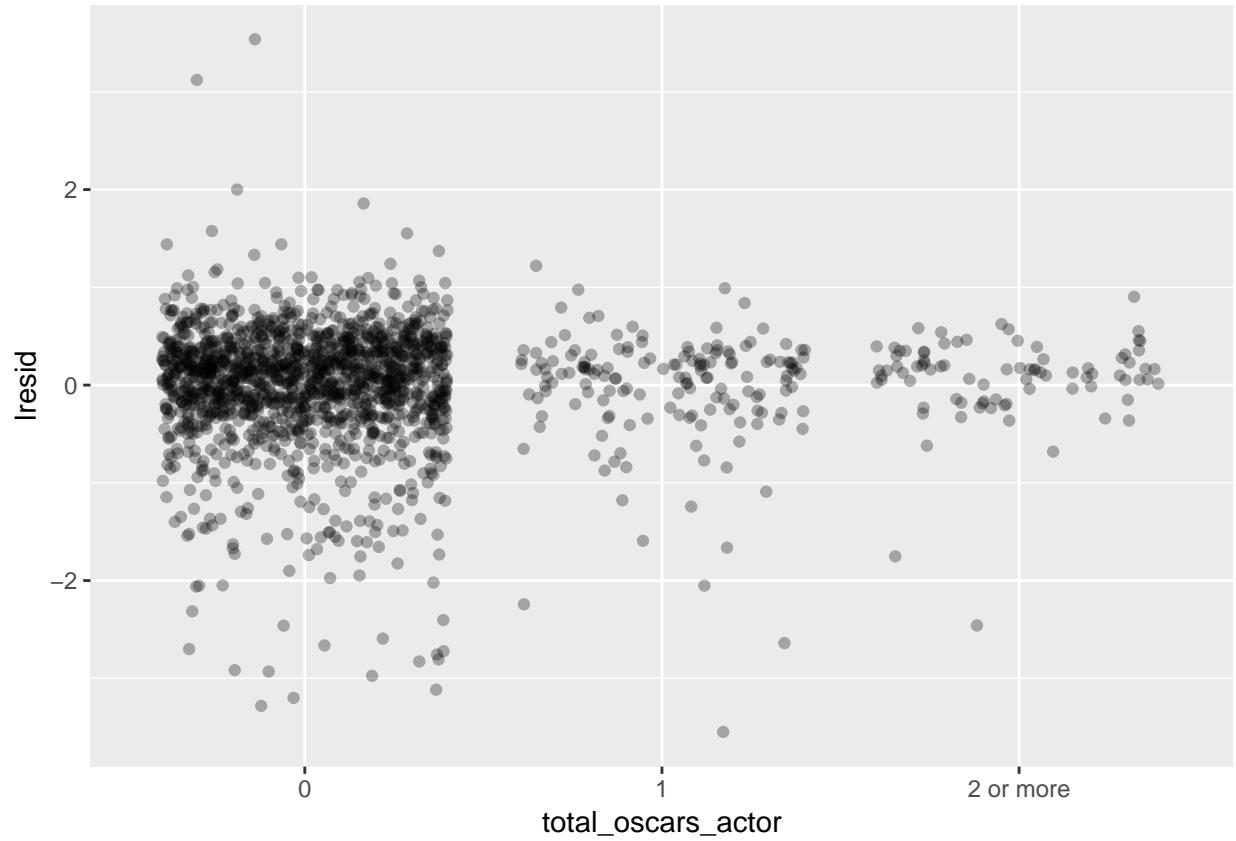
```
## Warning: Removed 94 rows containing missing values (geom_point).
```



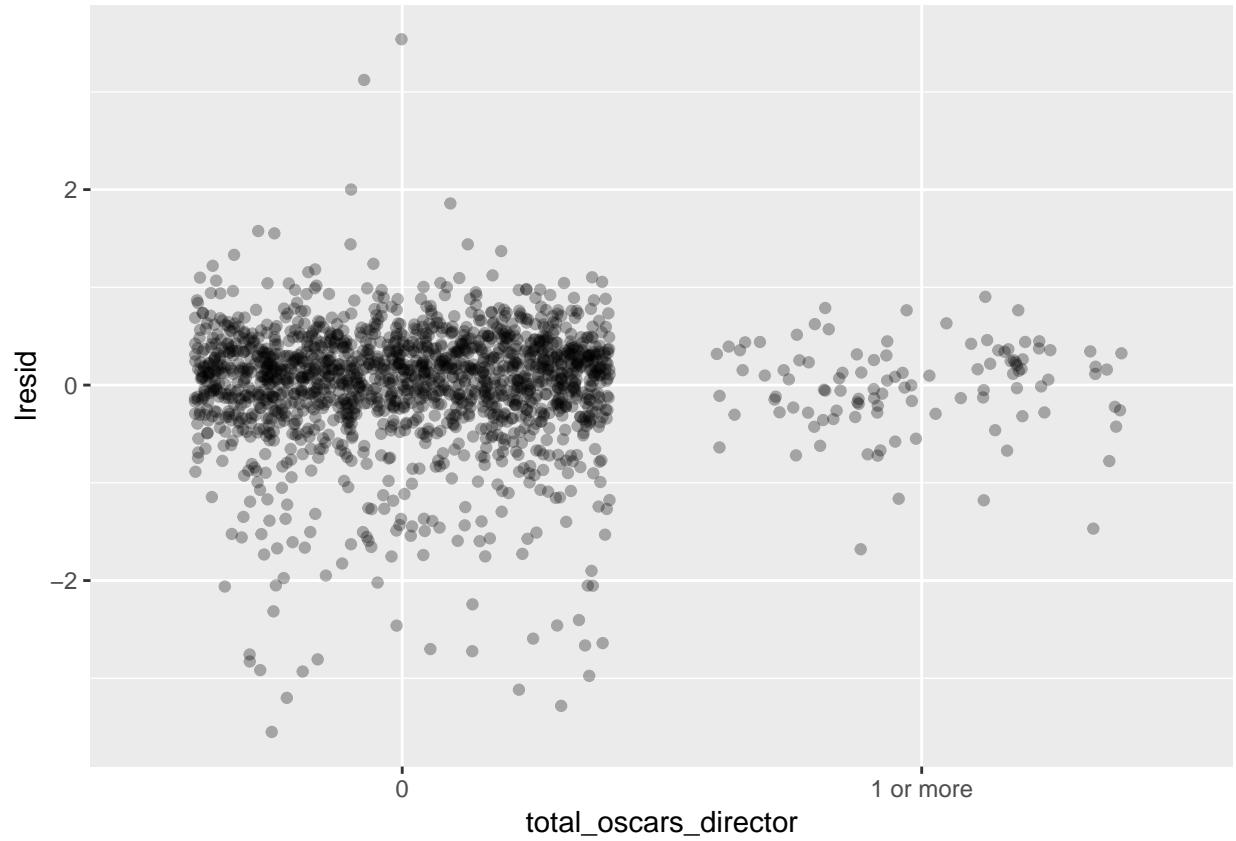
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



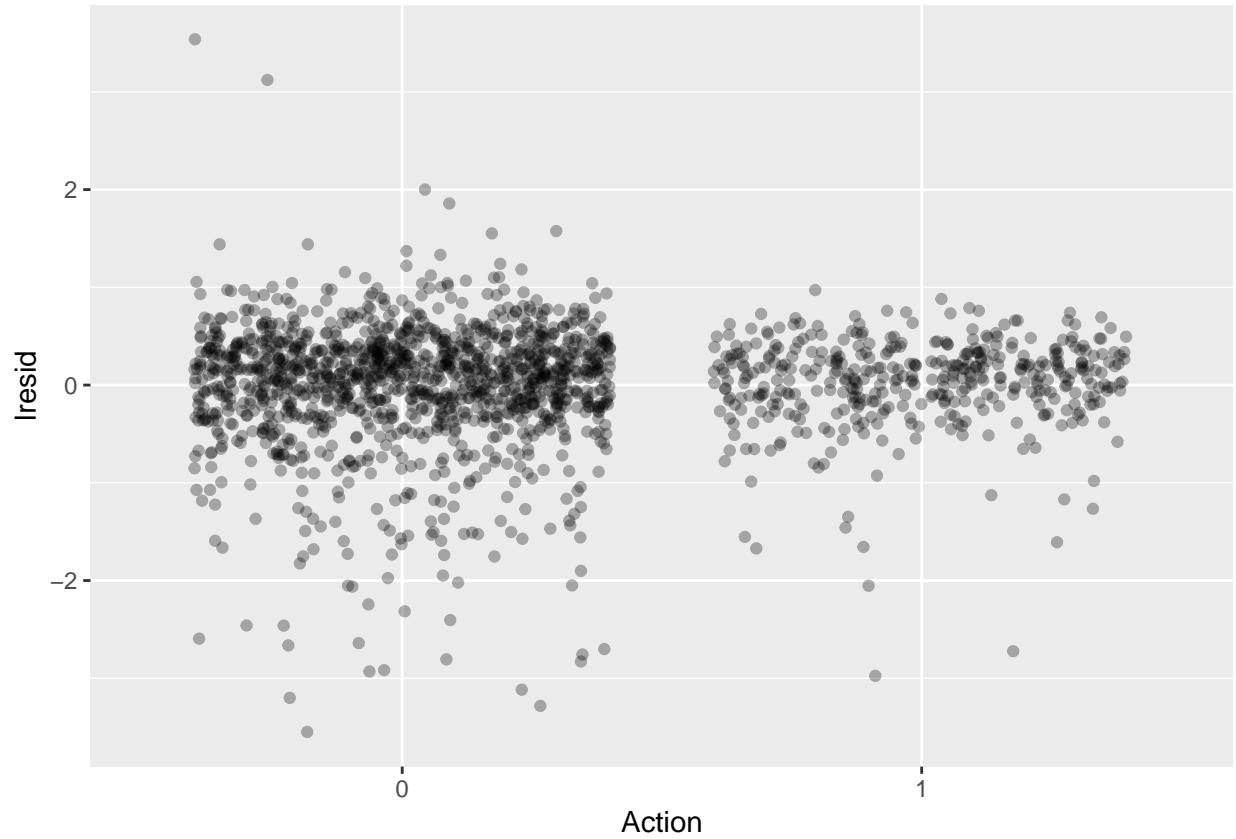
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



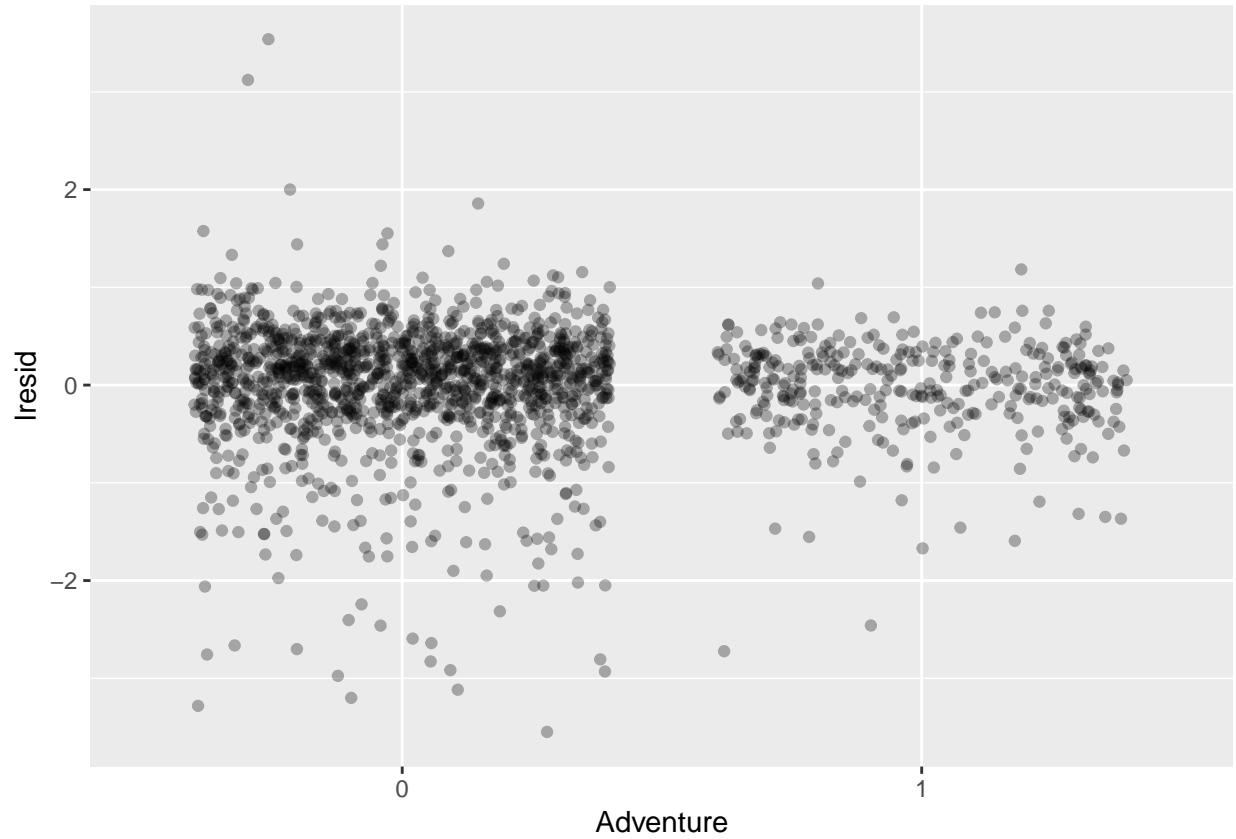
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



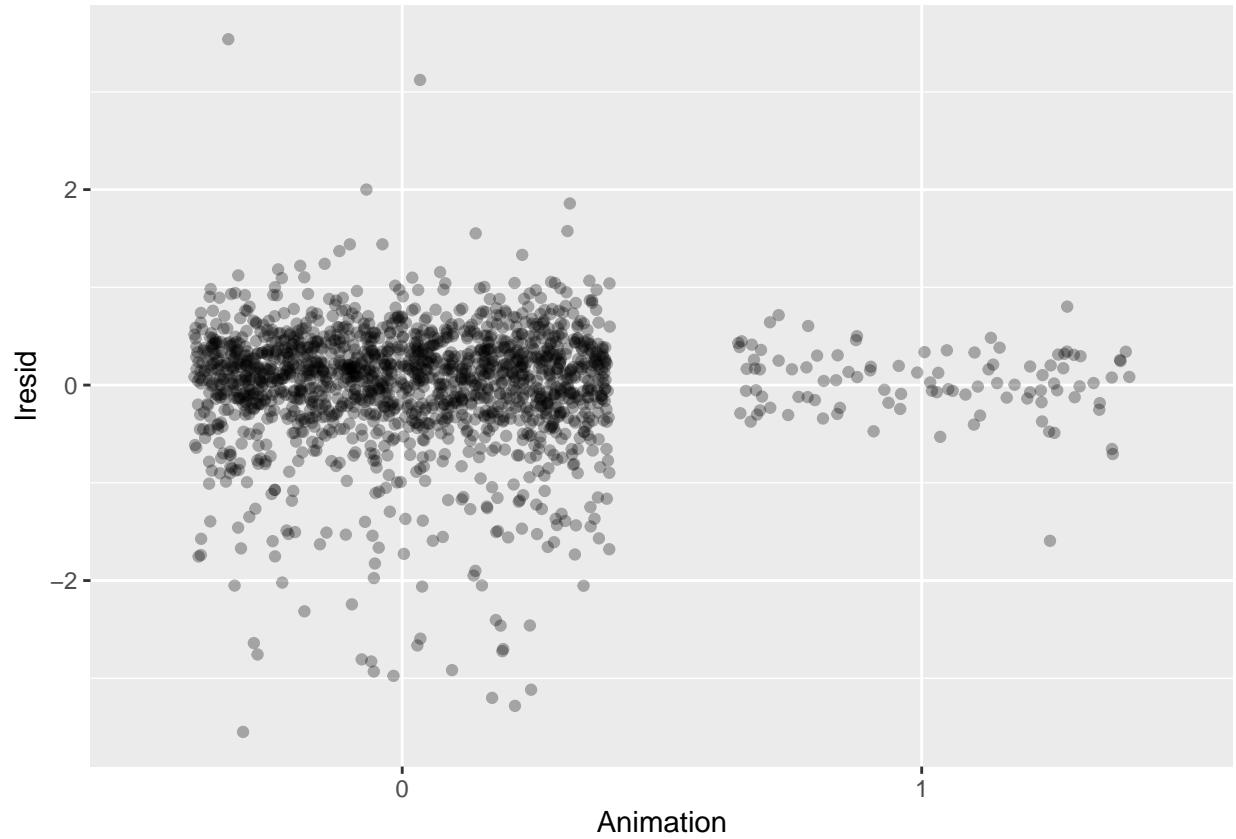
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



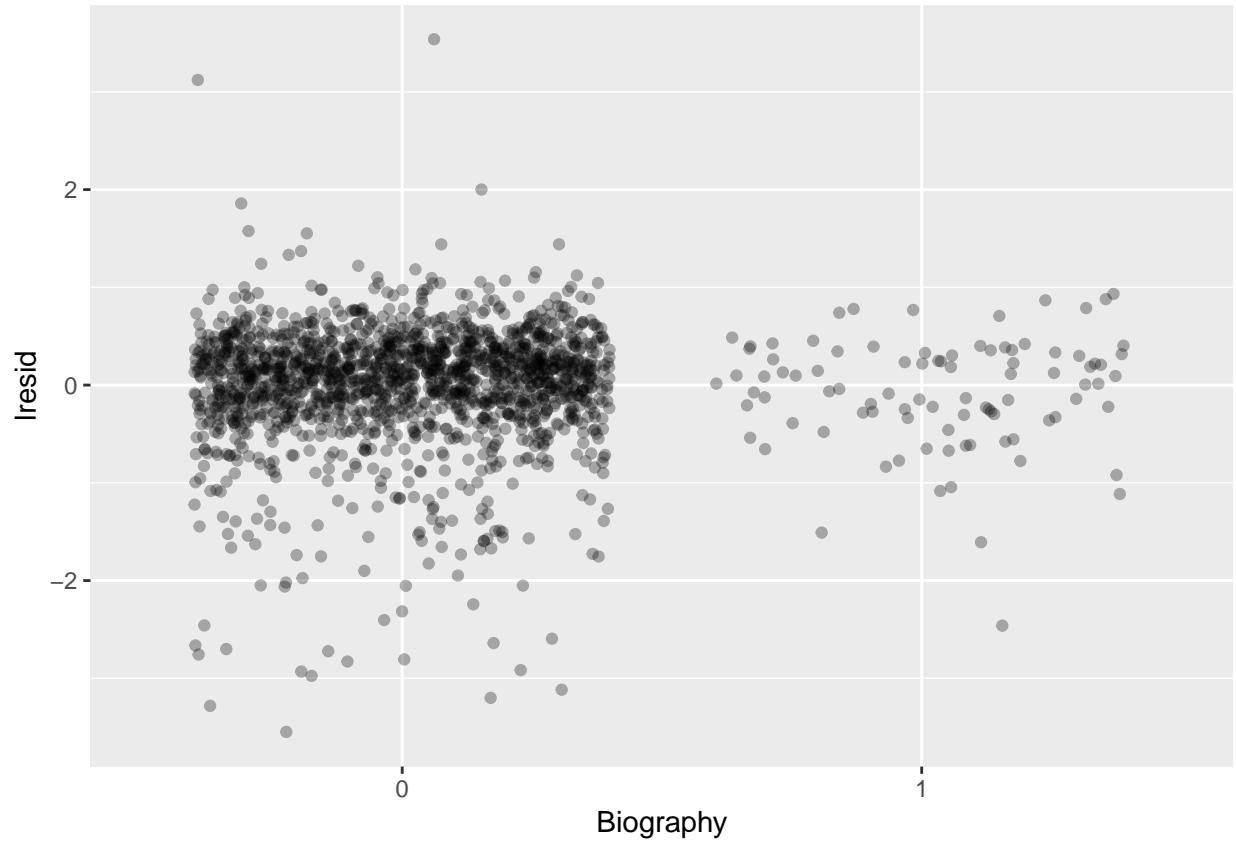
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



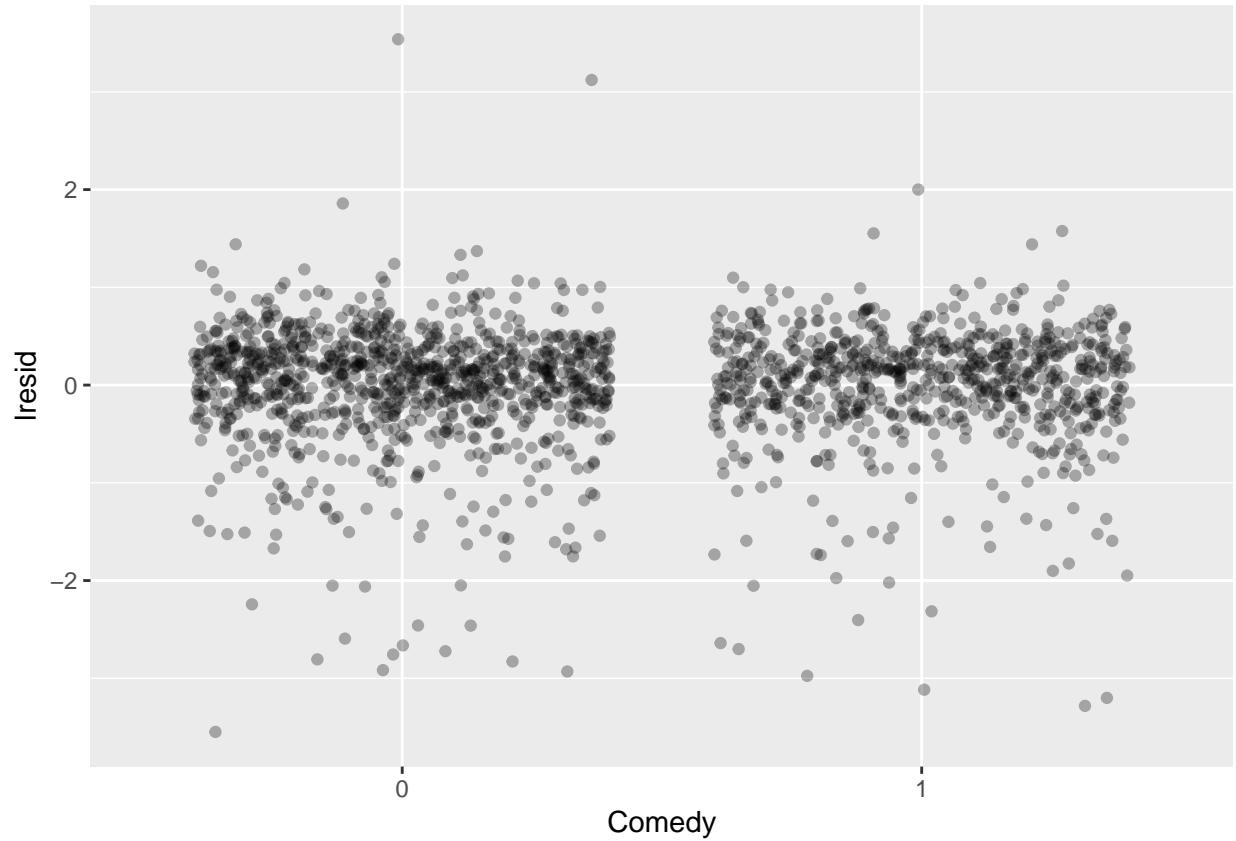
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



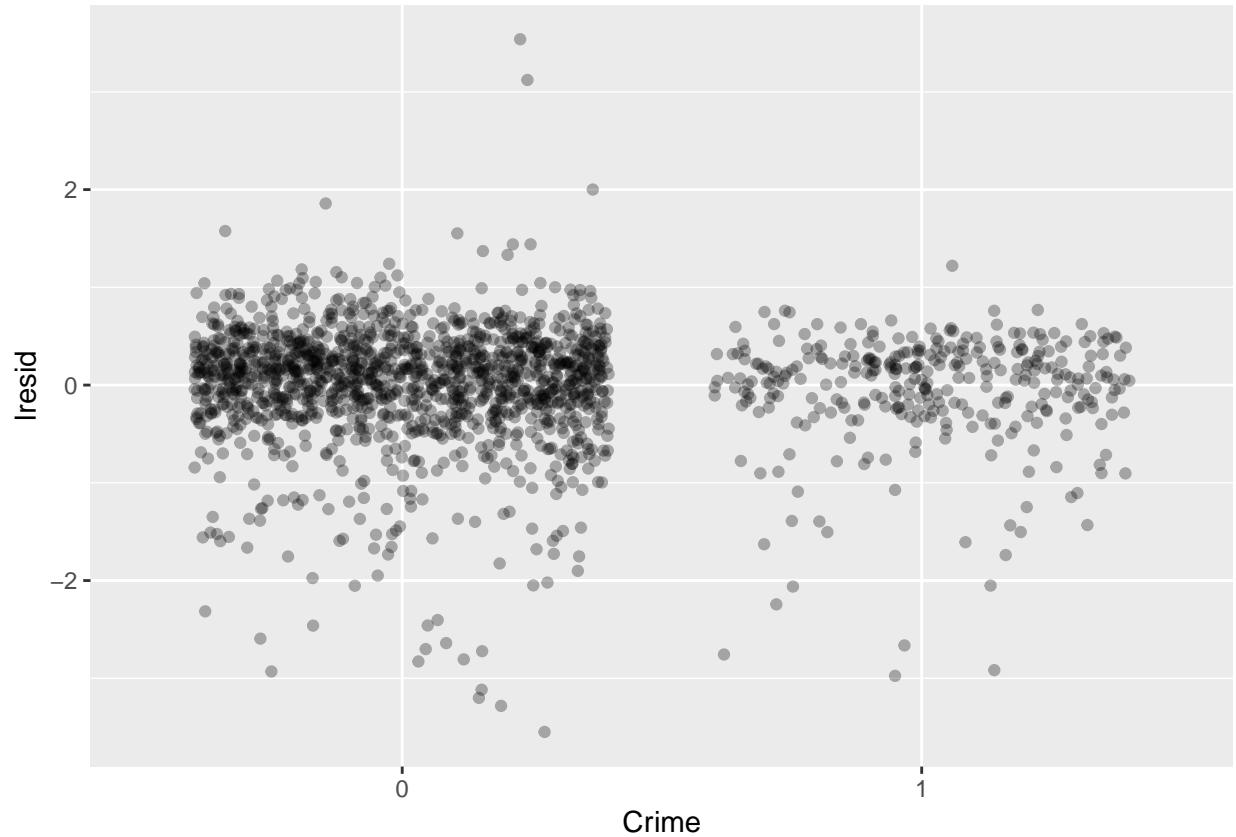
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



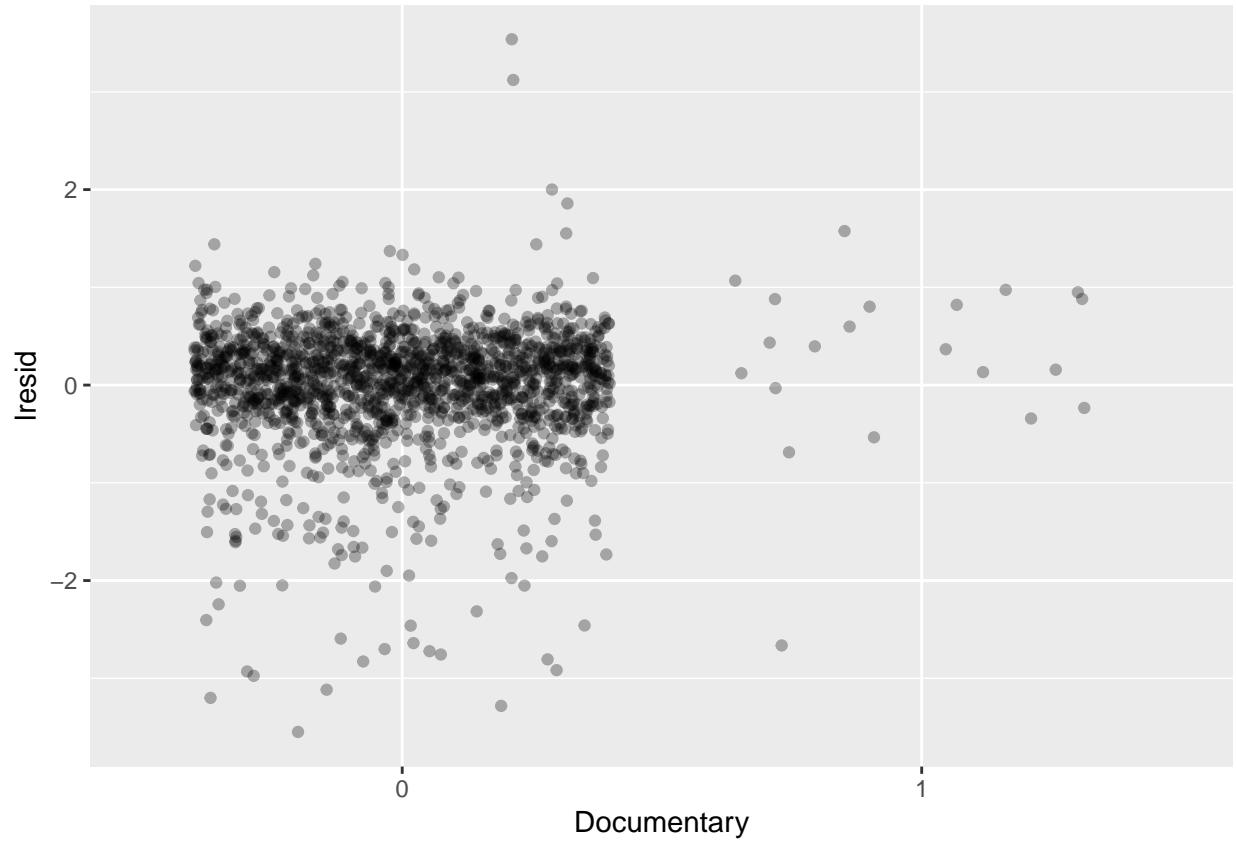
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



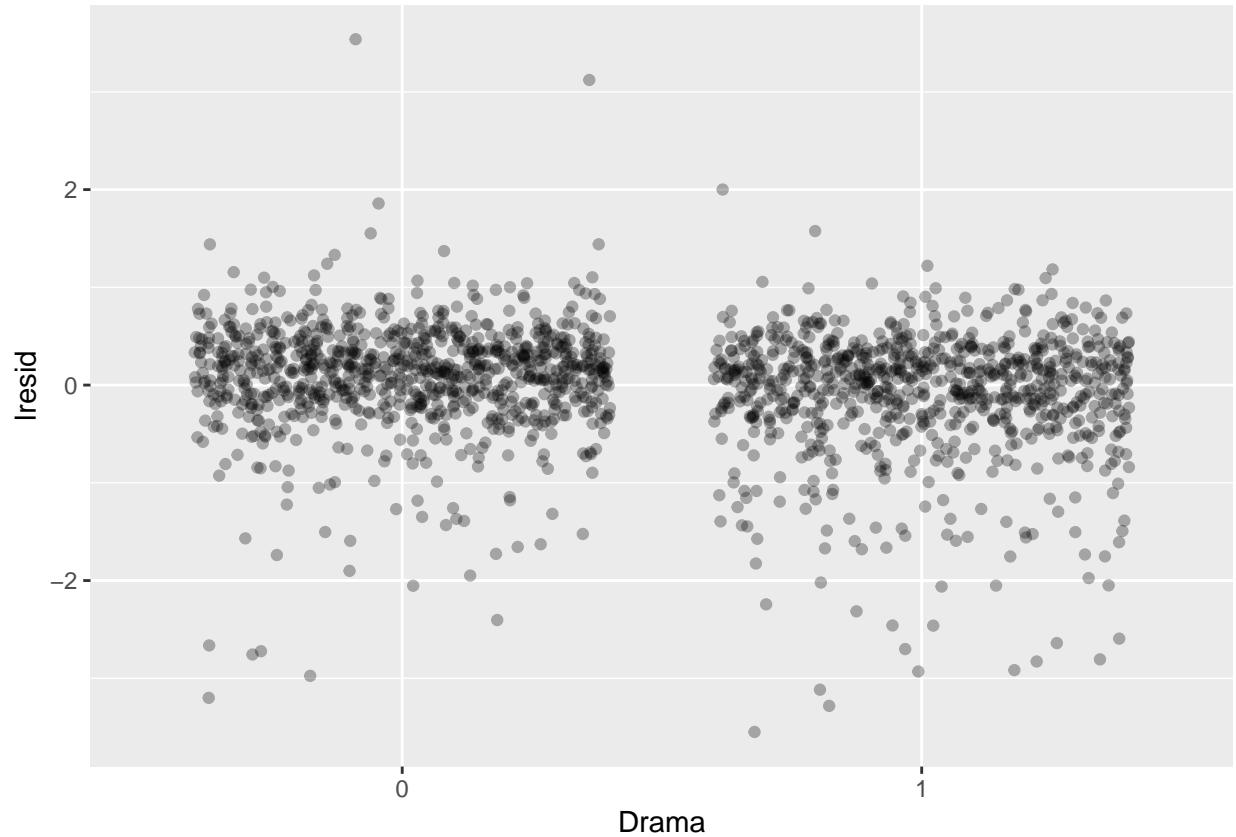
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



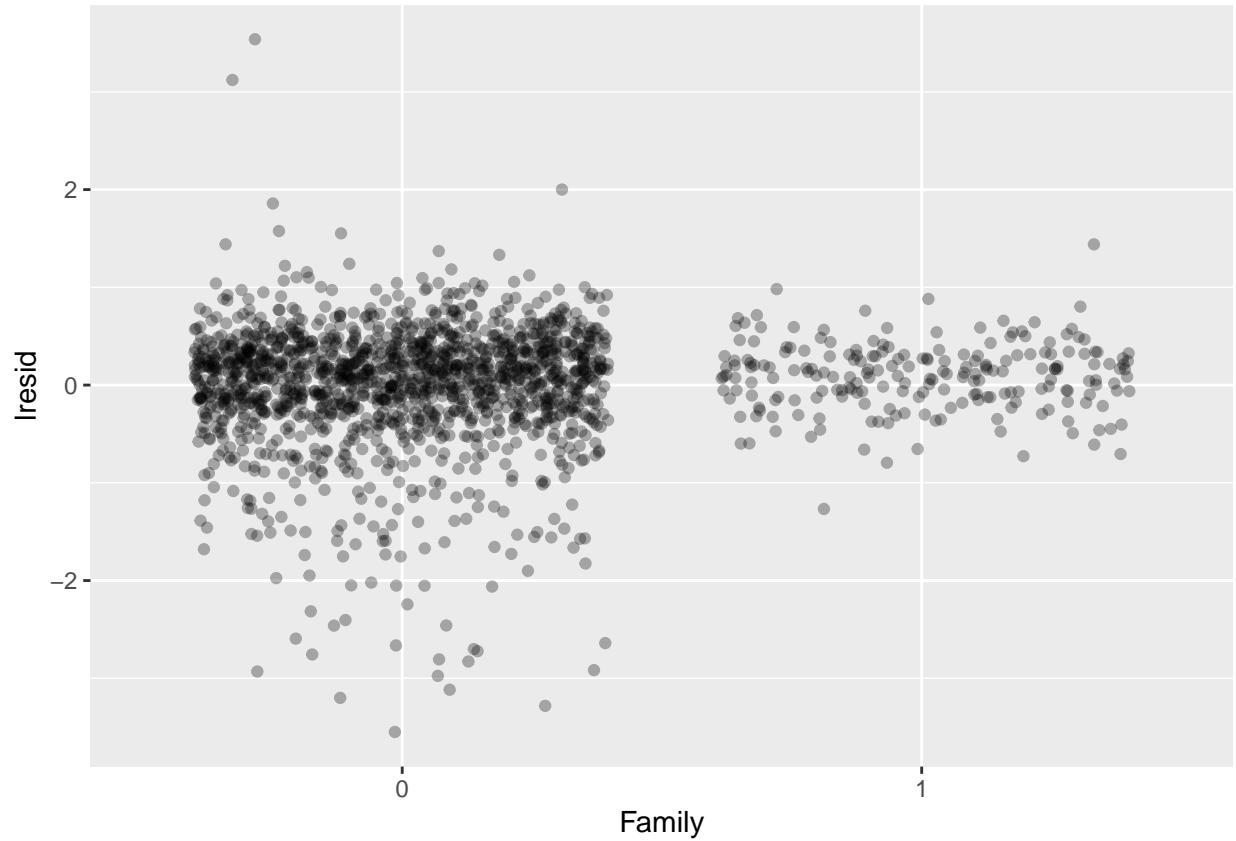
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



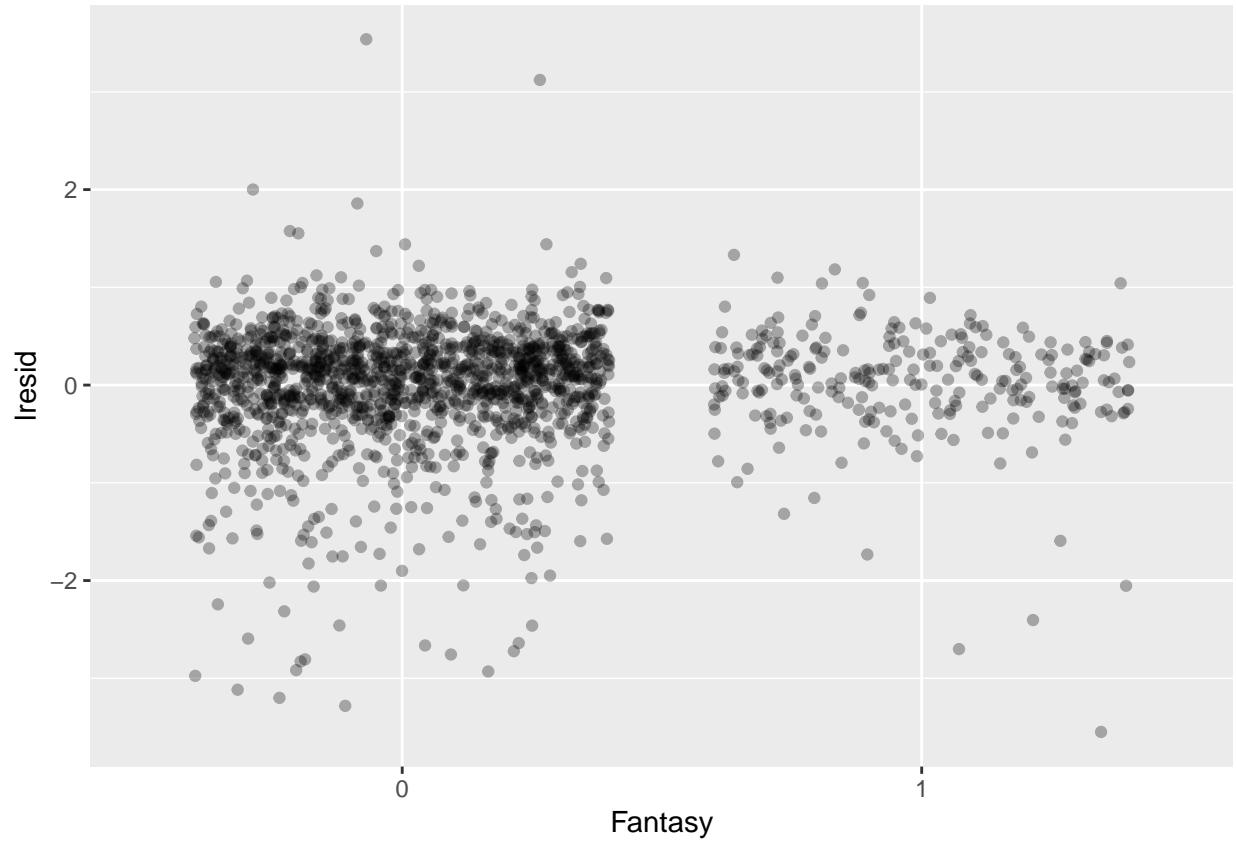
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



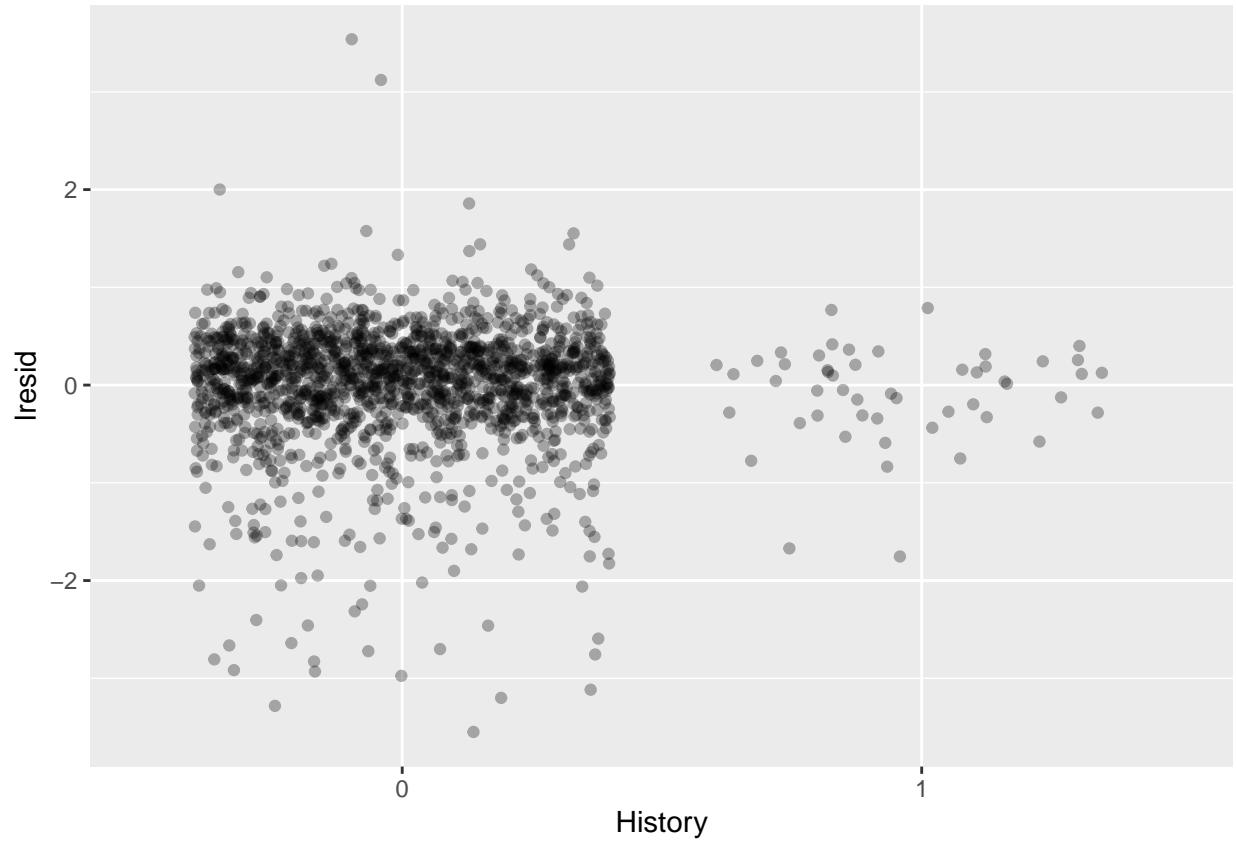
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



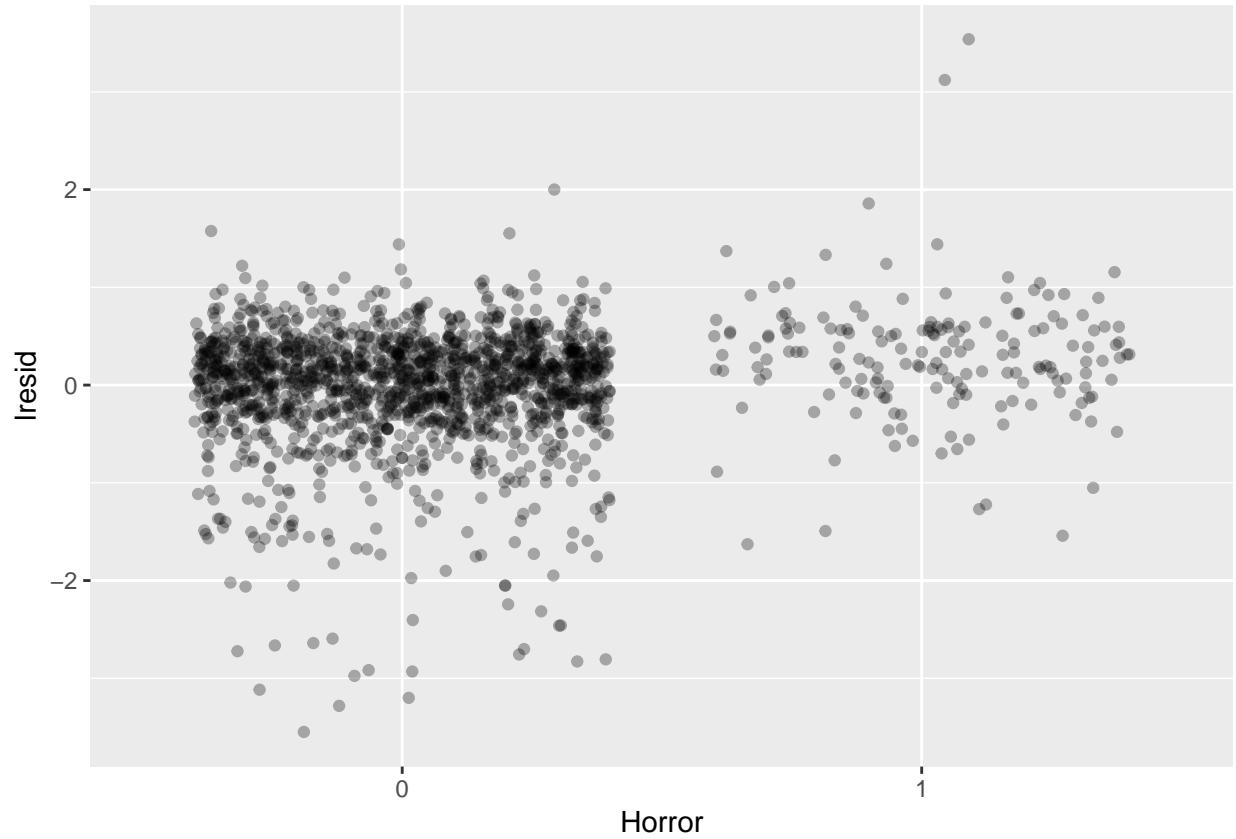
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



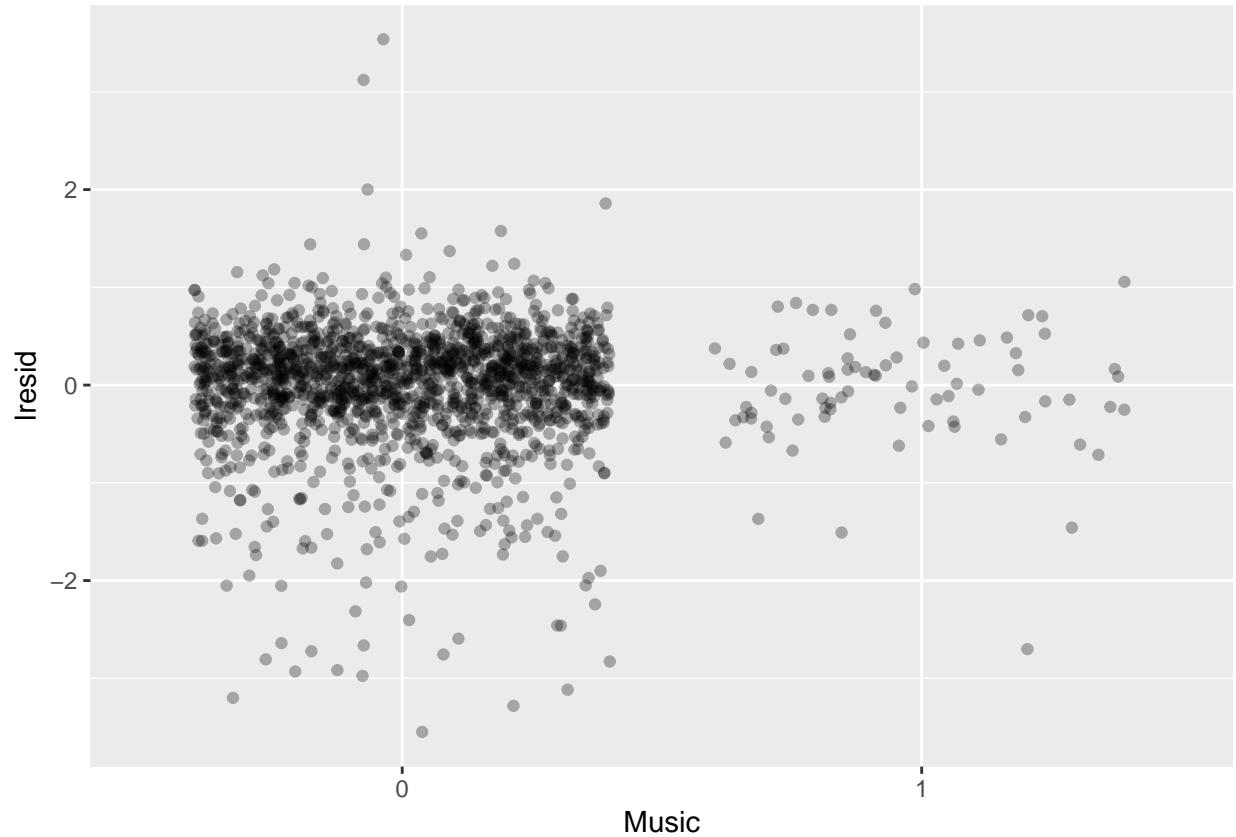
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



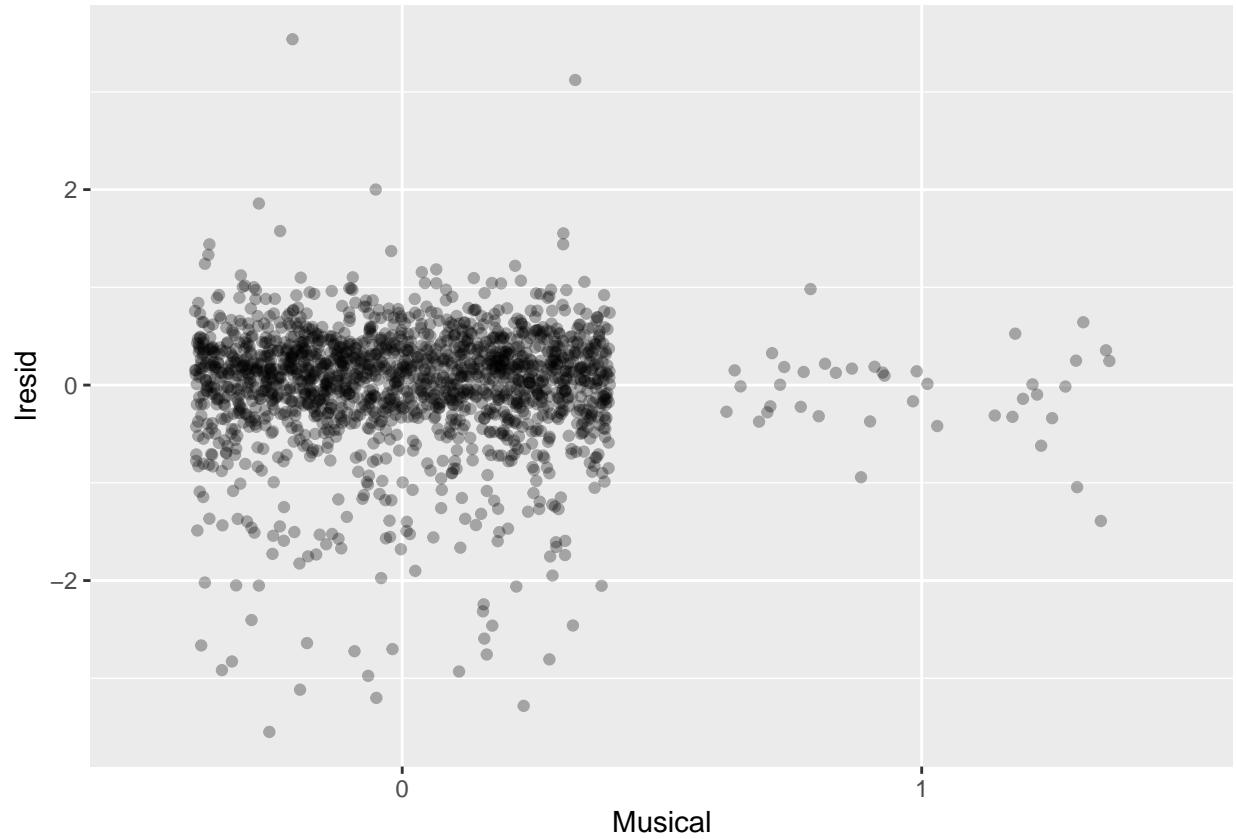
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



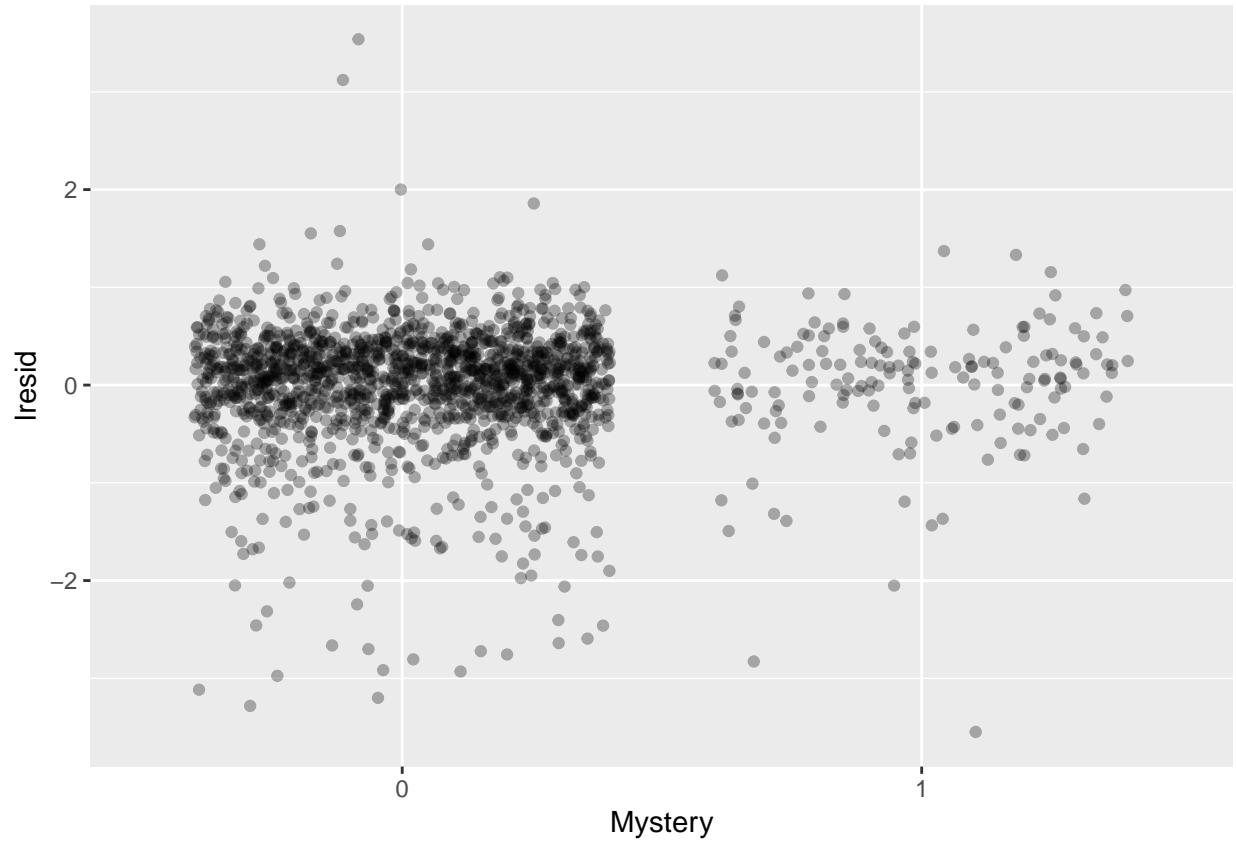
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



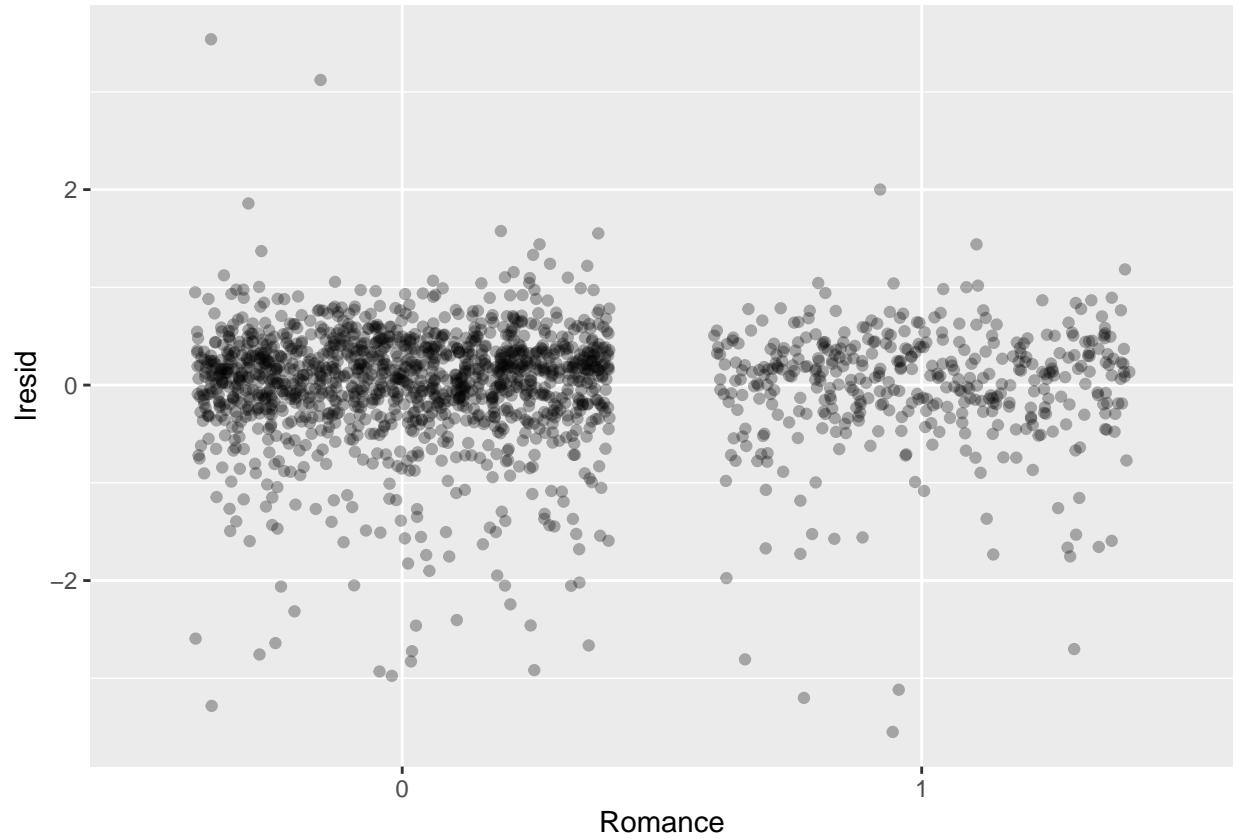
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



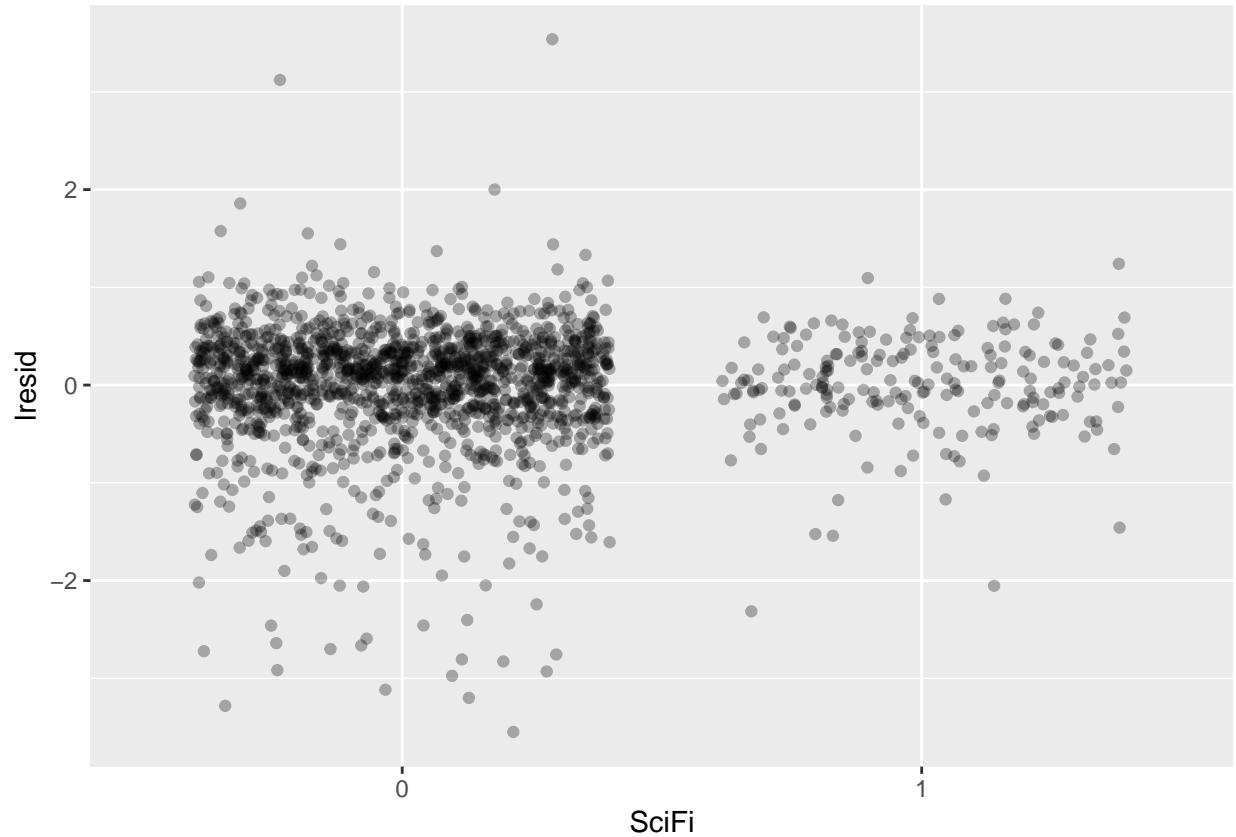
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



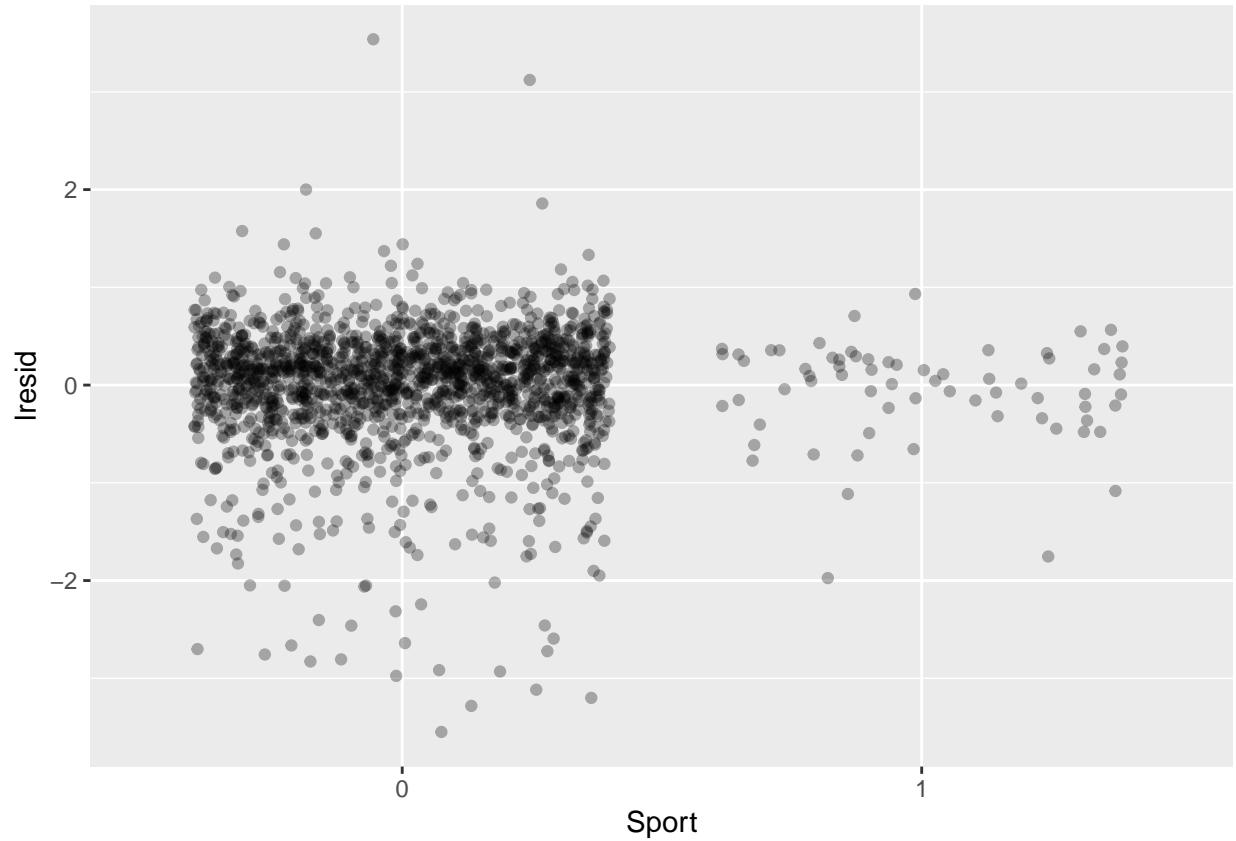
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



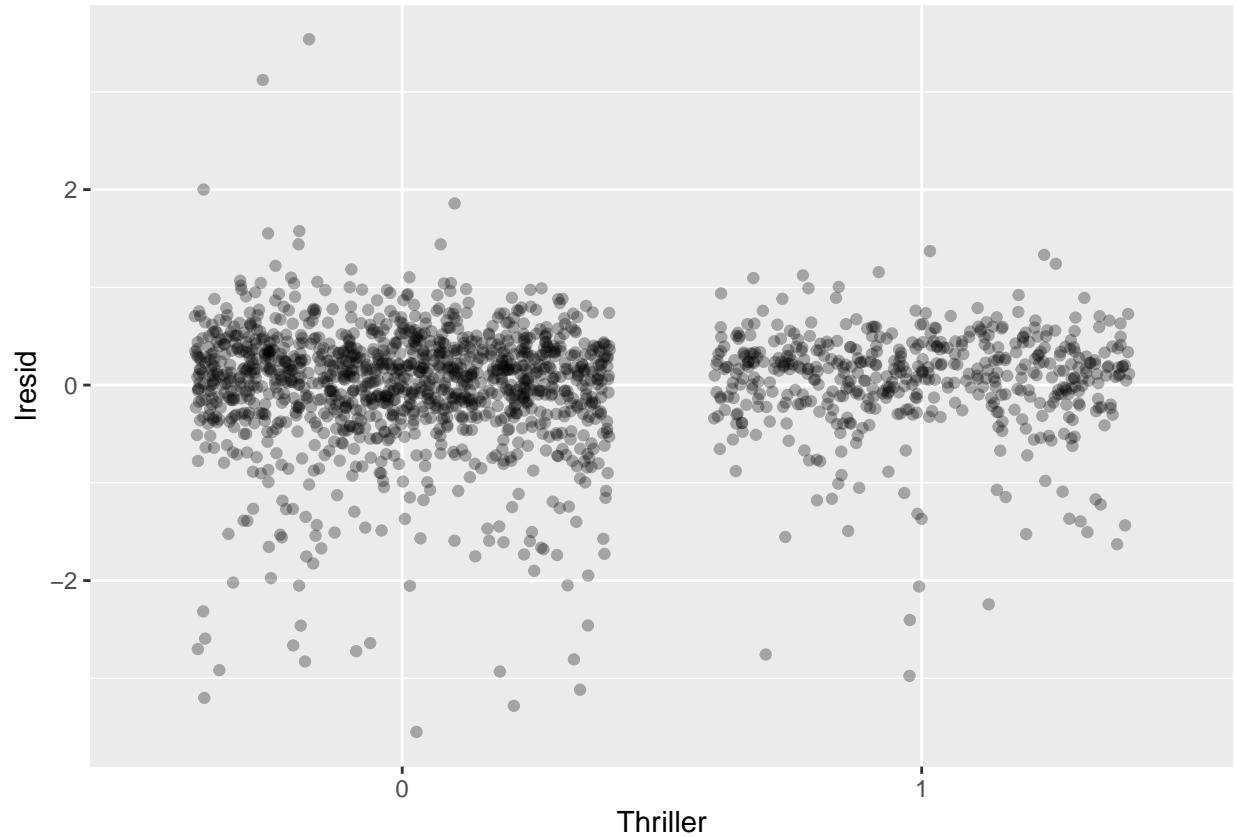
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



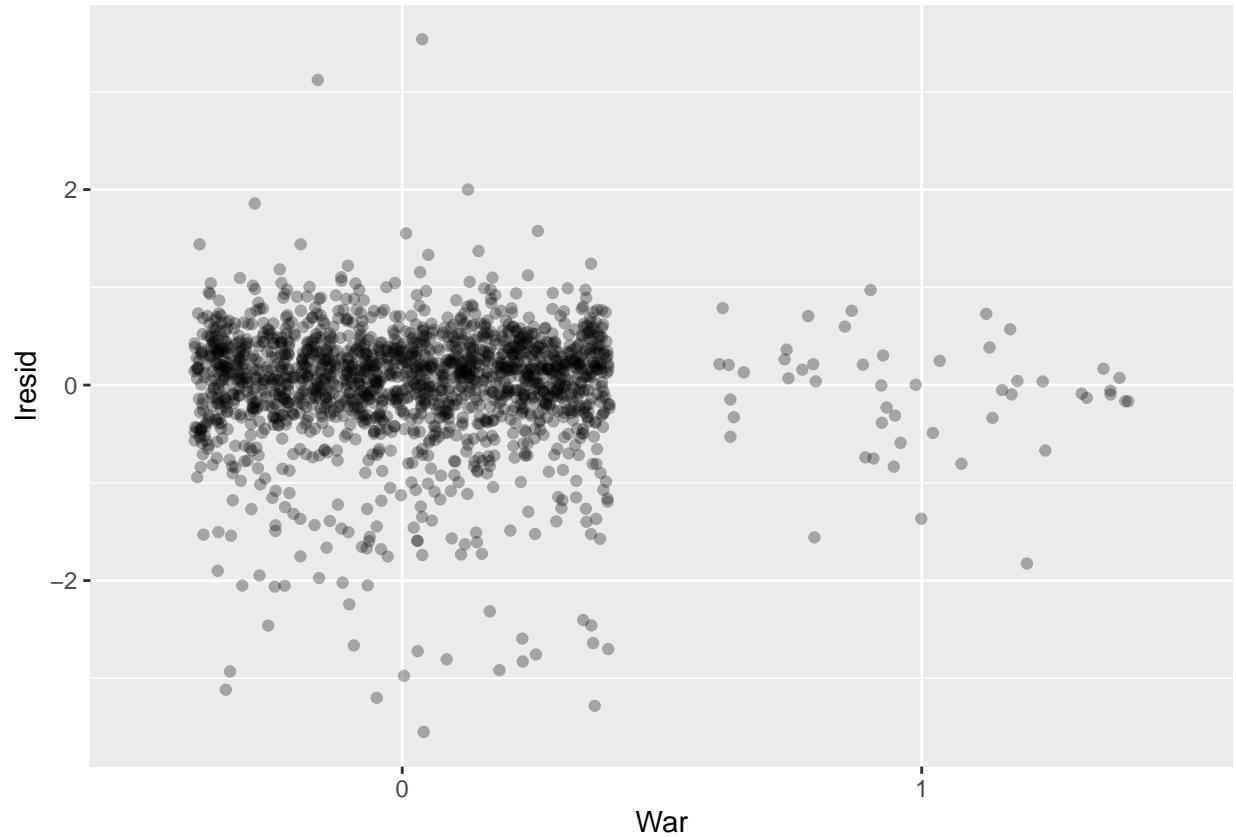
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



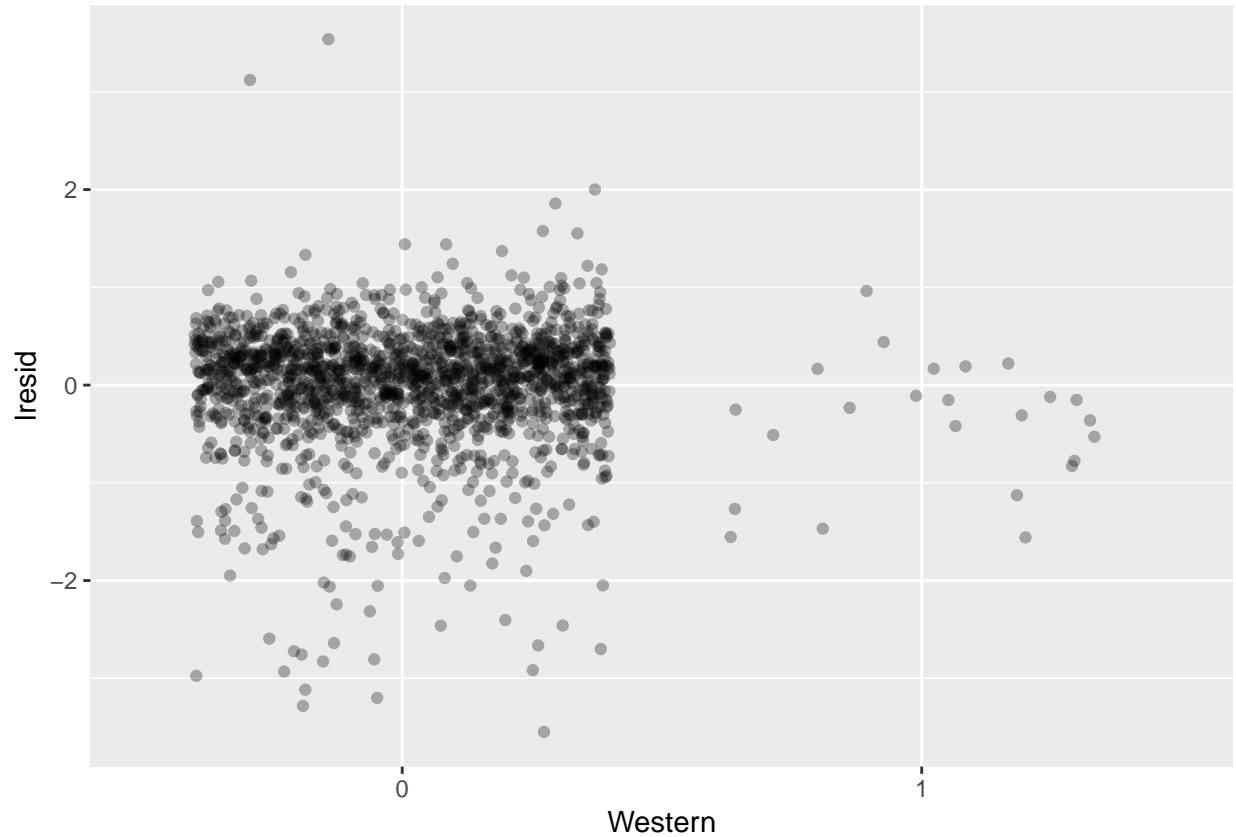
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



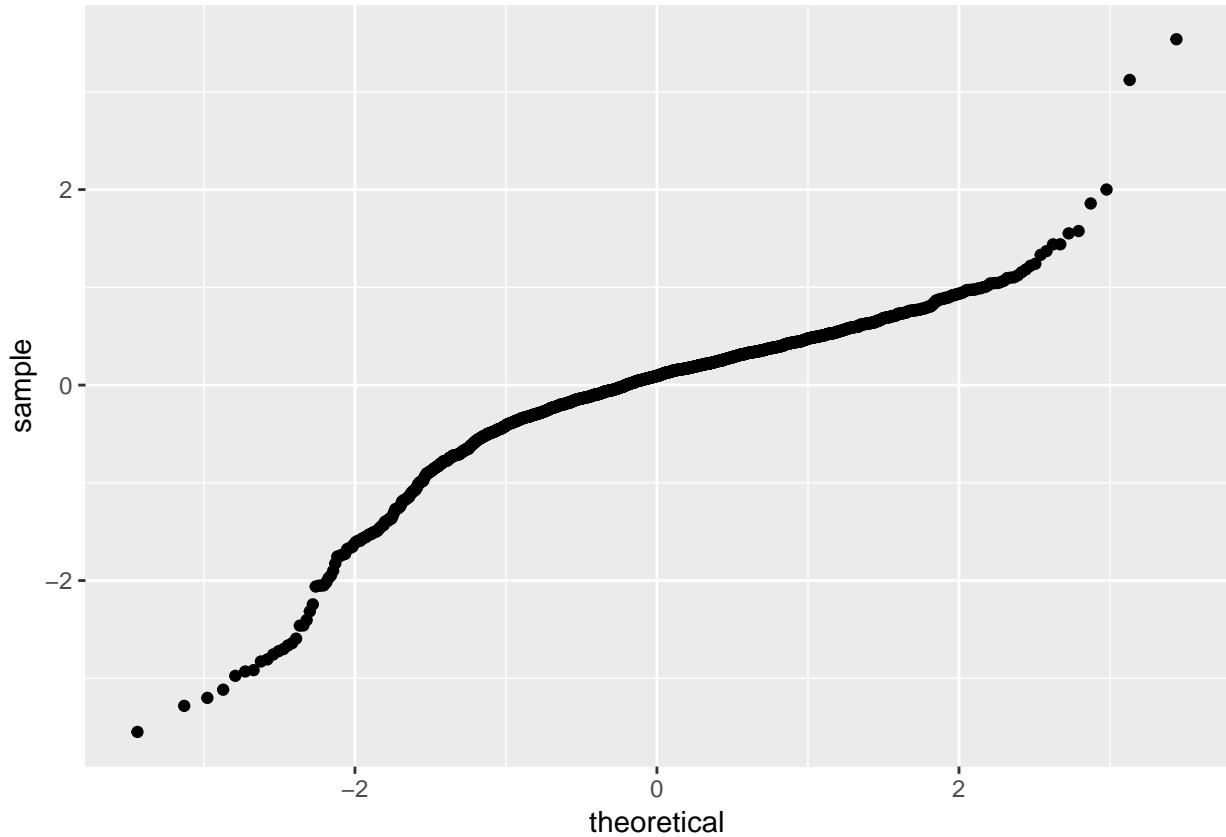
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



```
## Warning: Removed 132 rows containing missing values (geom_point).
```



```
## Warning: Removed 132 rows containing non-finite values (stat_qq).
```



```

# again blair and paranormal
train %>%
  add_residuals(mod_all2, 'lresid') %>%
  filter(lresid > 2.1)

## # A tibble: 2 x 58
##   director_name num_critic_for_~ director_facebook_likes actor_3_facebook_likes
##   <chr>           <dbl>                  <dbl>                  <dbl>
## 1 Daniel Myrick      360                   19                   39
## 2 Oren Peli          409                  110                   21
## # ... with 54 more variables: actor_2_name <chr>,
## #   actor_1_facebook_likes <dbl>, gross <dbl>, genres <chr>,
## #   actor_1_name <chr>, movie_title <chr>,
## #   cast_total_facebook_likes <dbl>, actor_3_name <chr>,
## #   num_user_for_reviews <dbl>, language <chr>, country <chr>,
## #   content_rating <fct>, budget <dbl>, year <fct>,
## #   actor_2_facebook_likes <dbl>, imdb_score <dbl>,
## #   movie_facebook_likes <dbl>, Action <fct>, Adventure <fct>,
## #   Animation <fct>, Biography <fct>, Comedy <fct>, Crime <fct>,
## #   Documentary <fct>, Drama <fct>, Family <fct>, Fantasy <fct>,
## #   History <fct>, Horror <fct>, Music <fct>, Musical <fct>,
## #   Mystery <fct>, Romance <fct>, SciFi <fct>, Sport <fct>,
## #   Thriller <fct>, War <fct>, Western <fct>, gdpd <dbl>,
## #   real_budget <dbl>, real_gross <dbl>, rev_budget <dbl>,
## #   profit_dum <dbl>, total_oscars_actor1 <dbl>,
## #   total_oscars_actor2 <dbl>, total_oscars_actor3 <dbl>,
## #   total_oscars_director <fct>, total_oscars_actor <fct>,

```

```

## #   real_gross_log <dbl>, real_budget_log <dbl>,
## #   director_facebook_likes_log <dbl>,
## #   cast_total_facebook_likes_log <dbl>, imdb_score_log <dbl>,
## #   lresid <dbl>
train %>%
  add_residuals(mod_all2, 'lresid') %>%
  filter(lresid < -2)

## # A tibble: 25 x 58
##   director_name num_critic_for_~ director_facebo~ actor_3_faceboo~
##   <chr>           <dbl>           <dbl>           <dbl>
## 1 Akiva Goldsm~        189          167          778
## 2 Nick Cassave~       177          415          545
## 3 John Hillcoat      355          214         3000
## 4 George Gallo       45           269          384
## 5 John Duigan        16           18           186
## 6 Daniel Algra~      42            3           442
## 7 Anand Tucker      134           14           356
## 8 Jake Paltrow       50           17            66
## 9 Timothy Hines        1            0            247
## 10 Michael Mere~     23            7           875
## # ... with 15 more rows, and 54 more variables: actor_2_name <chr>,
## #   actor_1_facebook_likes <dbl>, gross <dbl>, genres <chr>,
## #   actor_1_name <chr>, movie_title <chr>,
## #   cast_total_facebook_likes <dbl>, actor_3_name <chr>,
## #   num_user_for_reviews <dbl>, language <chr>, country <chr>,
## #   content_rating <fct>, budget <dbl>, year <fct>,
## #   actor_2_facebook_likes <dbl>, imdb_score <dbl>,
## #   movie_facebook_likes <dbl>, Action <fct>, Adventure <fct>,
## #   Animation <fct>, Biography <fct>, Comedy <fct>, Crime <fct>,
## #   Documentary <fct>, Drama <fct>, Family <fct>, Fantasy <fct>,
## #   History <fct>, Horror <fct>, Music <fct>, Musical <fct>,
## #   Mystery <fct>, Romance <fct>, SciFi <fct>, Sport <fct>,
## #   Thriller <fct>, War <fct>, Western <fct>, gdpd <dbl>,
## #   real_budget <dbl>, real_gross <dbl>, rev_budget <dbl>,
## #   profit_dum <dbl>, total_oscars_actor1 <dbl>,
## #   total_oscars_actor2 <dbl>, total_oscars_actor3 <dbl>,
## #   total_oscars_director <fct>, total_oscars_actor <fct>,
## #   real_gross_log <dbl>, real_budget_log <dbl>,
## #   director_facebook_likes_log <dbl>,
## #   cast_total_facebook_likes_log <dbl>, imdb_score_log <dbl>,
## #   lresid <dbl>

```

Prediction

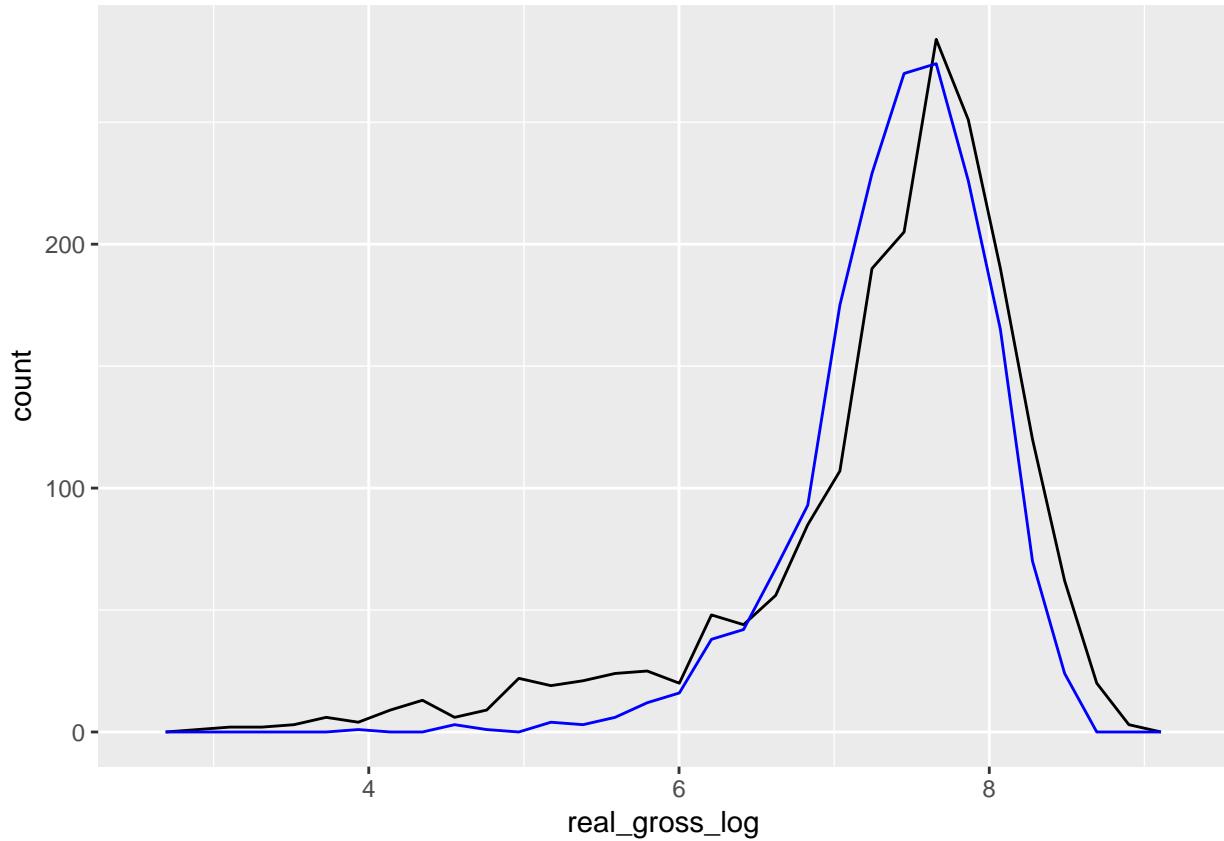
```

train %>%
  add_predictions(mod_all2, 'lpred') %>%
  ggplot() +
  geom_freqpoly(aes(x = real_gross_log)) +
  geom_freqpoly(aes(x = lpred), color = 'blue')

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

```

```
## Warning: Removed 132 rows containing non-finite values (stat_bin).
```



Fit really simple model: Anova compare

Budget always very significant and decreases RMSE significantly.

Use this to compare with other more complex models via arima

```
# filter so that same sample size as mod_all models so can compare
mod_simple <- lm(real_gross_log ~ real_budget_log, data = train %>% filter(!is.na(content_rating)))

# anova does show that the more complex models do explain real_gross_log better than simple
anova(mod_all, mod_simple)

## Analysis of Variance Table
##
## Model 1: real_gross_log ~ Adventure + Action + Family + Mystery + Documentary +
##           Drama + History + Romance + real_budget_log + imdb_score_log +
##           year + content_rating
## Model 2: real_gross_log ~ real_budget_log
##   Res.Df   RSS Df Sum of Sq    F    Pr(>F)
## 1   1668  614.85
## 2   1717  725.98 -49    -111.12 6.1524 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
anova(mod_all2, mod_simple)

## Analysis of Variance Table
```

```

## 
## Model 1: real_gross_log ~ real_budget_log + imdb_score_log + year + Comedy +
##           content_rating + Mystery
## Model 2: real_gross_log ~ real_budget_log
##   Res.Df   RSS Df Sum of Sq    F    Pr(>F)
## 1    1674 644.16
## 2    1717 725.98 -43   -81.819 4.9448 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Start with budget

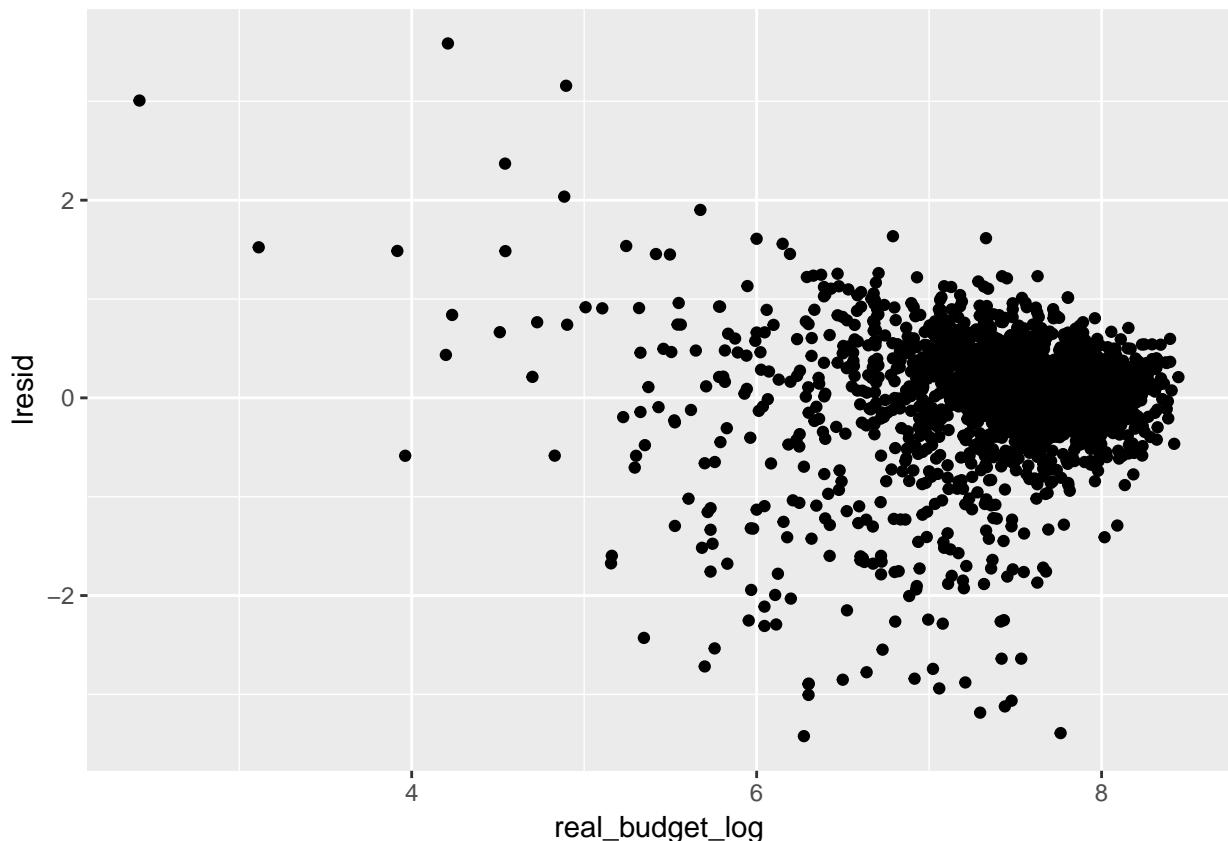
Look at residual plots and determine what other variables should be added

Budget has the most obvious relationship with revenue, both graphically and logically

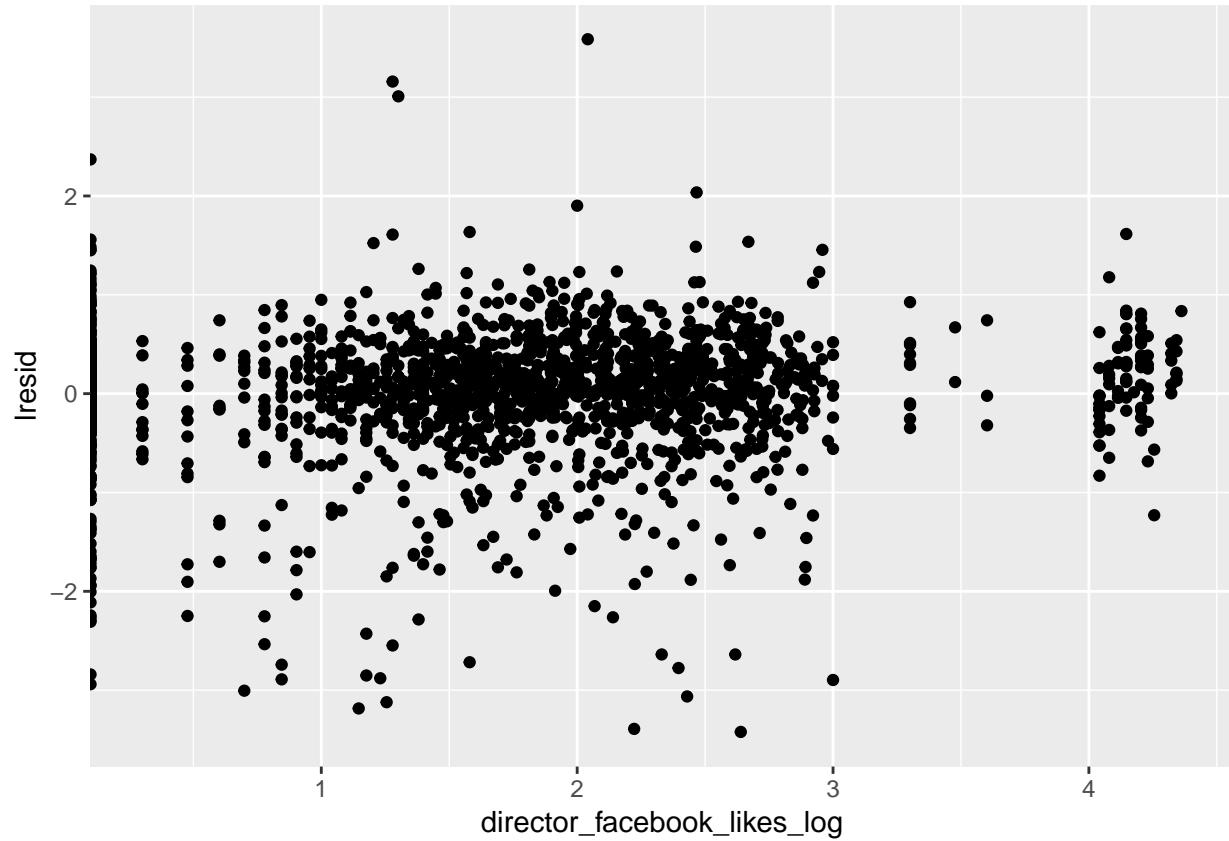
```
mod_simple <- lm(real_gross_log ~ real_budget_log, data = train)
```

```
gr_resid(mod_simple)
```

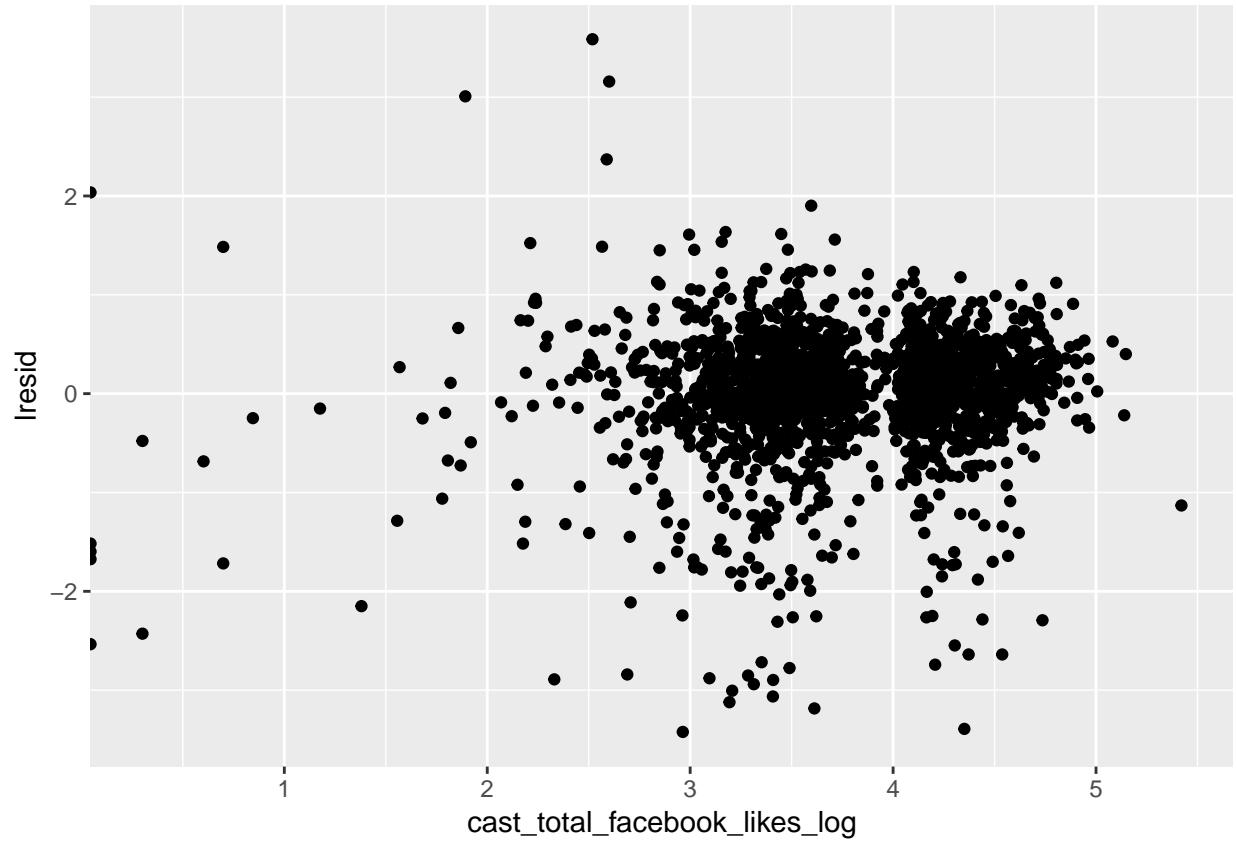
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



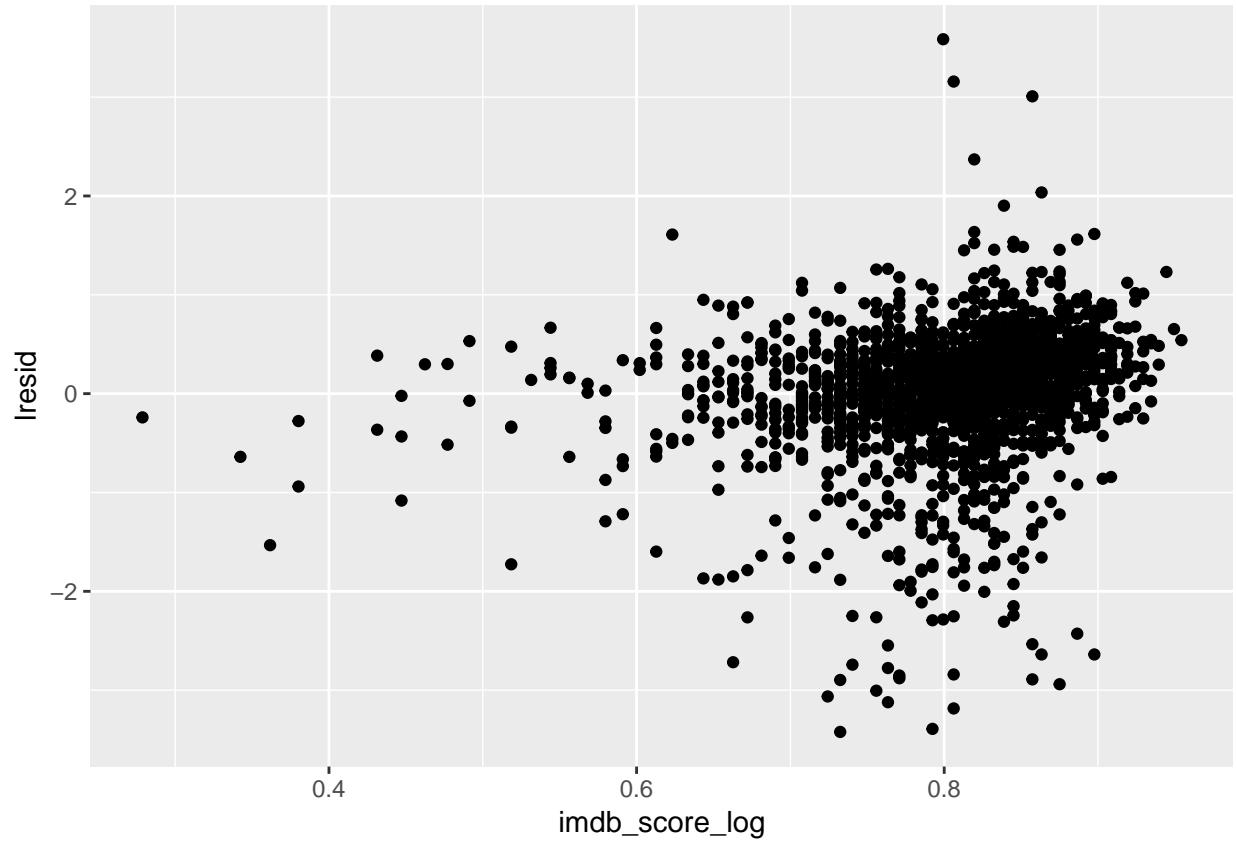
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



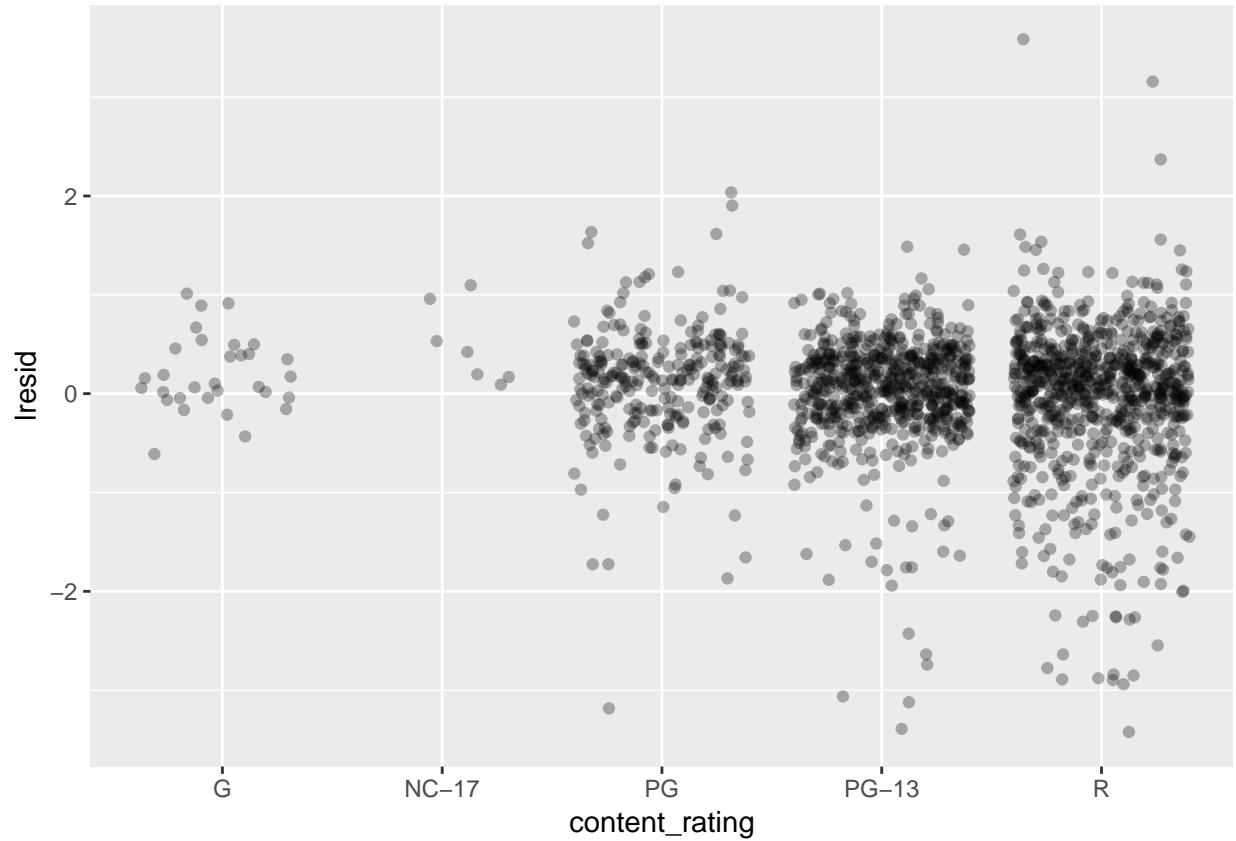
```
## Warning: Removed 102 rows containing missing values (geom_point).
```

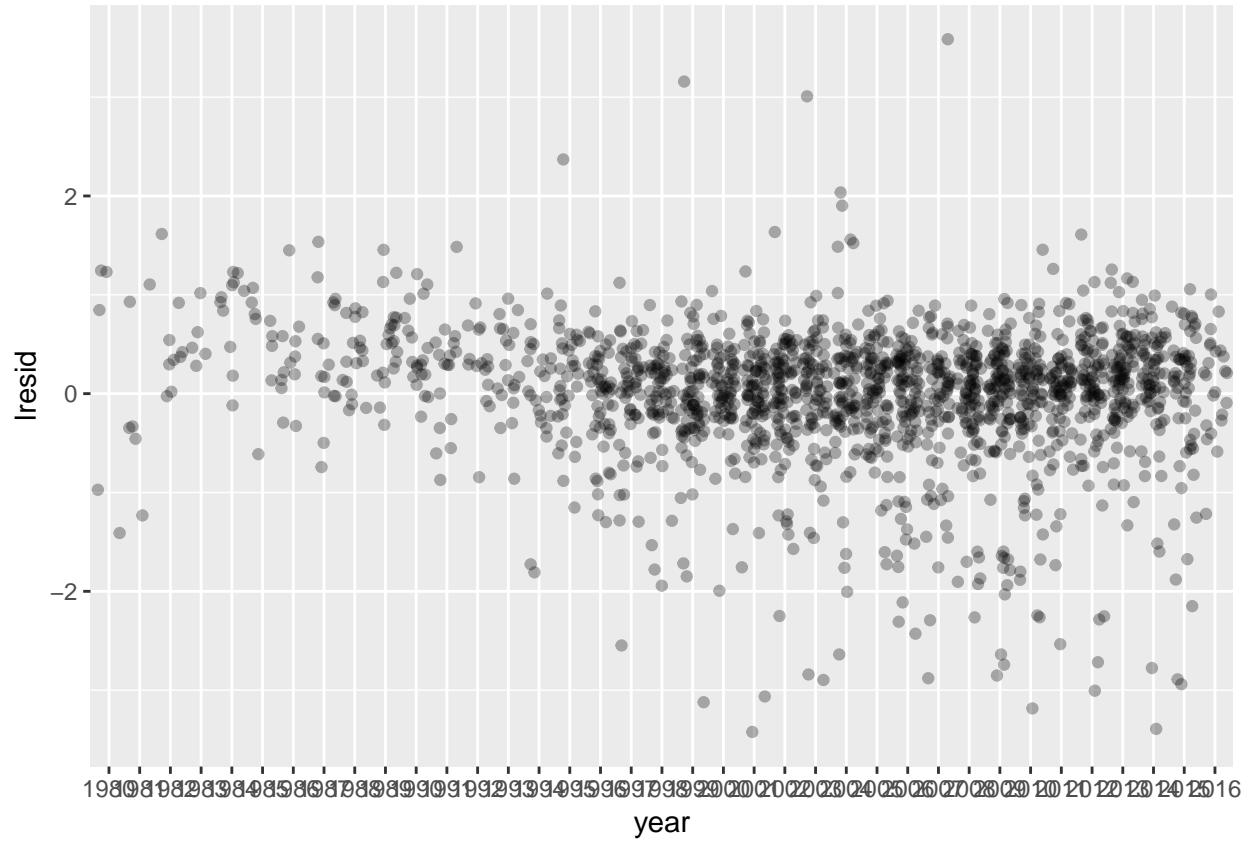


```
## Warning: Removed 102 rows containing missing values (geom_point).
```

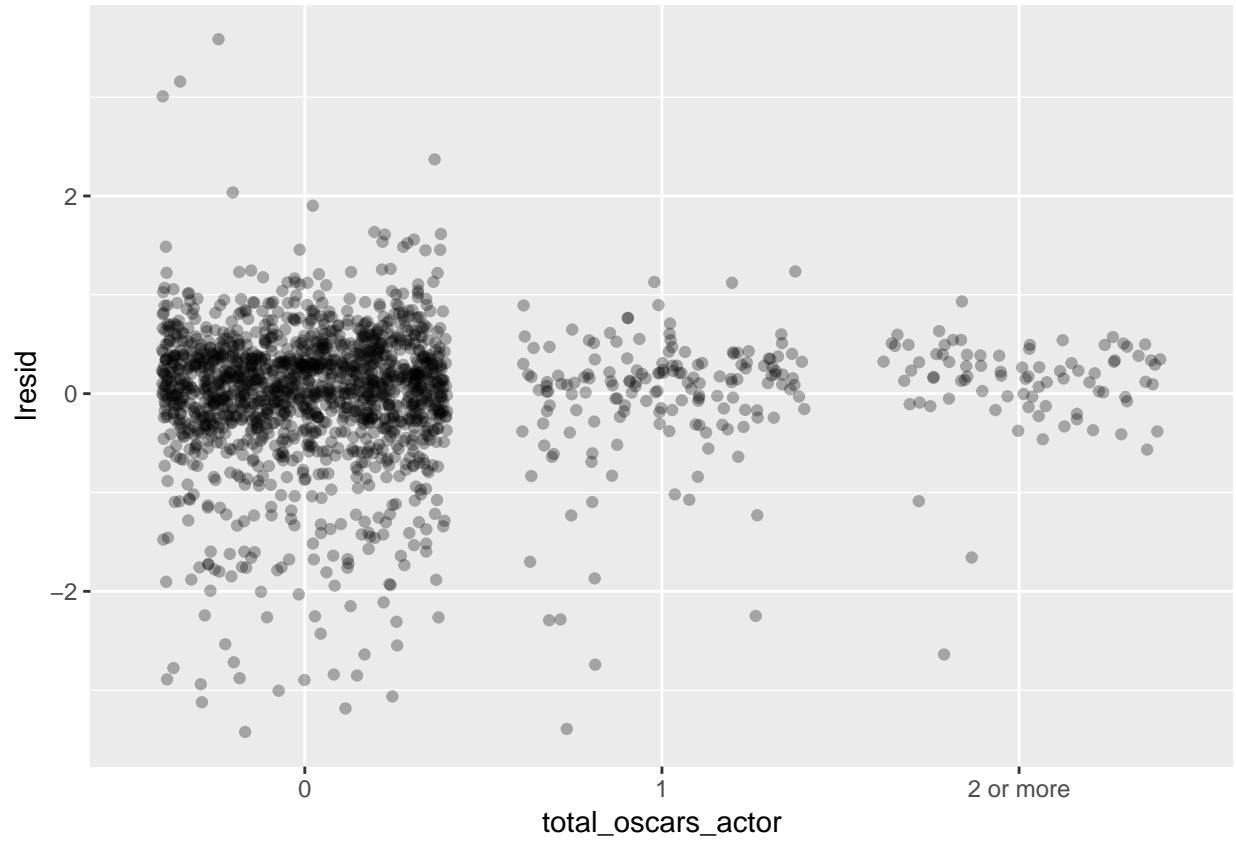


```
## Warning: Removed 94 rows containing missing values (geom_point).
```

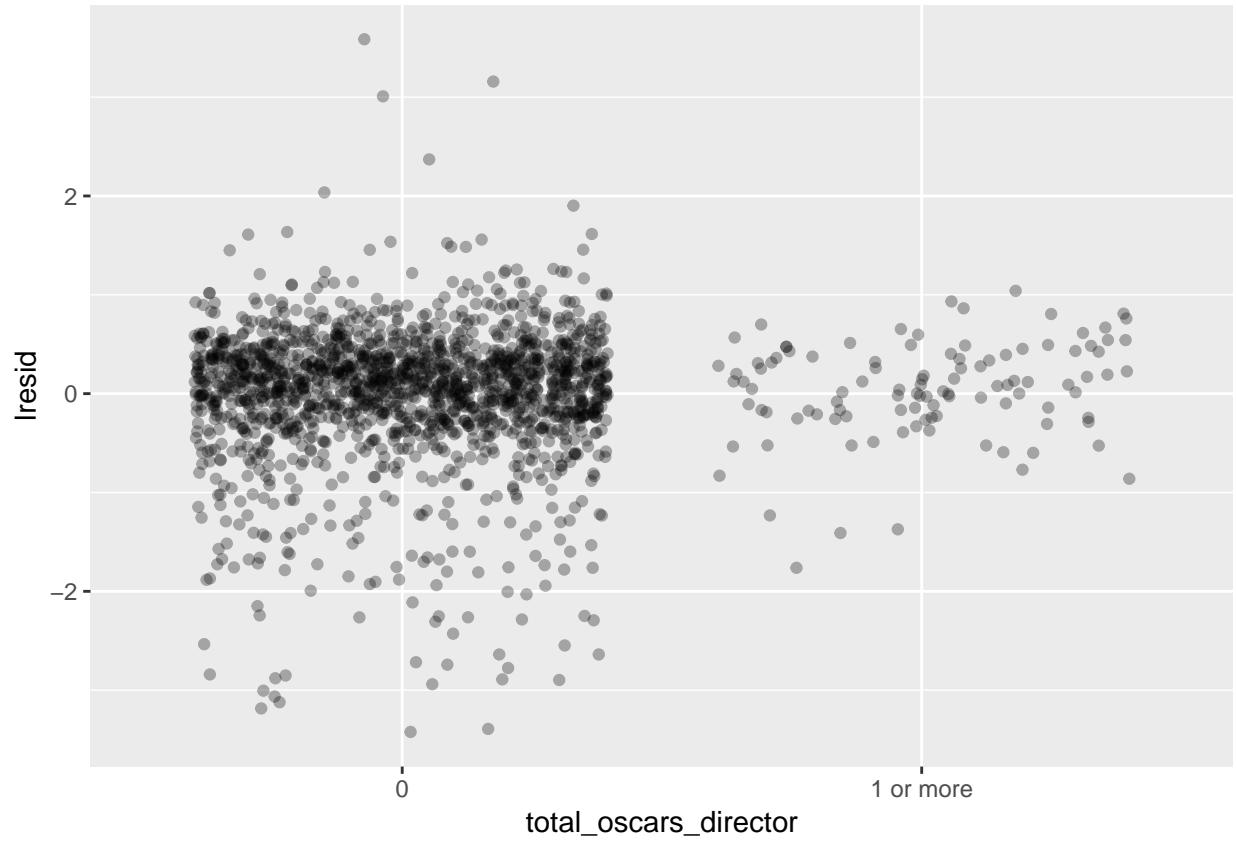




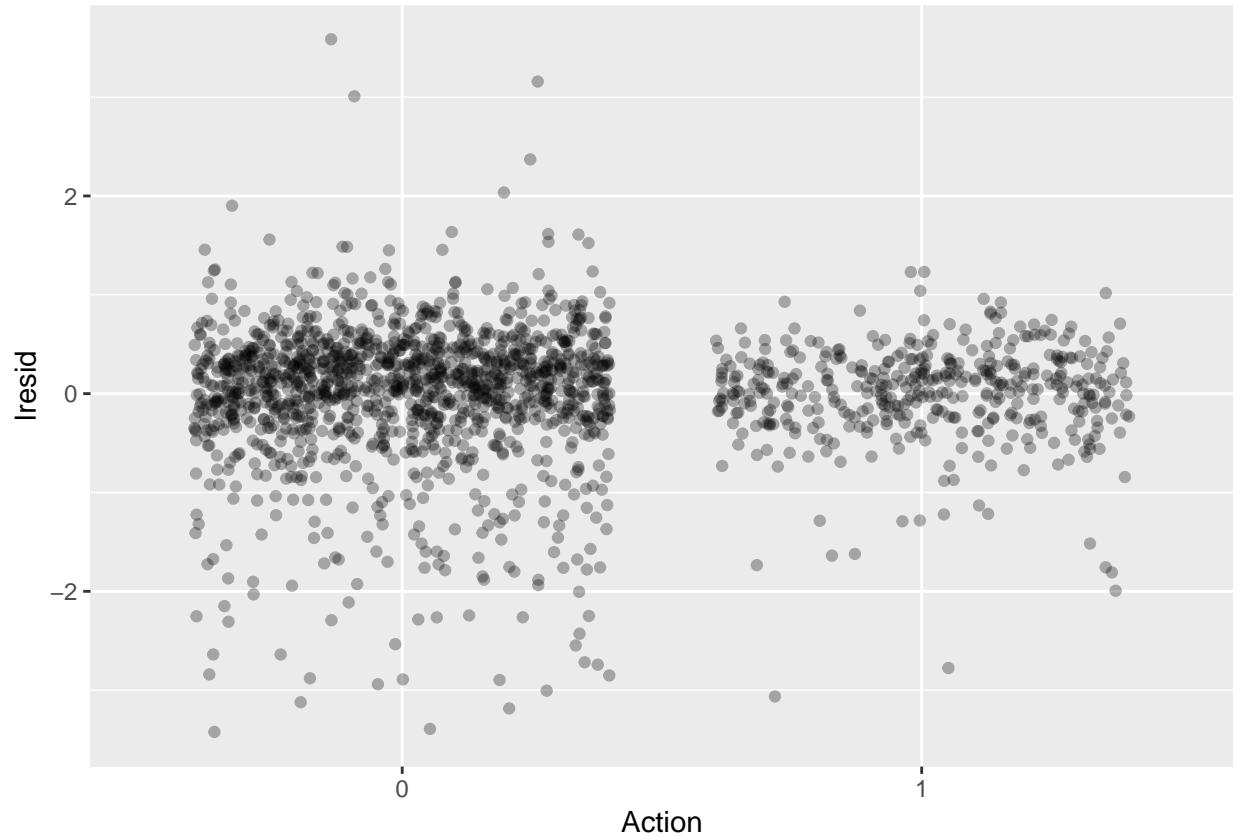
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



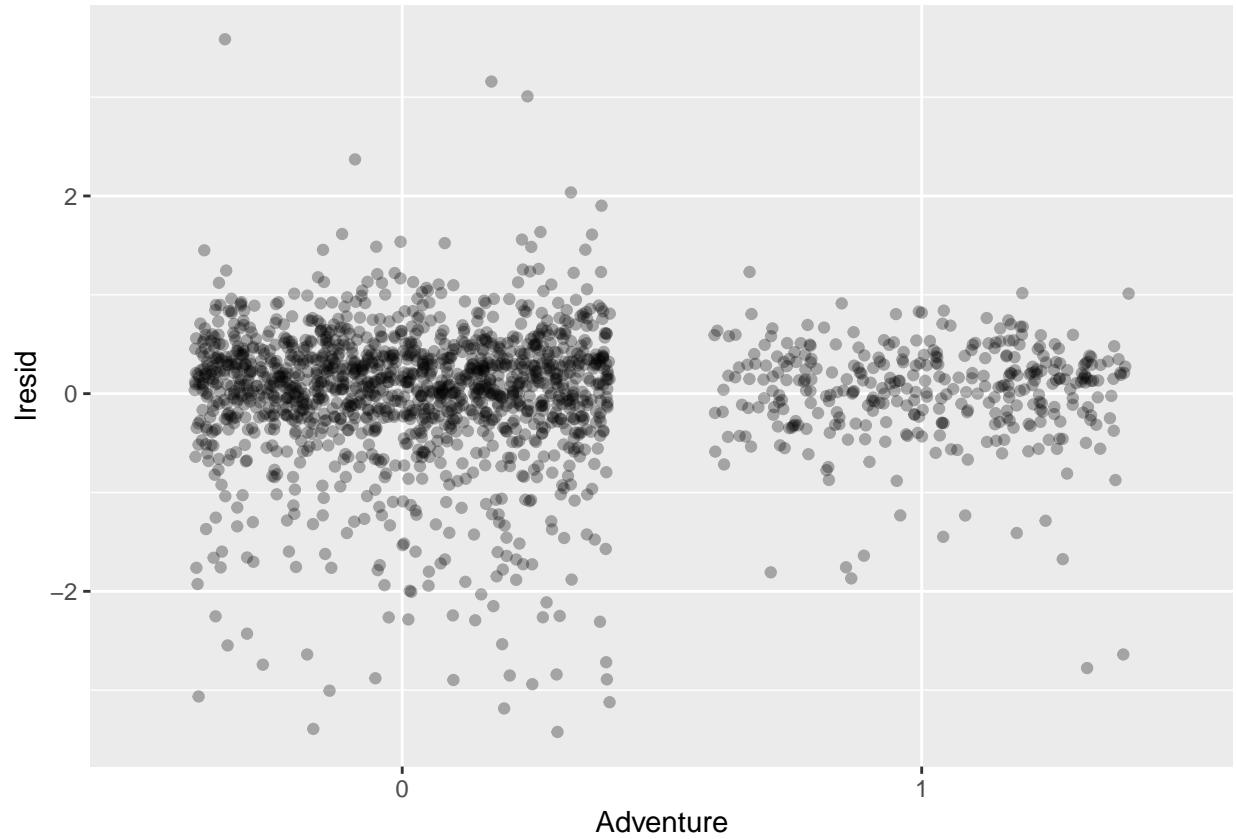
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



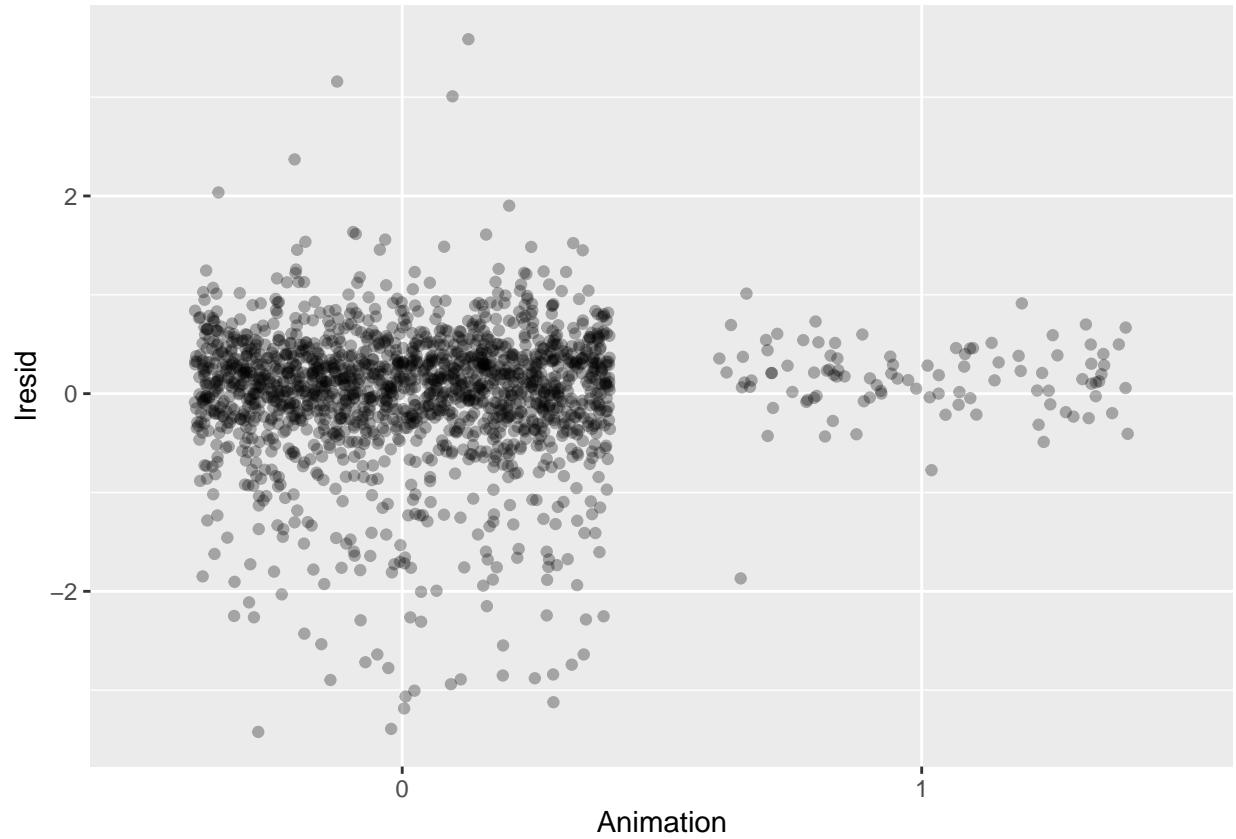
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



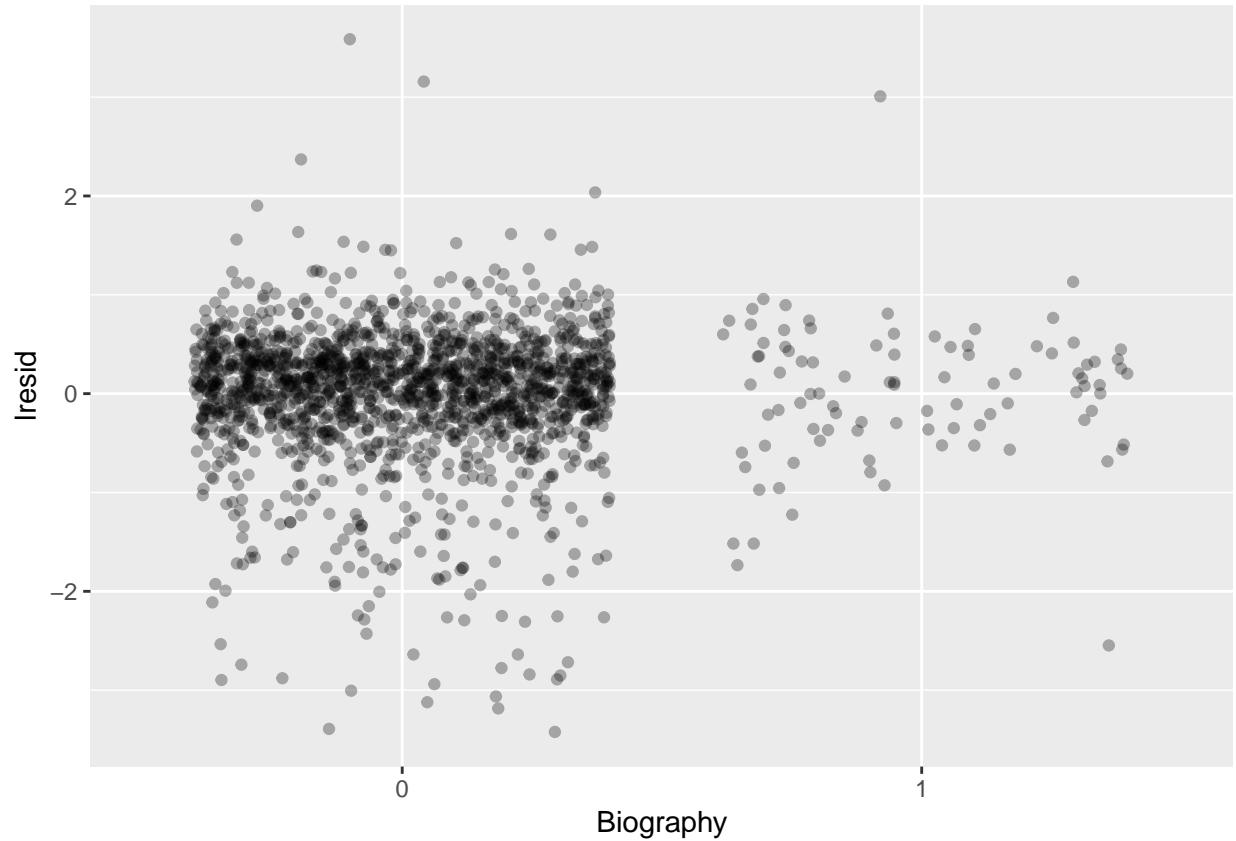
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



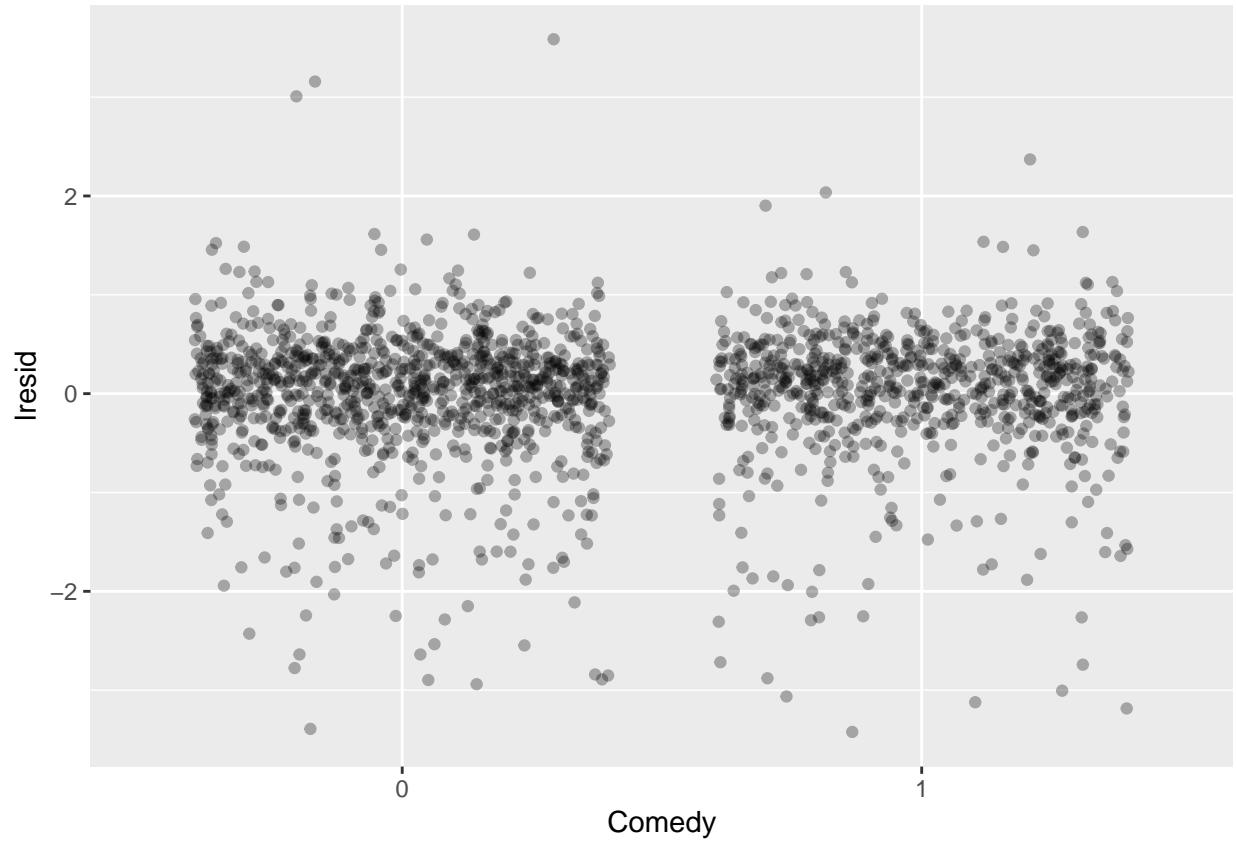
```
## Warning: Removed 102 rows containing missing values (geom_point).
```

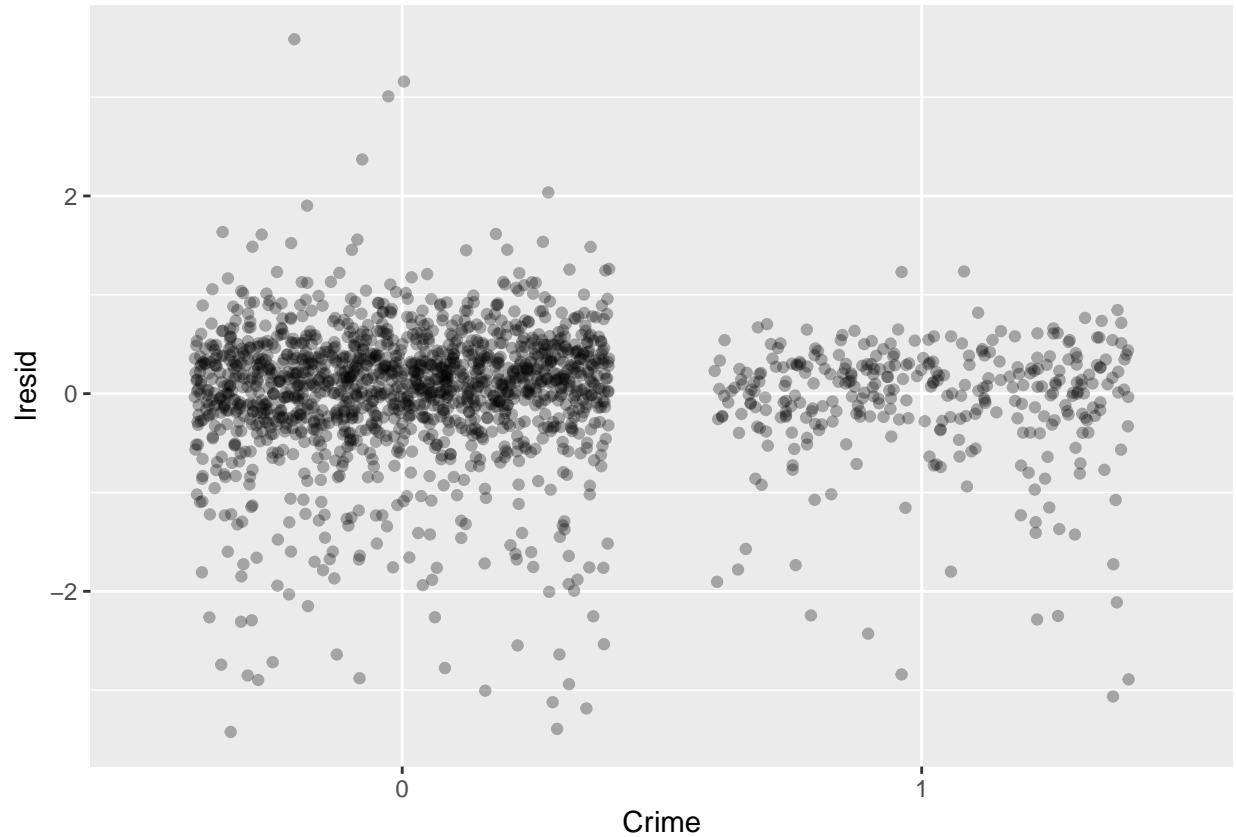


```
## Warning: Removed 102 rows containing missing values (geom_point).
```

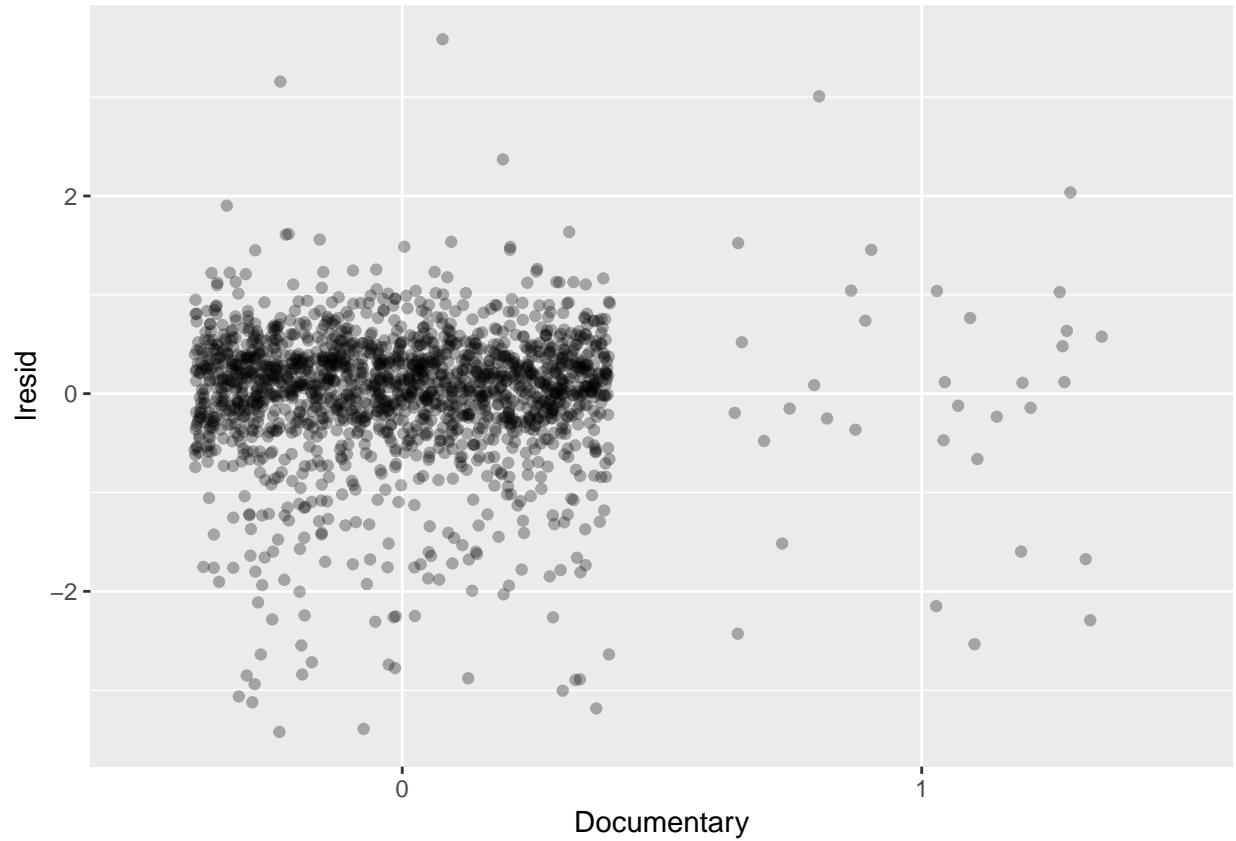


```
## Warning: Removed 102 rows containing missing values (geom_point).
```

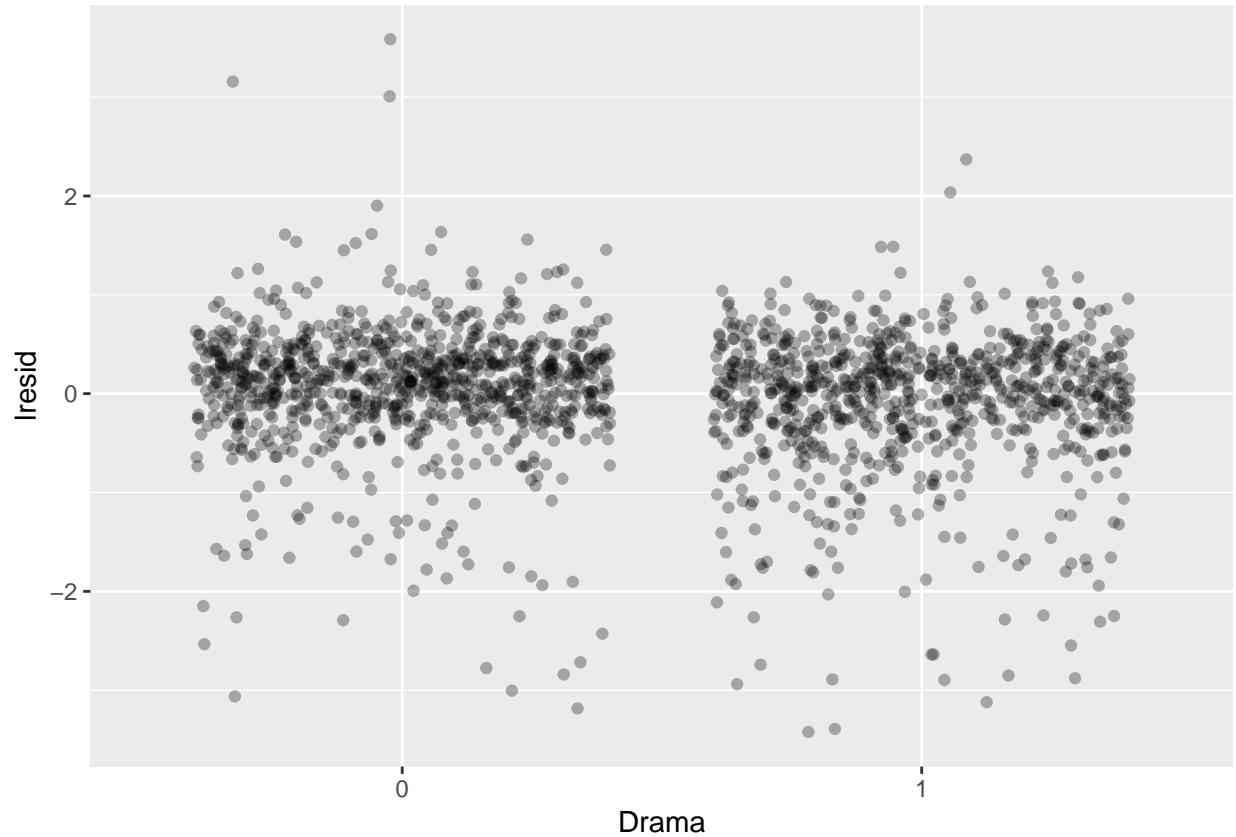




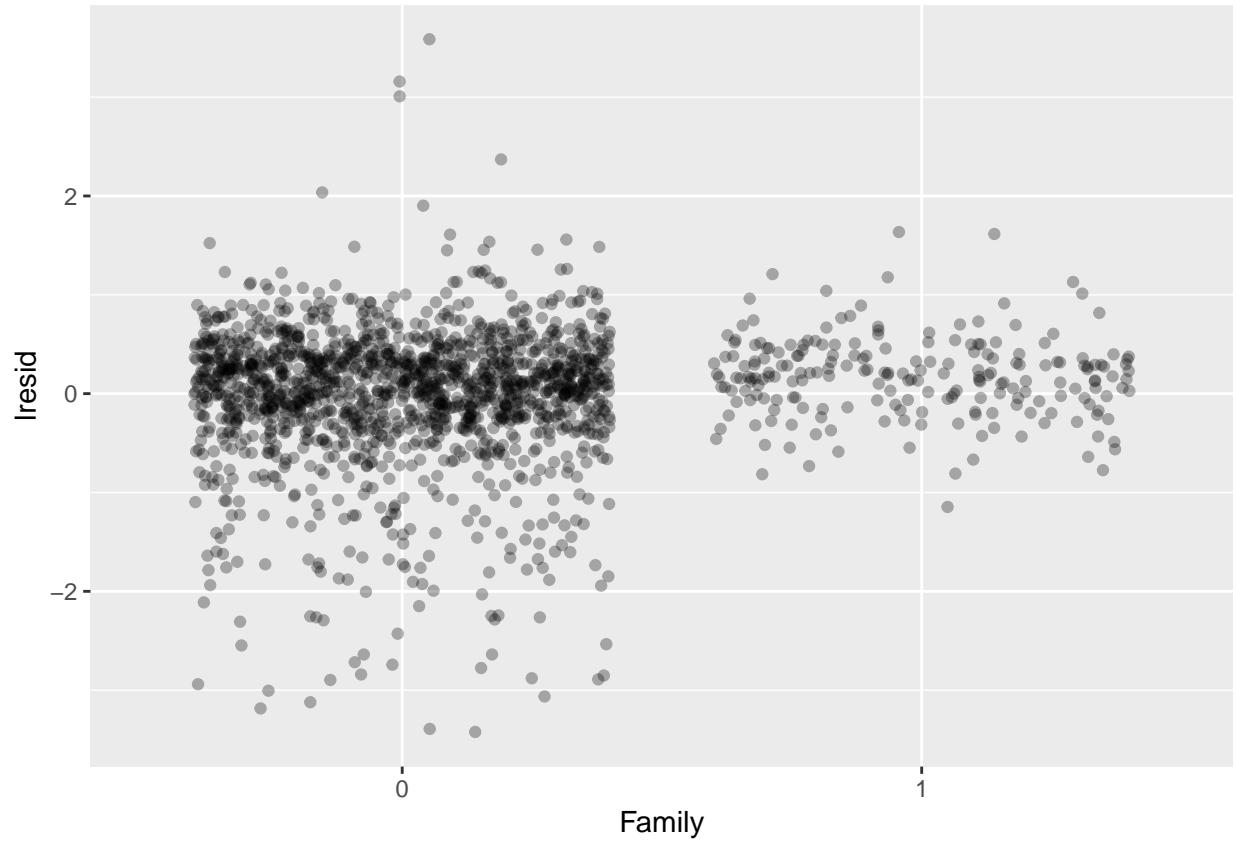
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



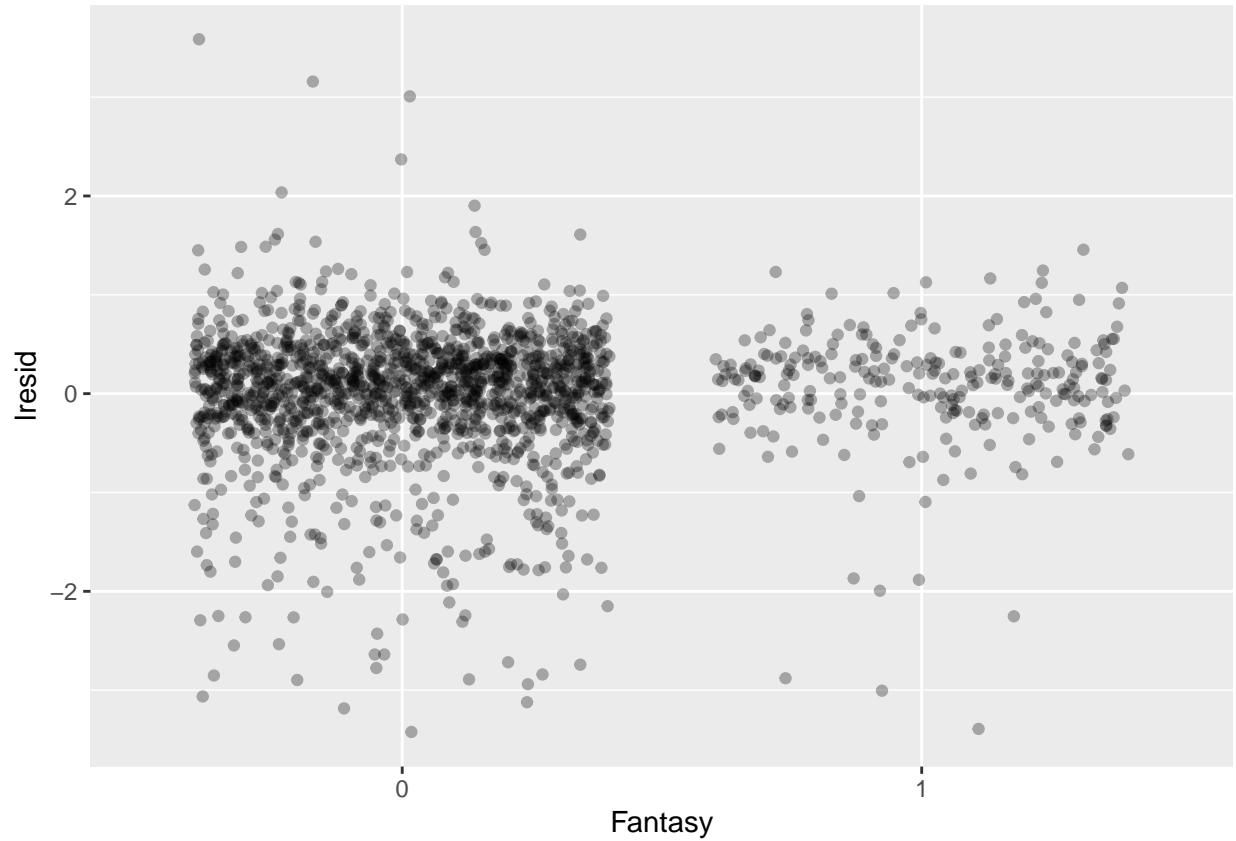
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



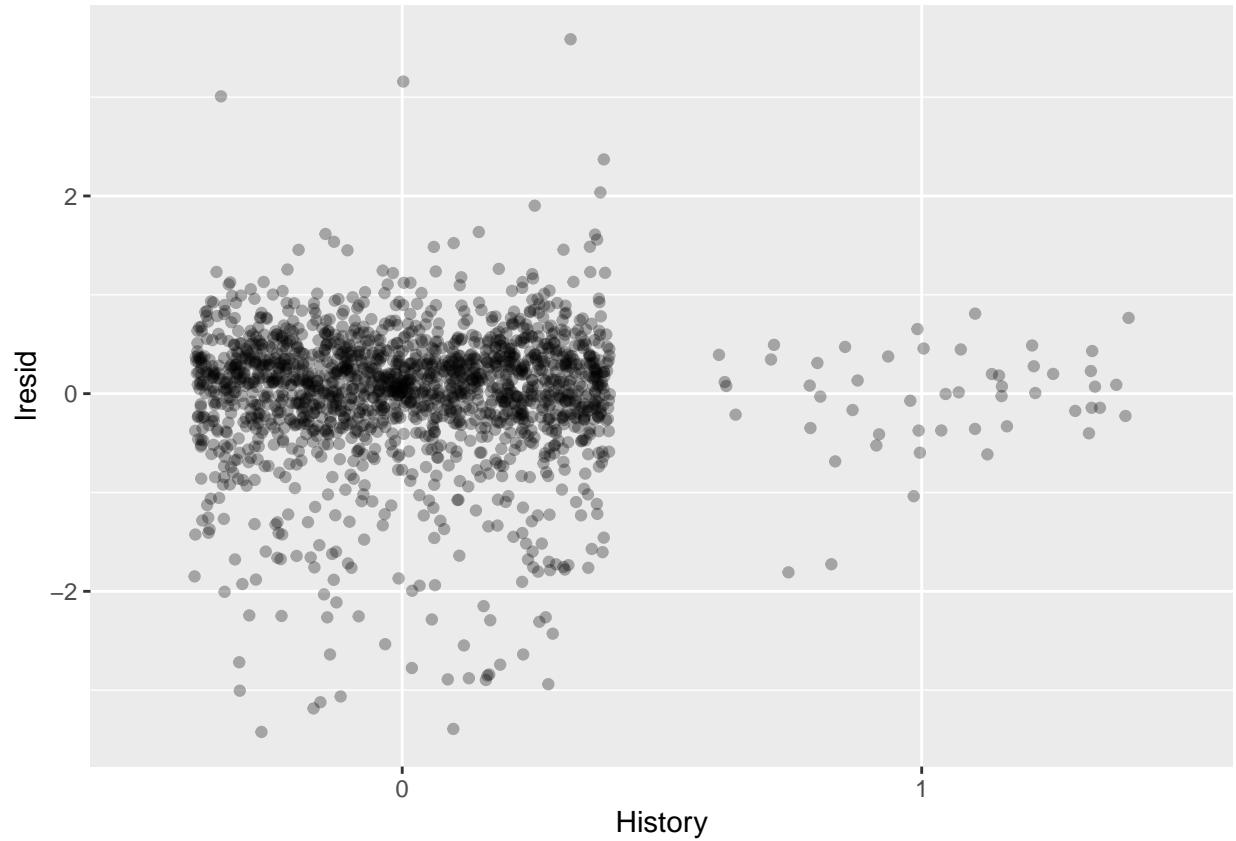
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



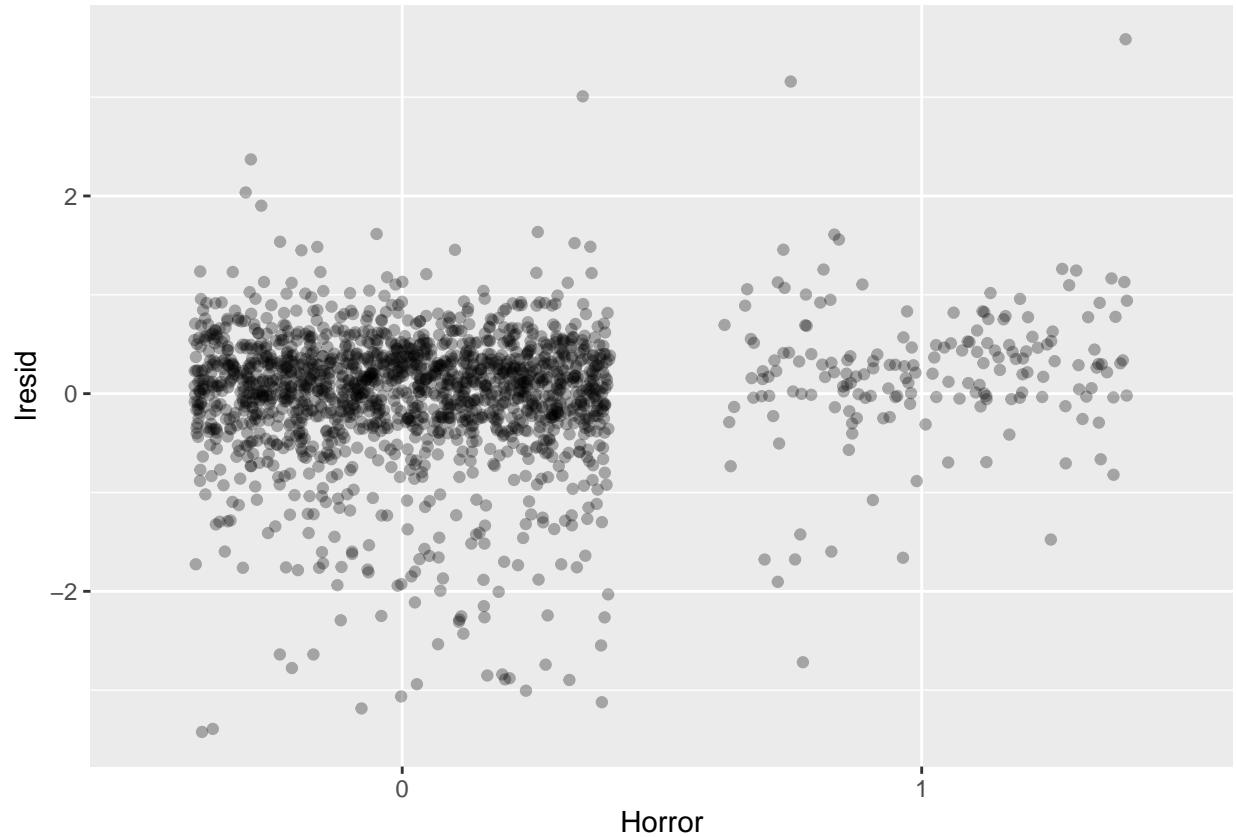
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



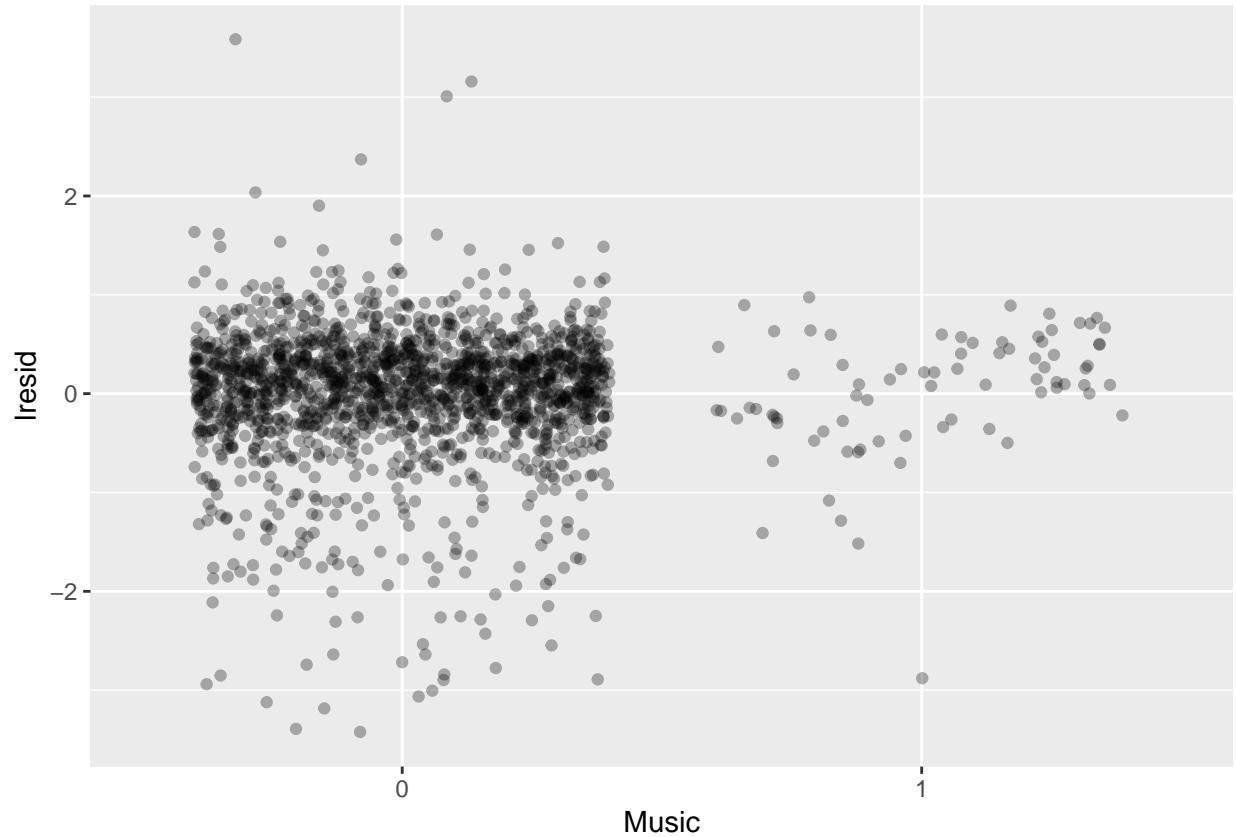
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



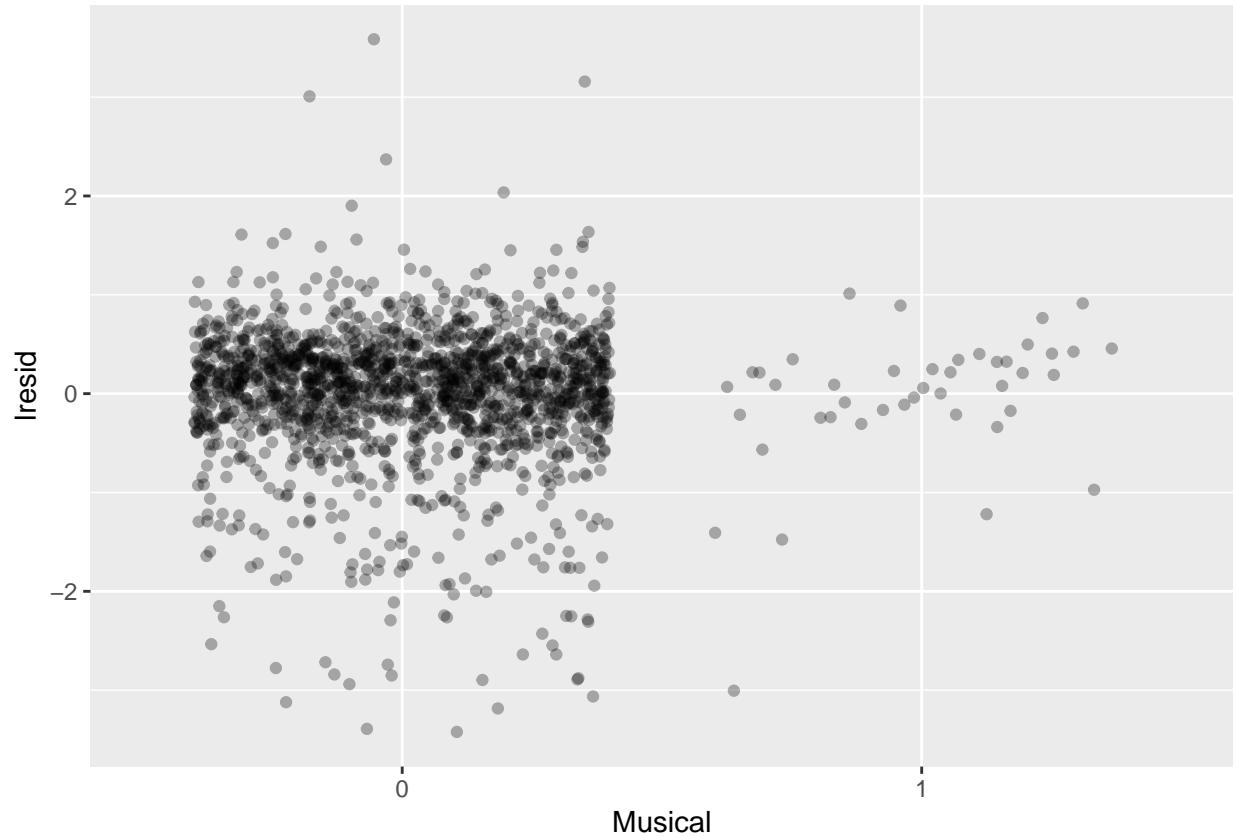
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



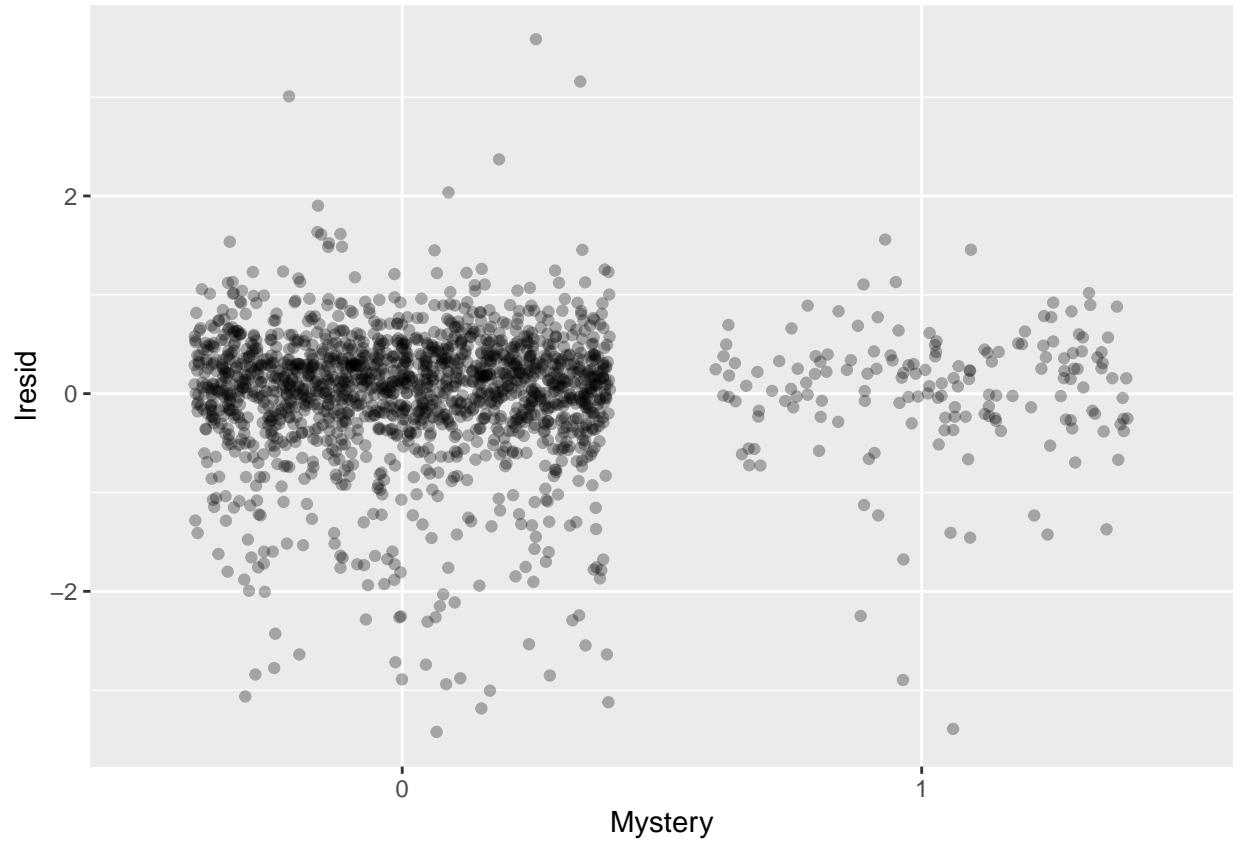
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



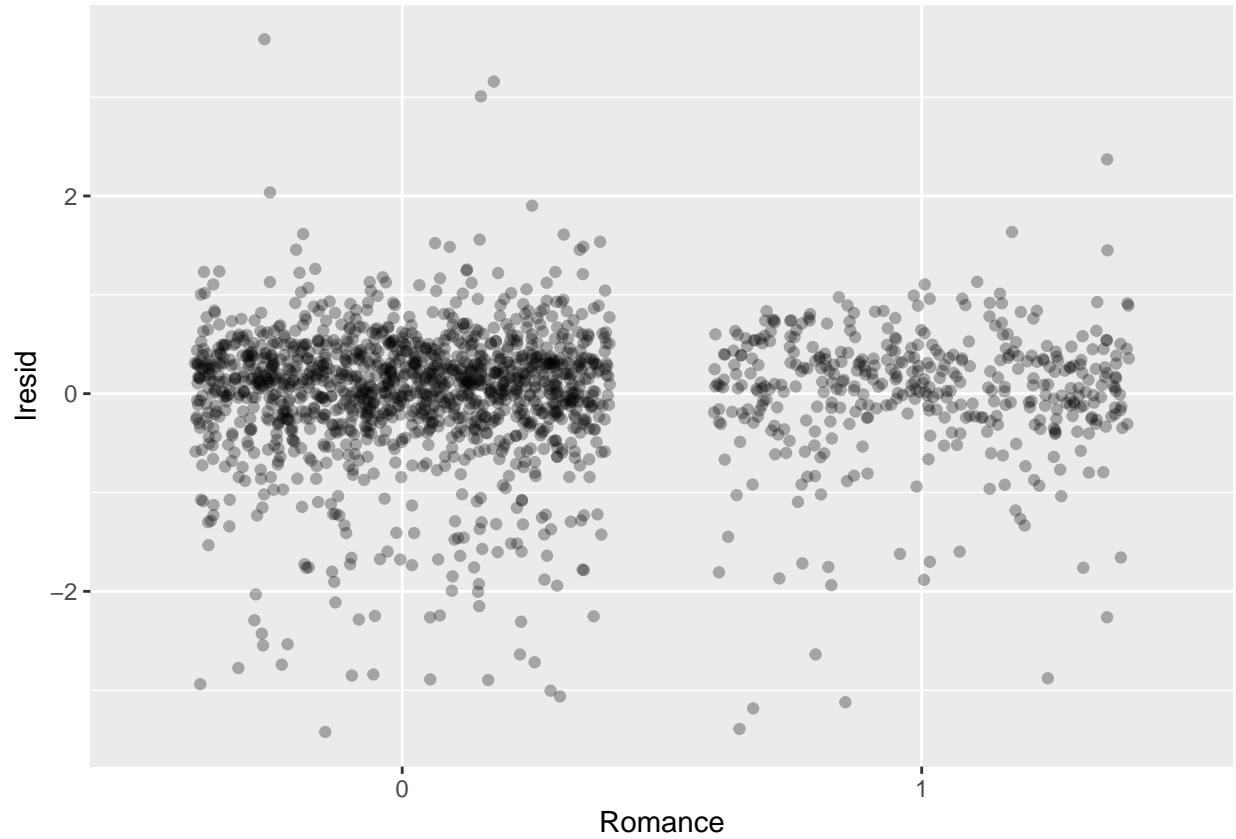
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



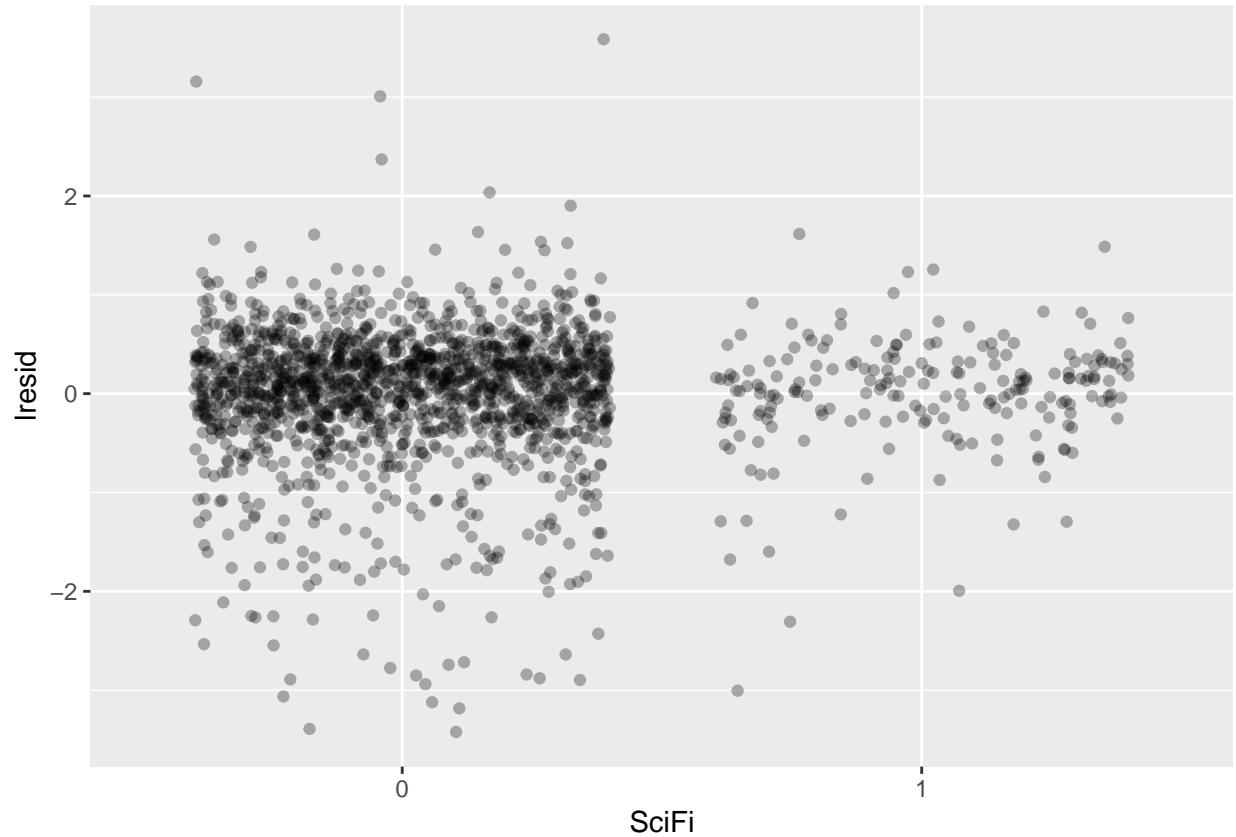
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



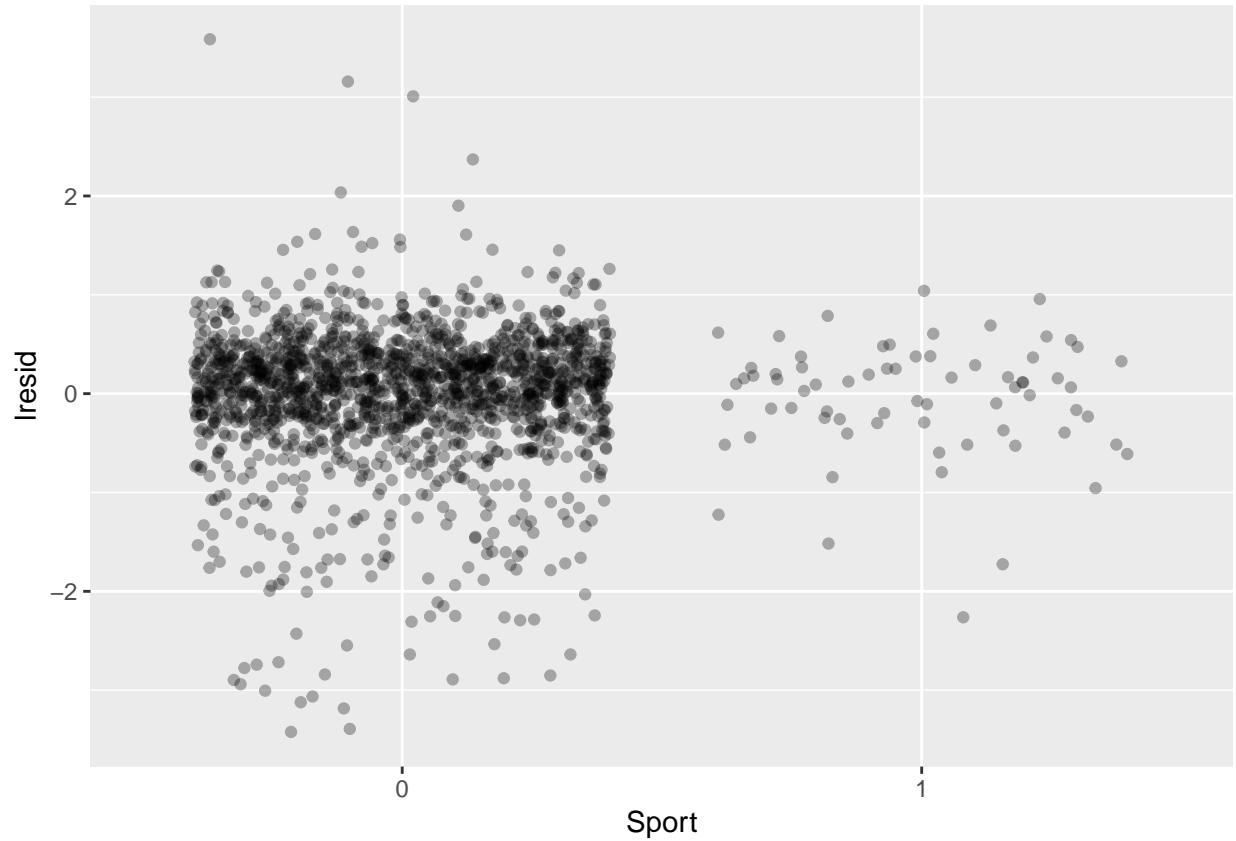
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



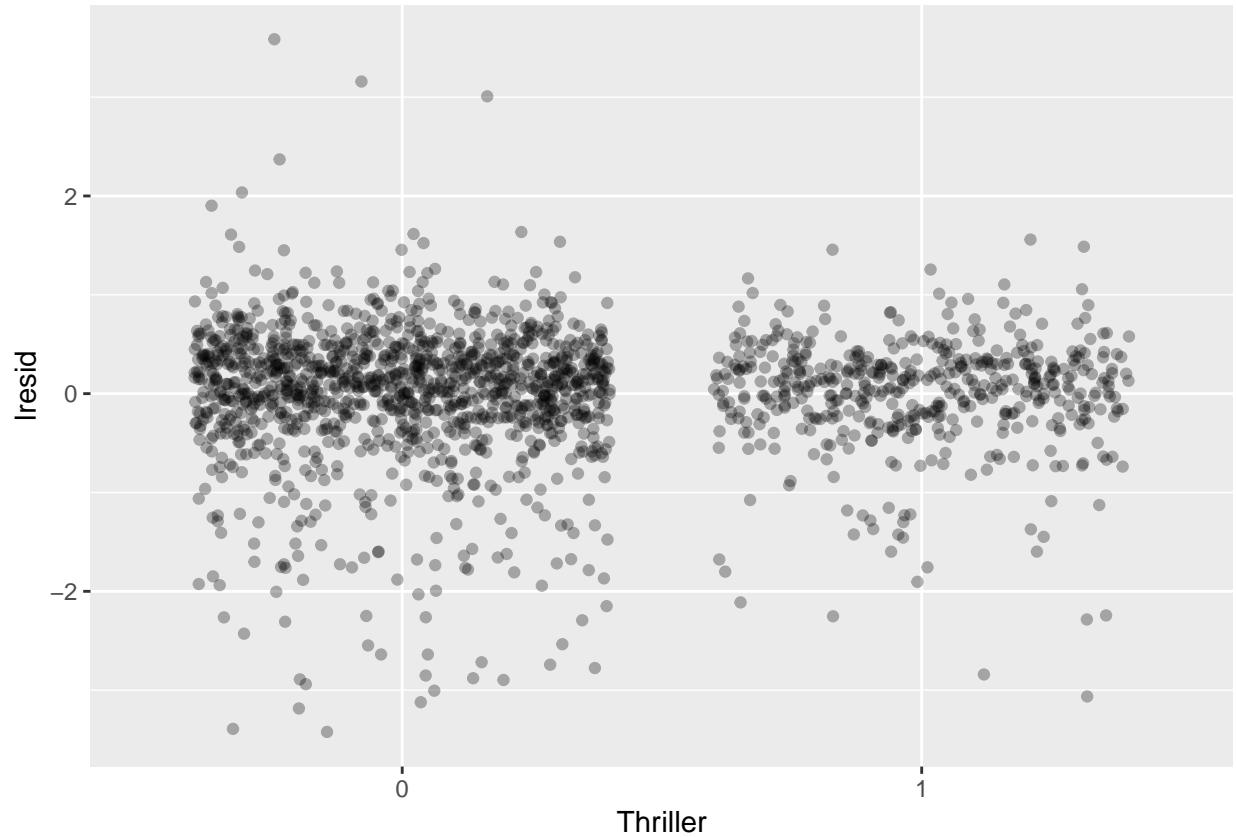
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



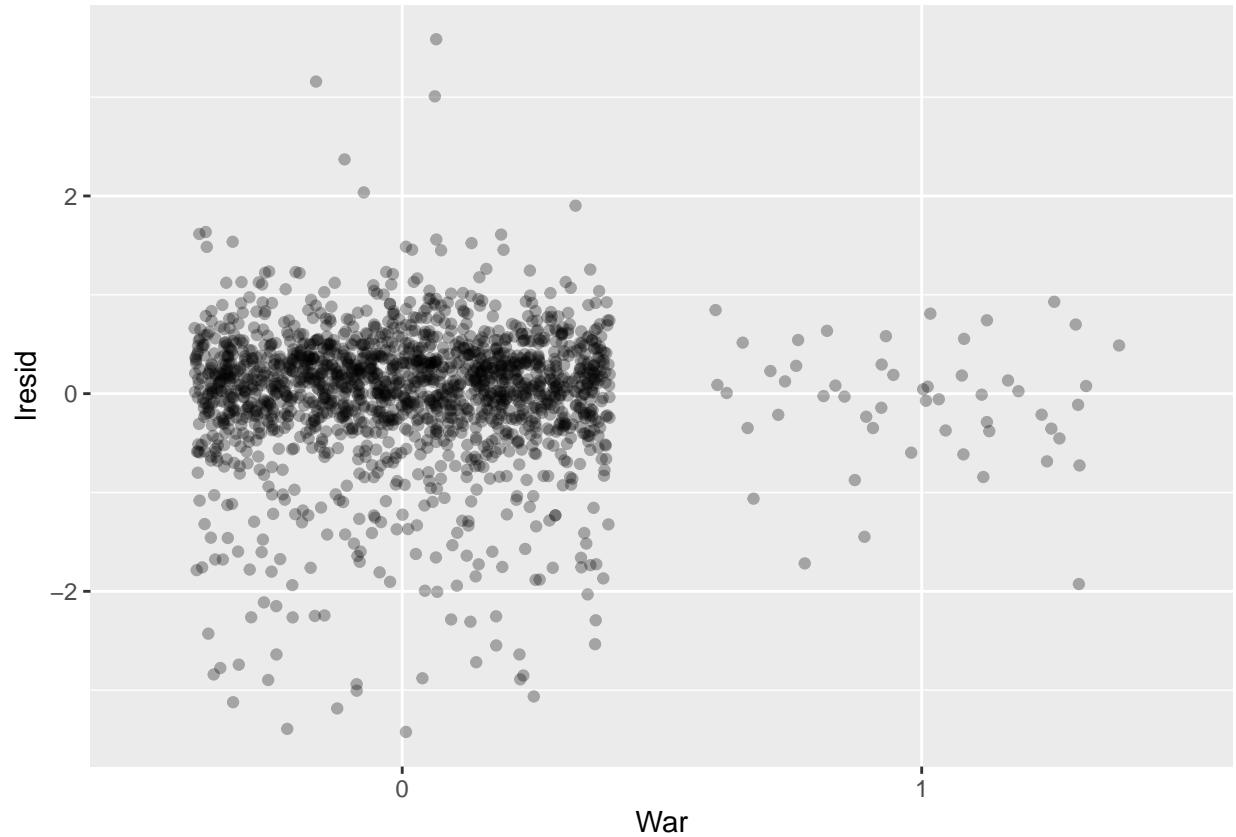
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



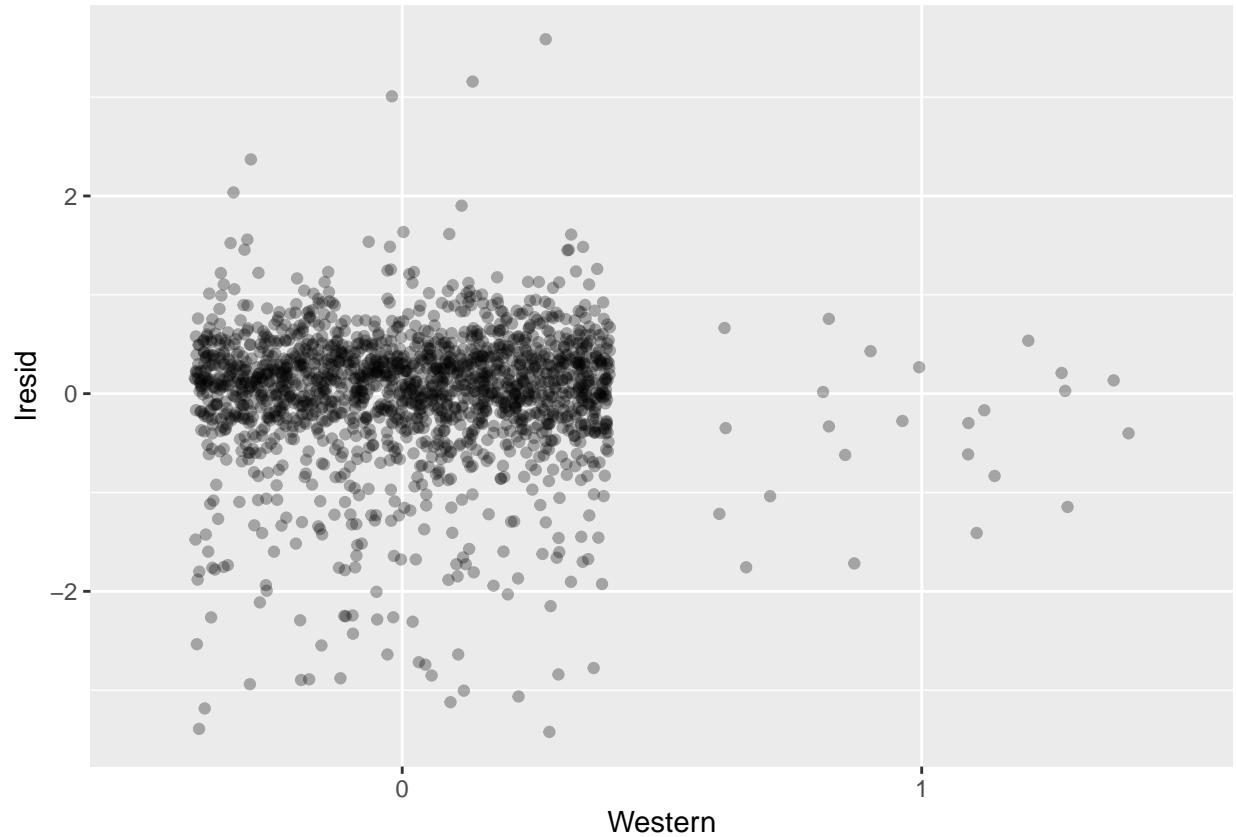
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



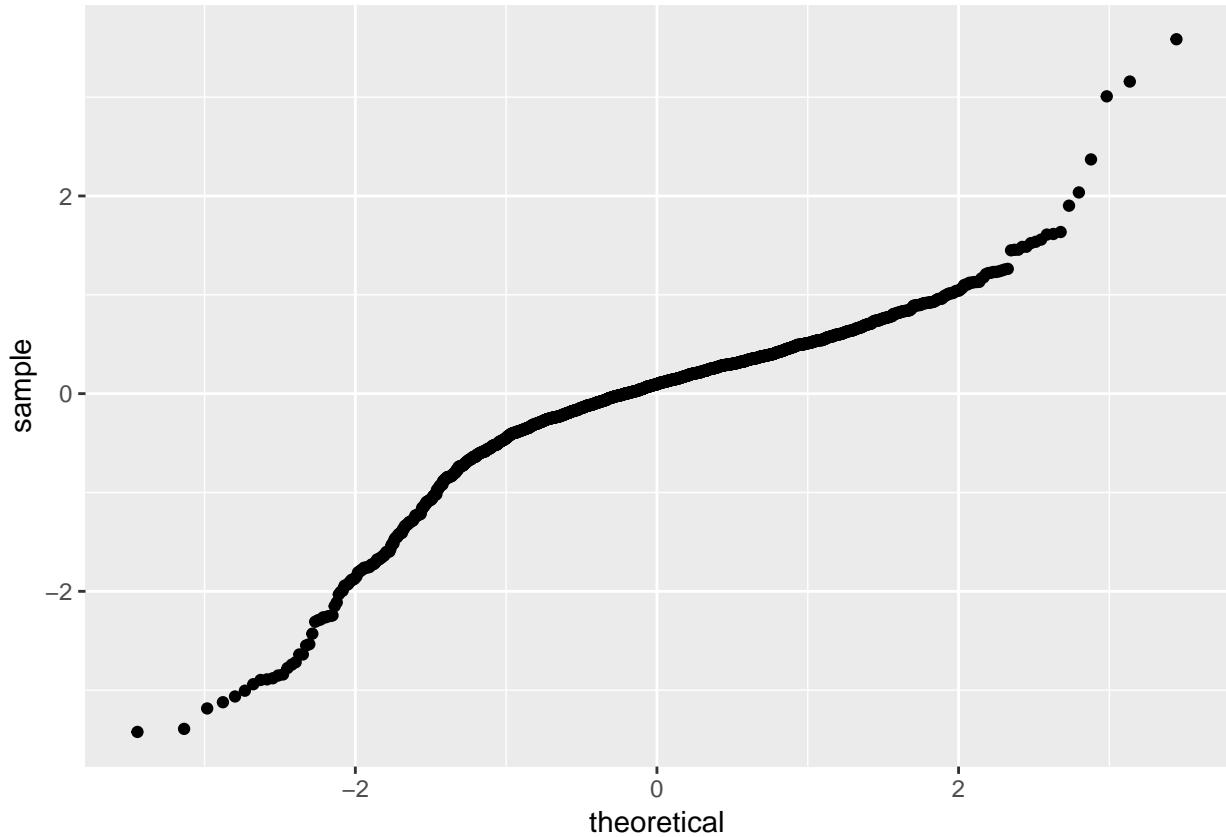
```
## Warning: Removed 102 rows containing missing values (geom_point).
```



```
## Warning: Removed 102 rows containing missing values (geom_point).
```



```
## Warning: Removed 102 rows containing non-finite values (stat_qq).
```



Use variables with non-random relationships to the residuals.

```

mod_simple_plus <- lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating +
                         Family + Horror, data = train)

summary(mod_simple_plus)

##
## Call:
## lm(formula = real_gross_log ~ real_budget_log + imdb_score_log +
##     year + content_rating + Family + Horror, data = train)
##
## Residuals:
##      Min      1Q  Median      3Q      Max 
## -3.4657 -0.2254  0.0845  0.3428  3.1910 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.01551   0.38954 -0.040   0.968    
## real_budget_log  0.81650   0.02597 31.440 < 2e-16 ***
## imdb_score_log  1.90618   0.19916  9.571 < 2e-16 ***
## year1981      -0.16844   0.37031 -0.455   0.649    
## year1982       0.21167   0.34124  0.620   0.535    
## year1983       0.21119   0.38653  0.546   0.585    
## year1984       0.51790   0.33037  1.568   0.117    
## year1985       0.26373   0.34127  0.773   0.440    
## year1986       0.06613   0.32617  0.203   0.839    
## 
```

```

## year1987          0.12252   0.31799   0.385   0.700
## year1988          0.09277   0.31192   0.297   0.766
## year1989          0.29569   0.30645   0.965   0.335
## year1990          0.07858   0.30783   0.255   0.799
## year1991          0.01181   0.31163   0.038   0.970
## year1992          -0.03313   0.31691   -0.105   0.917
## year1993          -0.05160   0.30813   -0.167   0.867
## year1994          -0.24165   0.30249   -0.799   0.424
## year1995          -0.08662   0.29423   -0.294   0.769
## year1996          -0.25017   0.28719   -0.871   0.384
## year1997          -0.15542   0.28722   -0.541   0.589
## year1998          -0.29490   0.28660   -1.029   0.304
## year1999          -0.26826   0.28327   -0.947   0.344
## year2000          -0.18925   0.28459   -0.665   0.506
## year2001          -0.28716   0.28262   -1.016   0.310
## year2002          -0.26763   0.28261   -0.947   0.344
## year2003          -0.22398   0.28405   -0.789   0.431
## year2004          -0.24424   0.28332   -0.862   0.389
## year2005          -0.27229   0.28249   -0.964   0.335
## year2006          -0.34943   0.28297   -1.235   0.217
## year2007          -0.34075   0.28438   -1.198   0.231
## year2008          -0.37503   0.28252   -1.327   0.185
## year2009          -0.37781   0.28163   -1.342   0.180
## year2010          -0.39954   0.28206   -1.417   0.157
## year2011          -0.28393   0.28449   -0.998   0.318
## year2012          -0.17336   0.28324   -0.612   0.541
## year2013          -0.10621   0.28265   -0.376   0.707
## year2014          -0.16106   0.28447   -0.566   0.571
## year2015          -0.29355   0.28662   -1.024   0.306
## year2016          -0.16115   0.30283   -0.532   0.595
## content_ratingNC-17 -0.18595   0.27781   -0.669   0.503
## content_ratingPG    -0.02590   0.11953   -0.217   0.829
## content_ratingPG-13  0.12934   0.13640   0.948   0.343
## content_ratingR     -0.07096   0.13545   -0.524   0.600
## Family1            0.33683   0.07868   4.281 1.97e-05 ***
## Horror1            0.34213   0.05192   6.590 5.89e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6106 on 1674 degrees of freedom
##   (132 observations deleted due to missingness)
## Multiple R-squared:  0.4874, Adjusted R-squared:  0.4739
## F-statistic: 36.17 on 44 and 1674 DF,  p-value: < 2.2e-16
rmse(mod_simple_plus, data = valid)

## [1] 0.6293731
anova(mod_simple_plus)

## Analysis of Variance Table
##
## Response: real_gross_log
##              Df Sum Sq Mean Sq  F value    Pr(>F)
## real_budget_log  1 491.67 491.67 1318.5462 < 2.2e-16 ***

```

```

## imdb_score_log      1  28.18   28.18   75.5846 < 2.2e-16 ***
## year                36  36.51    1.01    2.7201 2.232e-07 ***
## content_rating       4   14.13   3.53    9.4767 1.417e-07 ***
## Family               1    6.74   6.74   18.0651 2.252e-05 ***
## Horror               1   16.19   16.19   43.4245 5.886e-11 ***
## Residuals          1674 624.21    0.37
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Stepwise just with those variables with non-random relationships with residual

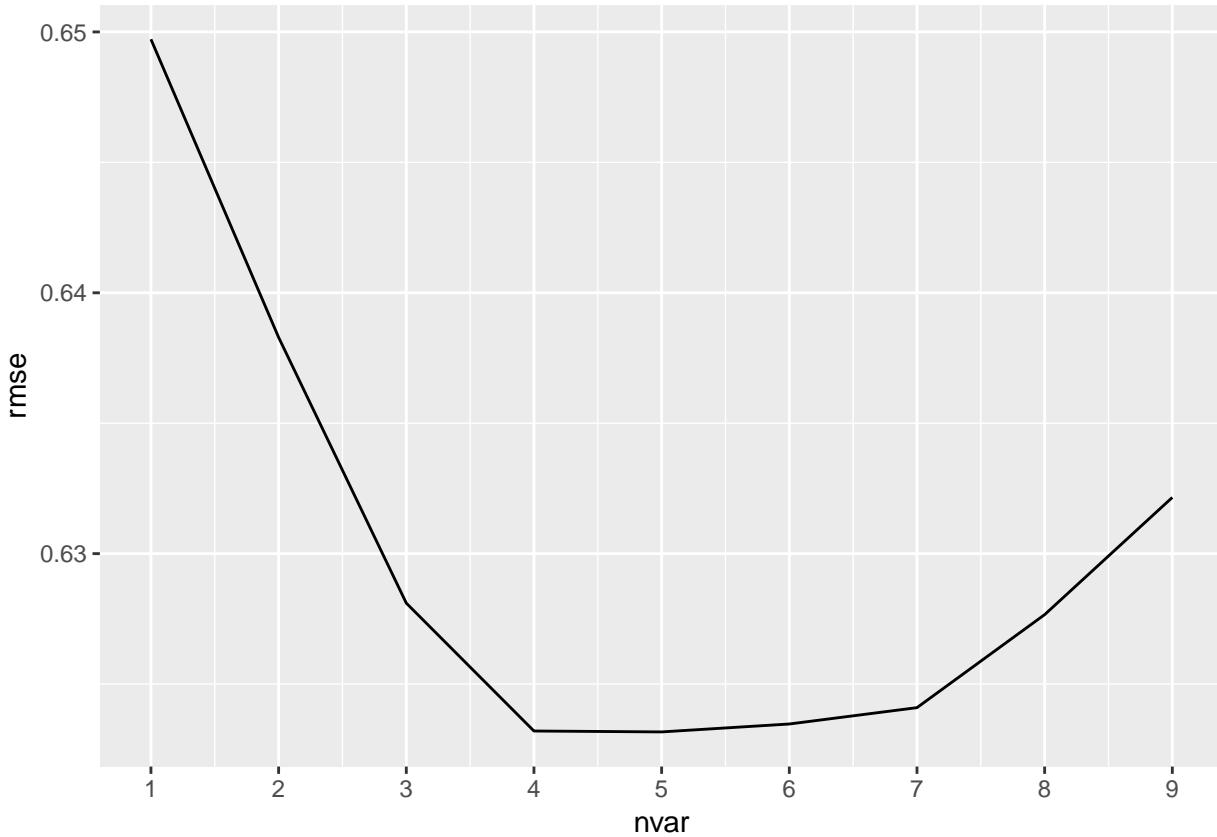
```

# stepwise
# ALL potentially relevant variables
rmse_lst <- step_wise_loop(df = train %>% select(real_budget_log, imdb_score_log, year,
                                                 content_rating, Family, Western, Fantasy, Horror,
                                                 Documentary))

## real_budget_log
##           0.6497192
## [1] 1
## imdb_score_log
##           0.6382915
## [1] 2
##       year
## 0.6281035
## [1] 3
## content_rating
##           0.623201
## [1] 4
## Documentary
## 0.6231664
## [1] 5
##     Fantasy
## 0.6234698
## [1] 6
##     Family
## 0.6240969
## [1] 7
##     Western
## 0.6276628
## [1] 8
##     Horror
## 0.6321549

# graph RMSE vs number of variables
fit_rmse <- tibble(nvar = 1:length(rmse_lst),
                     rmse = rmse_lst)
ggplot(fit_rmse) + geom_line(aes(x = nvar, y = rmse))+
  scale_x_continuous(breaks = seq(1, length(rmse_lst), by = 1))

```



```
# after var 4, decreases too small or increase

mod_simple_plus2 <- lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating,
                        data = train)

summary(mod_simple_plus2)

##
## Call:
## lm(formula = real_gross_log ~ real_budget_log + imdb_score_log +
##     year + content_rating, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -3.4983 -0.2230  0.0973  0.3455  3.4758 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 0.444778  0.391868  1.135  0.25653    
## real_budget_log 0.821725  0.026058 31.534 < 2e-16 *** 
## imdb_score_log 1.626517  0.199024  8.172 5.89e-16 *** 
## year1981    -0.238907  0.376557 -0.634  0.52587    
## year1982     0.297881  0.346750  0.859  0.39043    
## year1983     0.235272  0.393319  0.598  0.54981    
## year1984     0.504909  0.336175  1.502  0.13331    
## year1985     0.335770  0.346892  0.968  0.33322
```

```

## year1986      0.111891  0.331663  0.337  0.73589
## year1987      0.212045  0.322869  0.657  0.51143
## year1988      0.226528  0.316488  0.716  0.47424
## year1989      0.285472  0.311721  0.916  0.35991
## year1990      0.178997  0.312725  0.572  0.56714
## year1991      0.014859  0.316338  0.047  0.96254
## year1992      0.045736  0.321585  0.142  0.88692
## year1993     -0.003696  0.312675 -0.012  0.99057
## year1994     -0.184559  0.306846 -0.601  0.54761
## year1995     -0.042381  0.298579 -0.142  0.88714
## year1996     -0.208917  0.291490 -0.717  0.47365
## year1997     -0.131576  0.291460 -0.451  0.65173
## year1998     -0.246108  0.290619 -0.847  0.39721
## year1999     -0.232583  0.287331 -0.809  0.41837
## year2000     -0.141456  0.288697 -0.490  0.62421
## year2001     -0.245089  0.286678 -0.855  0.39271
## year2002     -0.222859  0.286430 -0.778  0.43665
## year2003     -0.176337  0.287847 -0.613  0.54022
## year2004     -0.193313  0.287370 -0.673  0.50123
## year2005     -0.210828  0.286412 -0.736  0.46177
## year2006     -0.308781  0.286785 -1.077  0.28177
## year2007     -0.267002  0.288303 -0.926  0.35452
## year2008     -0.323941  0.286439 -1.131  0.25825
## year2009     -0.325844  0.285759 -1.140  0.25434
## year2010     -0.338852  0.286066 -1.185  0.23637
## year2011     -0.217375  0.288402 -0.754  0.45112
## year2012     -0.110811  0.287004 -0.386  0.69948
## year2013     -0.037216  0.286569 -0.130  0.89669
## year2014     -0.103257  0.288515 -0.358  0.72047
## year2015     -0.216276  0.290823 -0.744  0.45718
## year2016     -0.069711  0.307146 -0.227  0.82048
## content_ratingNC-17 -0.197326  0.272730 -0.724  0.46946
## content_ratingPG -0.123677  0.119665 -1.034  0.30150
## content_ratingPG-13 -0.175106  0.115920 -1.511  0.13109
## content_ratingR   -0.346131  0.115941 -2.985  0.00287 **

## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6214 on 1676 degrees of freedom
##   (132 observations deleted due to missingness)
## Multiple R-squared:  0.4685, Adjusted R-squared:  0.4552
## F-statistic: 35.18 on 42 and 1676 DF,  p-value: < 2.2e-16

rmse(mod_simple_plus2, data = valid)

## [1] 0.623201

# when consider factors as one variable, they are significant
anova(mod_simple_plus2)

## Analysis of Variance Table
##
## Response: real_gross_log
##              Df Sum Sq Mean Sq  F value    Pr(>F)
## real_budget_log  1 491.67 491.67 1273.3487 < 2.2e-16 ***

```

```

## imdb_score_log      1  28.18    28.18   72.9937 < 2.2e-16 ***
## year                 36  36.51     1.01    2.6269 6.334e-07 ***
## content_rating       4   14.13     3.53    9.1519 2.589e-07 ***
## Residuals          1676 647.14     0.39
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
# number of observations
nobs(mod_simple_plus2)

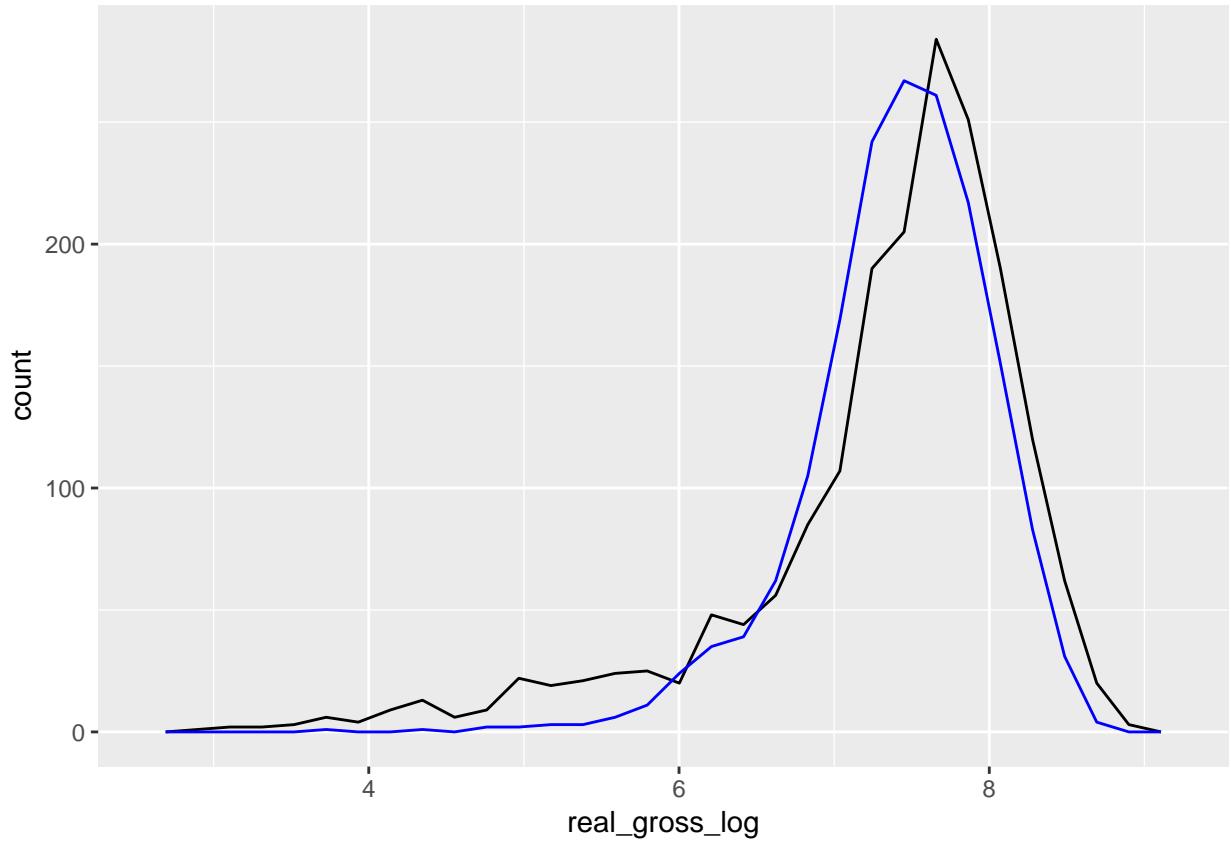
## [1] 1719
# compare with anova
anova(mod_simple_plus, mod_simple_plus2)

## Analysis of Variance Table
##
## Model 1: real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating +
##           Family + Horror
## Model 2: real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating
##   Res.Df   RSS Df Sum of Sq    F   Pr(>F)
## 1    1674 624.21
## 2    1676 647.14 -2   -22.929 30.745 7.71e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Predictions
train %>%
  add_predictions(mod_simple_plus, 'lpred') %>%
  ggplot() +
  geom_freqpoly(aes(x = real_gross_log)) +
  geom_freqpoly(aes(x = lpred), color = 'blue')

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 132 rows containing non-finite values (stat_bin).

```

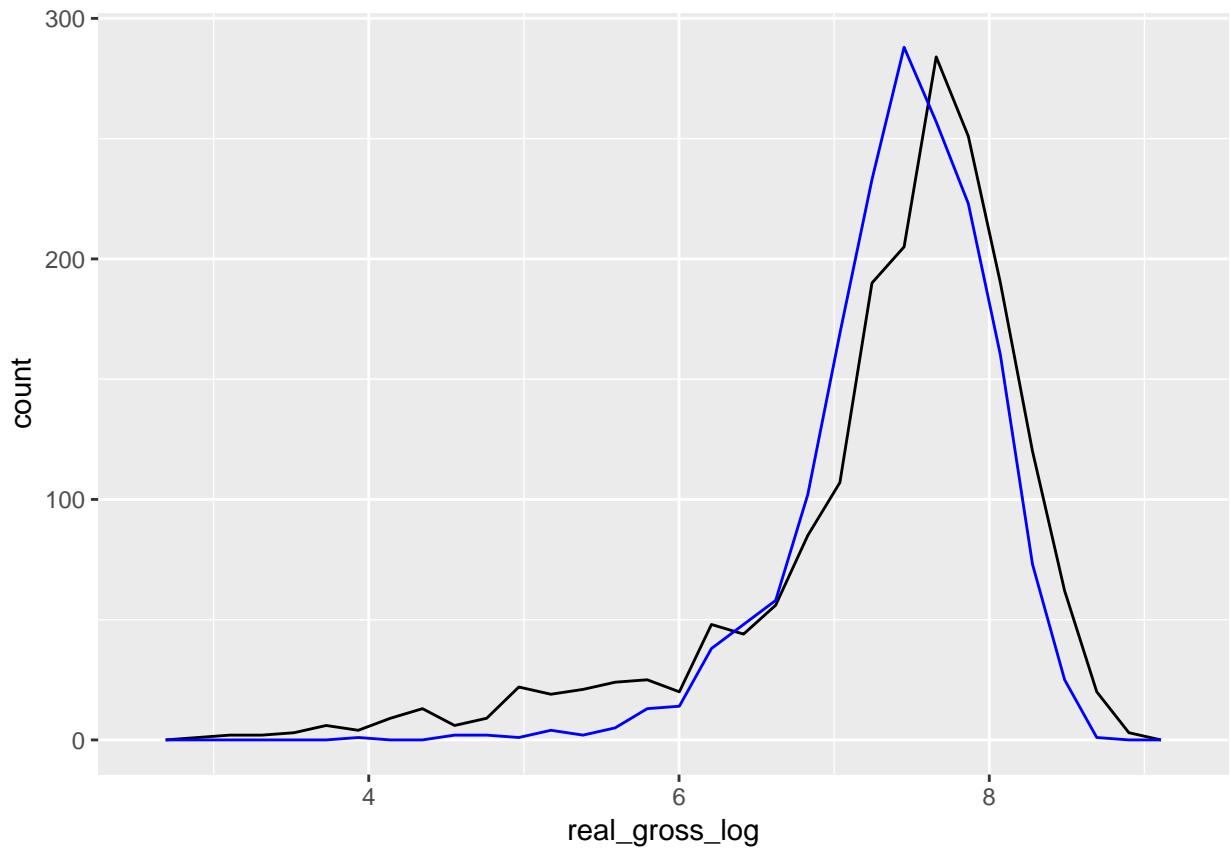


```

train %>%
  add_predictions(mod_simple_plus2, 'lpred') %>%
  ggplot() +
  geom_freqpoly(aes(x = real_gross_log)) +
  geom_freqpoly(aes(x = lpred), color = 'blue')

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 132 rows containing non-finite values (stat_bin).

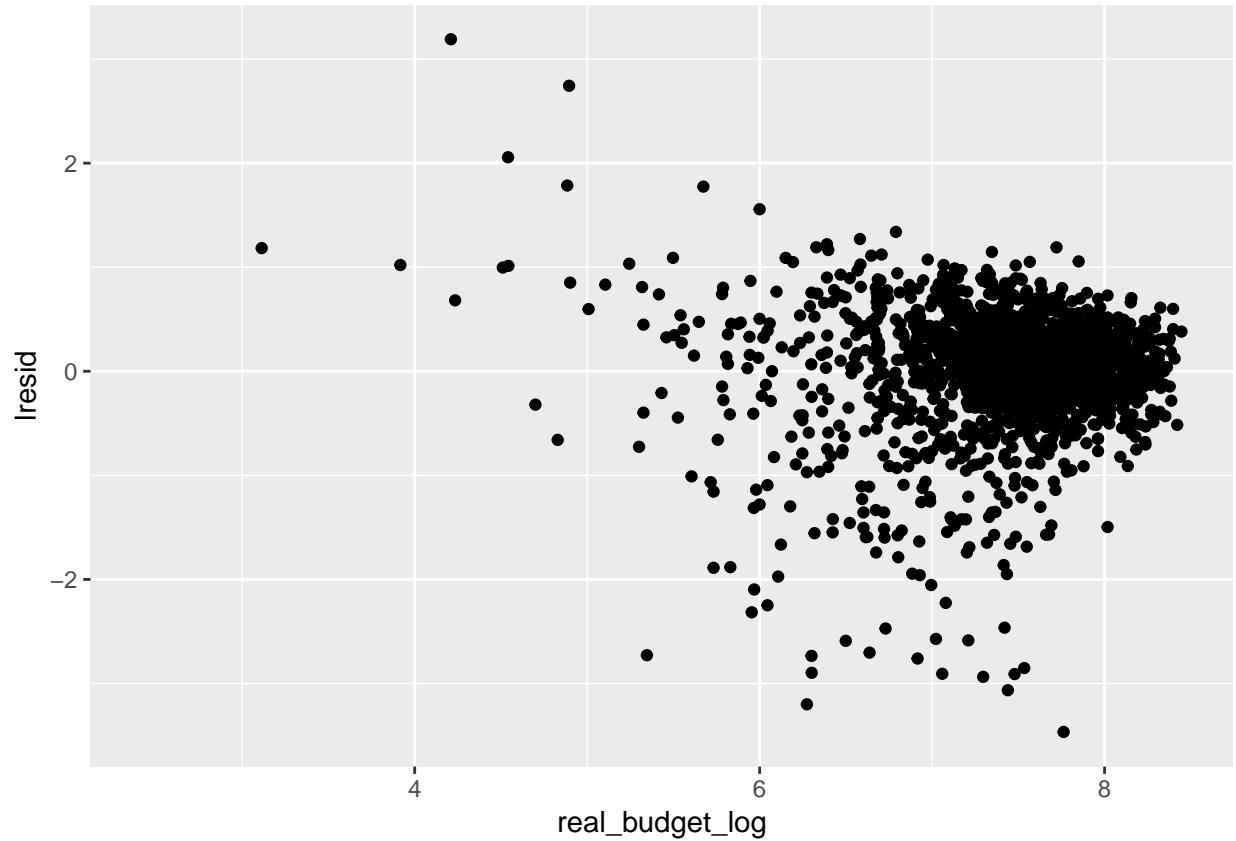
```



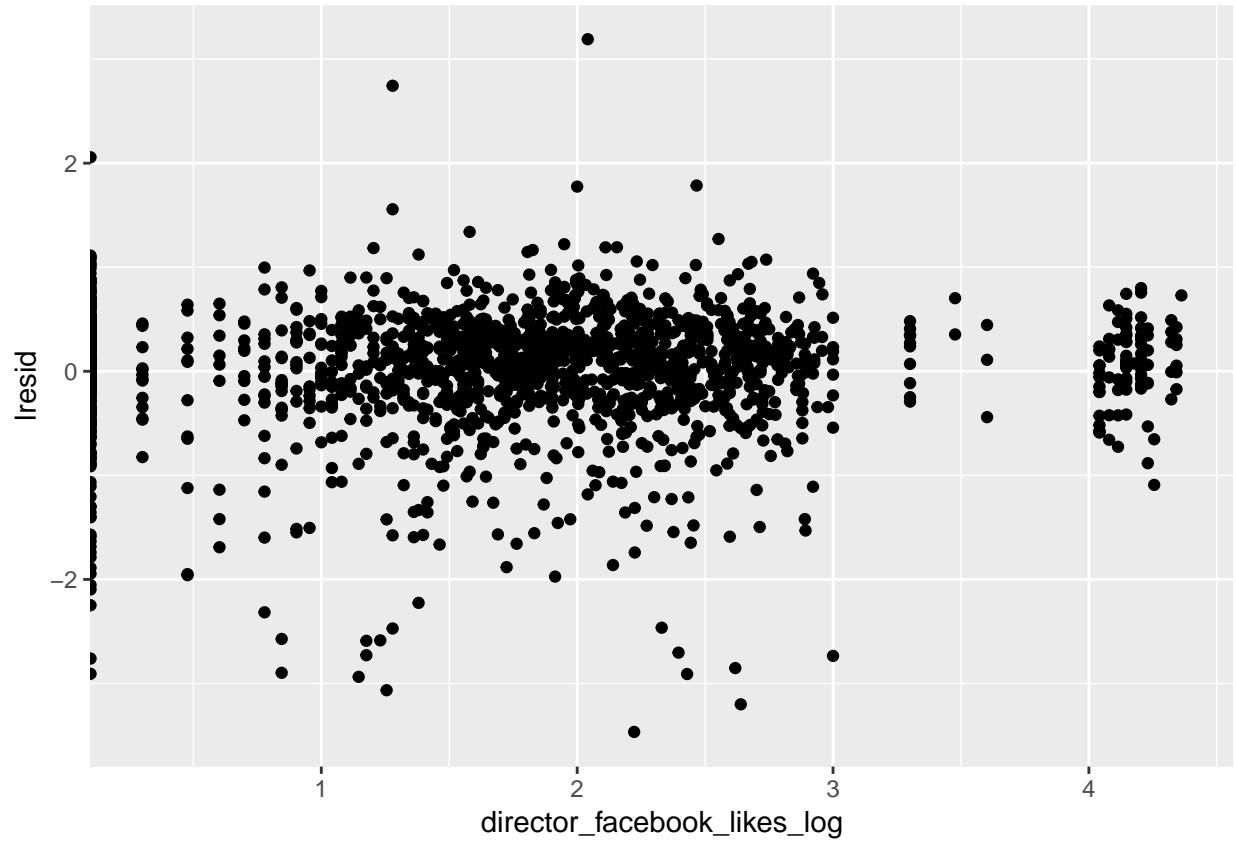
Residuals

```
gr_resid(mod_simple_plus)
```

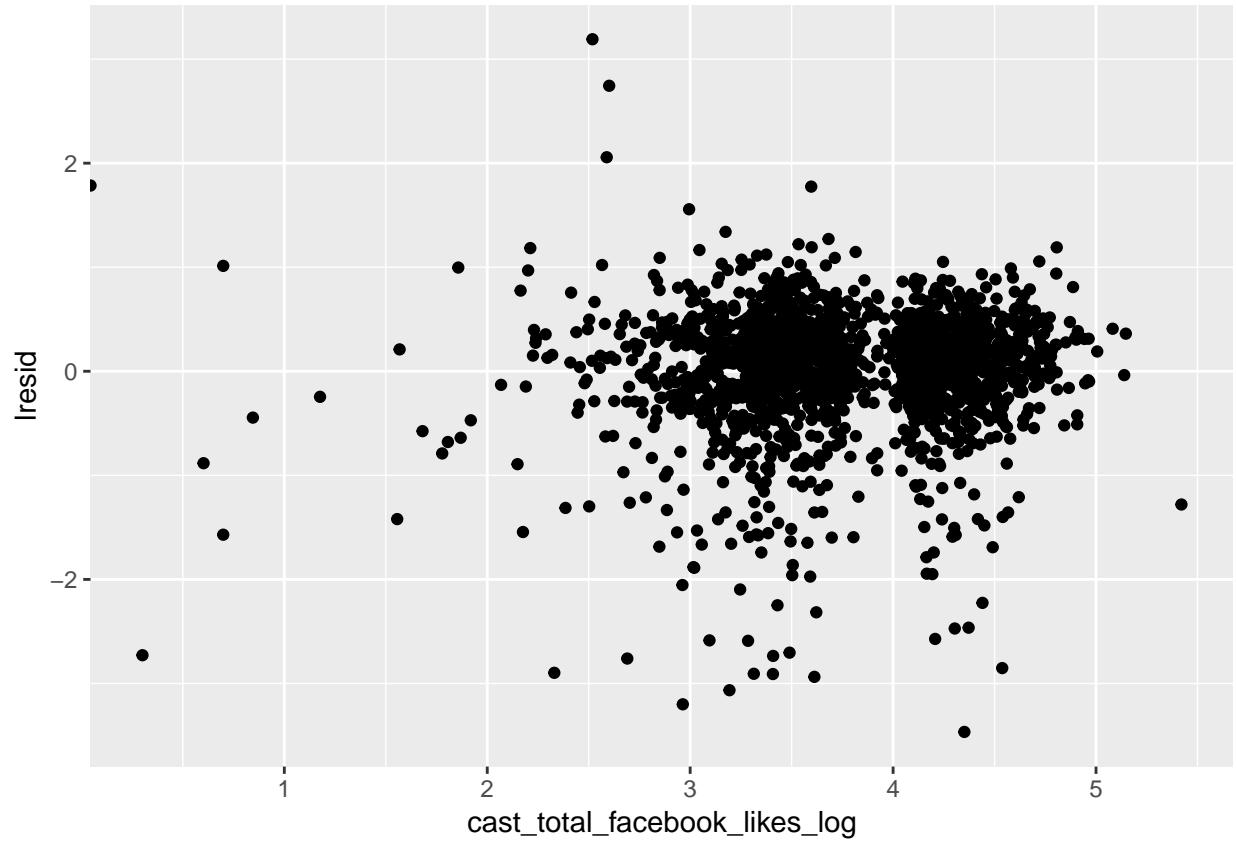
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



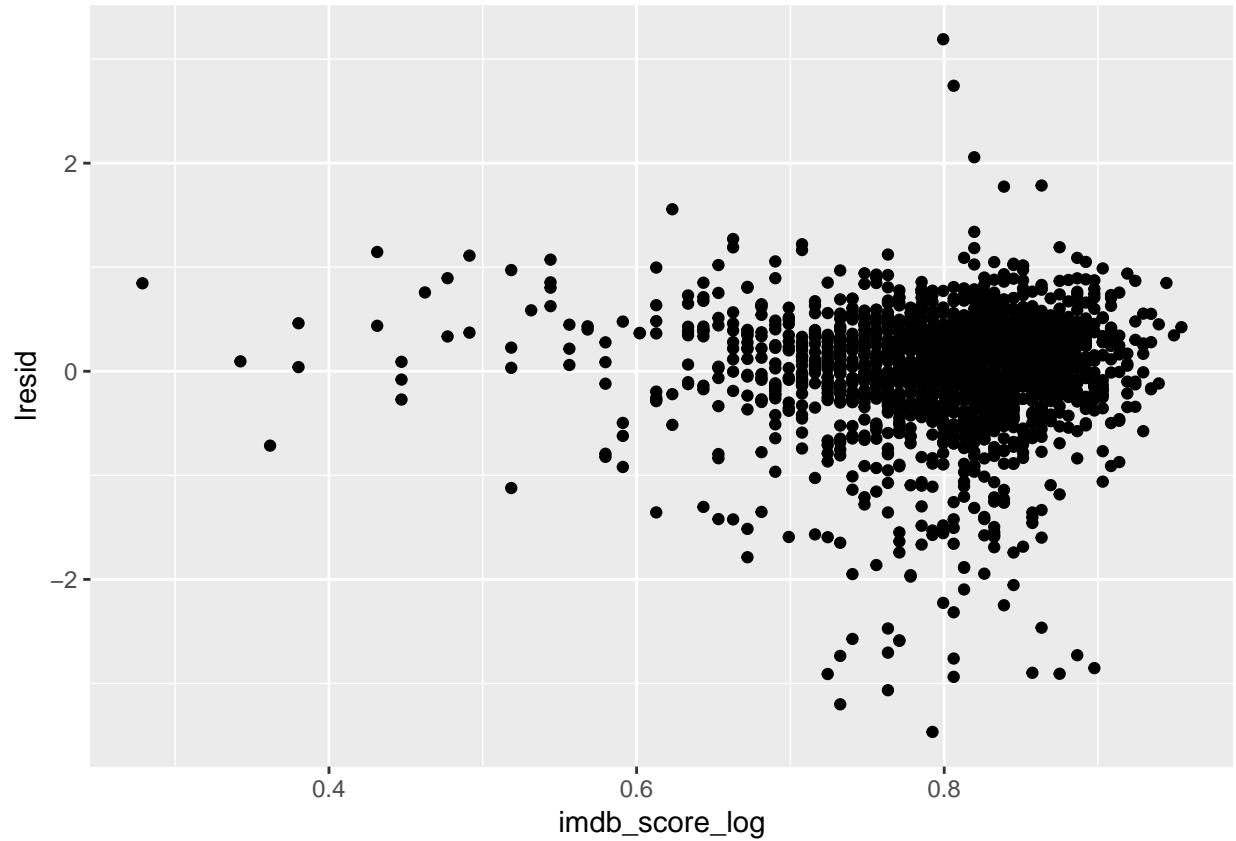
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



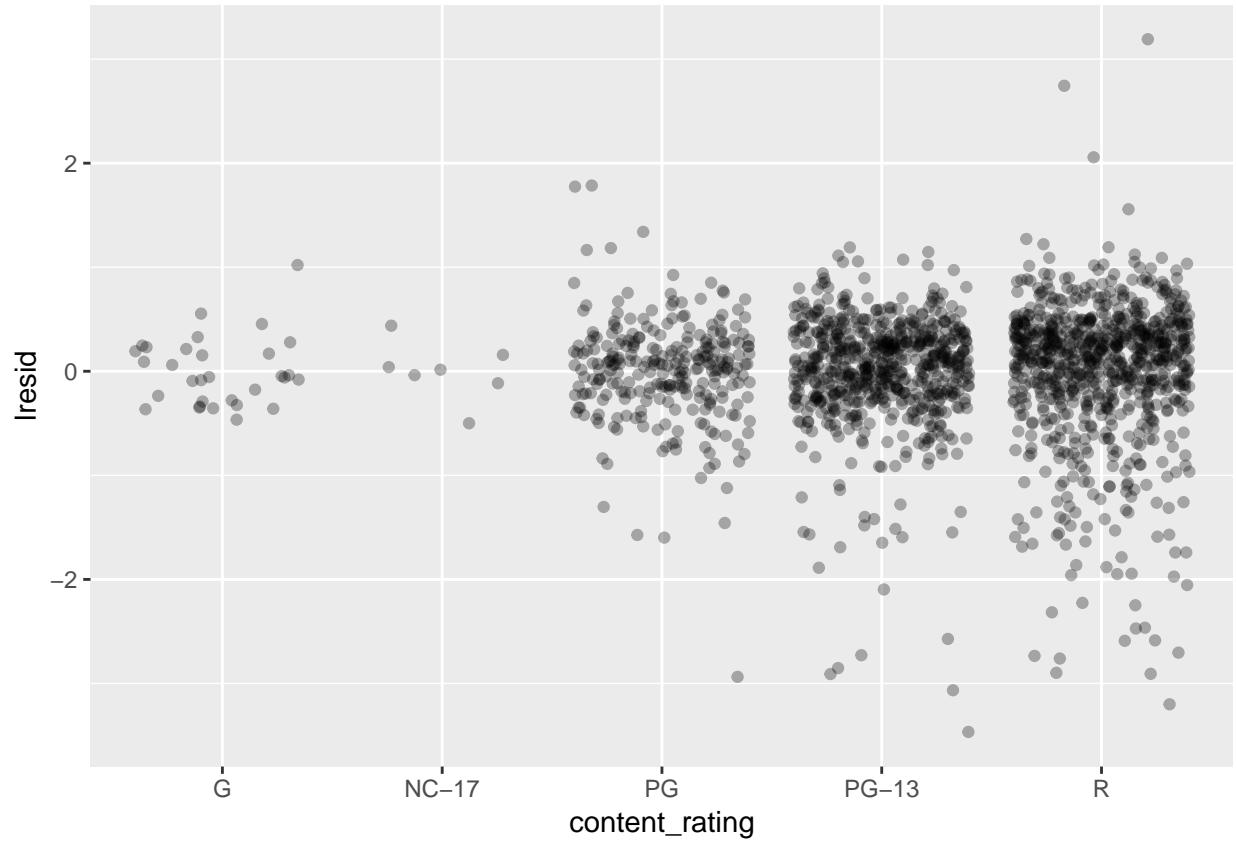
```
## Warning: Removed 132 rows containing missing values (geom_point).
```

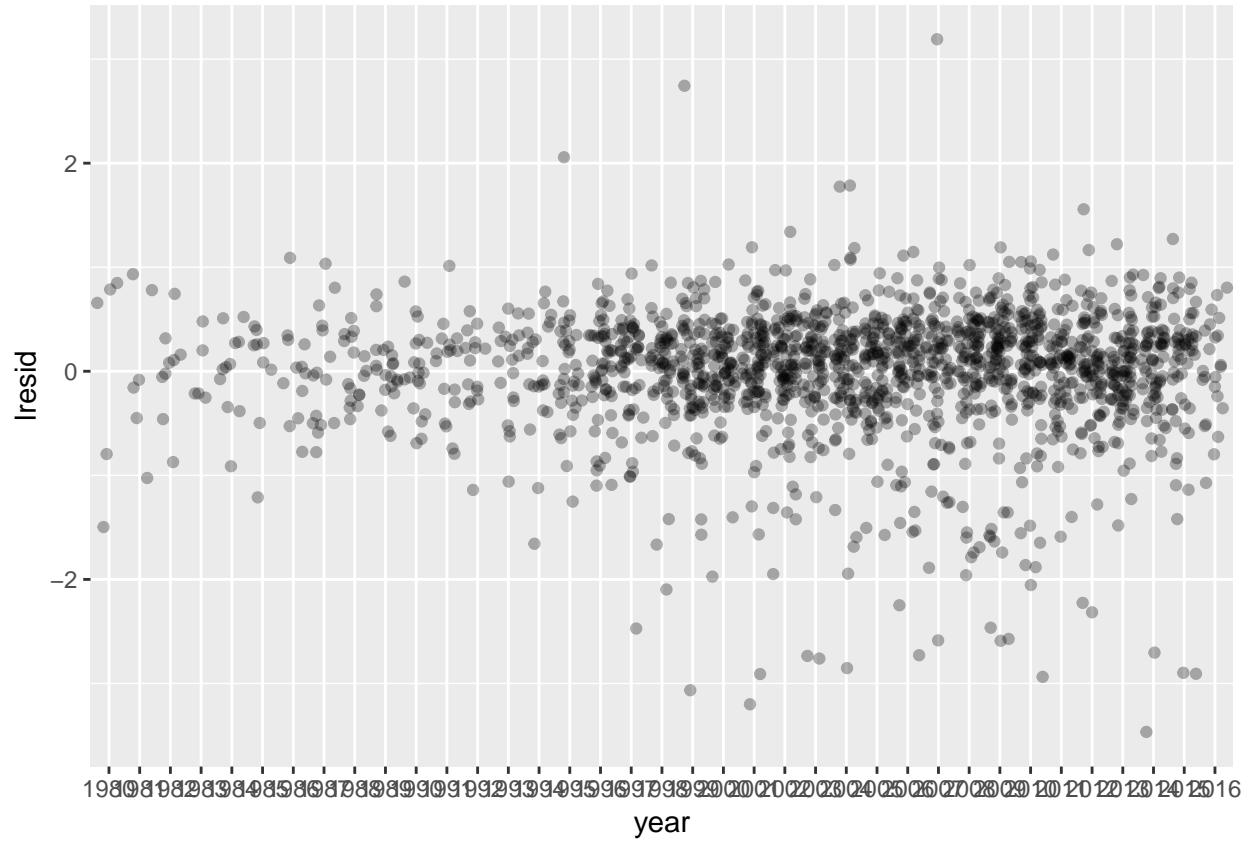


```
## Warning: Removed 132 rows containing missing values (geom_point).
```

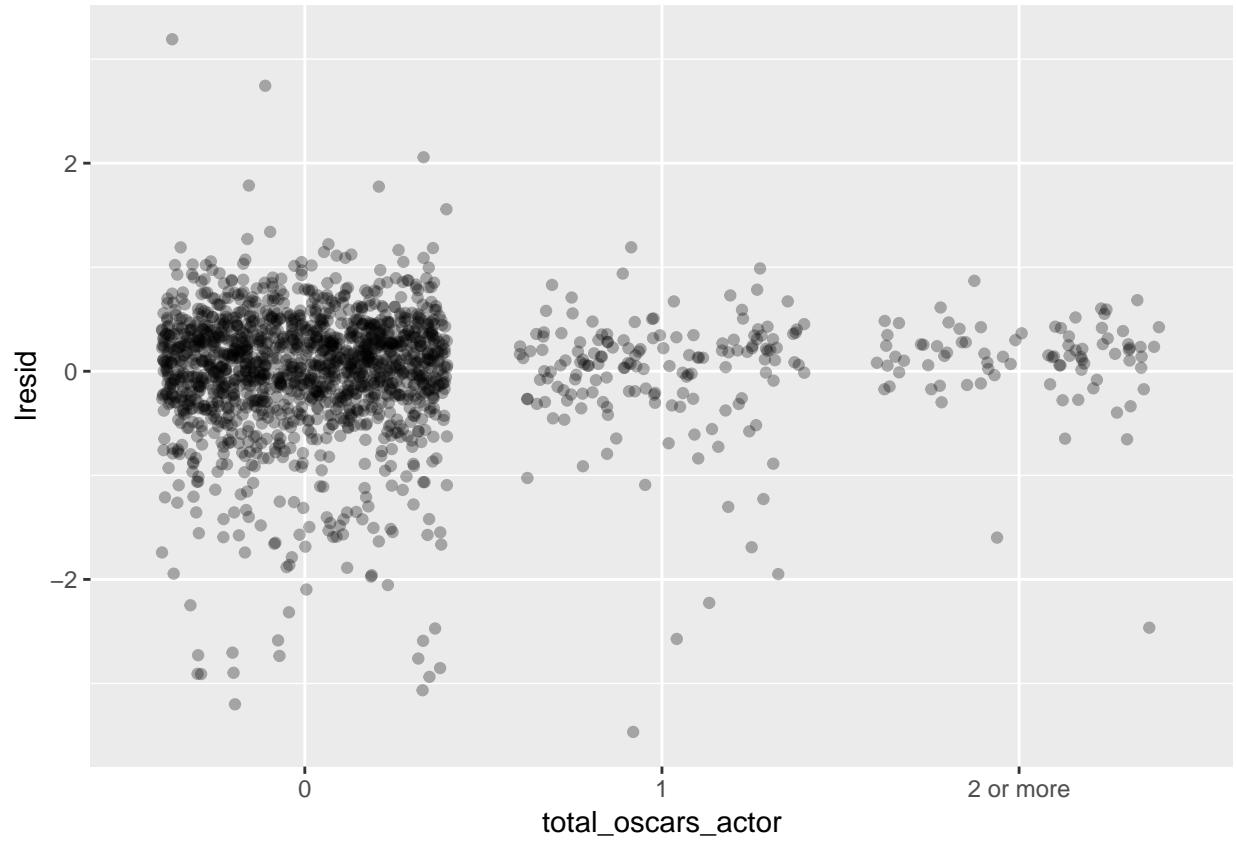


```
## Warning: Removed 94 rows containing missing values (geom_point).
```

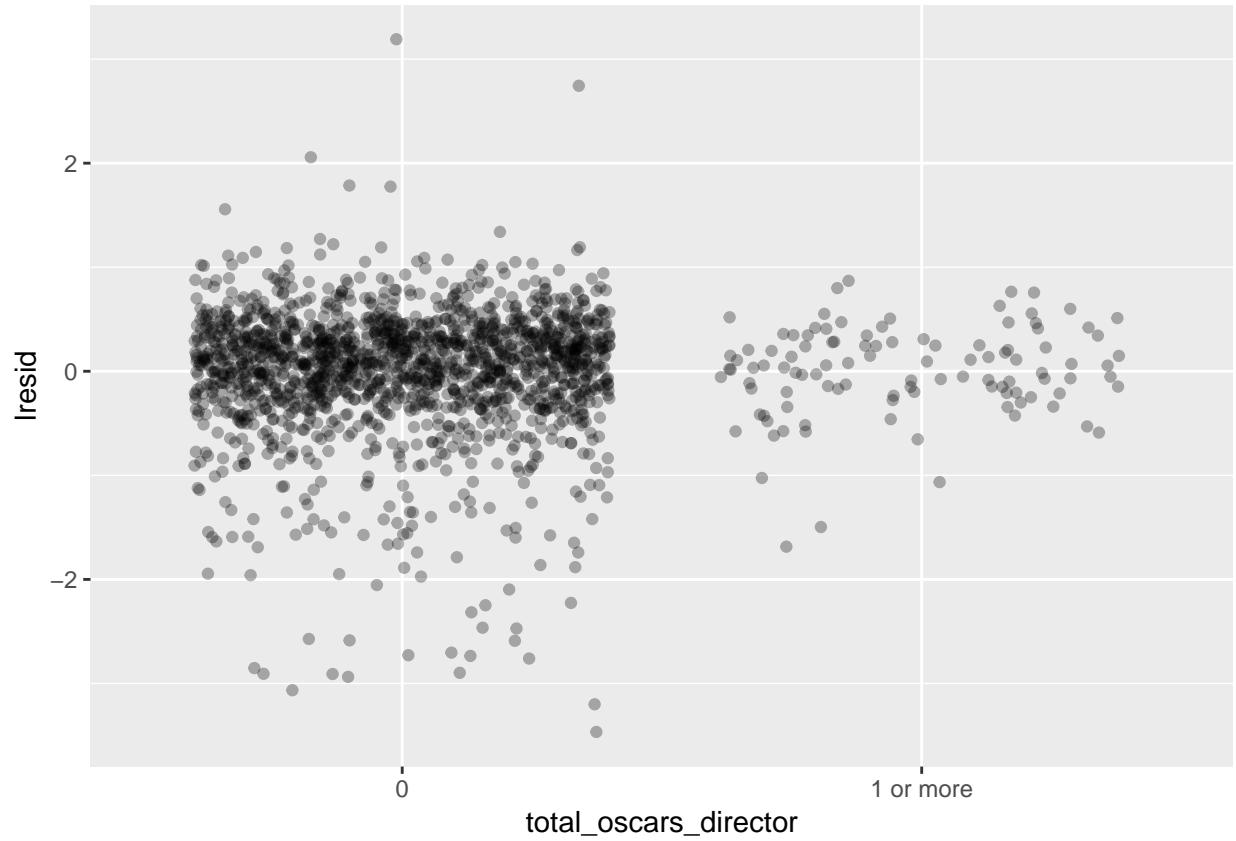




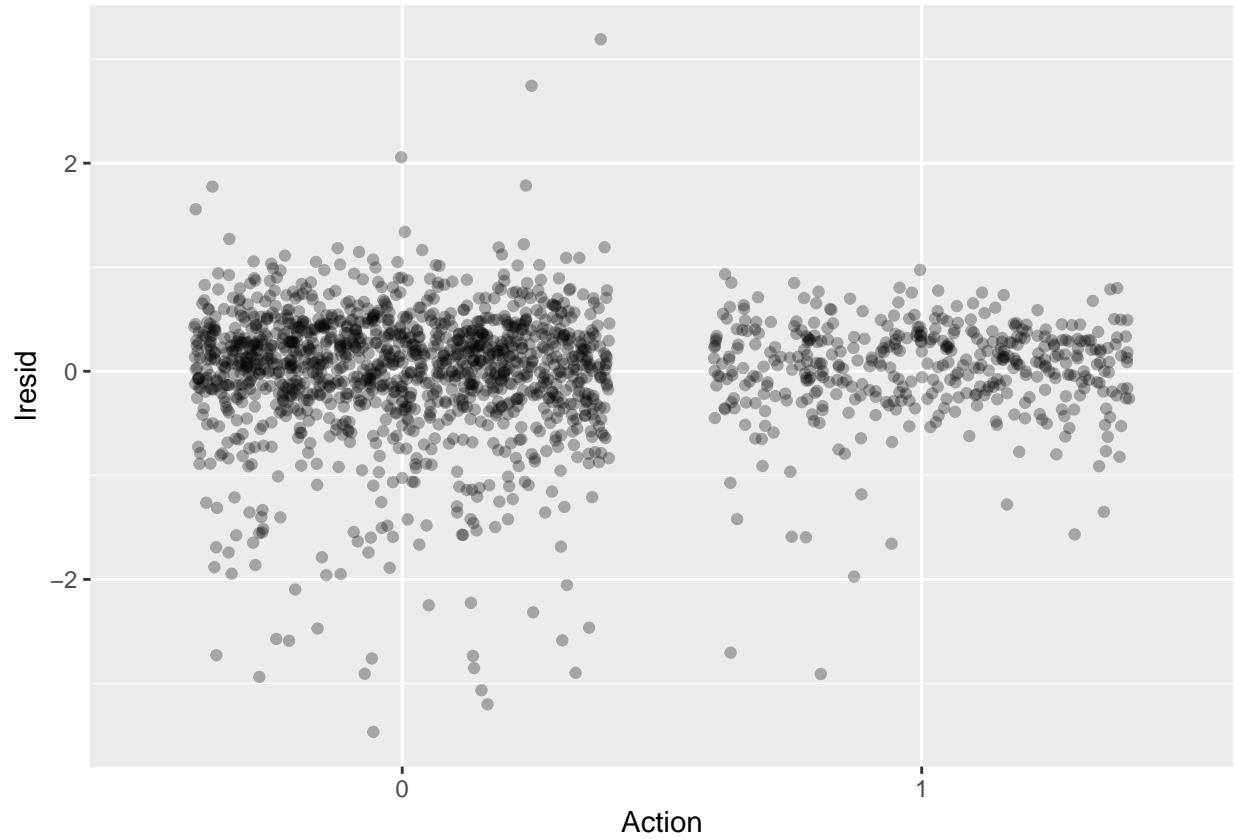
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



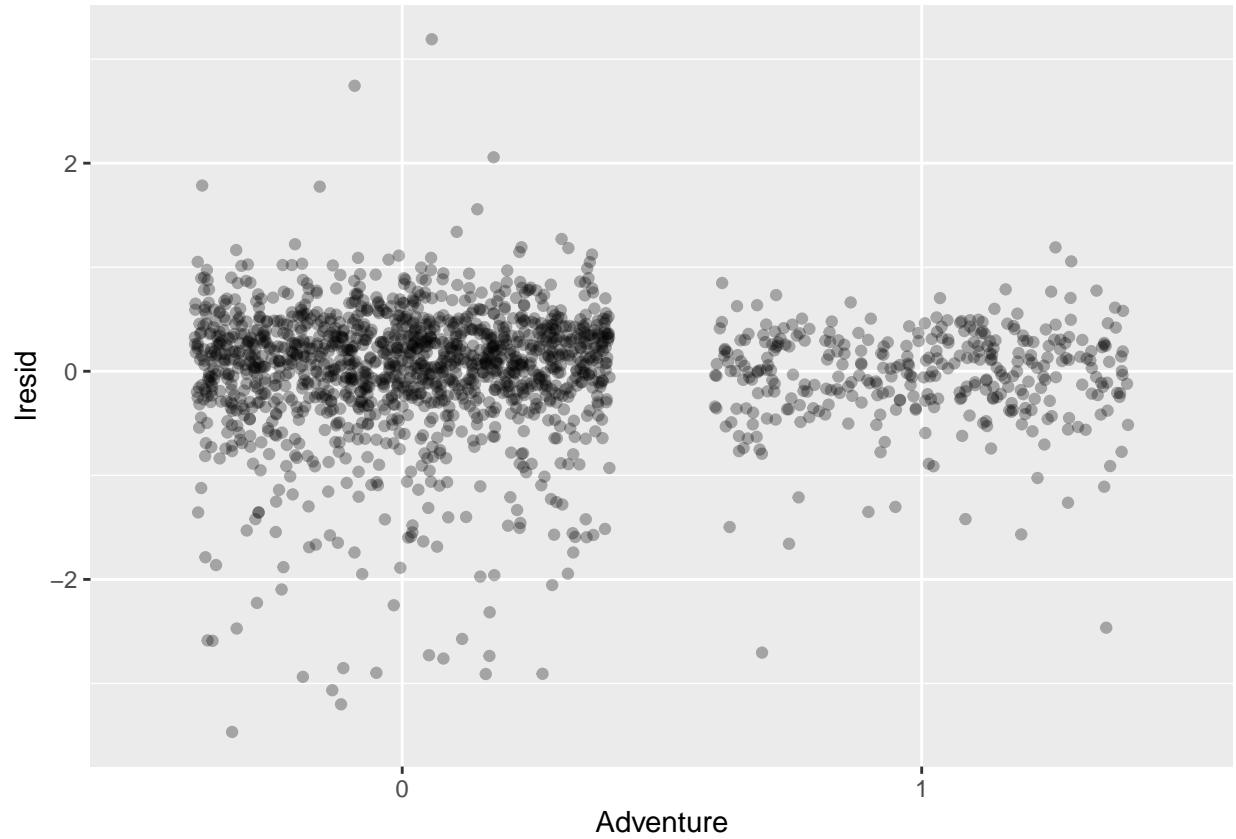
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



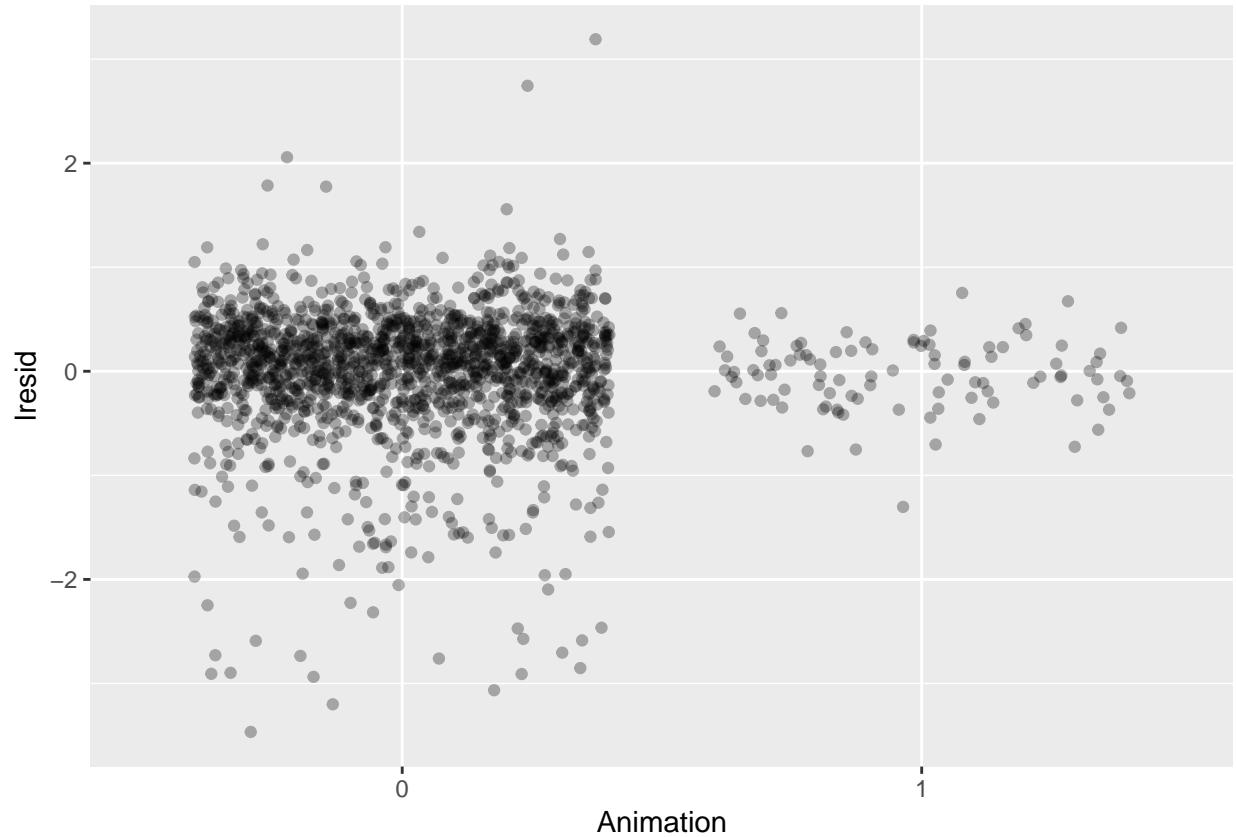
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



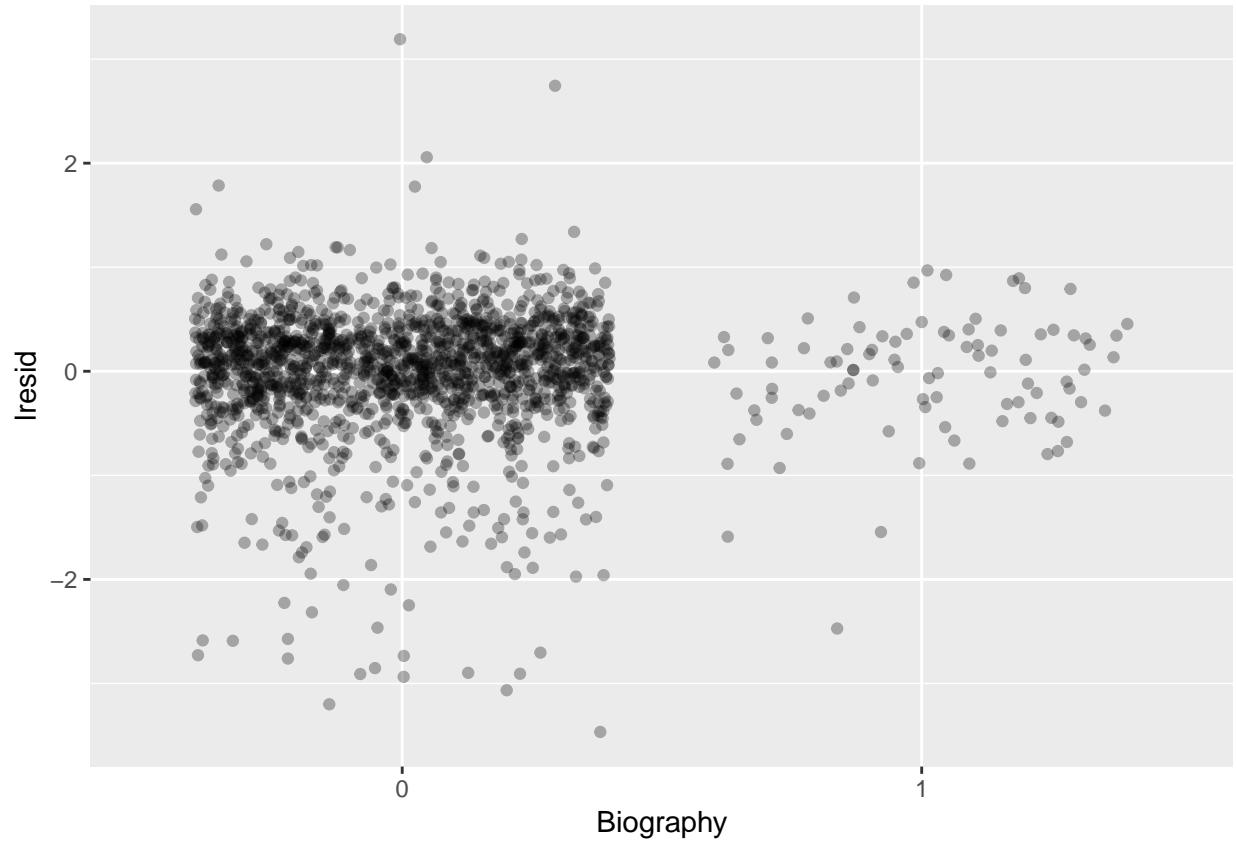
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



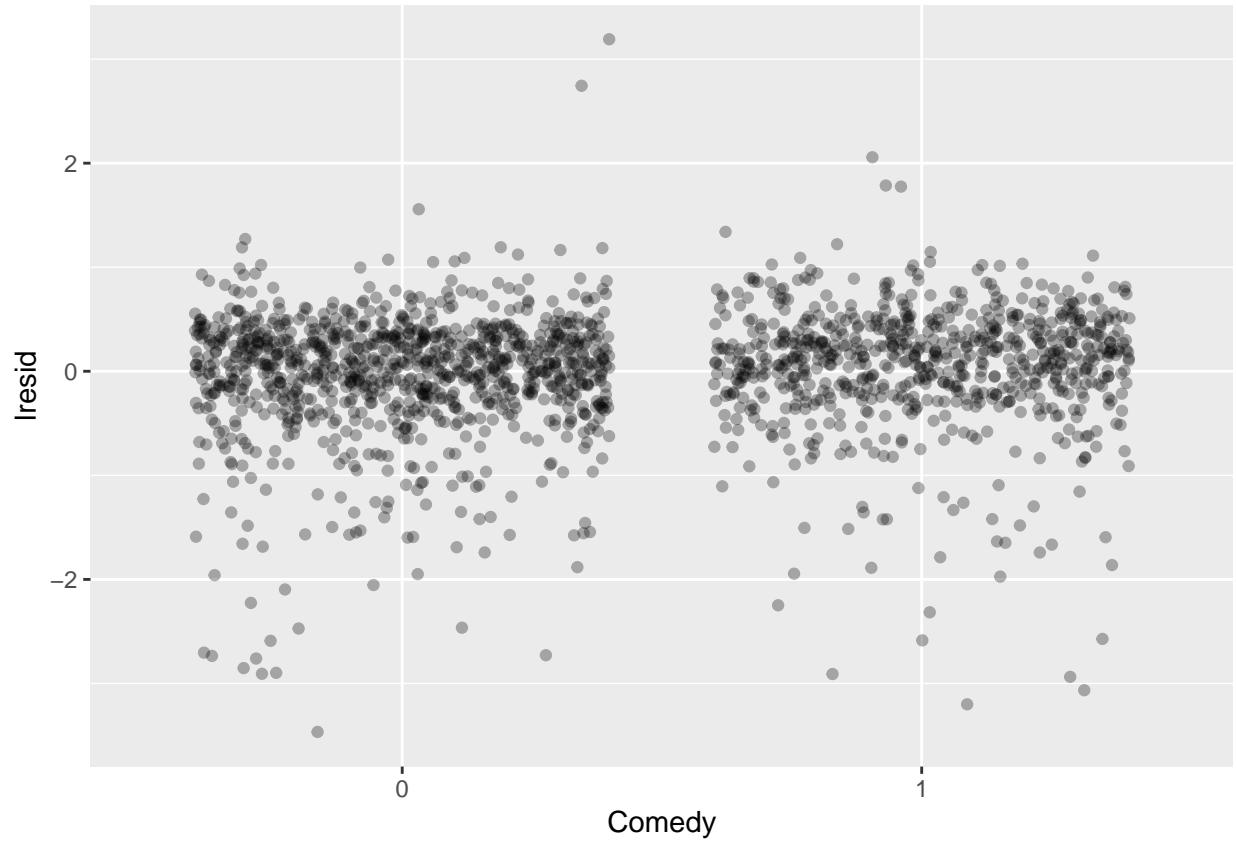
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



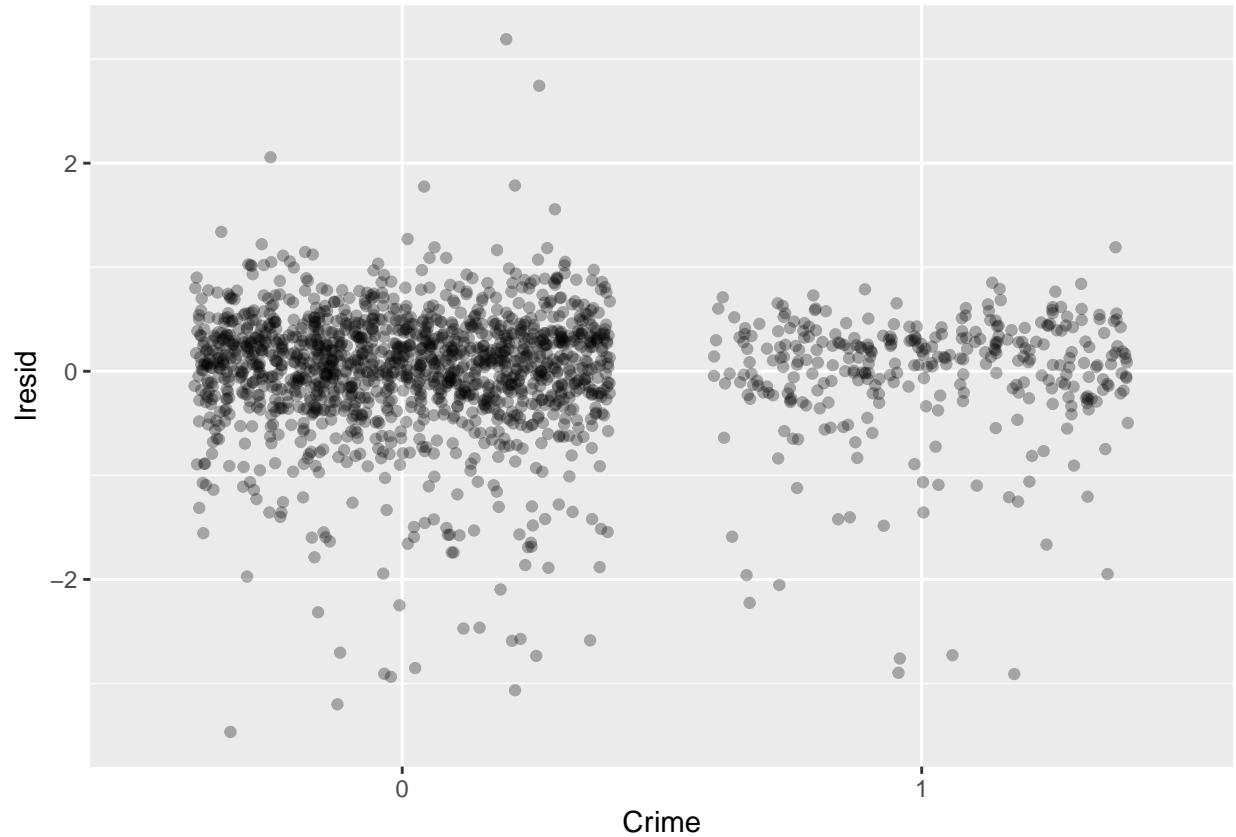
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



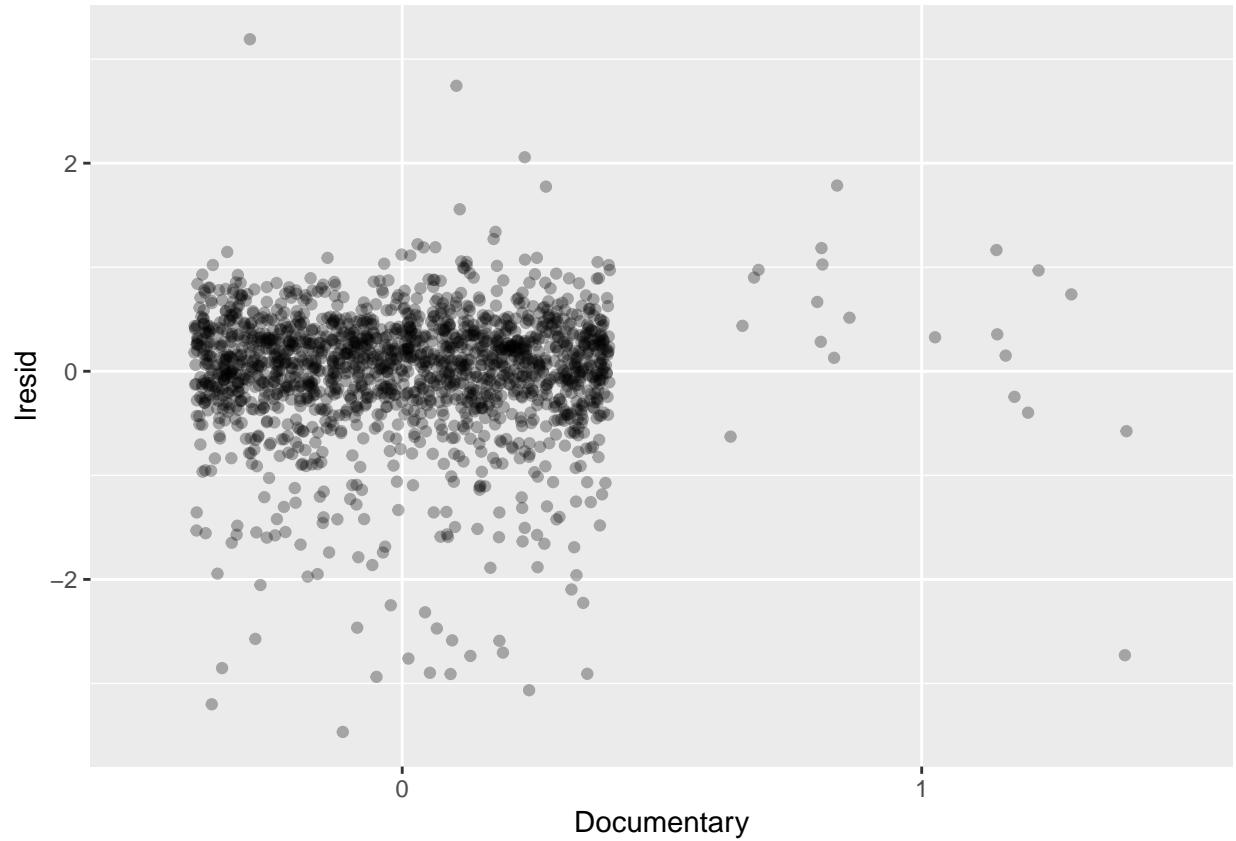
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



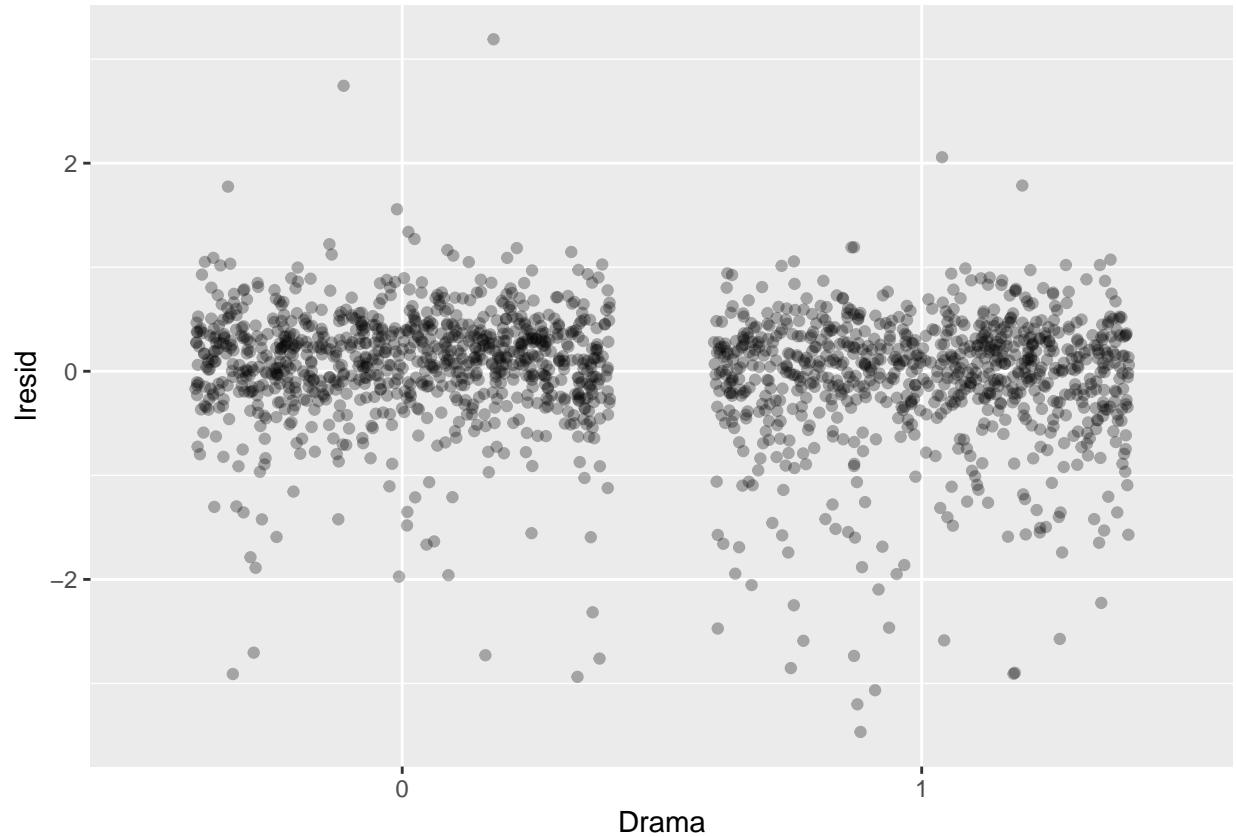
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



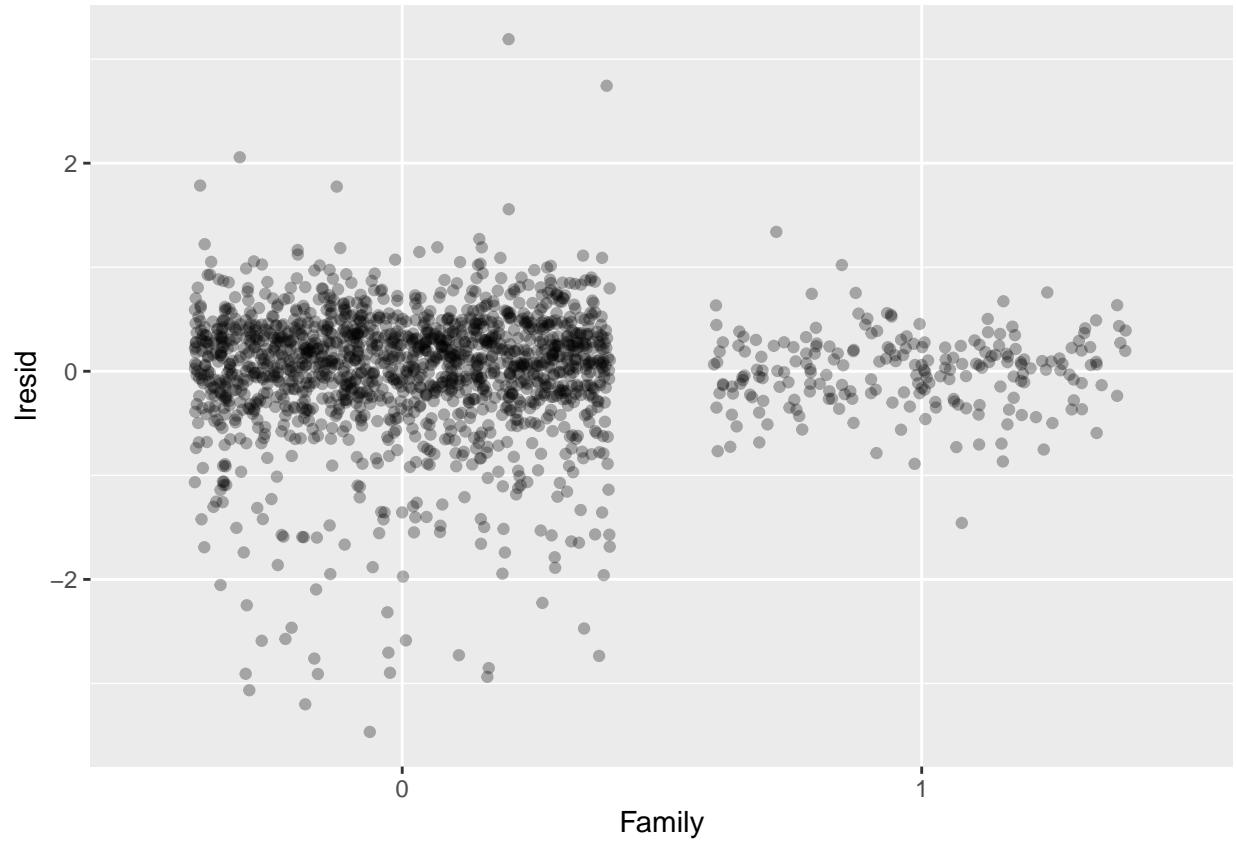
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



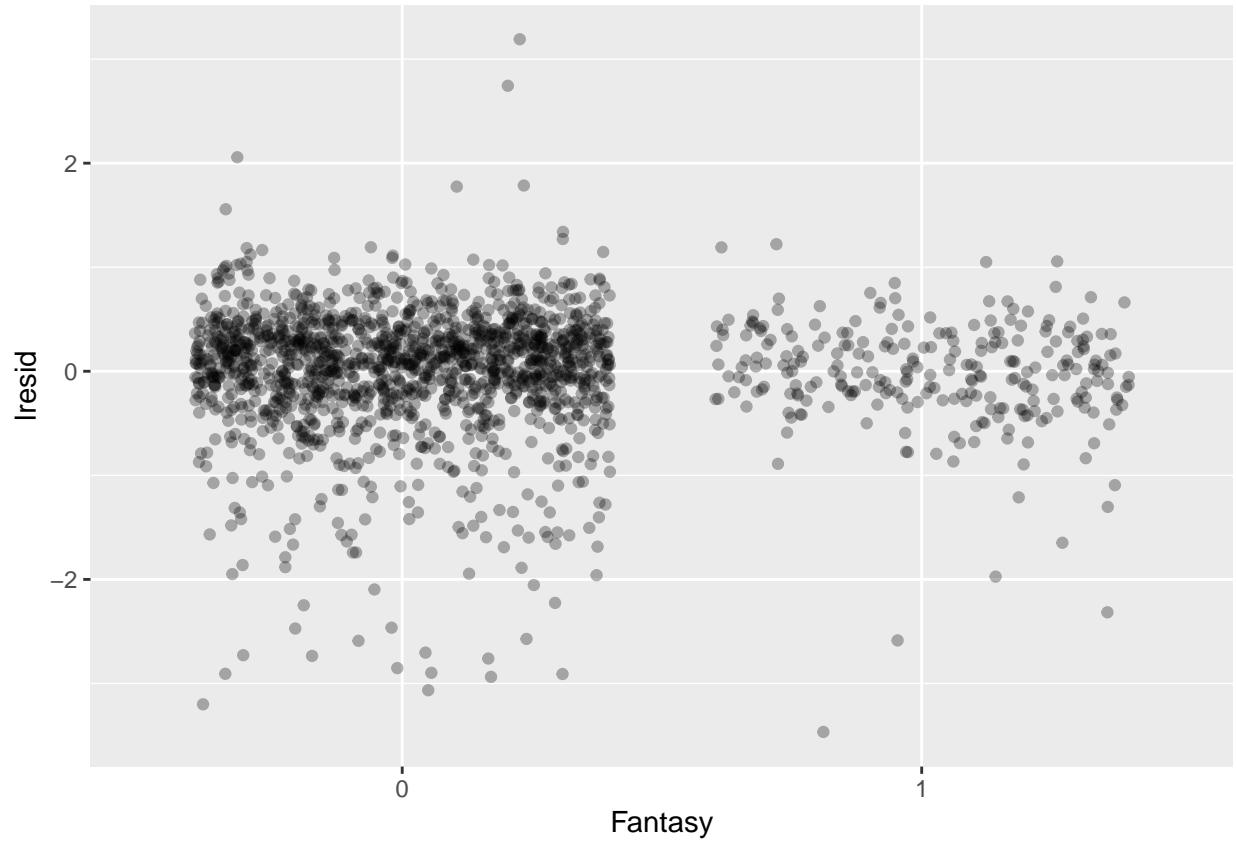
```
## Warning: Removed 132 rows containing missing values (geom_point).
```

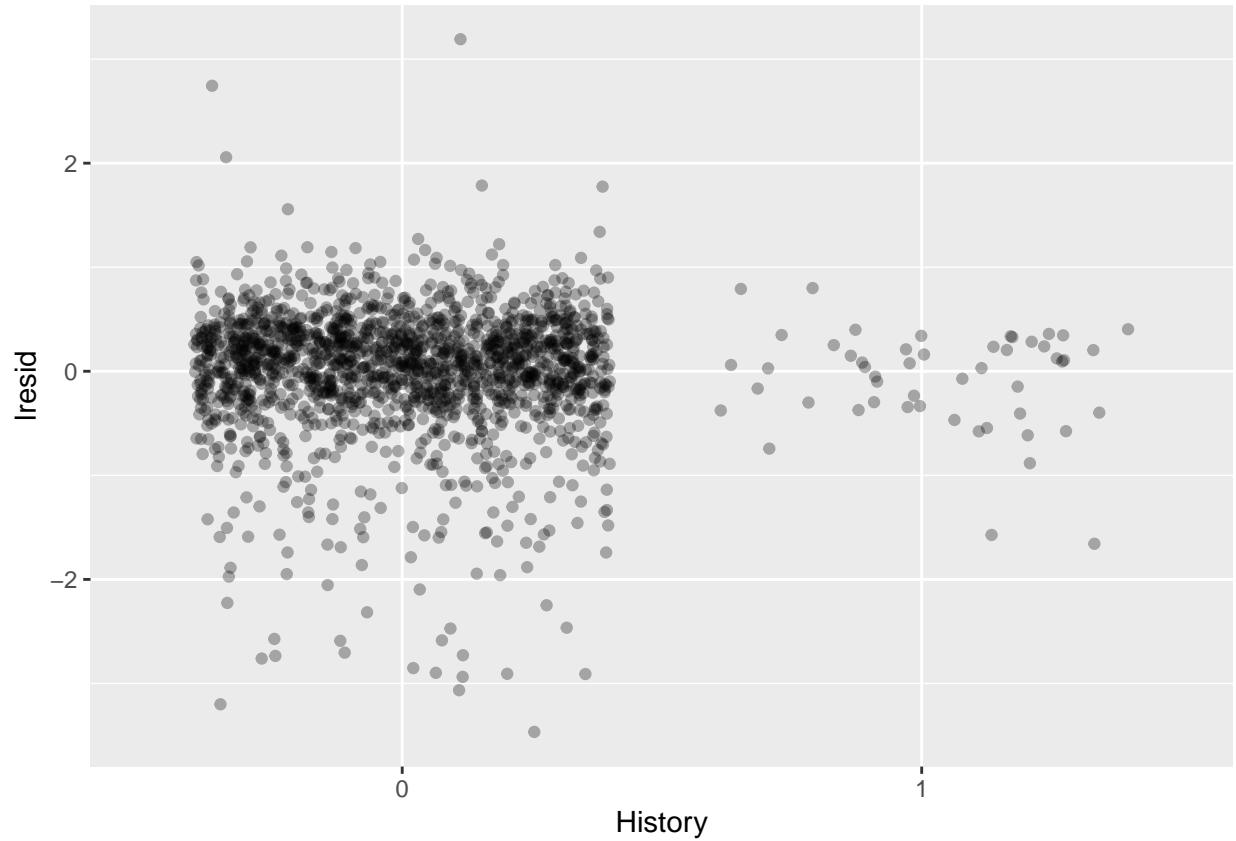


```
## Warning: Removed 132 rows containing missing values (geom_point).
```

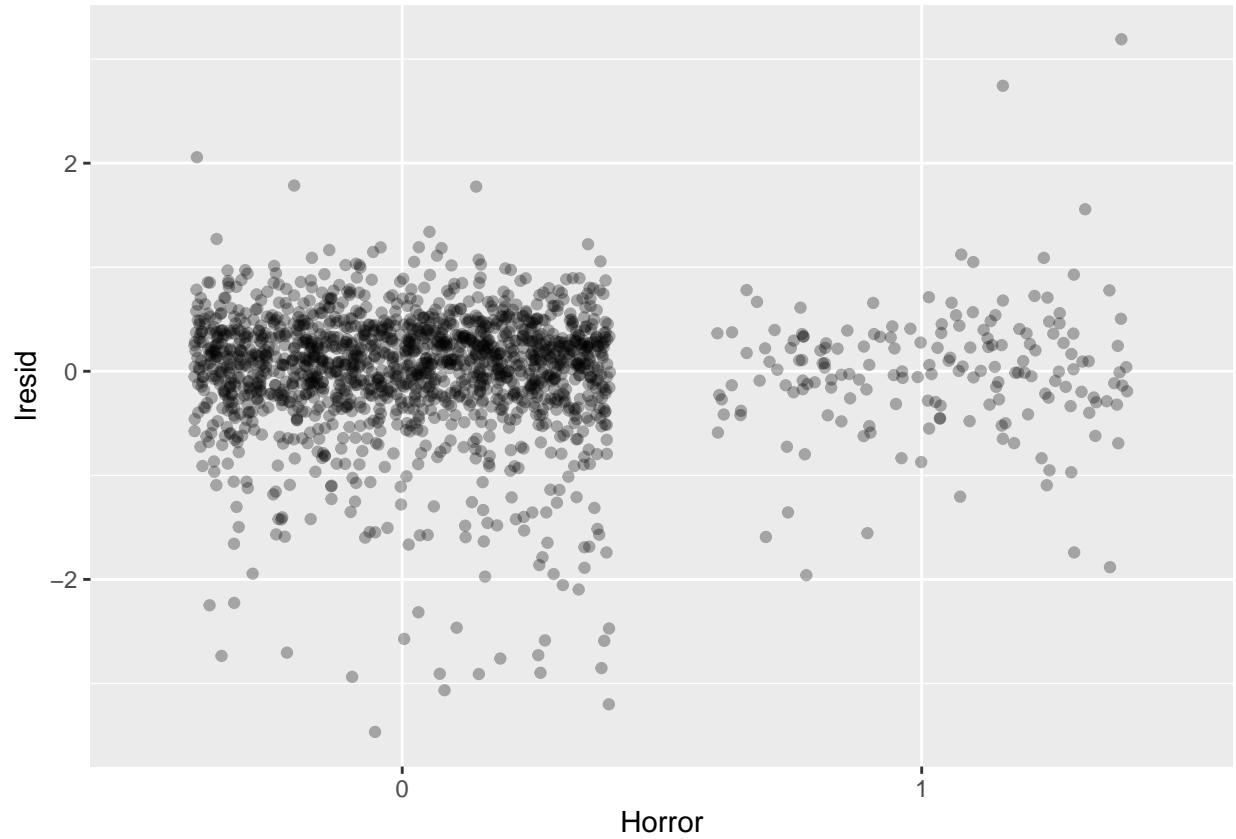


```
## Warning: Removed 132 rows containing missing values (geom_point).
```

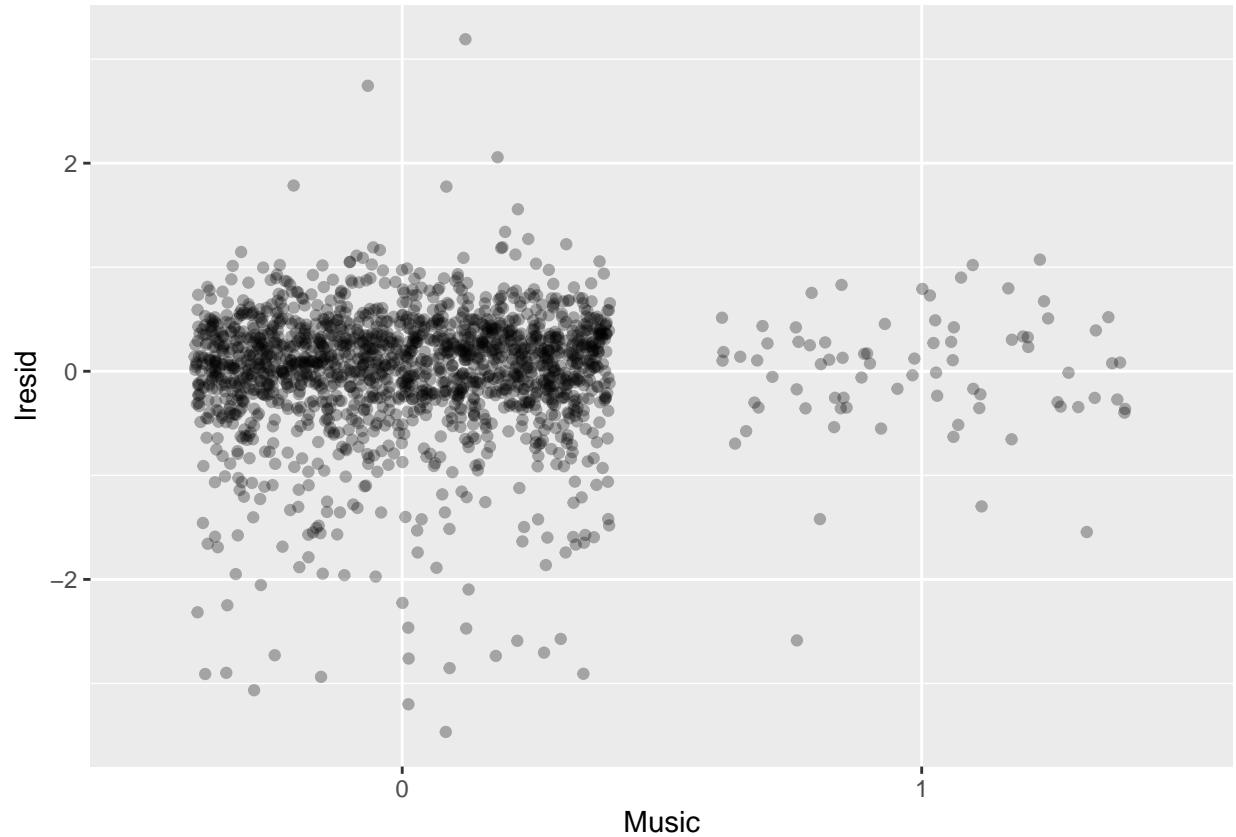




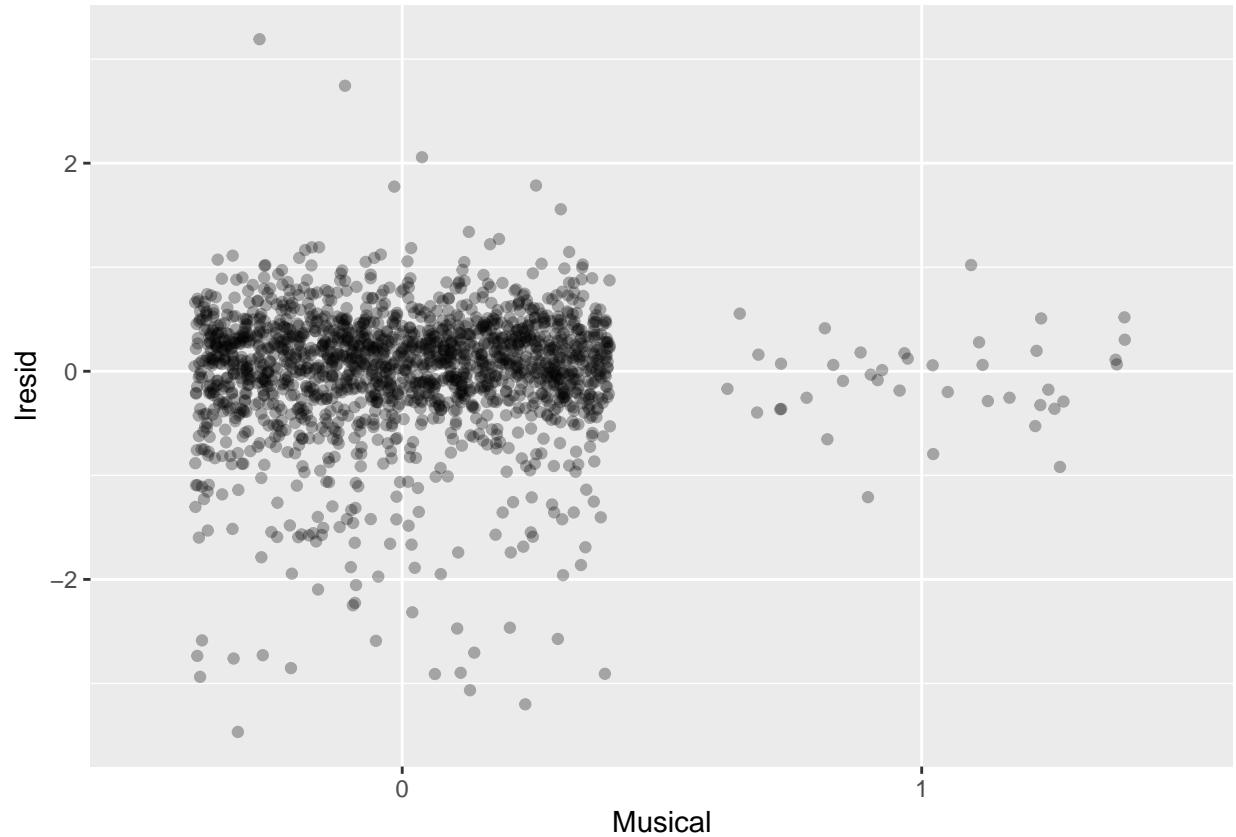
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



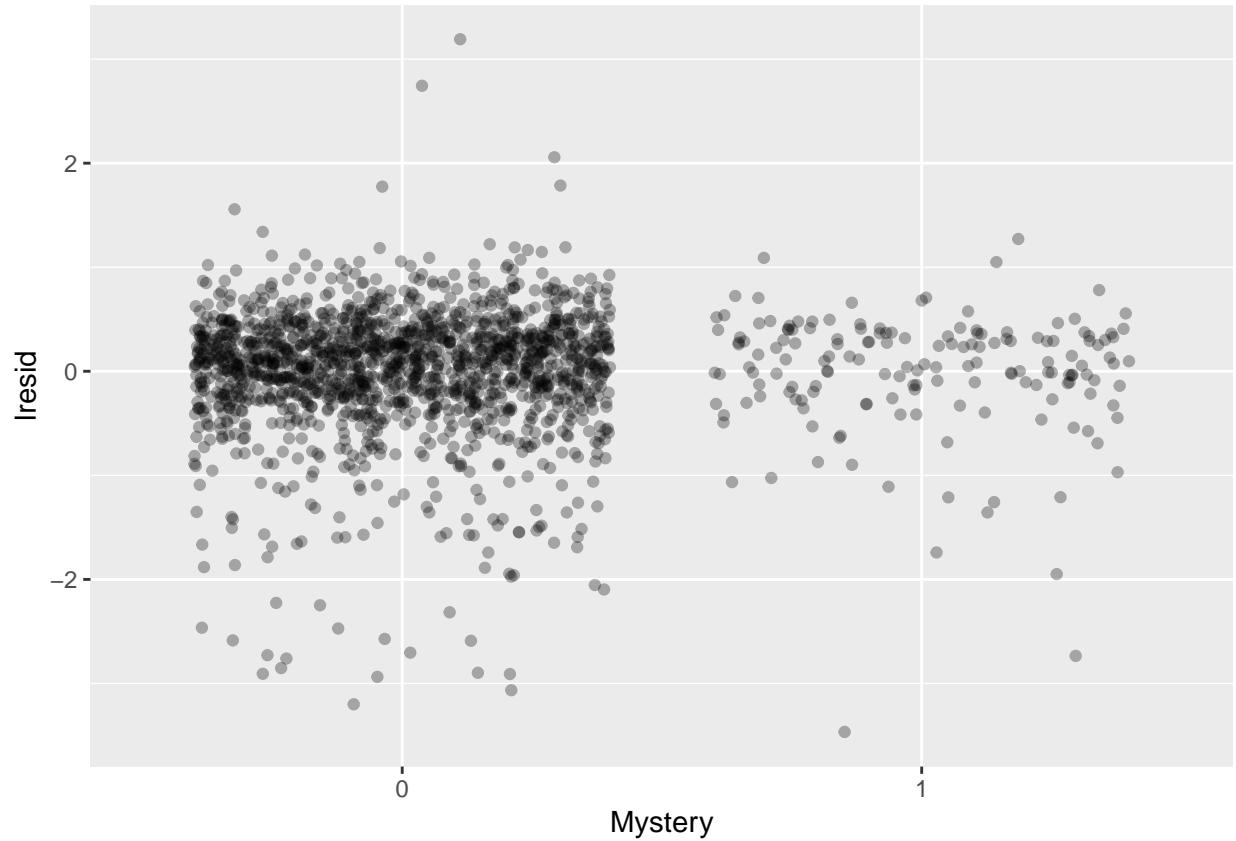
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



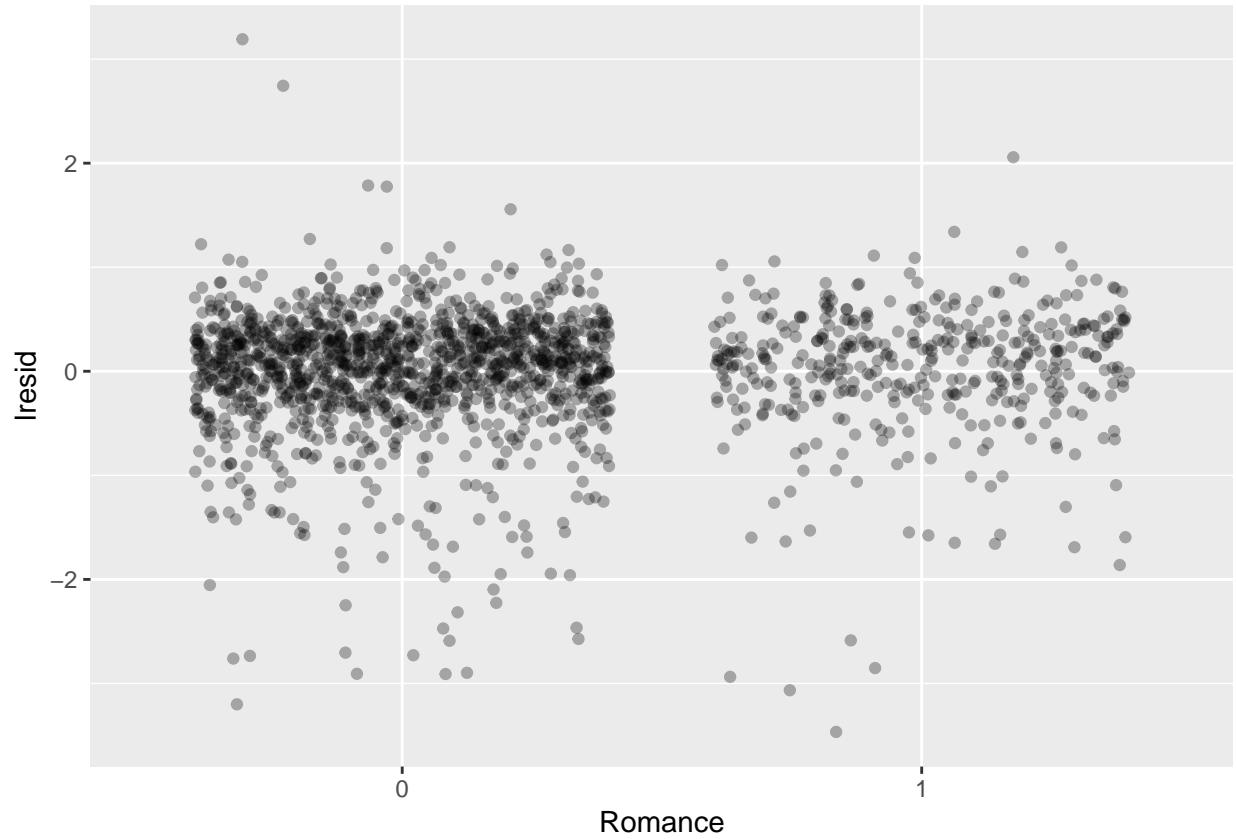
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



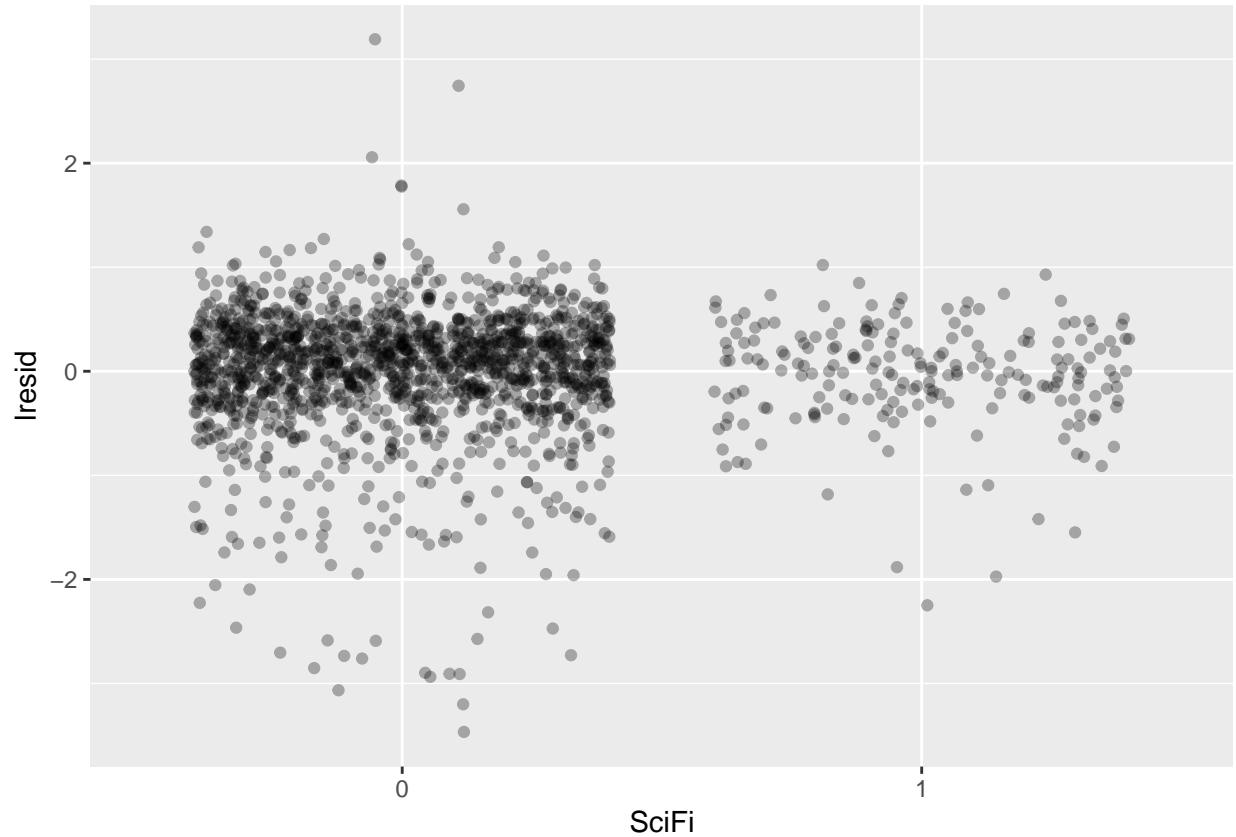
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



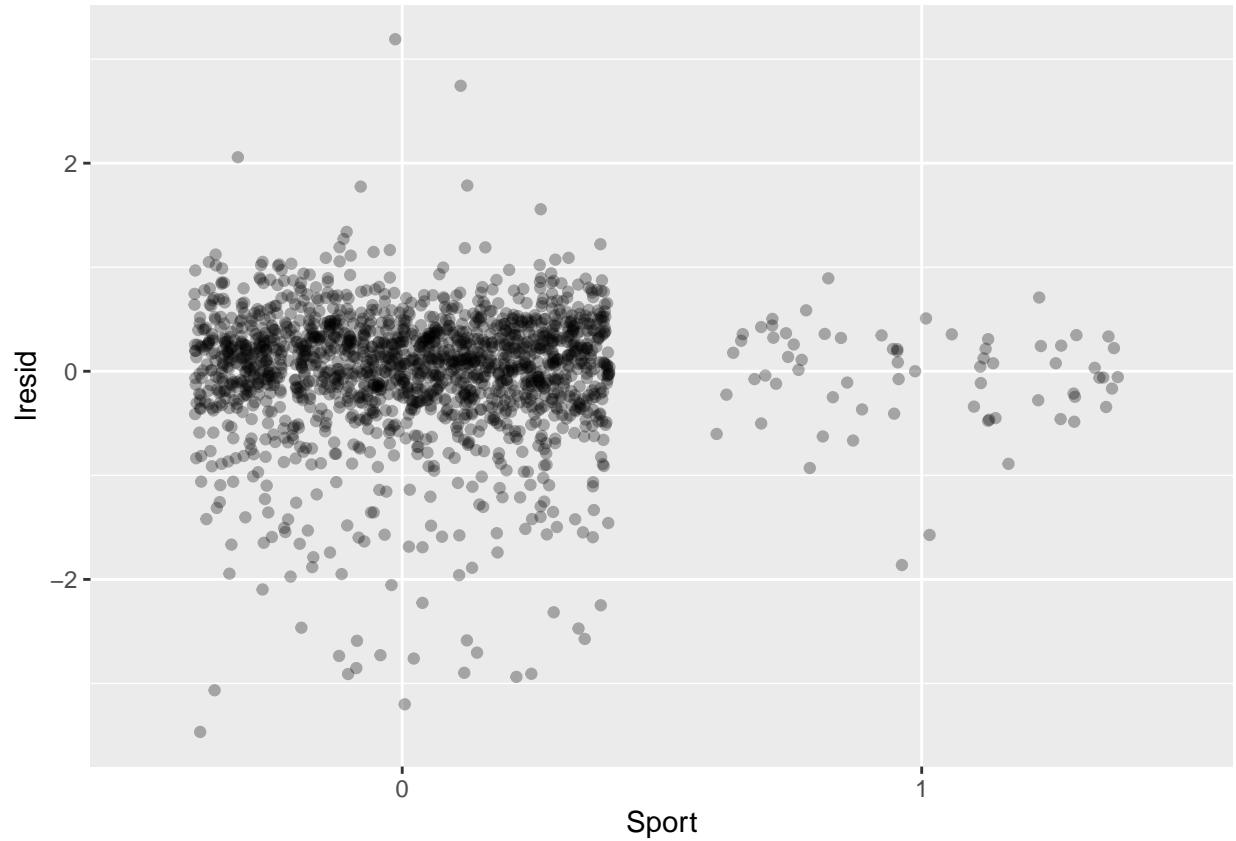
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



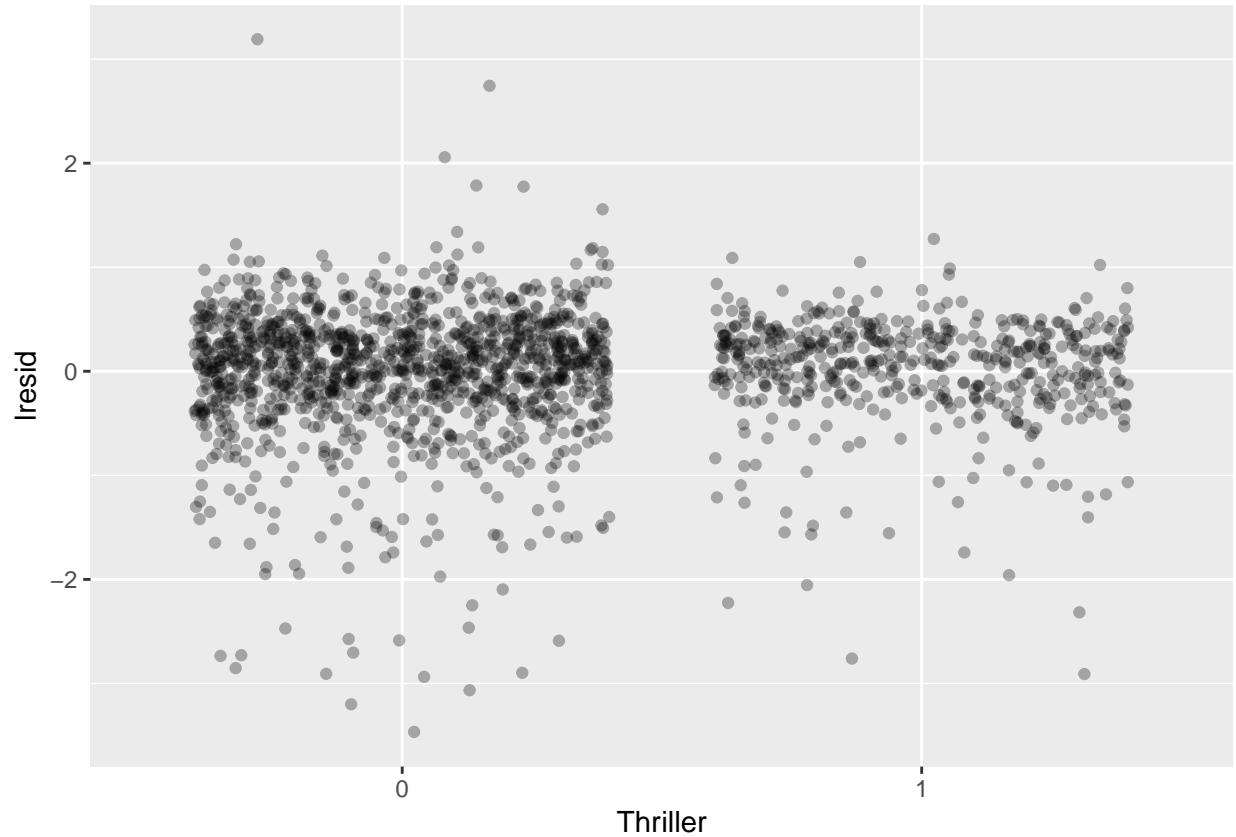
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



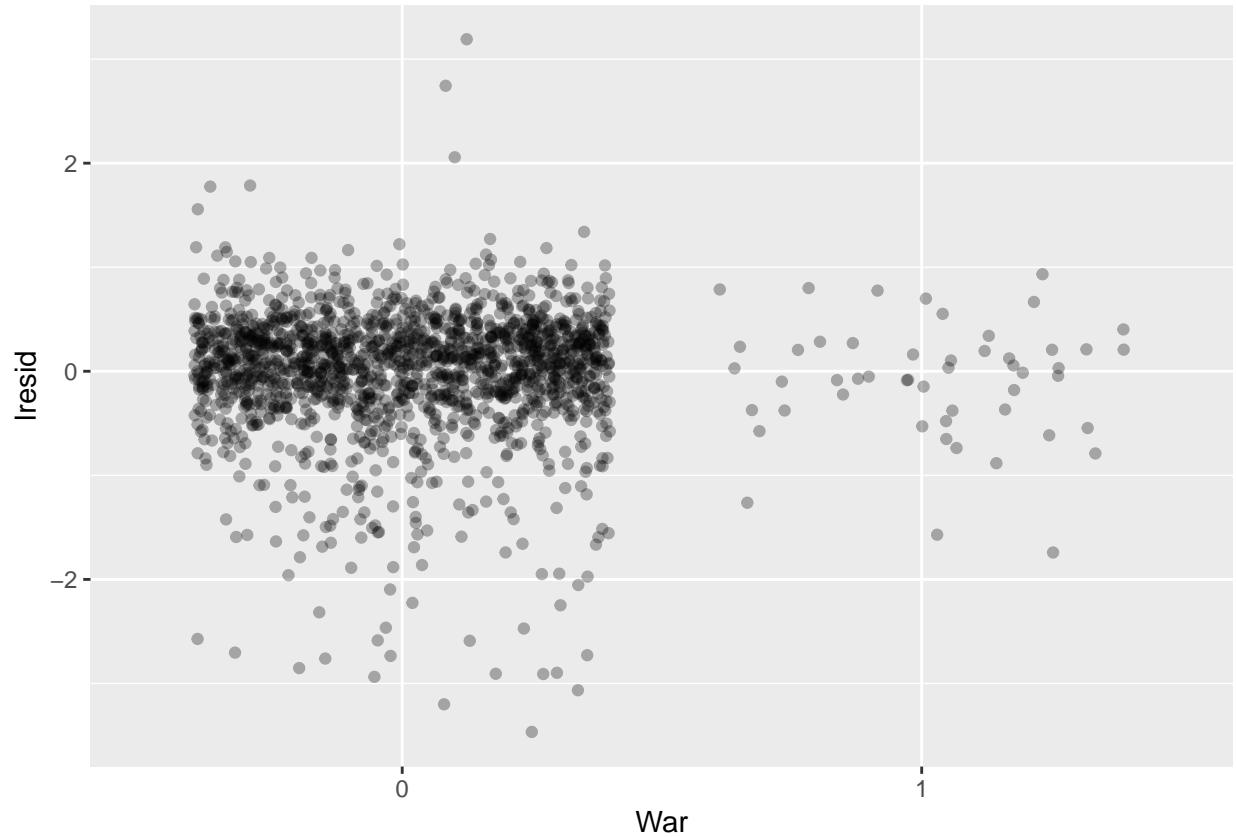
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



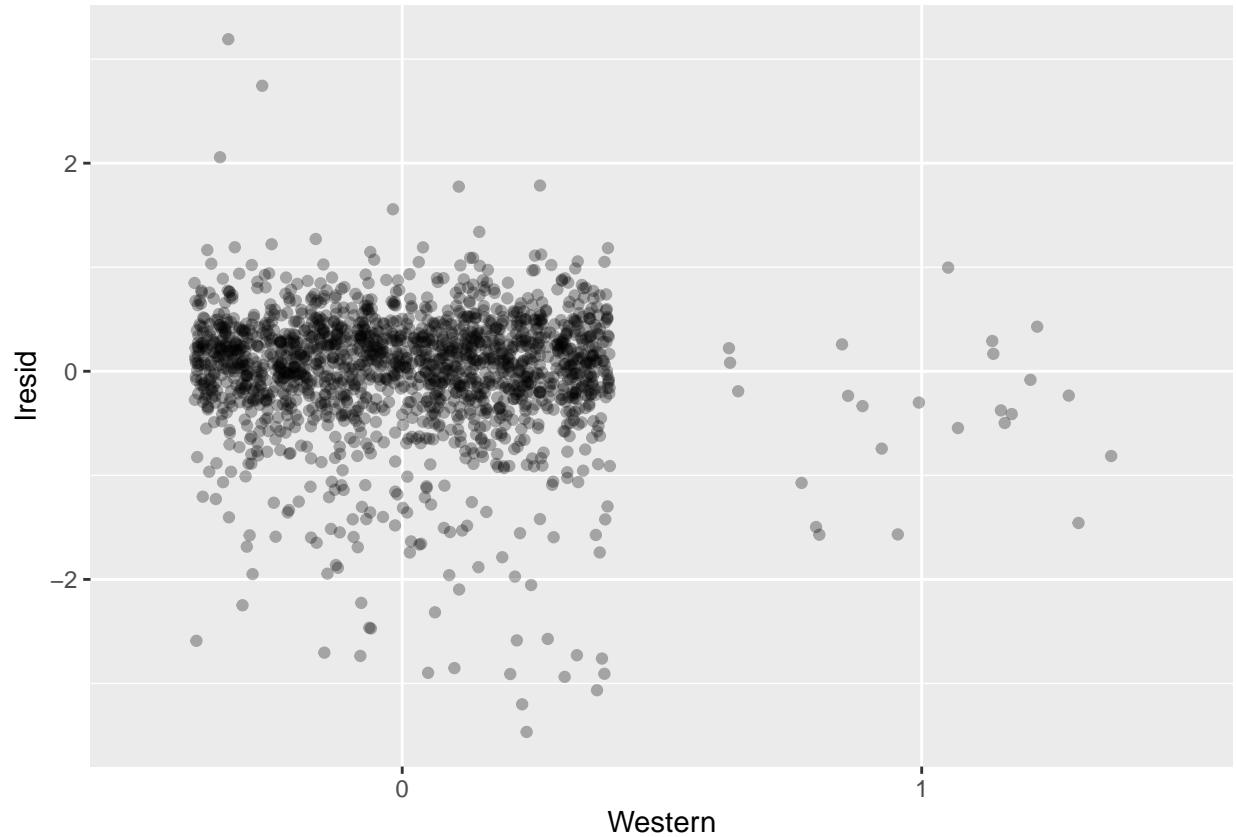
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



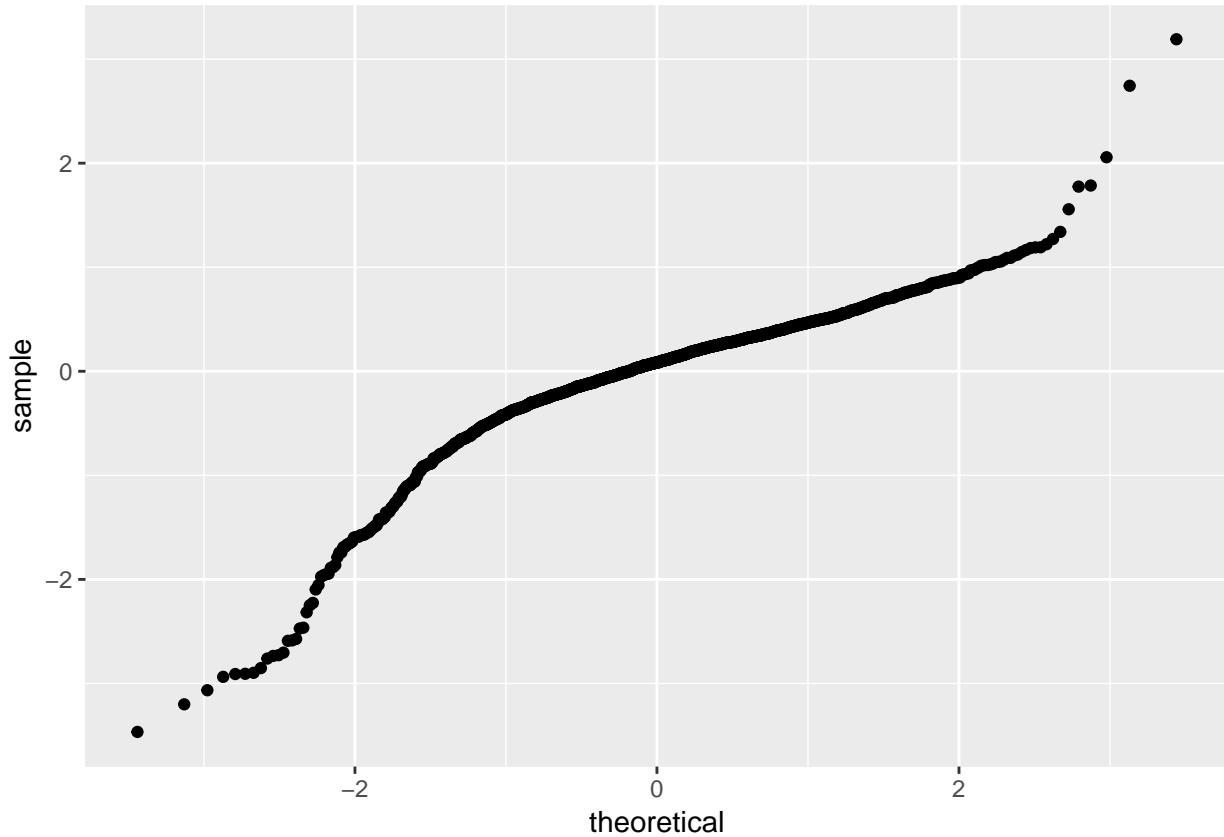
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



```
## Warning: Removed 132 rows containing missing values (geom_point).
```

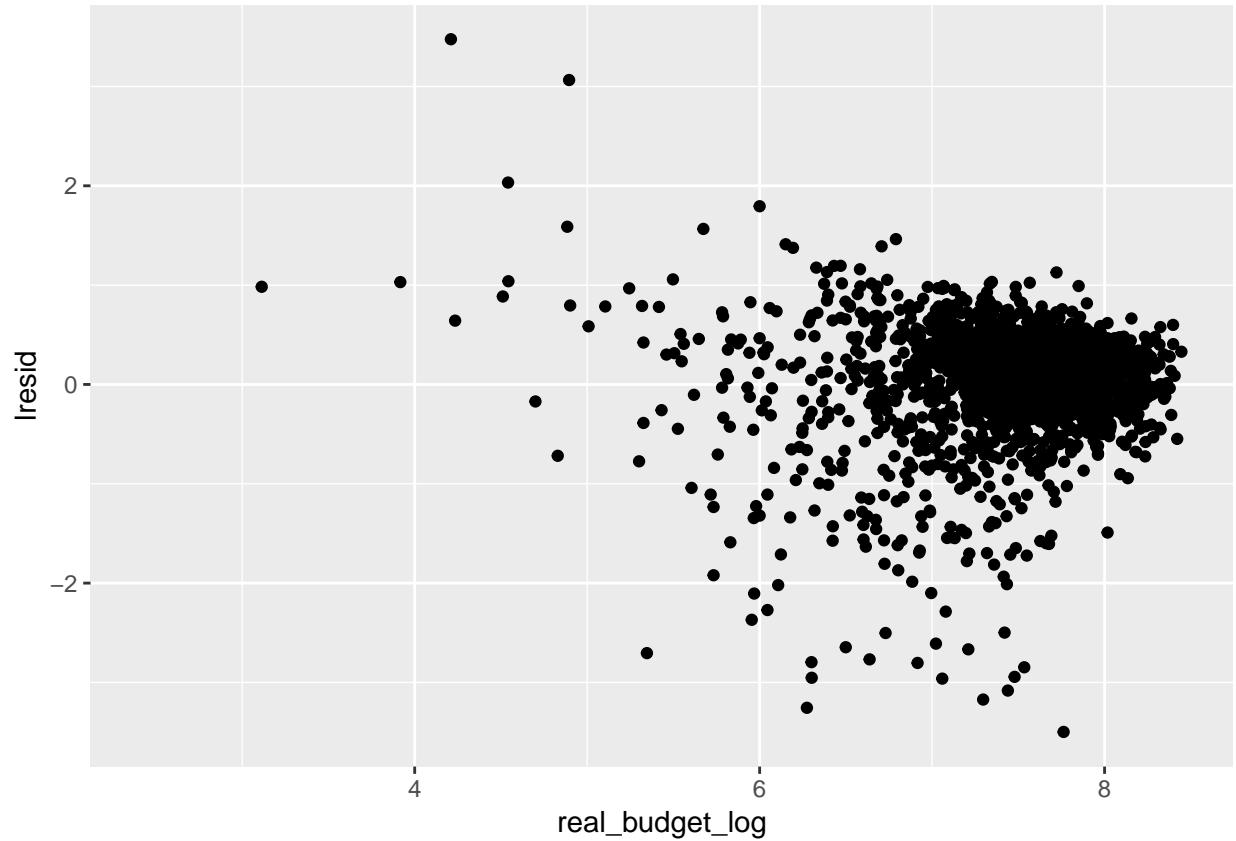


```
## Warning: Removed 132 rows containing non-finite values (stat_qq).
```

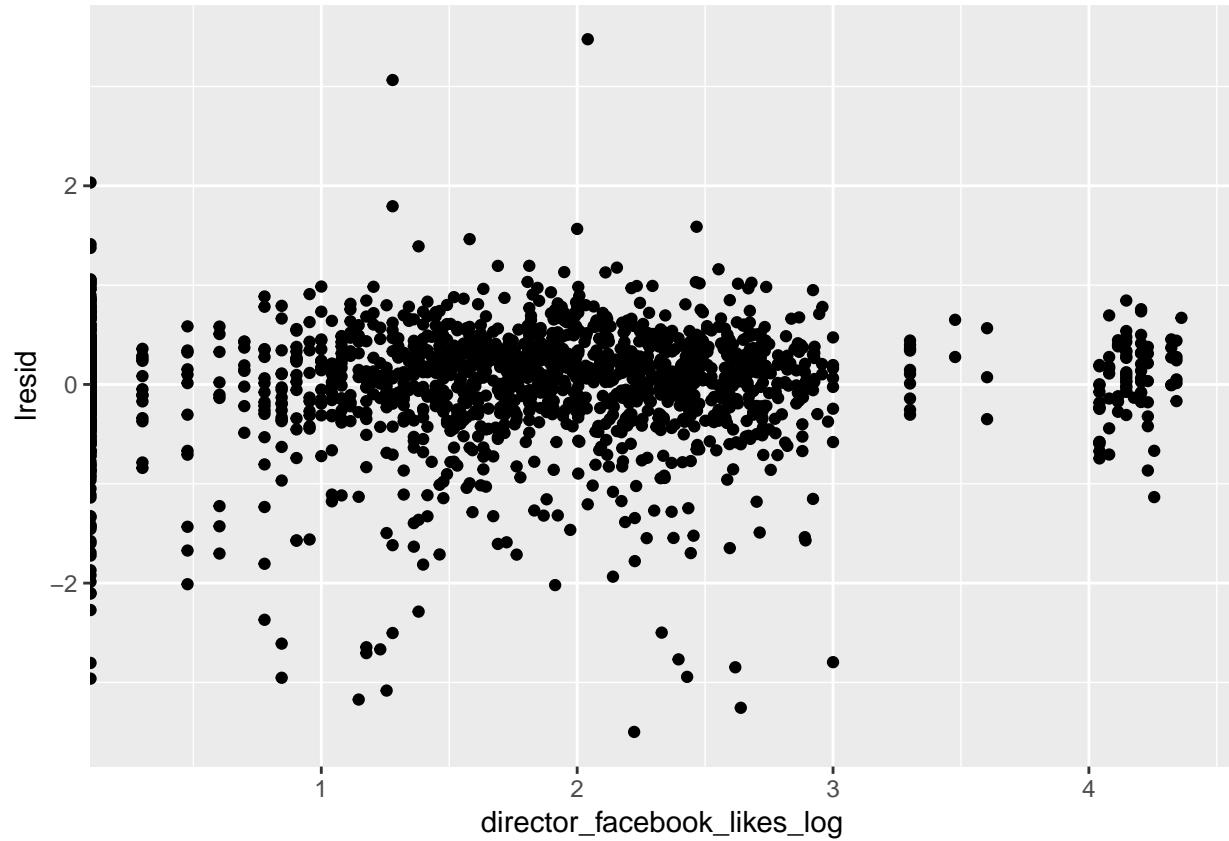


```
gr_resid(mod_simple_plus2)
```

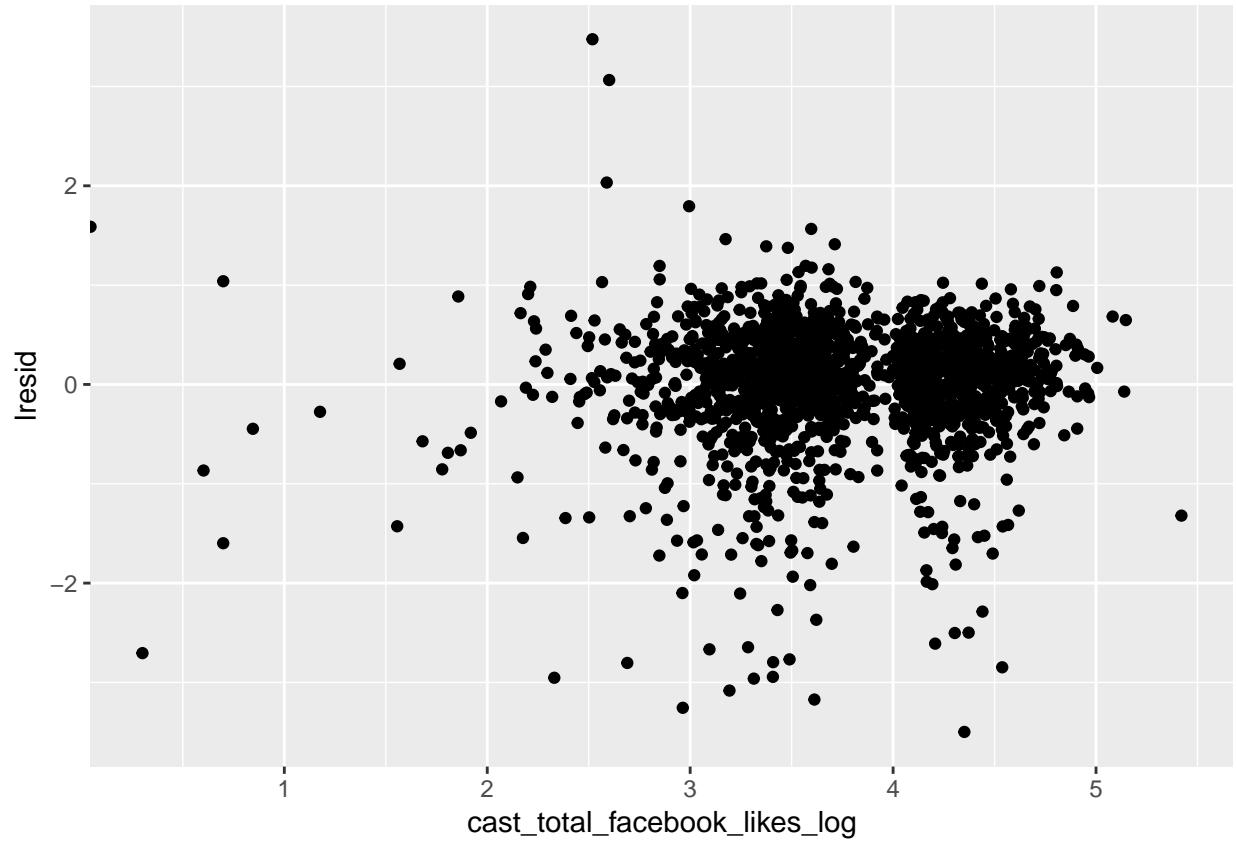
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



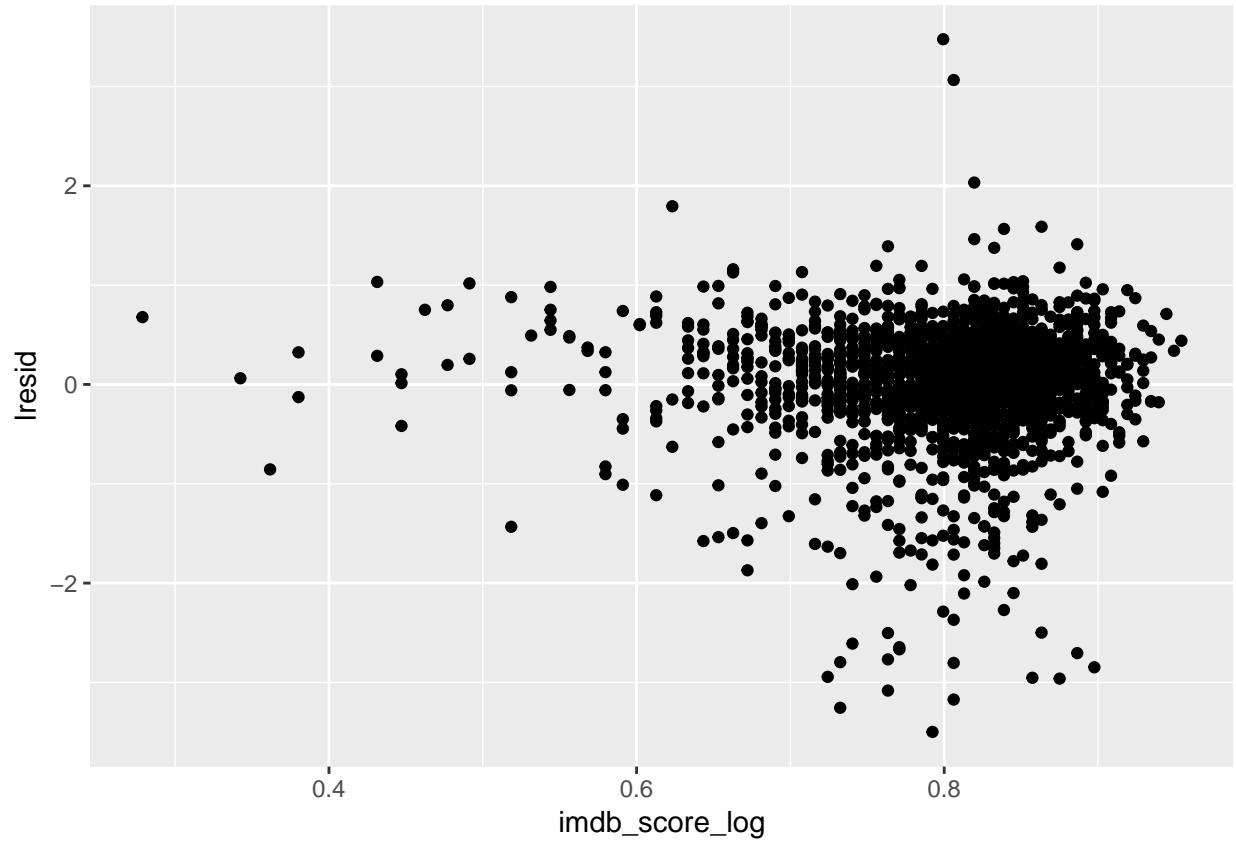
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



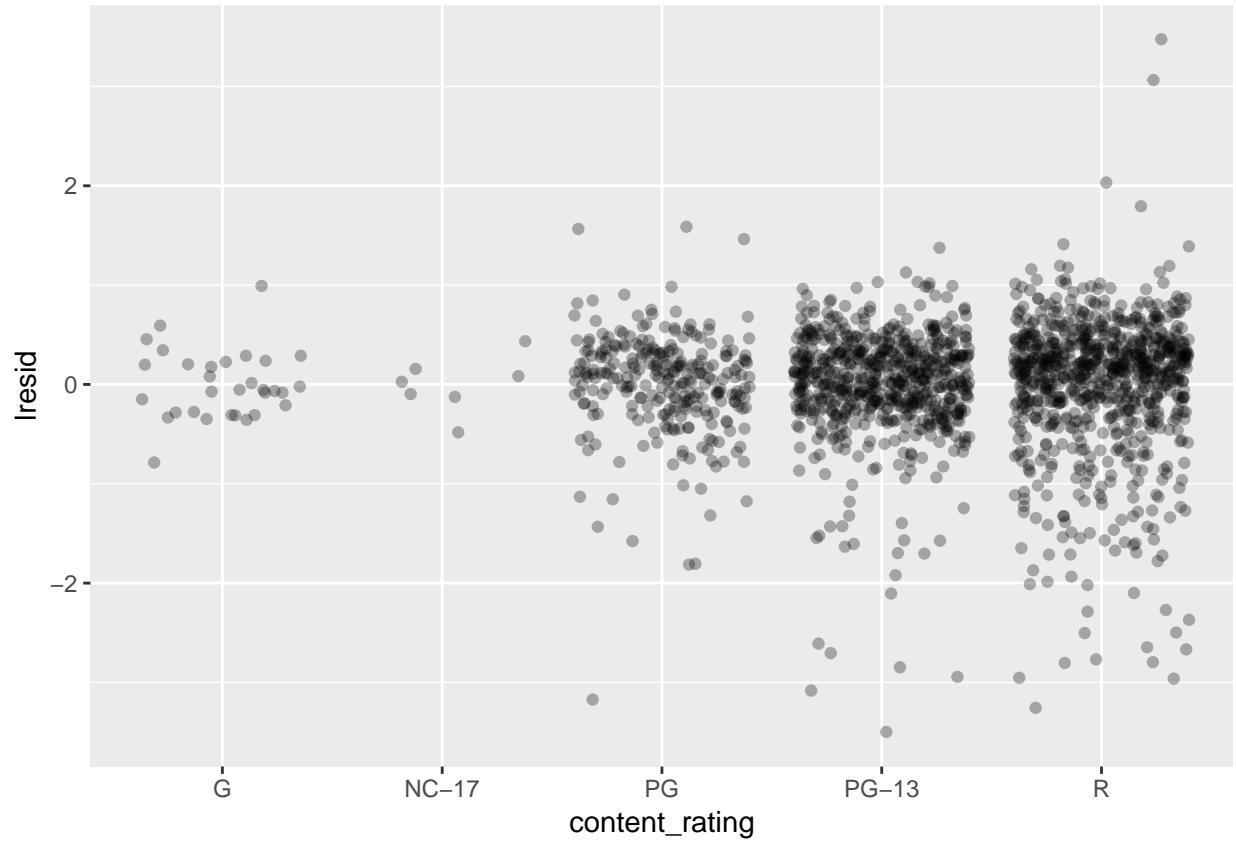
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



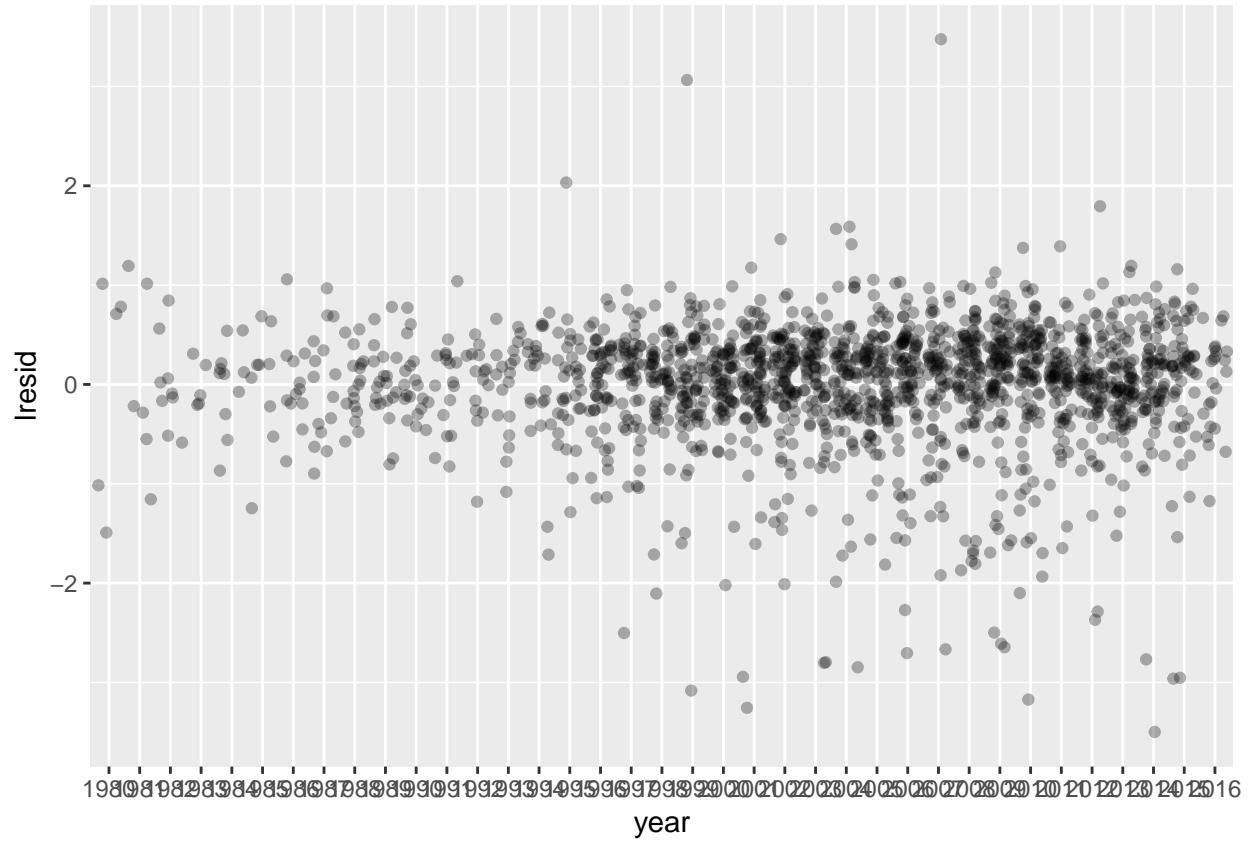
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



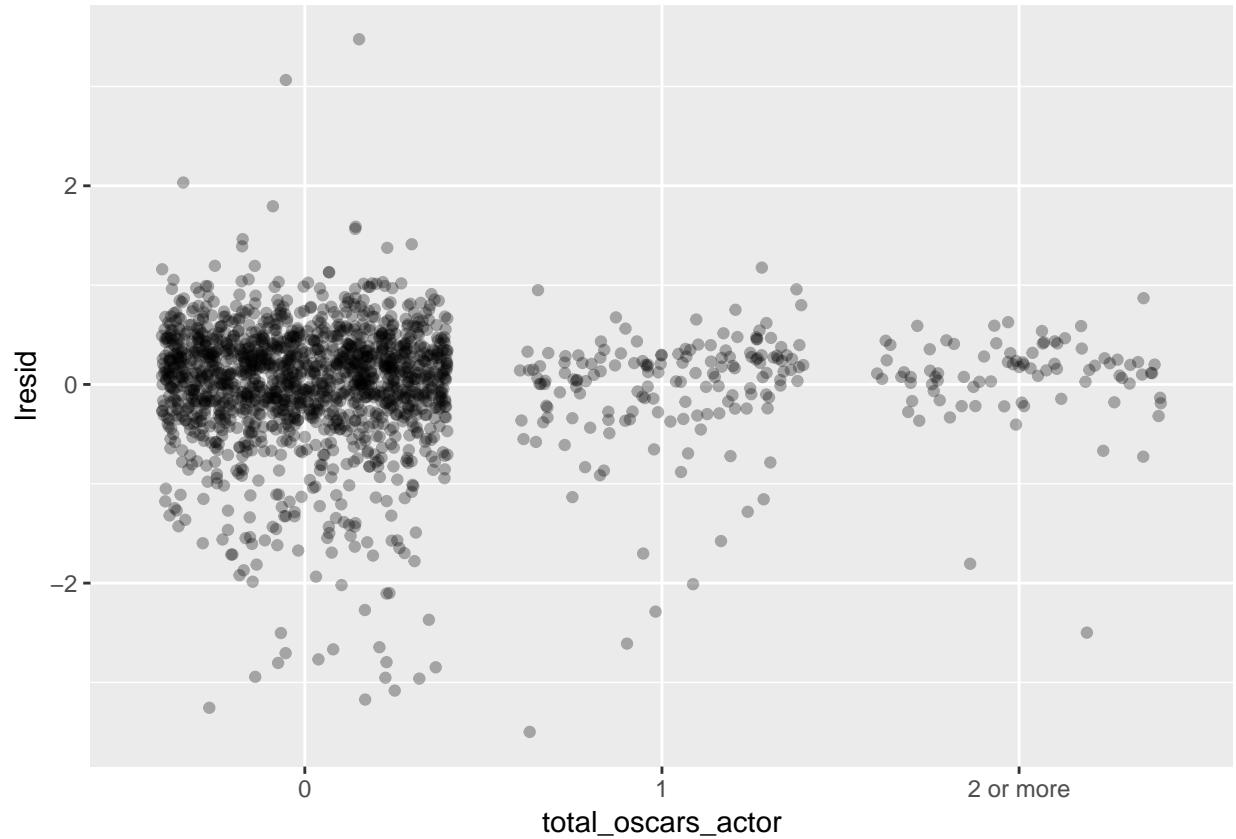
```
## Warning: Removed 94 rows containing missing values (geom_point).
```



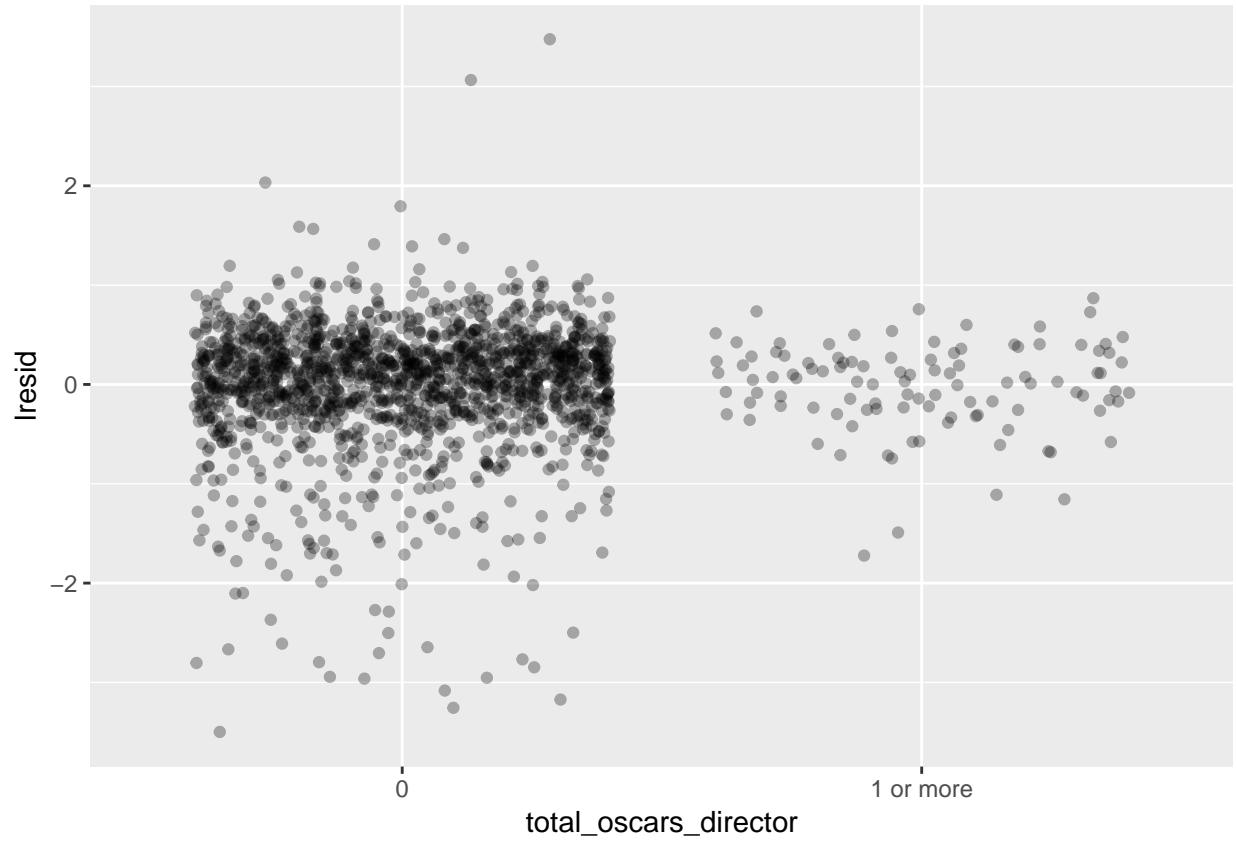
```
## Warning: Removed 132 rows containing missing values (geom_point).
```

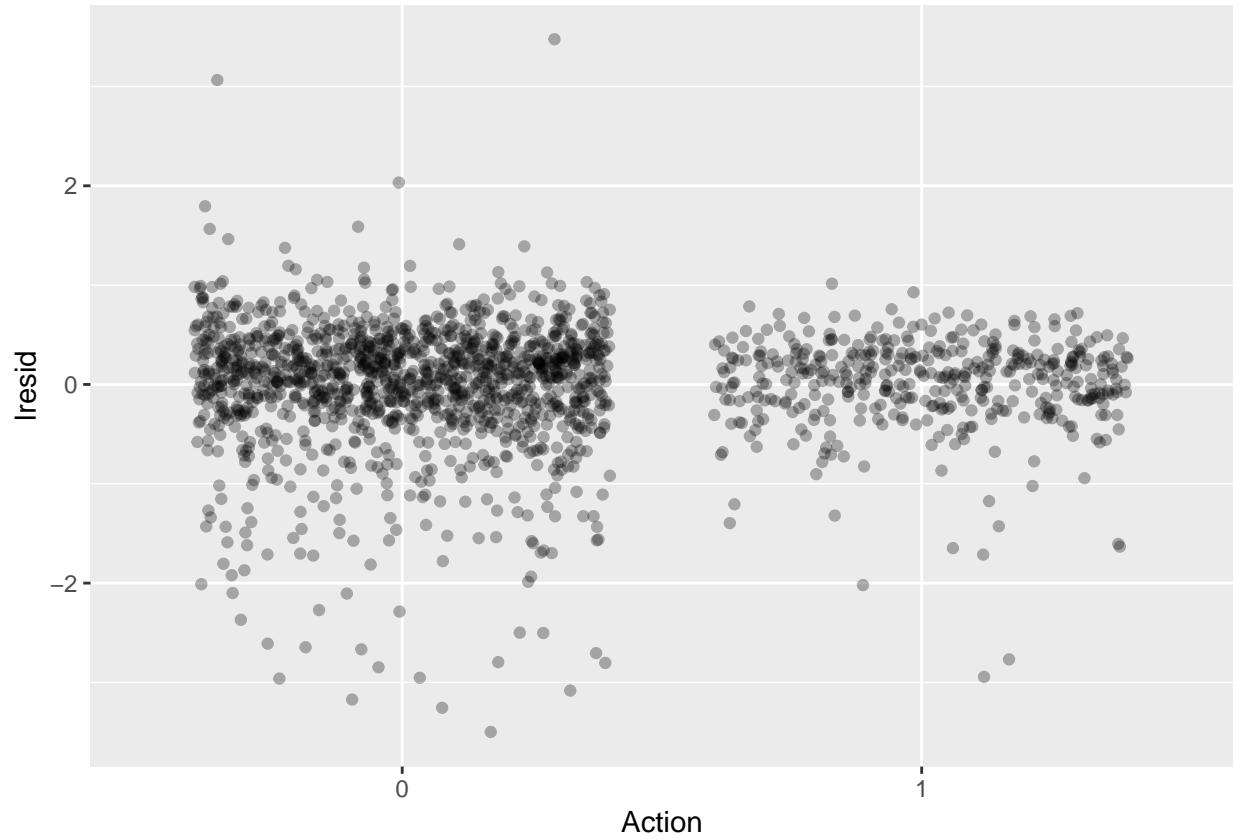


```
## Warning: Removed 132 rows containing missing values (geom_point).
```

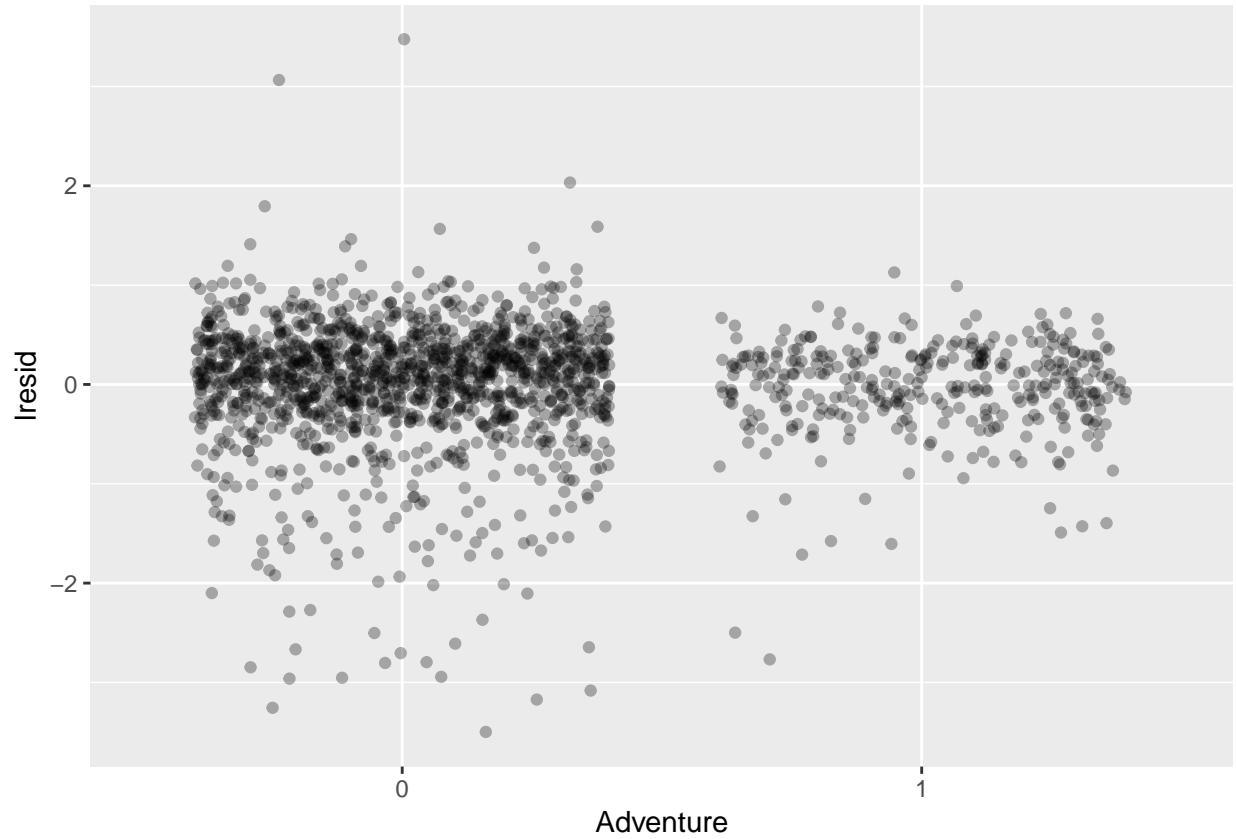


```
## Warning: Removed 132 rows containing missing values (geom_point).
```

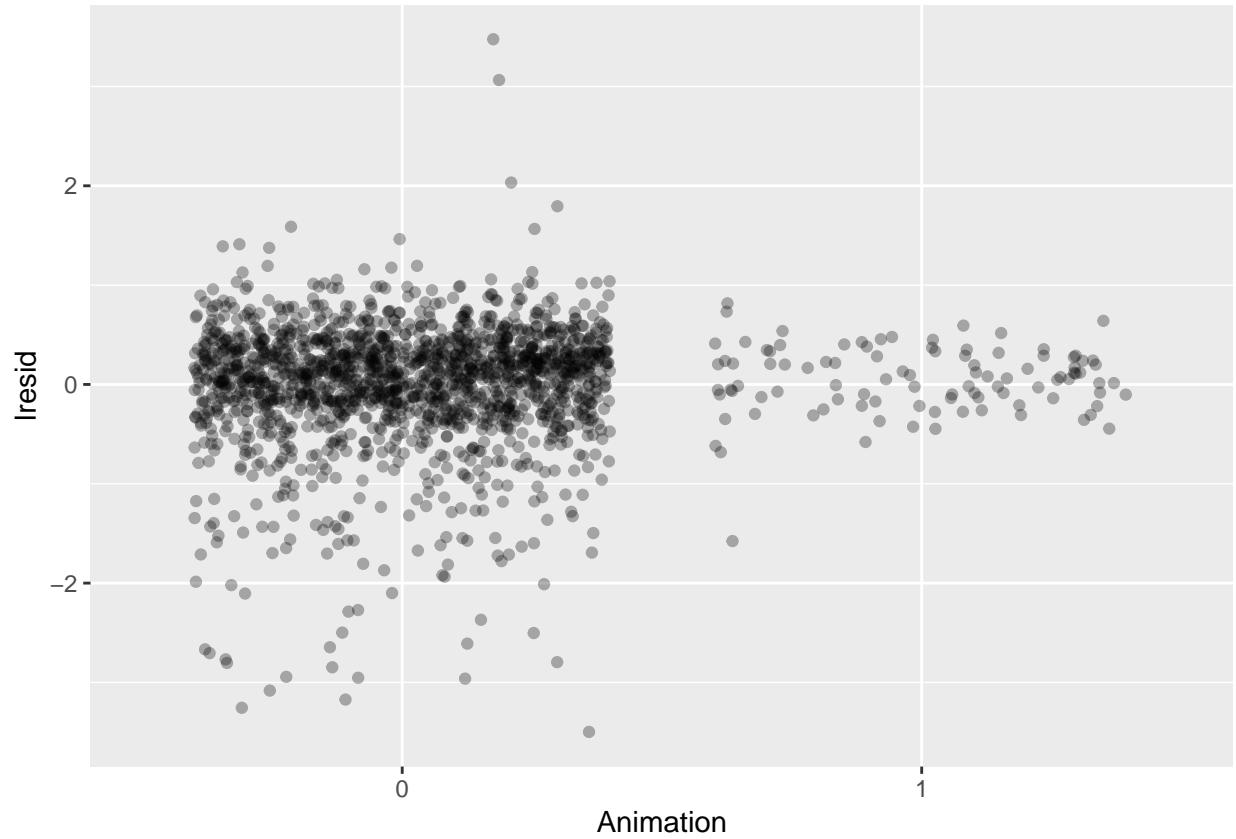




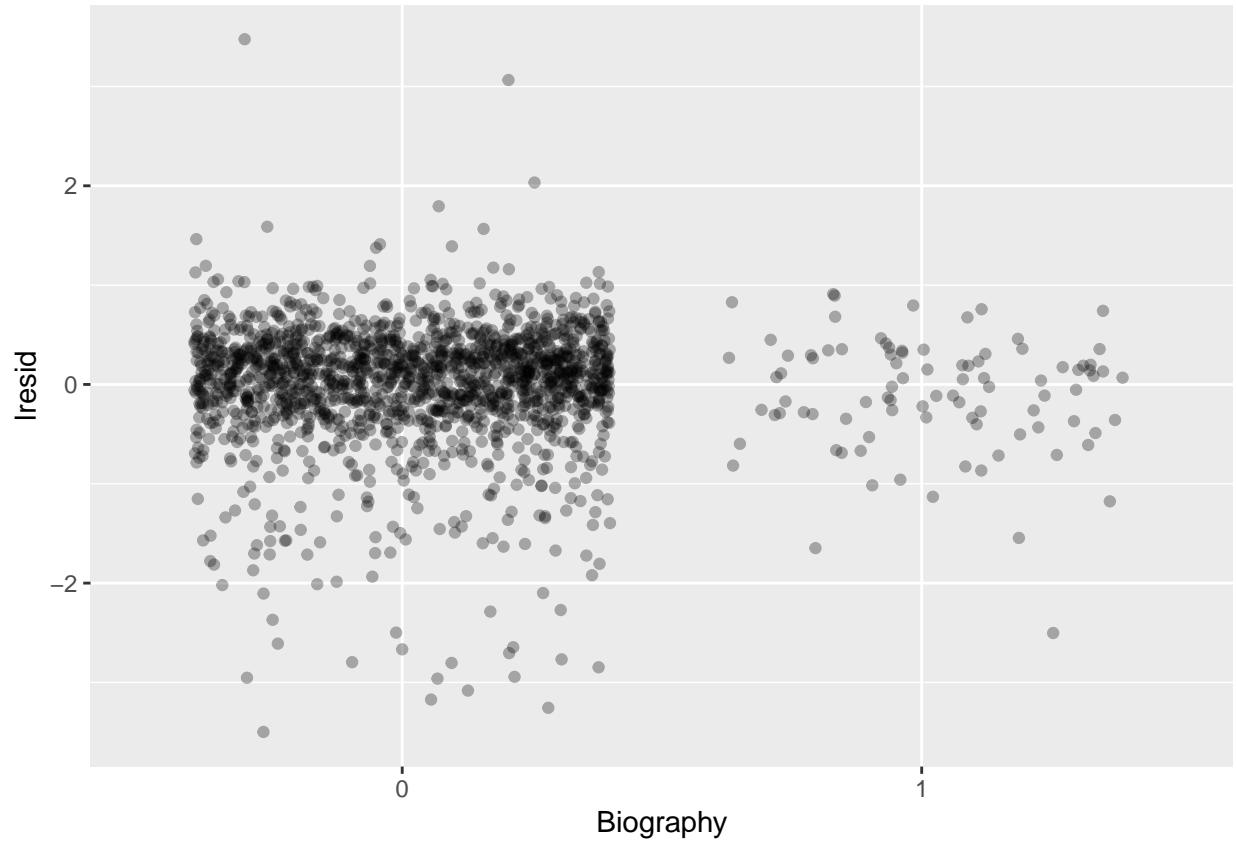
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



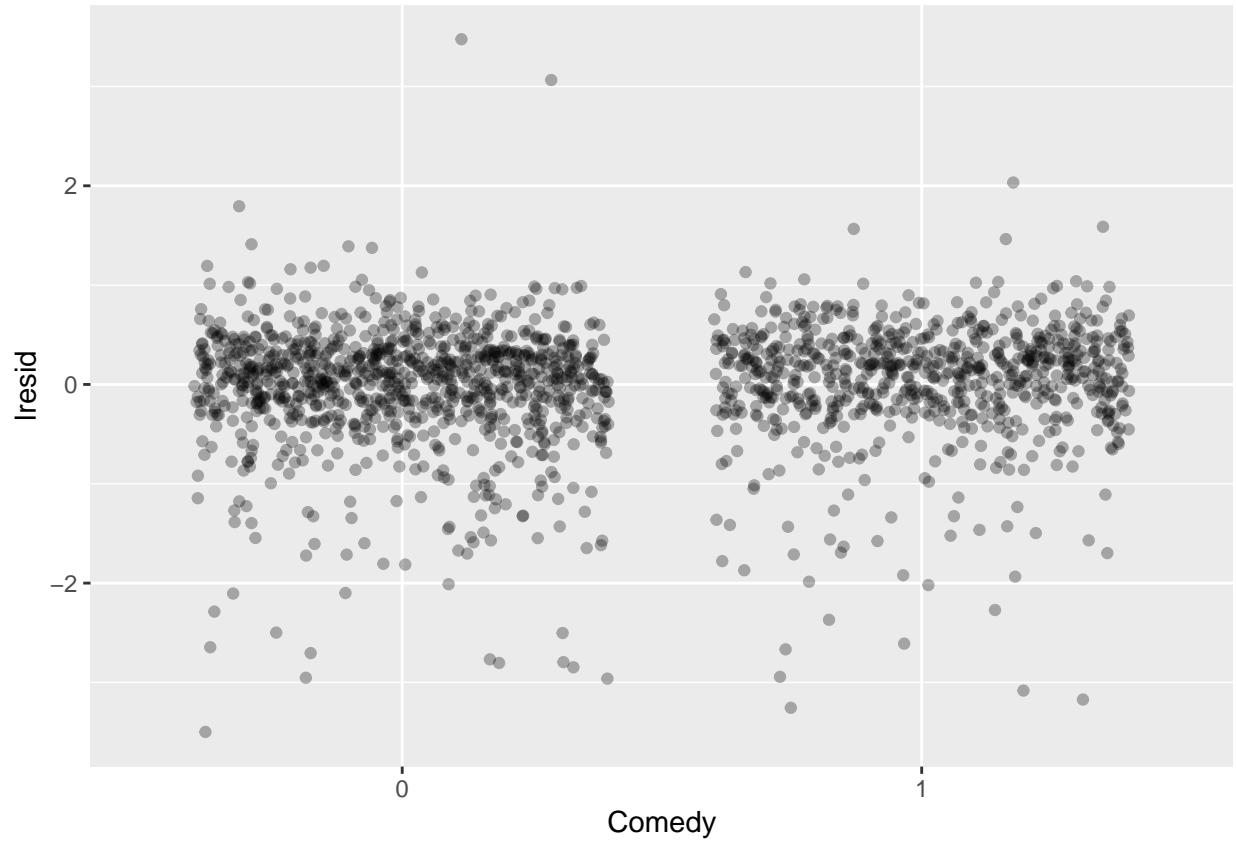
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



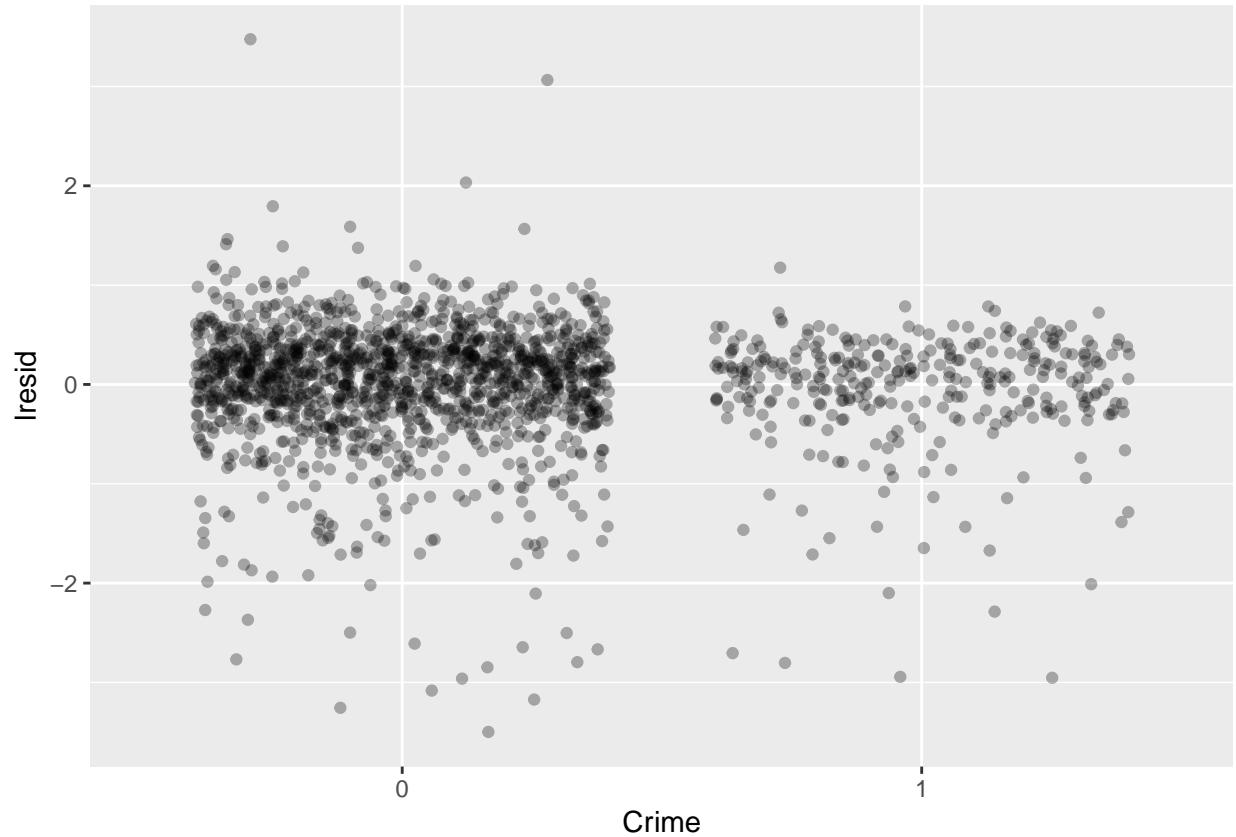
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



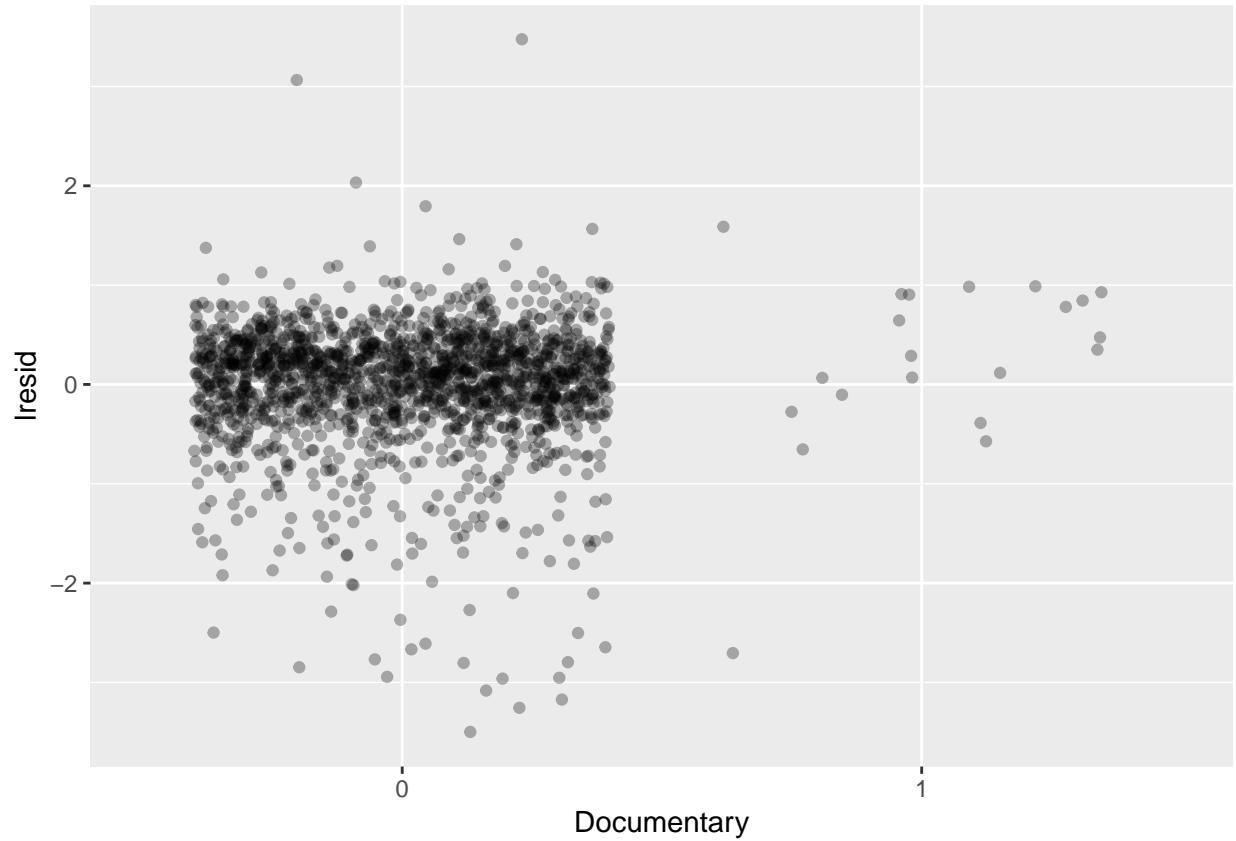
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



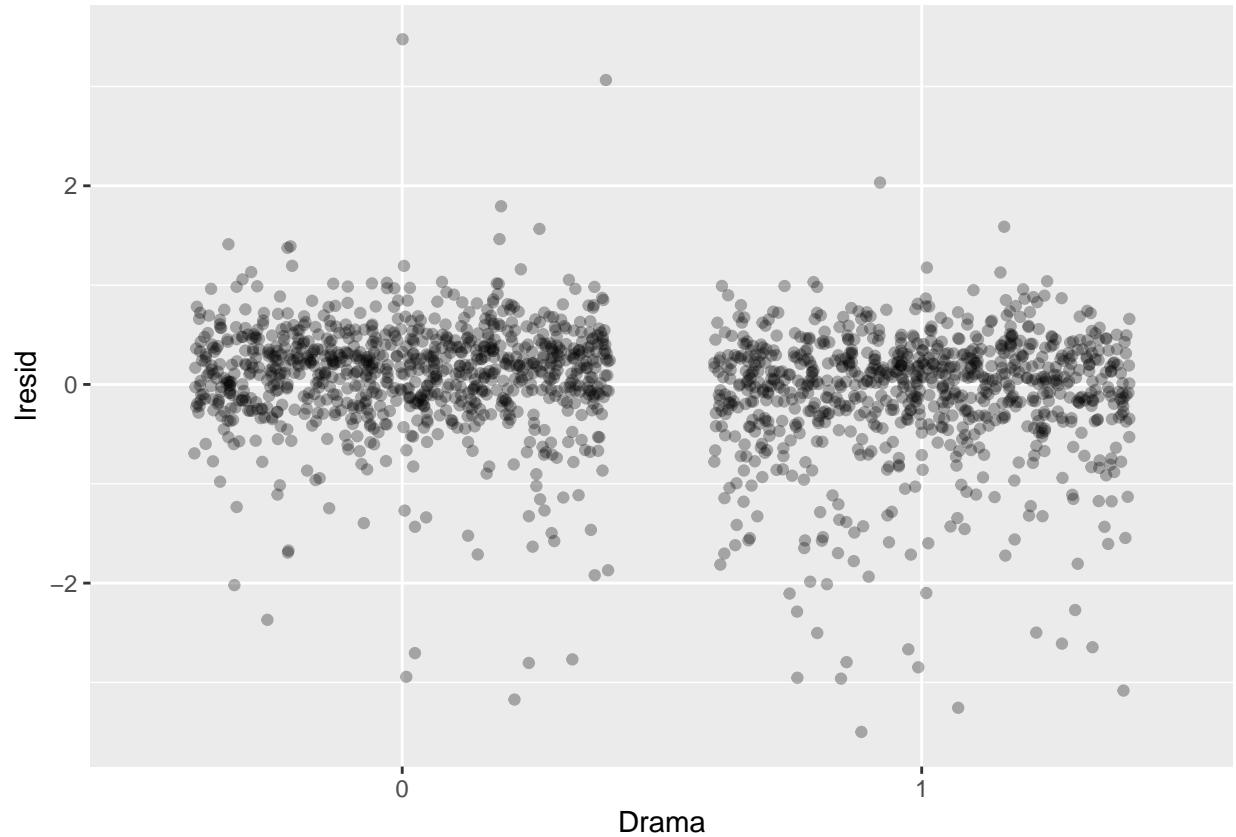
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



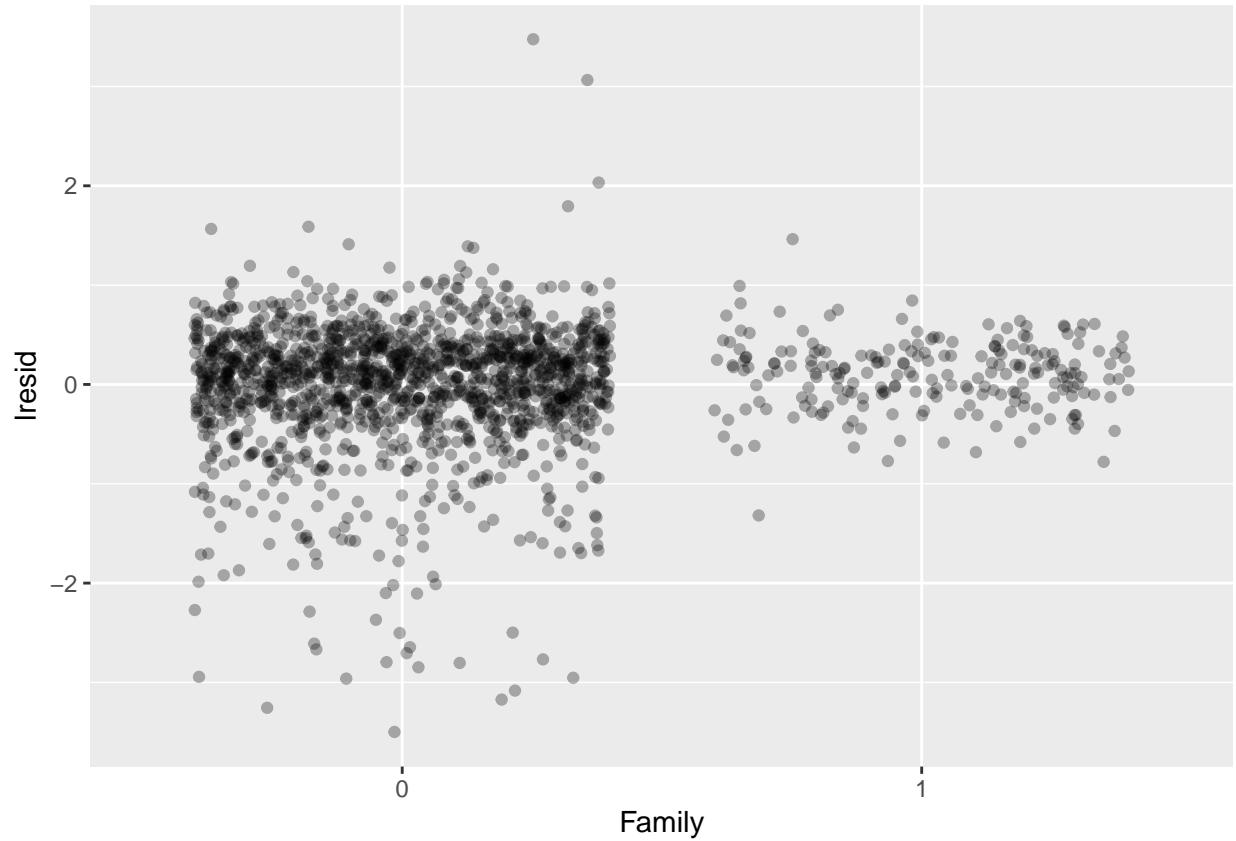
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



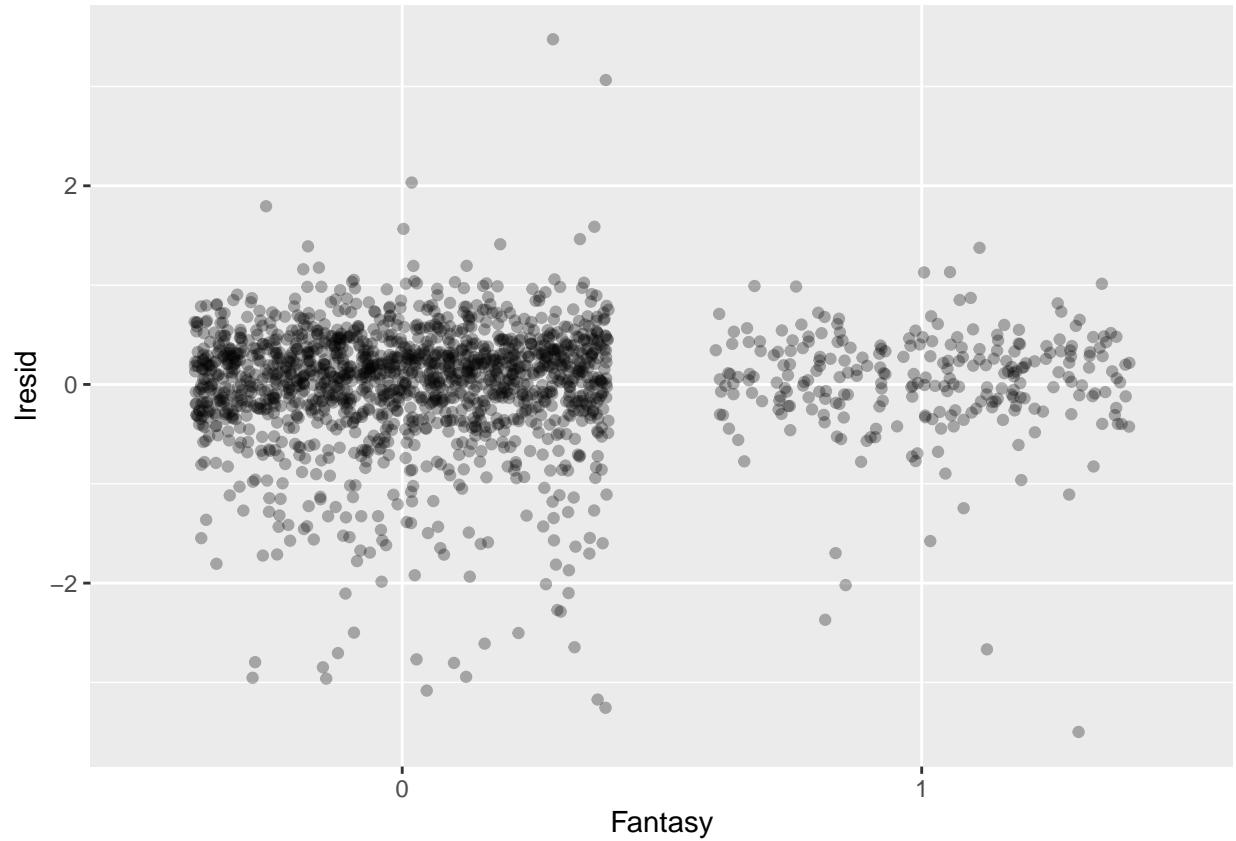
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



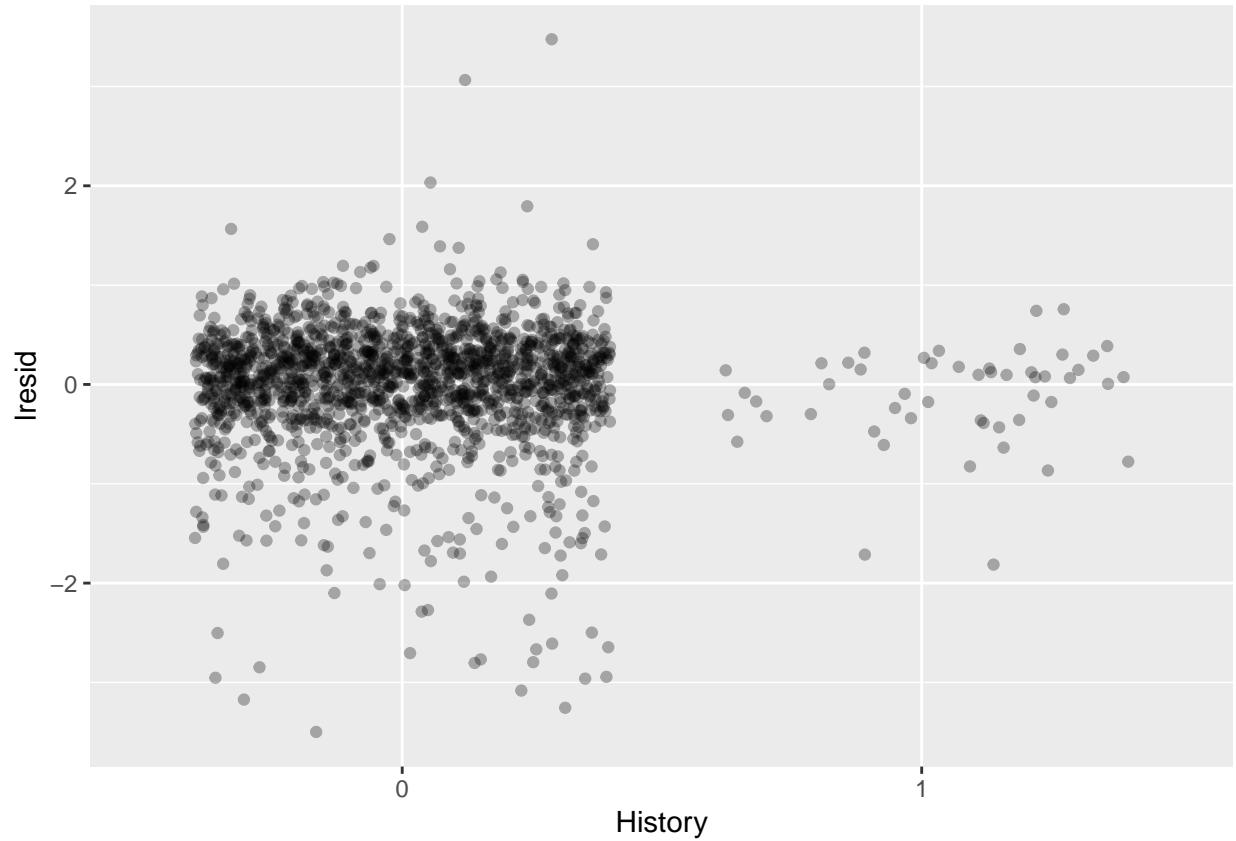
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



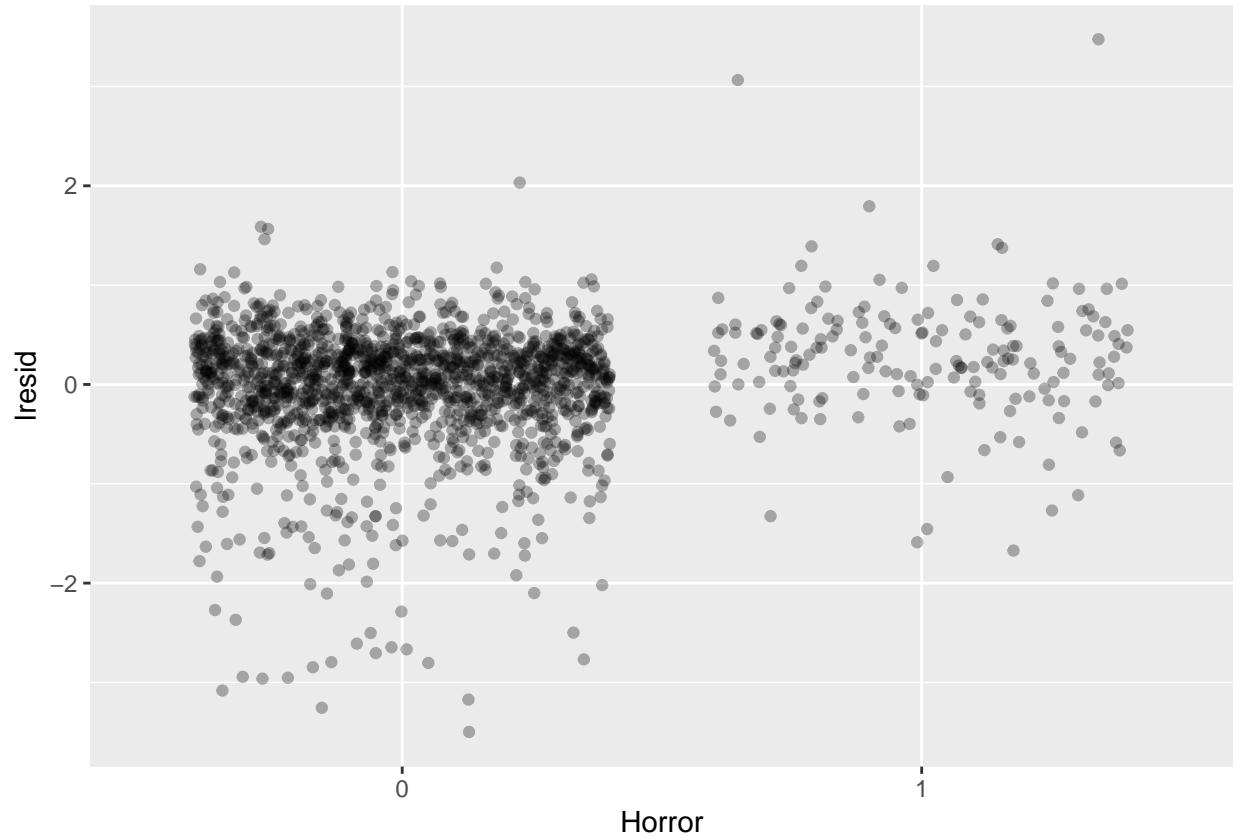
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



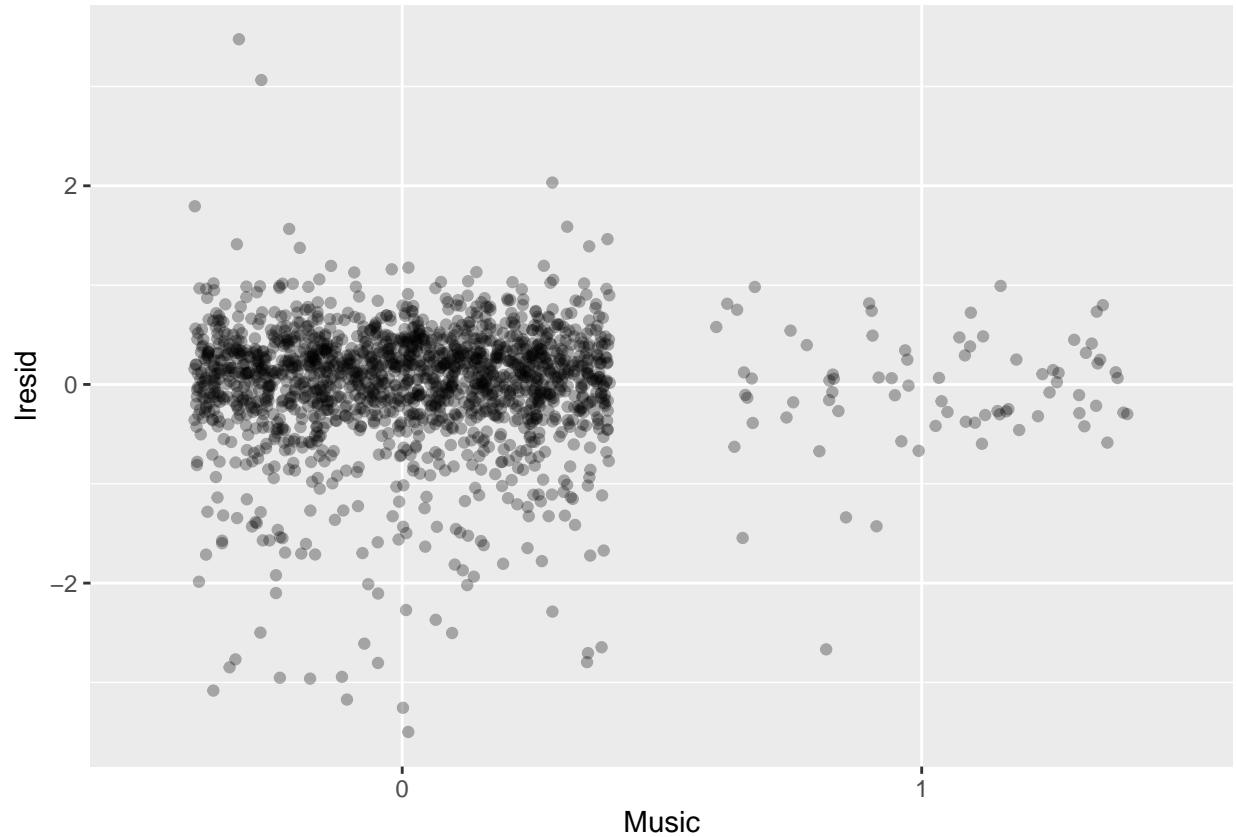
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



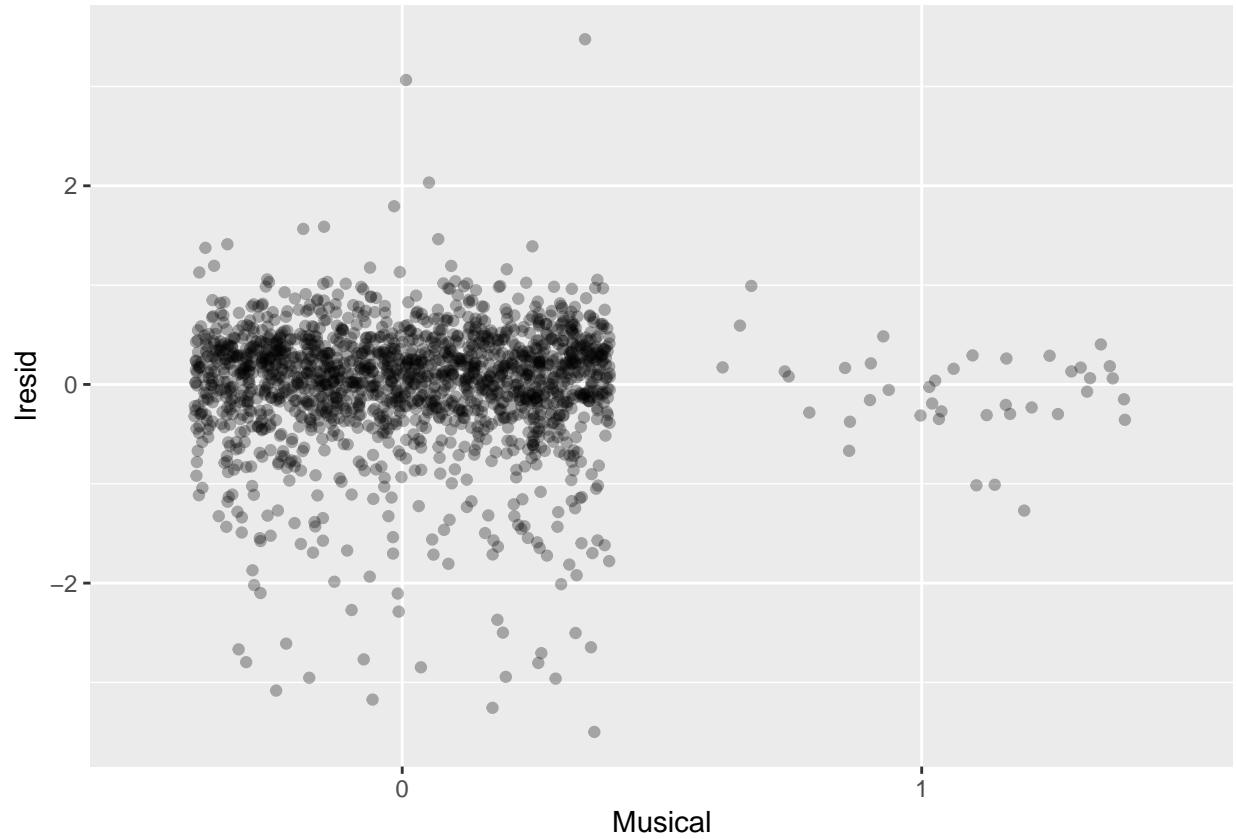
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



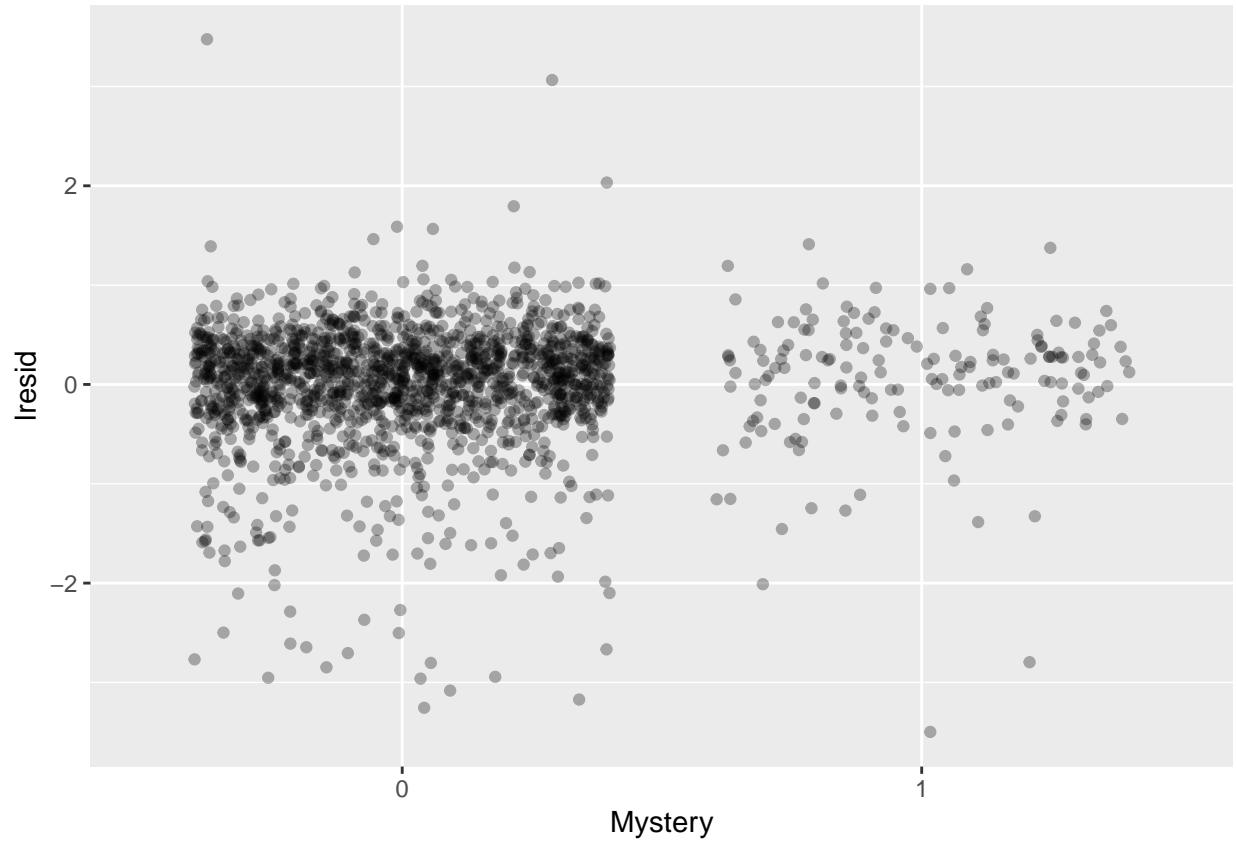
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



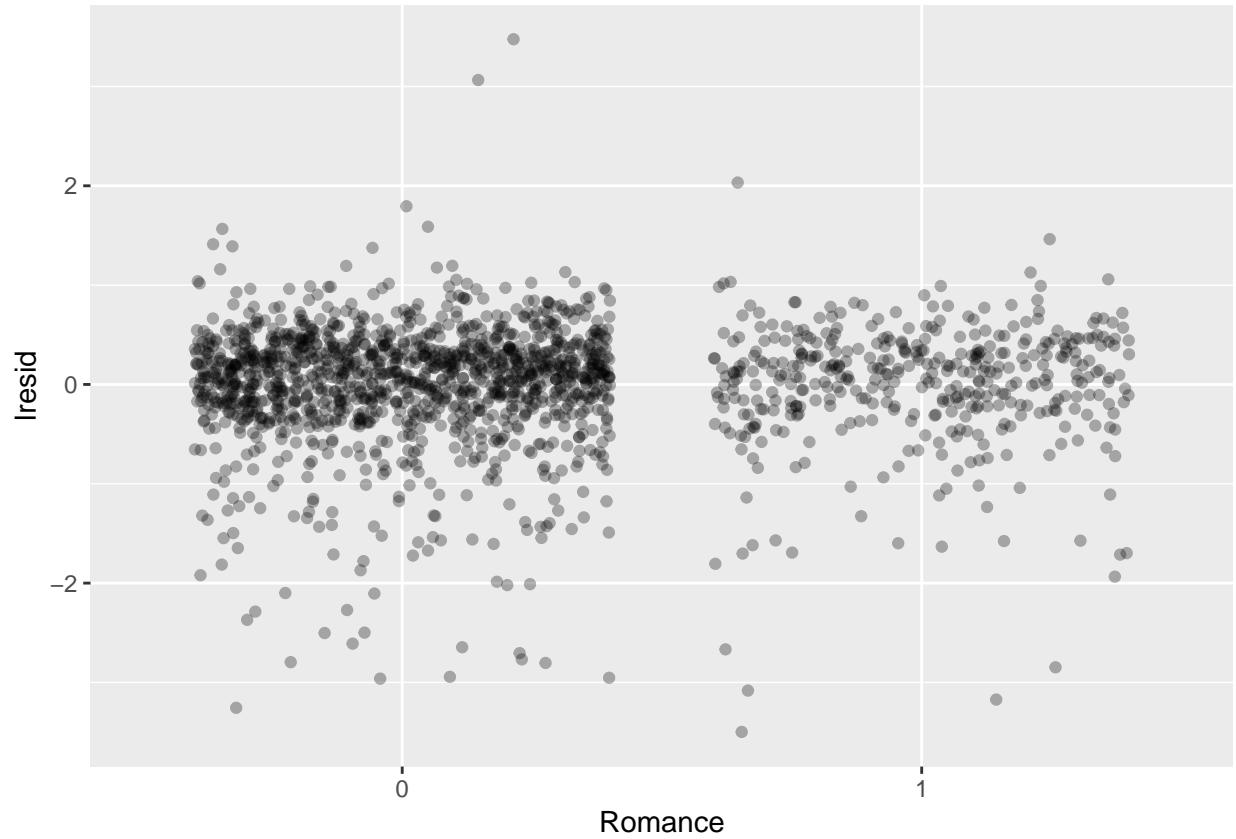
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



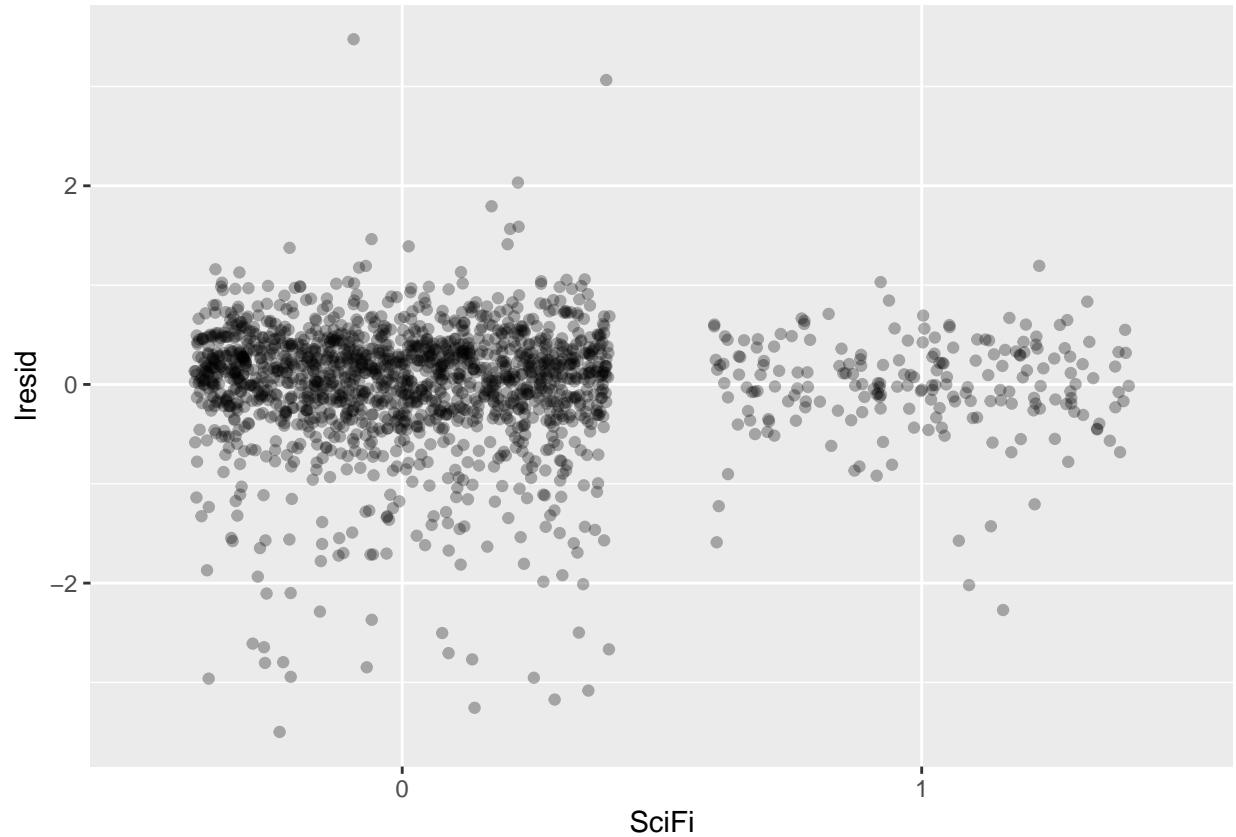
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



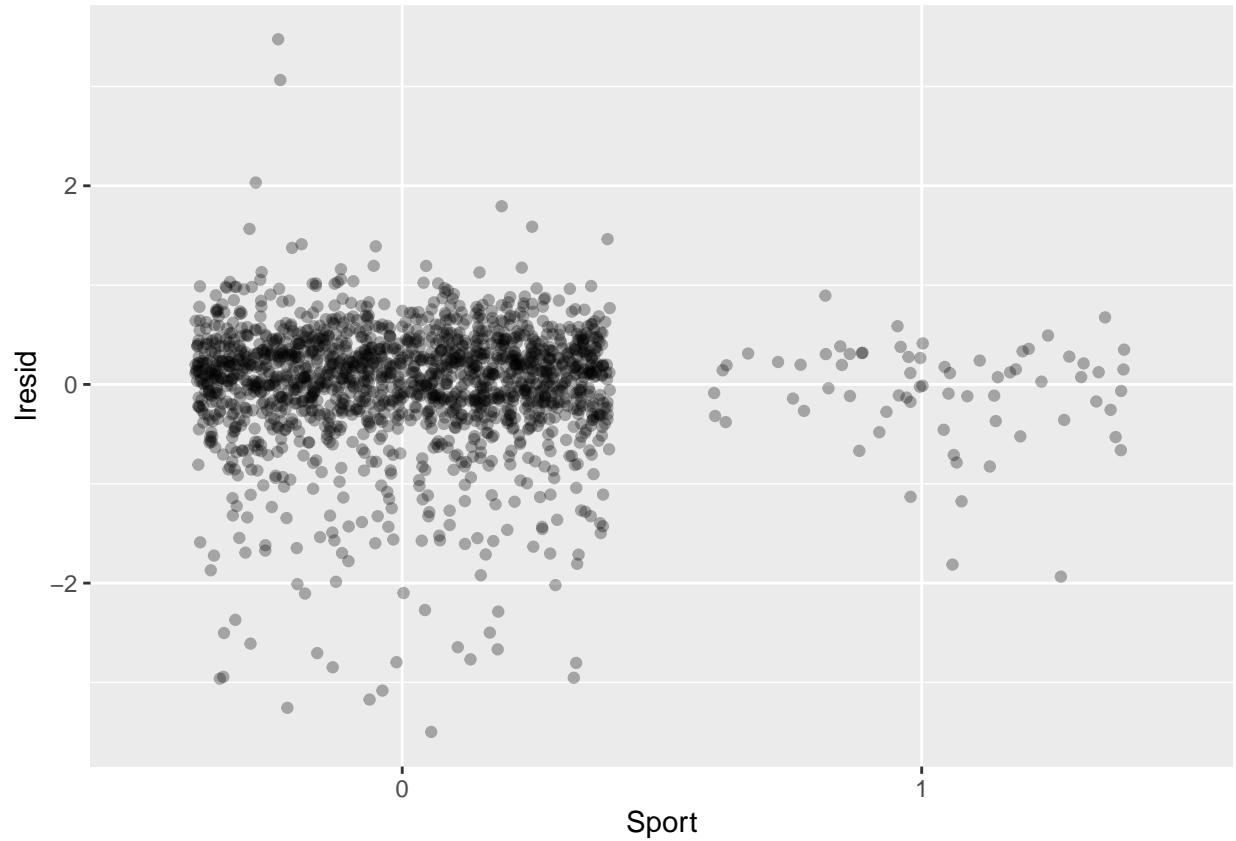
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



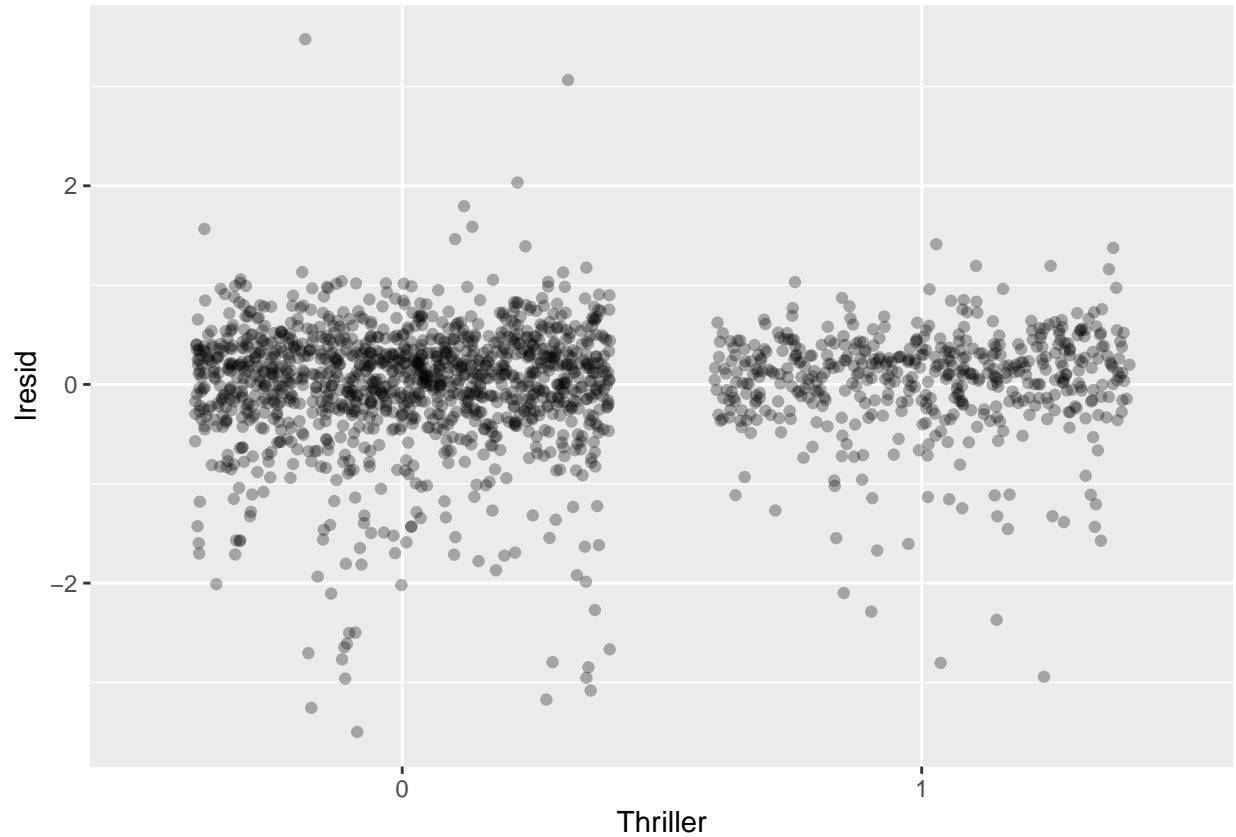
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



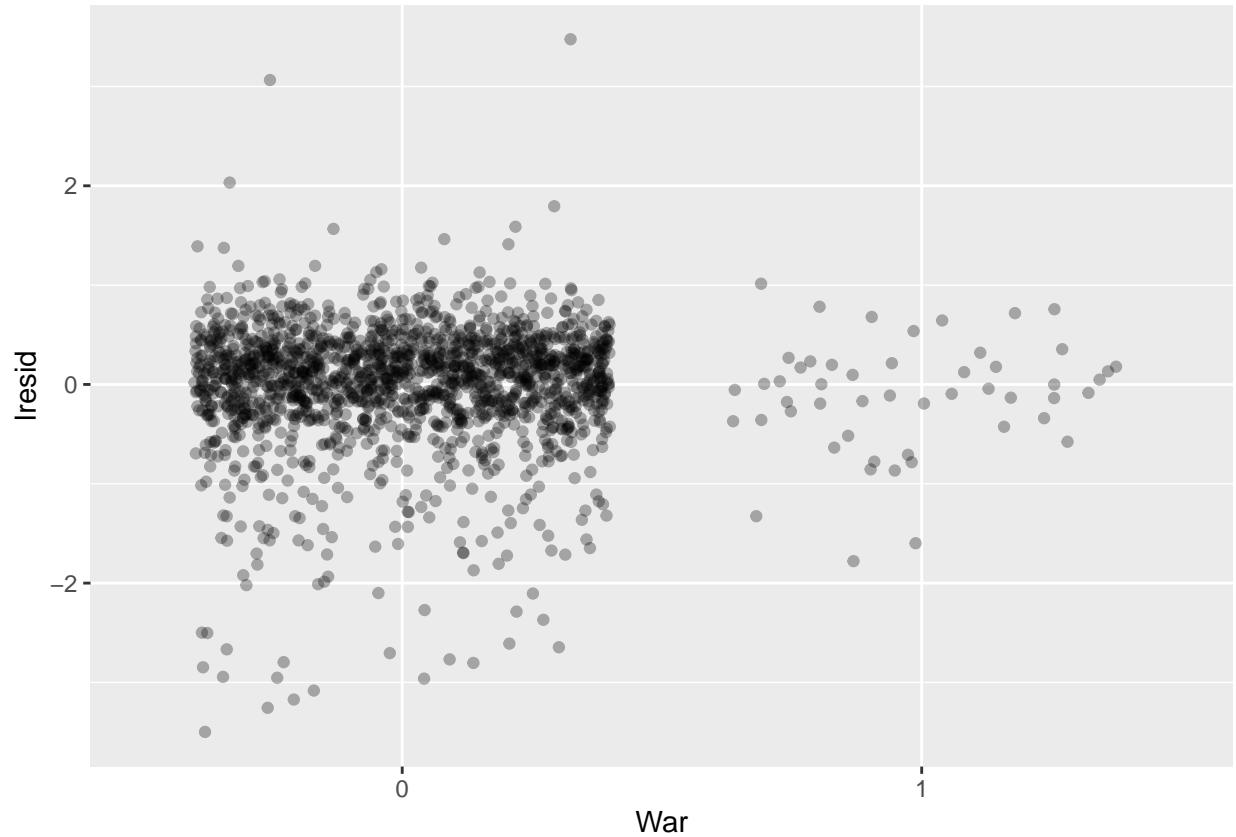
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



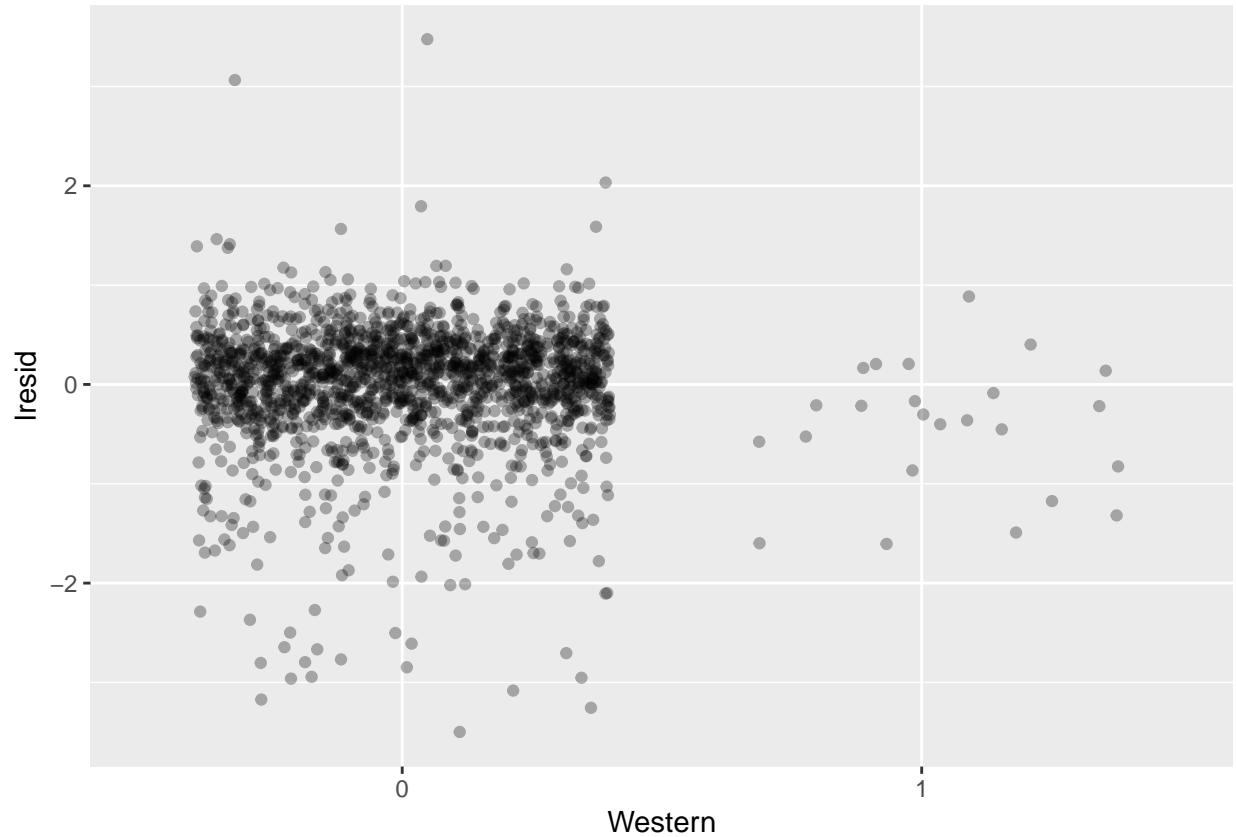
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



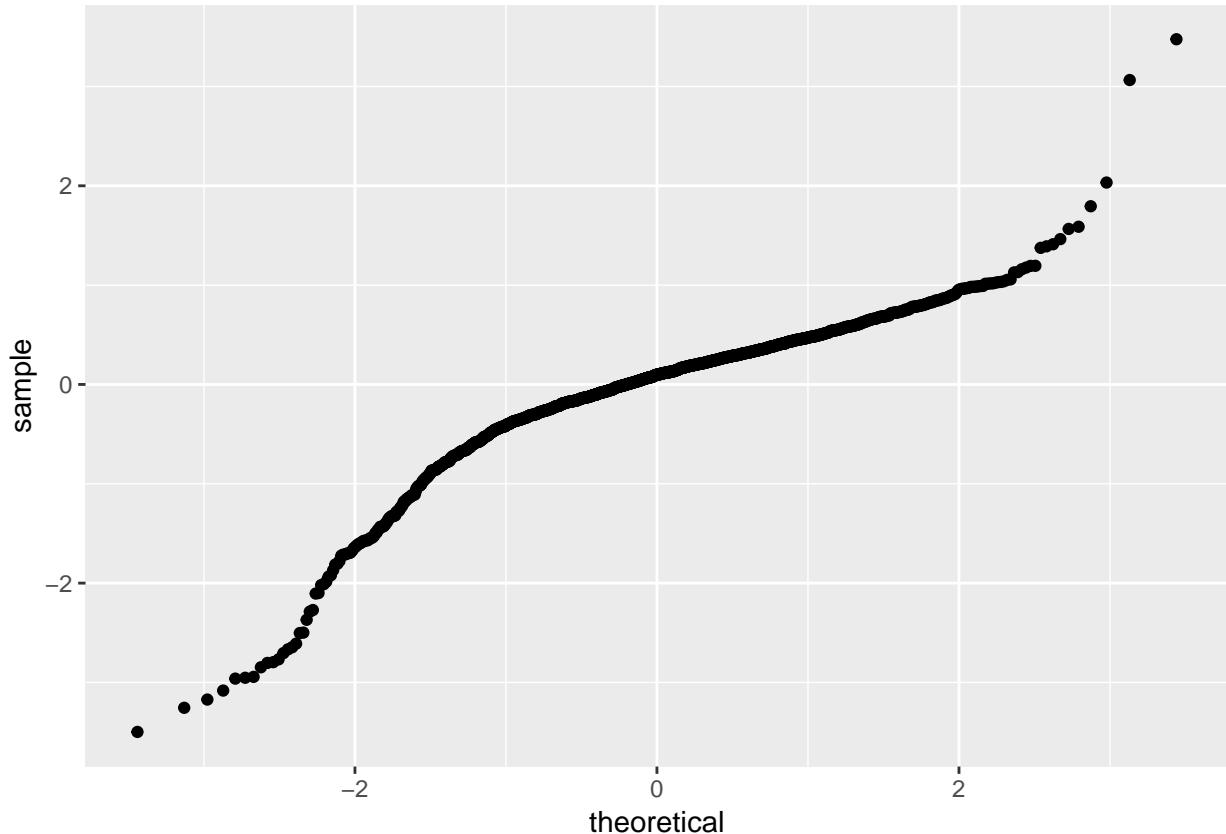
```
## Warning: Removed 132 rows containing missing values (geom_point).
```



```
## Warning: Removed 132 rows containing missing values (geom_point).
```



```
## Warning: Removed 132 rows containing non-finite values (stat_qq).
```



```
# slightly non-random relationships with some of the variables that are included in mod_simple_plus.
# but those variables don't improve model fit enough
```

Try treating genre as one variable: they all need to be added at once

Step wise. Genre should not be included.

```
rmse(lm(real_gross_log ~ Action + Adventure + Animation + Biography + Comedy + Crime + Documentary
        + Drama + Family + Fantasy + History + Horror + Music + Musical + Mystery + Romance +
        SciFi + Sport + Thriller + War + Western, data = train), data = valid)

## [1] 0.8780216

rmse(lm(real_gross_log ~ real_budget_log, data = train), data = valid) # .649

## [1] 0.6497192

rmse(lm(real_gross_log ~ imdb_score_log, data = train), data = valid)

## [1] 0.9273344

rmse(lm(real_gross_log ~ year, data = train), data = valid)

## [1] 0.9297676

rmse(lm(real_gross_log ~ content_rating, data = train), data = valid)

## [1] 0.8789942
```

```

rmse(lm(real_gross_log ~ cast_total_facebook_likes, data = train), data = valid)

## [1] 0.96301

rmse(lm(real_gross_log ~ director_facebook_likes, data = train), data = valid)

## [1] 0.9360115

rmse(lm(real_gross_log ~ total_oscars_actor, data = train), data = valid)

## [1] 0.9406675

rmse(lm(real_gross_log ~ total_oscars_director, data = train), data = valid)

## [1] 0.9352206

rmse(lm(real_gross_log ~ real_budget_log + Action + Adventure + Animation + Biography + Comedy + Crime +
      + Drama + Family + Fantasy + History + Horror + Music + Musical + Mystery + Romance +
      + SciFi + Sport + Thriller + War + Western, data = train), data = valid)

## [1] 0.6605224

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log, data = train), data = valid) # .638

## [1] 0.6382915

rmse(lm(real_gross_log ~ real_budget_log + year, data = train), data = valid)

## [1] 0.6395231

rmse(lm(real_gross_log ~ real_budget_log + content_rating, data = train), data = valid)

## [1] 0.6445119

rmse(lm(real_gross_log ~ real_budget_log + cast_total_facebook_likes, data = train), data = valid)

## [1] 0.6516661

rmse(lm(real_gross_log ~ real_budget_log + director_facebook_likes, data = train), data = valid)

## [1] 0.6501299

rmse(lm(real_gross_log ~ real_budget_log + total_oscars_actor, data = train), data = valid)

## [1] 0.6510423

rmse(lm(real_gross_log ~ real_budget_log + total_oscars_director, data = train), data = valid)

## [1] 0.6494547

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + Action + Adventure + Animation + Biography +
      + SciFi + Sport + Thriller + War + Western, data = train), data = valid)

## [1] 0.6371965

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year, data = train), data = valid) # .628

## [1] 0.6281035

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + content_rating, data = train), data = valid)

## [1] 0.631455

```

```

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + cast_total_facebook_likes, data = train), data = valid)
## [1] 0.6390821

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + director_facebook_likes, data = train), data = valid)
## [1] 0.6389138

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + total_oscars_actor, data = train), data = valid)
## [1] 0.6385564

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + total_oscars_director, data = train), data = valid)
## [1] 0.6383809

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + Action + Adventure + Animation + Biography + Crime + Drama + Family + History + Mystery + Romance + SciFi + Sport + Thriller + War + Western, data = train), data = valid)
## [1] 0.62748

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating, data = train), data = valid)
## [1] 0.623201

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + cast_total_facebook_likes, data = train), data = valid)
## [1] 0.629357

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + director_facebook_likes, data = train), data = valid)
## [1] 0.6284716

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + total_oscars_actor, data = train), data = valid)
## [1] 0.6286905

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + total_oscars_director, data = train), data = valid)
## [1] 0.6281841

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating + Action + Adventure + Crime + Drama + Family + History + Mystery + Romance + SciFi + Sport + Thriller + War + Western, data = train), data = valid)
## [1] 0.6249983

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating + cast_total_facebook_likes, data = train), data = valid)
## [1] 0.6234161

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating + director_facebook_likes, data = train), data = valid)
## [1] 0.6235301

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating + total_oscars_actor, data = train), data = valid)
## [1] 0.6240377

rmse(lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating + total_oscars_director, data = train), data = valid)
## [1] 0.623331

# RMSE has leveled off. not even including genre yet

```