

Modeling_Katrina_Cleaned

Katrina Truebebach

March 18, 2019

```
rm(list = ls())
```

Load cleaned data

```
load(file = '~/DS5110/data/proj_cleaned_dta.RData')
```

```
# need to drop years before 1980: too sparse
# most of those years have 1 or 0 observations. If including year in the model, we aren't getting any
# only necessary when including year in model, but hard to compare different models then b/c data differ
train <- train %>% filter(as.integer(as.character(year)) >= 1980)
valid <- valid %>% filter(as.integer(as.character(year)) >= 1980)
test <- test %>% filter(as.integer(as.character(year)) >= 1980)

# calculate logs
train <- train %>%
  mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
    'imdb_score', 'real_gross'), funs(log = log10(.)))
valid <- valid %>%
  mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
    'imdb_score', 'real_gross'), funs(log = log10(.)))
test <- test %>%
  mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
    'imdb_score', 'real_gross'), funs(log = log10(.)))
```

Write Functions to Automate

Write function to automate stepwise

Note: not using the step() function because can't fit and find RMSE on different datasets (train, valid)

```
# function to automate each step of stepwise variable selection
# df_vars is the dataset with only the relevant variables
# var_lst is the list of variables that are in the base model
# formula is the formula with those variables besides the y variable
step_wise_step <- function(df_vars, var_lst = NULL, formula = NULL) {
  # if first step
  if (length(var_lst) == 0) {
    # rmse with each variable against real_gross
    rmse_vars <- sapply(names(df_vars), function(var) {
      # rmse of model
      rmse(lm(as.formula(str_c('real_gross_log ~', var)), data = train), data = valid)
    })
    # if > first step: exclude variables from var_lst from data and include in model formula
  } else {
    rmse_vars <- sapply(names(df_vars %>% select(-var_lst)), function(var) {
      # rmse of model
      rmse(lm(as.formula(str_c('real_gross_log ~', formula, ' + ', var)),
        data = train), data = valid)
    })
  }
}
```

```

    })
  }
  # return the name and value of the genre that resulted in the lowest RMSE
  return(rmse_vars[which.min(rmse_vars)])
}

# function to loop through each step wise loop
# adding optional starting vars and formula in case want to build off of an existing formula
step_wise_loop <- function(df_vars, starting_vars = NULL, starting_formula = NULL) {
  # list to store min RMSE from each step in
  rmse_lst <- c()

  # first step: no genre_lst or formula (default values NULL)
  min_rmse_var <- step_wise_step(df = df_vars, var_lst = starting_vars, formula = starting_formula)
  print(min_rmse_var)

  # add to list of genres, formula, and min RMSE list
  var_lst <- c(starting_vars, names(min_rmse_var))
  formula <- str_c(starting_formula, '+', names(min_rmse_var))
  rmse_lst <- c(rmse_lst, min(min_rmse_var))

  # if have starting variables, take those out of the number we are iterating through
  if (!is.null(starting_vars)) {
    df_vars_seq <- df_vars %>% select(-starting_vars)
  } else {
    df_vars_seq <- df_vars
  }
  # loop through until have considered every variable
  for (i in seq(1:(ncol(df_vars_seq)-1))) {
    print(i)
    # step
    min_rmse_var <- step_wise_step(df = df_vars, var_lst = var_lst, formula = formula)
    print(min_rmse_var)

    # add to lists
    var_lst <- c(var_lst, names(min_rmse_var))
    formula <- str_c(formula, ' + ', names(min_rmse_var))
    rmse_lst <- c(rmse_lst, min(min_rmse_var))
  }
  return(rmse_lst)
}

```

Function to graph the residuals from a model against all potential variables (included and excluded)

```

gr_resid <- function(mod, df = train) {
  # graph residuals
  # get log versions of variables since residuals are log: same scale
  df_resid <- df %>%
    add_residuals(mod, 'lresid') %>%
    mutate_at(vars('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
                  'imdb_score'), funs(log = log10(.)))

  # graph each against log residual: continuous
  lapply(c('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',

```

```

    'imdb_score'), function(var) {
print(df_resid %>%
  ggplot() +
  geom_point(aes_string(str_c(var, '_log'), y = 'lresid')))
})

# categorical
# can't log categorical variables
lapply(c('content_rating', 'year', 'total_oscars_actor', 'total_oscars_director', all_genre_vars), function(var) {
  print(df_resid %>%
    filter(!is.na(!!rlang::sym(var))) %>%
    ggplot() +
    geom_jitter(aes_string(var, 'lresid'), alpha = .3))
})
}

```

Fit Model with Genre Variables vs Real Revenue

Start with genre because it looks like it has a strong relationship with revenue (different genres have different average revenues) and because this was our original hypothesis.

However, based on our EDA, the average revenue of some genres vs not genre is almost the same. Thus we want to start by determining which genres should be included by fitting a model of only genre variables.

Note this is not as simple as a categorical variable because one movie can have multiple genres (adventure, action, comedy).

Step Wise Selection

Dependent variable is $\log(\text{real_gross})$. Makes model look better *and* a lot of the relationships with other variables are more linear with log, so we will need to use this as y variable in the main model.

```

# version of train set with just genre columns to loop through
all_genre_vars <- c('Action', 'Adventure', 'Animation', 'Biography', 'Comedy', 'Crime', 'Documentary',
  'Drama', 'Family', 'Fantasy', 'History', 'Horror', 'Music', 'Musical', 'Mystery',
  'Romance', 'SciFi', 'Sport', 'Thriller', 'War', 'Western')


train_genre_only <- train %>% select(all_genre_vars)

# step wise implement
# return list of all min RMSE from each step -> graph
rmse_lst <- step_wise_loop(df = train_genre_only)

# graph RMSE at each step
fit_rmse <- tibble(nvar = 1:length(rmse_lst),
  rmse = rmse_lst)
ggplot(fit_rmse, aes(x = nvar, y = rmse)) + geom_line() +
  scale_x_continuous(breaks = seq(1, length(rmse_lst), by = 1)) +
  geom_vline(xintercept = 8, linetype = 'dashed') +
  geom_vline(xintercept = 10, linetype = 'dashed') +
  labs(x = 'Number of Variables', y = 'RMSE',
    title = 'RMSE vs Number of Variables: Include 8 or 10')

```

RMSE vs Number of Variables: Include 8 or 10



after var 8, decreases too small or increase (debatably 10?)

model based off of step wise

```
mod_genre <- lm(real_gross_log ~ Adventure + Action + Family + Mystery + Romance + Drama + History + Documentary, data = train)
mod_genre10 <- lm(real_gross_log ~ Adventure + Action + Family + Mystery + Romance + Drama + History + Documentary, data = train)
```

```
summary(mod_genre)
```

```
##
## Call:
## lm(formula = real_gross_log ~ Adventure + Action + Family + Mystery +
##      Romance + Drama + History + Documentary, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.0859 -0.3211  0.1680  0.5636  1.7697
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   7.21040    0.04041  178.418 < 2e-16 ***
## Adventure1    0.25572    0.05959   4.291 1.87e-05 ***
## Action1       0.37493    0.05337   7.025 2.99e-12 ***
## Family1       0.46184    0.06733   6.859 9.43e-12 ***
## Mystery1      0.19739    0.07021   2.811 0.004985 **
## Romance1      0.10130    0.04886   2.073 0.038297 *
```

```
## Drama1      -0.23074    0.04398  -5.247 1.73e-07 ***
## History1    0.45334    0.12298   3.686 0.000234 ***
## Documentary1 -1.10142    0.13044  -8.444 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8582 on 1842 degrees of freedom
## Multiple R-squared:  0.1648, Adjusted R-squared:  0.1612
## F-statistic: 45.44 on 8 and 1842 DF,  p-value: < 2.2e-16
rmse(mod_genre, data = valid)

## [1] 0.8729182
# two additional variables Musical and War are insignificant and RMSE goes from .873 to .872
summary(mod_genre10)

##
## Call:
## lm(formula = real_gross_log ~ Adventure + Action + Family + Mystery +
##      Romance + Drama + History + Documentary + Musical + War,
##      data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.0873 -0.3245  0.1659  0.5635  1.7803
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   7.21185    0.04047  178.215 < 2e-16 ***
## Adventure1    0.25575    0.05958   4.292 1.86e-05 ***
## Action1       0.36445    0.05378   6.776 1.66e-11 ***
## Family1       0.46793    0.06792   6.890 7.65e-12 ***
## Mystery1      0.20093    0.07027   2.859 0.00429 **
## Romance1      0.10540    0.04894   2.154 0.03140 *
## Drama1       -0.23707    0.04420  -5.364 9.18e-08 ***
## History1      0.38804    0.13002   2.984 0.00288 **
## Documentary1 -1.11351    0.13064  -8.524 < 2e-16 ***
## Musical1     -0.06322    0.13019  -0.486 0.62730
## War1         0.19858    0.12895   1.540 0.12375
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8581 on 1840 degrees of freedom
## Multiple R-squared:  0.166, Adjusted R-squared:  0.1615
## F-statistic: 36.63 on 10 and 1840 DF,  p-value: < 2.2e-16
rmse(mod_genre10, data = valid)

## [1] 0.8715984
# list of these variables for future use
genre_xvar <- c('Adventure', 'Action', 'Family', 'Mystery',
               'Documentary', 'Drama', 'History', 'Romance')
```

This model selection by and large makes sense. All included variables are significant at some level. However, according to Qiang's graphs in EDA, some of the included genres do not make a real difference to

real_gross. Especially History. Also, some genres that look like they would make a significant difference are not included. For example, Animation.

Thoughts:

- There are a few genres that define almost all of the movies (For example, almost 80% of the movies are either Adventure, Action, Romance, or Drama). Thus, the relationship between revenue and some genres can be explained by other genres. For example, 93 out of 101 Animation movies are also Family. So Animation's effect on revenue may already be captured by Family, which is included in the model.
- On the flip side, History is included even though it seems to have a negligible effect on revenue based on the EDA bar graphs. I don't have a great explanation for this other than it was close to the cutoff RMSE for being included. 53 out of 55 History movies are also Drama. So unclear why included.

```
train %>% filter(Animation == 1, Family == 1) %>% count() # 93
train %>% filter(Animation == 1) %>% count() # 101

train %>% filter(History == 1) %>% count() # 52
train %>% filter(History == 1, Drama == 1) %>% count() # 51
```

Residuals graph


I like this geom_jitter view better. Can see individual points. Most movies have some outliers where actual makes less money than predicted based on the included genres. But tricky because movies are multiple genres. So could be because that movie is also another genre that makes less money. Bulk of observations around zero.


```
train_resid <- train %>%
  add_residuals(mod_genres, 'lresid')


# graph residuals against each variable included in the model
# most look random except Adventure
lapply(genres_xvar, function(var) {
  train_resid %>%
    ggplot() +
    geom_jitter(aes_string(var, y = 'lresid'), alpha = .3)
})

# graph residuals against each genre not included in the model
lapply(names(train_genres_only %>% select(-genres_xvar)), function(var) {
  train_resid %>%
    ggplot() +
    geom_jitter(aes_string(var, y = 'lresid'), alpha = .3)
})
```

Display a couple of plots for presentation: some diversity. But show all around 0








QQ-Plot Not great.

```
# residuals themselves are NOT normally distributed
# qq plot
train_resid %>% ggplot(aes(sample = lresid)) +
  geom_qq() +
  geom_qq_line() +
  labs(title = 'Residual QQPlot: deviations at tails',
        x = 'Theoretical Quantiles', y = 'Sample Quantiles')
```

Residual QQPlot: deviations at tails



Plot Predictions

Plot prediction for mean real revenue against each genre included in the model.

AND genres not included in the model: still pretty good with some exceptions. Evidence that these genres are not useful for prediction/are covered by other genres.

```
train_pred <- train %>%
  add_predictions(mod_genre, 'lpred') %>%
  mutate(pred = 10^lpred)

lapply(genre_xvar, function(var) {
  gr <- train_pred %>%
    ggplot(aes_string(x = var)) +
    geom_boxplot(aes(y = real_gross)) +
    # include mean
    stat_summary(aes(y = real_gross), fun.y = mean, geom = 'point', size = 4) +
    geom_point(data = train_pred %>% group_by(!!rlang::sym(var)) %>% summarize(mean = mean(pred)),
              aes(y = mean), color = 'red', size = 2) +
    scale_y_log10() +
    labs(y = 'Log Real Gross Revenue', title = 'Revenue Actual vs Predicted: Log Scale \n Successful Pro
})

# predictions against other genres
lapply(names(train_genre_only %>% select(-genre_xvar)), function(var) {
  train_pred %>%
    ggplot(aes_string(x = var)) +
    geom_boxplot(aes(y = real_gross)) +
```


```

# include mean
stat_summary(aes(y = real_gross), fun.y = mean, geom = 'point', size = 4) +
geom_point(data = train_pred %>% group_by(!rlang::sym(var)) %>% summarize(mean = mean(pred)),
          aes(y = mean), color = 'red', size = 2) +
scale_y_log10() +
labs(y = 'Log Real Gross Revenue', title = 'Revenue Actual vs Predicted: Log Scale \n Successful Pr
})

```

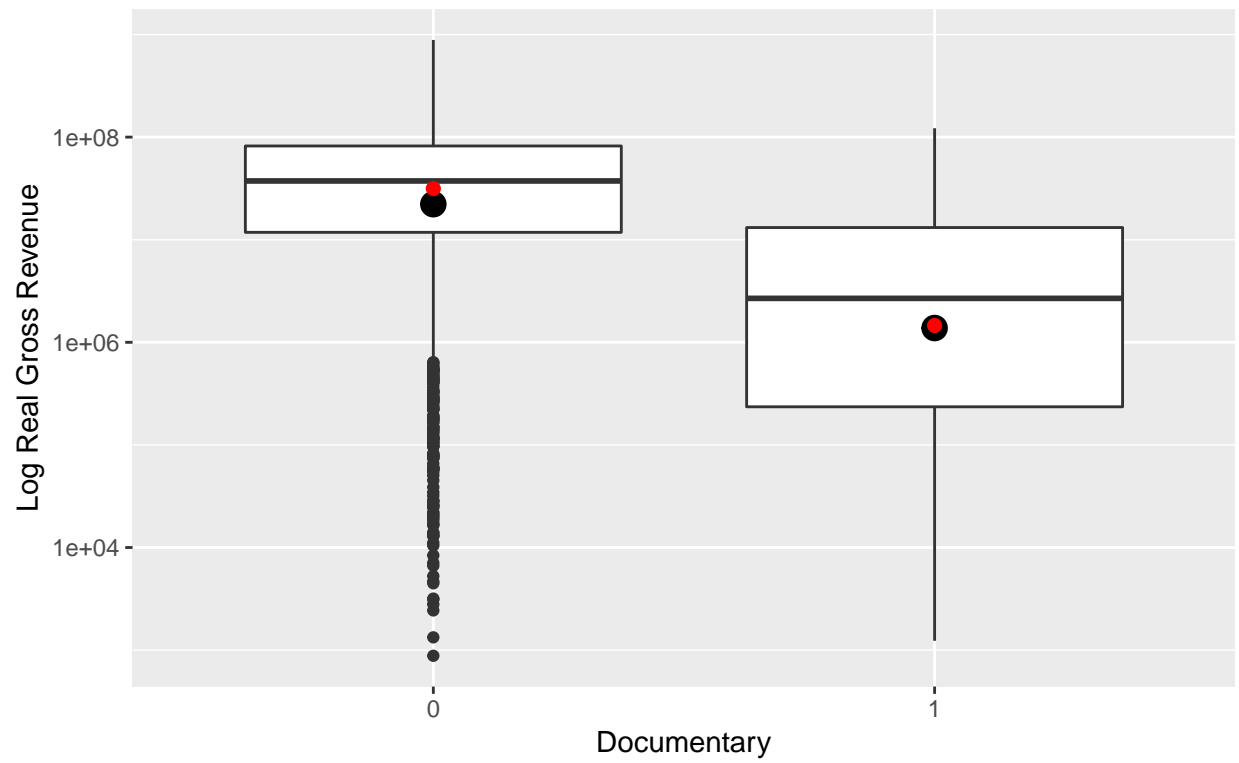
Display a couple of plots for presentation. In slides, report: note about black dot is true mean, red dot is pre-

Revenue Actual vs Predicted: Log Scale Successful Prediction




dicted


Revenue Actual vs Predicted: Log Scale
Successful Prediction



Revenue Actual vs Predicted: Log Scale
Successful Prediction Even With Genres Not Included




Revenue Actual vs Predicted: Log Scale Successful Prediction Even With Genres Not Included



Overall predictions: clearly not enough to just specify genres

```
train %>%
  add_predictions(mod_genre, 'lpred') %>%
  mutate(pred = 10^lpred) %>%
  ggplot() +
  geom_freqpoly(aes(x = real_gross, color = 'Actual')) +
  geom_freqpoly(aes(x = pred, color = 'Predicted')) +
  scale_x_log10() +
  labs(x = 'Log Real Gross Revenue', y = 'Count') +
  scale_color_manual(name = '', values = c(Actual = 'black', Predicted = 'blue'))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




Glmnet: sparse


Try to eliminate most of the genre variables using glmnet and see if results are similar to stepwise. Can't do statistical tests, so not useful for analysis, but can use to aid justification.

```
# matrix of x and y variables
x <- as.matrix(train_genre_only %>% mutate_all(funs(as.numeric(as.character(.)))))
y <- as.matrix(train$real_gross_log)

# glmnet process form class
mod_sparse <- glmnet(x, y, family = 'gaussian')
plot(mod_sparse, xvar = 'lambda', label = TRUE)
```



```
mod_sparse <- cv.glmnet(x, y)
plot(mod_sparse)
```

```
coef(mod_sparse, s = 'lambda.min') # use min lambda
```

```
## 22 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  7.157540227
## Action      0.317543067
## Adventure    0.236828217
## Animation    0.135506319
## Biography    0.145534271
## Comedy      -0.010964561
## Crime       -0.014150867
## Documentary -1.084241783
## Drama       -0.237845227
## Family      0.405574661
## Fantasy     0.112498696
## History     0.321743675
## Horror      .
## Music       0.172500588
## Musical     -0.074590785
## Mystery     0.108157227
## Romance     0.123131935
## SciFi       0.003661742
## Sport       0.150296957
## Thriller    0.166927824
## War         0.211110797
## Western     -0.216137862
```

```
coef(mod_sparse, s = 'lambda.1se') # use most sparse
```

```
## 22 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  7.24955074
## Action      0.21115913
## Adventure   0.18312646
## Animation   .
## Biography   .
## Comedy      .
## Crime       .
## Documentary -0.54632201
## Drama       -0.06351752
## Family      0.22984009
## Fantasy     .
## History     .
## Horror      .
## Music       .
## Musical     .
## Mystery     .
## Romance     .
## SciFi       .
## Sport       .
## Thriller    .
## War         .
## Western     .
```

Use genre model as a base

Next, we started with the genre model and added in additional variables: clear that genre is not sufficient. Several steps:

- Plot other variables (budget, facebook likes, total number of Oscars, imdb score, number of oscars, content rating, year) against the residuals from the genre model.
 - All of these have somewhat non-random relationships with residuals. Especially budget and IMDB score. Movies with higher budgets make more revenue than predicted by genre (positive residual)
- Based on these non-random relationships, used them all in step wise, but started with the genre variables from the genre model as a base.
 - Log budget and IMDB score had the strongest relationships with log revenue in EDA. Facebook likes and Oscars had especially weak relationships. Thus the final variables included are not surprising.
- Looked at new residuals of included and excluded relationships. All fairly random.
- Looked at predictions for each individual variable. Good job even for variables not included in the model.
 - Note that these predictions are not simple lines based on one coefficient because each movie's predicted revenue is based on several variables.
 - Year isn't great. But in most years, other factors would be more important. But some years, such as recessions, it is very important which is why it is important to include.
- Looked at predictions overall. Looks much better.

Overall, this is a pretty good model. Prediction is pretty good. Residuals are more normal. A lot of the variables are not significant, but they are significant when considered together in anova.


However, still not convinced this is the best model. I worry that it is overfitted and the genre variables are


not necessary. Maybe it is not a valid assumption that we should start with genre, especially since a lot did come in as insignificant. Perhaps these effects are being captured by other variables, so we should not start with the best assumption that we include genre.

```
train_resid <- train %>%
  add_residuals(mod_genre, 'lresid')

# graph each against log residual: continuous (log scale)
lapply(c('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes',
        'imdb_score'), function(var) {
  train_resid %>%
    ggplot() +
    geom_point(aes_string(str_c(var, '_log'), y = 'lresid'))
})

# categorical
# can't log categorical variables
lapply(c('content_rating', 'year', 'total_oscars_actor', 'total_oscars_director'), function(var) {
  train_resid %>%
    filter(!is.na(!rlang::sym(var))) %>%
    ggplot() +
    geom_jitter(aes_string(var, 'lresid'), alpha = .3)
})
```






```
# For some log, need to turn -Inf from log(0) to NA
train <- train %>%
  mutate_at(vars(contains('log')), funs(ifelse(is.infinite(.), NA, .)))
valid <- valid %>%
  mutate_at(vars(contains('log')), funs(ifelse(is.infinite(.), NA, .)))

# starting formula: genre
starting_formula = 'Adventure + Action + Family + Mystery + Documentary + Drama + History + Romance'

# stepwise starting with genre
rmse_lst <- step_wise_loop(df = train %>% select(genre_xvar, content_rating, real_budget, year,
                                                total_oscars_actor, total_oscars_director,
                                                imdb_score_log, real_budget_log,
                                                director_facebook_likes_log,
                                                cast_total_facebook_likes_log),
                          starting_vars = genre_xvar,
                          starting_formula = starting_formula)

# graph RMSE vs number of variables
fit_rmse <- tibble(nvar = 1:length(rmse_lst),
                  rmse = rmse_lst)
ggplot(fit_rmse, aes(x = nvar, y = rmse)) + geom_line() +
  scale_x_continuous(breaks = seq(1, length(rmse_lst), by = 1)) +
  geom_vline(xintercept = 4, linetype = 'dashed') +
  labs(x = 'Number of Variables', y = 'RMSE',
       title = 'RMSE vs Number of Variables: Include 4')
```

RMSE vs Number of Variables: Include 4



after var 4, decreases too small or increase

model with extra 4 variables

```
mod_all <- lm(real_gross_log ~ Adventure + Action + Family + Mystery +
              Documentary + Drama + History + Romance +
              real_budget_log + imdb_score_log + year + content_rating,
              data = train)
```

```
summary(mod_all)
```

```
##
## Call:
## lm(formula = real_gross_log ~ Adventure + Action + Family + Mystery +
##     Documentary + Drama + History + Romance + real_budget_log +
##     imdb_score_log + year + content_rating, data = train)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-3.4485	-0.2238	0.0878	0.3407	3.3377

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.016486	0.394298	-0.042	0.96665
Adventure1	-0.142890	0.045332	-3.152	0.00165 **
Action1	-0.011306	0.040640	-0.278	0.78090
Family1	0.311938	0.079689	3.914	9.43e-05 ***

```

## Mystery1          0.042585    0.051695    0.824    0.41019
## Documentary1      0.163354    0.141606    1.154    0.24884
## Drama1            -0.257951    0.034509   -7.475    1.24e-13 ***
## History1          -0.017739    0.089809   -0.198    0.84345
## Romance1          0.015177    0.037044    0.410    0.68208
## real_budget_log    0.812038    0.029073   27.931    < 2e-16 ***
## imdb_score_log     2.160658    0.206270   10.475    < 2e-16 ***
## year1981          -0.209854    0.368926   -0.569    0.56955
## year1982           0.182927    0.339404    0.539    0.58998
## year1983           0.247796    0.384325    0.645    0.51917
## year1984           0.500267    0.329092    1.520    0.12866
## year1985           0.288020    0.339707    0.848    0.39665
## year1986           0.065028    0.324926    0.200    0.84140
## year1987           0.172547    0.316476    0.545    0.58568
## year1988           0.147460    0.310528    0.475    0.63494
## year1989           0.279267    0.305019    0.916    0.36002
## year1990           0.062553    0.306795    0.204    0.83846
## year1991           0.044401    0.310206    0.143    0.88620
## year1992          -0.007279    0.315891   -0.023    0.98162
## year1993           0.025560    0.307311    0.083    0.93372
## year1994          -0.222206    0.301314   -0.737    0.46095
## year1995          -0.050133    0.293633   -0.171    0.86445
## year1996          -0.217635    0.286206   -0.760    0.44712
## year1997          -0.171164    0.286195   -0.598    0.54988
## year1998          -0.272722    0.285619   -0.955    0.33980
## year1999          -0.270022    0.282260   -0.957    0.33889
## year2000          -0.173448    0.283525   -0.612    0.54078
## year2001          -0.287595    0.281570   -1.021    0.30721
## year2002          -0.262480    0.281560   -0.932    0.35135
## year2003          -0.250251    0.282916   -0.885    0.37653
## year2004          -0.266357    0.282475   -0.943    0.34585
## year2005          -0.255937    0.281463   -0.909    0.36332
## year2006          -0.346753    0.281963   -1.230    0.21895
## year2007          -0.347511    0.283316   -1.227    0.22015
## year2008          -0.374356    0.281500   -1.330    0.18375
## year2009          -0.371088    0.280547   -1.323    0.18611
## year2010          -0.402398    0.281138   -1.431    0.15253
## year2011          -0.279867    0.283485   -0.987    0.32367
## year2012          -0.184045    0.281991   -0.653    0.51406
## year2013          -0.120662    0.281346   -0.429    0.66807
## year2014          -0.154492    0.283282   -0.545    0.58558
## year2015          -0.267933    0.285375   -0.939    0.34793
## year2016          -0.176387    0.301510   -0.585    0.55862
## content_ratingNC-17 -0.036289    0.275191   -0.132    0.89510
## content_ratingPG    -0.035081    0.119686   -0.293    0.76948
## content_ratingPG-13  0.139695    0.137054    1.019    0.30822
## content_ratingR     -0.052452    0.136983   -0.383    0.70184
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6071 on 1668 degrees of freedom
## (132 observations deleted due to missingness)
## Multiple R-squared:  0.495, Adjusted R-squared:  0.4799
## F-statistic: 32.71 on 50 and 1668 DF, p-value: < 2.2e-16

```

```
rmse(mod_all, data = valid)
```

```
## [1] 0.6206177
```

```
# when consider the factors as one variable, they are significant  
anova(mod_all)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: real_gross_log
```

```
##           Df Sum Sq Mean Sq F value    Pr(>F)        
## Adventure    1  75.64   75.64 205.2075 < 2.2e-16 ***  
## Action       1  24.10   24.10  65.3761 1.178e-15 ***  
## Family       1  32.07   32.07  86.9917 < 2.2e-16 ***  
## Mystery      1   2.40    2.40   6.5082 0.01083 *  
## Documentary  1   7.70    7.70  20.8945 5.210e-06 ***  
## Drama        1  14.60   14.60  39.6072 3.960e-10 ***  
## History      1   5.95    5.95  16.1431 6.134e-05 ***  
## Romance      1   1.86    1.86   5.0542 0.02470 *  
## real_budget_log 1 347.64  347.64 943.0834 < 2.2e-16 ***  
## imdb_score_log 1  46.38   46.38 125.8221 < 2.2e-16 ***  
## year        36  33.63    0.93   2.5341 1.764e-06 ***  
## content_rating 4  10.83    2.71   7.3440 7.330e-06 ***  
## Residuals   1668 614.85    0.37
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# number of observations
```

```
nobs(mod_all)
```

```
## [1] 1719
```

```
gr_resid(mod_all)
```

```
# two points with consistently really high residuals (greater than 2) i.e. based on all of their factors
```

```
### make a graph labeling these and pointing out
```

```
# often many points with large negative results less than -2: often get movies that are a flop. High budget
```

```
train %>%
```

```
  add_residuals(mod_all, 'lresid') %>%
```

```
  filter(lresid > 2.1)
```

```
train %>%
```

```
  add_residuals(mod_all, 'lresid') %>%
```

```
  filter(lresid < -2)
```

```
train %>%
```

```
  add_residuals(mod_all, 'lresid') %>%
```

```
  ggplot(aes(sample = lresid)) +
```

```
    geom_qq() +
```

```
    geom_qq_line() +
```


```
    labs(title = 'Residual QQQPlot: Improved, still Deviation at Tails',
```

```
          x = 'Theoretical Quantiles', y = 'Sample Quantiles')
```

```
## Warning: Removed 132 rows containing non-finite values (stat_qq).
```

```
## Warning: Removed 132 rows containing non-finite values (stat_qq_line).
```

Residual QQPlot: Improved, still Deviation at Tails

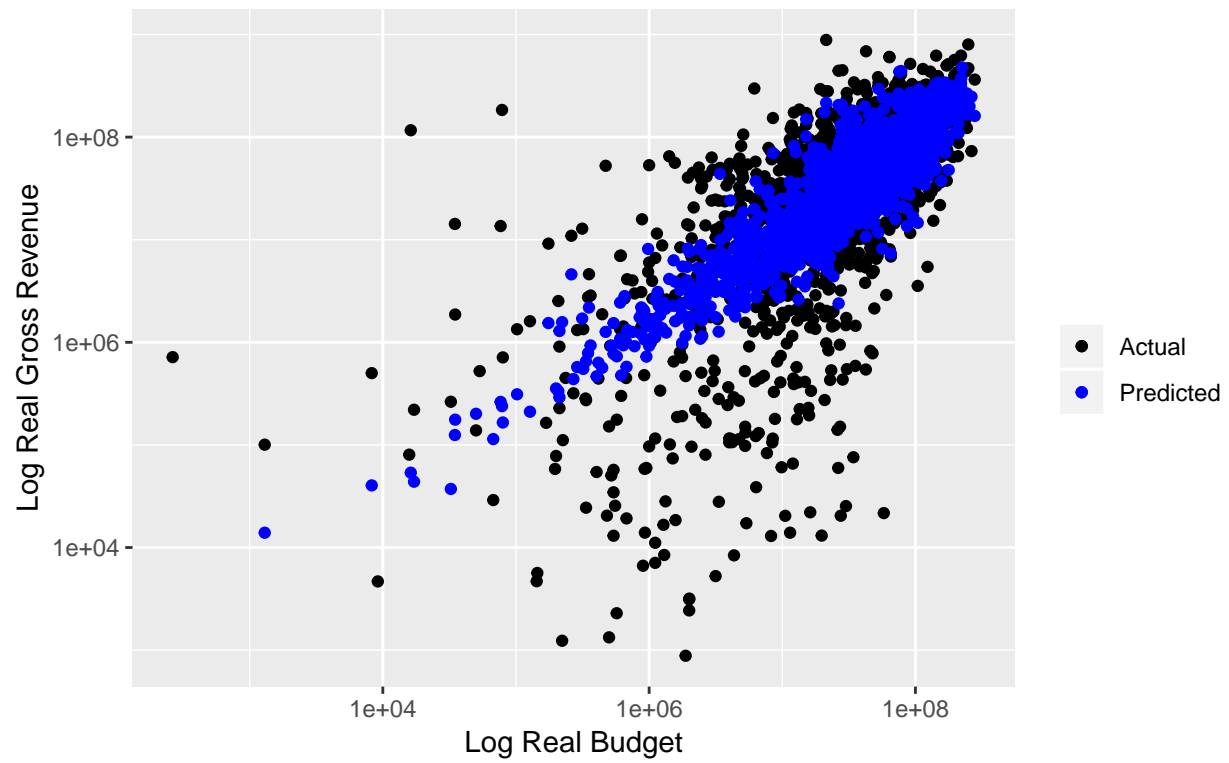


```
train_pred_all <- train %>%
  add_predictions(mod_all, 'lpred') %>%
  mutate(pred = 10^lpred)

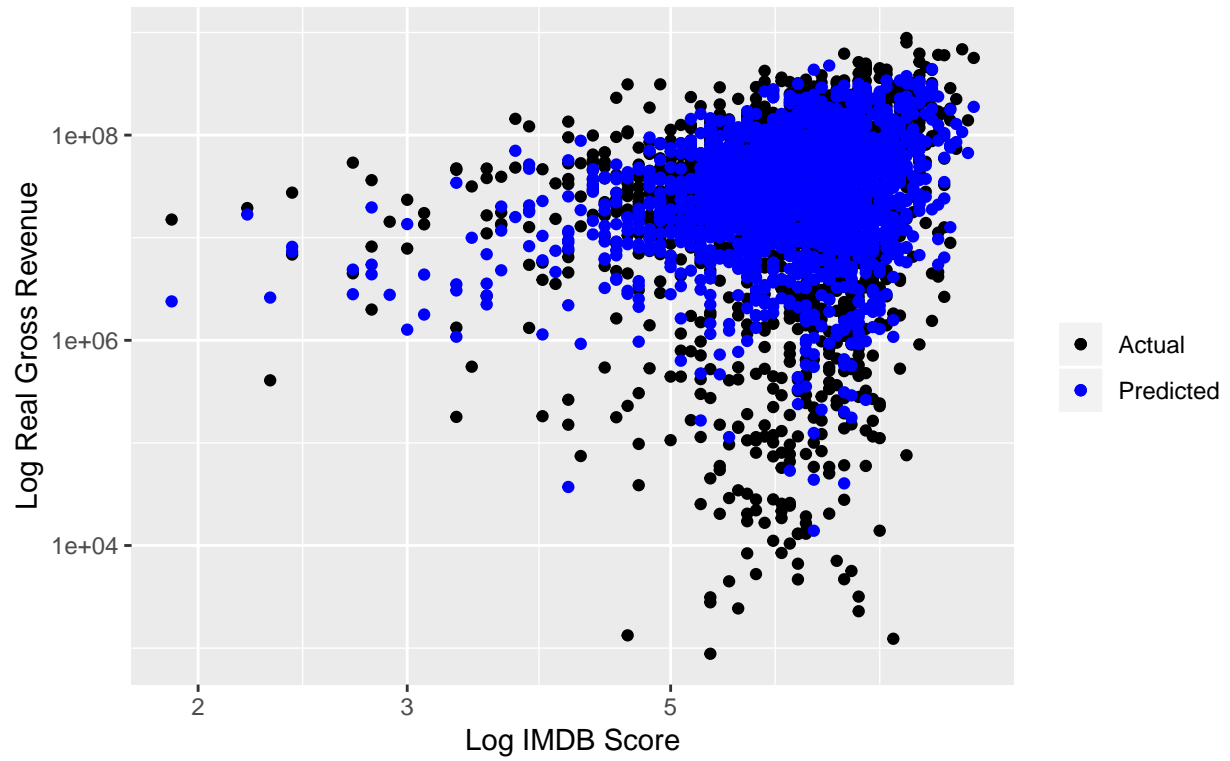
# point because many factors. won't just be a single coefficient determining
# SEE BOOK where did prediction based on a bunch of variables
lapply(c('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes', 'imdb_score'), function(var) {
  train_pred_all %>%
    ggplot(aes_string(x = var)) +
    geom_point(aes(y = real_gross, color = 'Actual')) +
    geom_point(aes(y = pred, color = 'Predicted')) +
    scale_y_log10() + scale_x_log10()
})

# predictions against other genres
lapply(c('content_rating', 'year', 'total_oscars_actor', 'total_oscars_director', all_genre_vars), function(var) {
  train_pred_all %>%
    ggplot(aes_string(x = var)) +
    geom_boxplot(aes(y = real_gross)) +
    # include mean
    stat_summary(aes(y = real_gross), fun.y = mean, geom = 'point', size = 4) +
    geom_point(data = train_pred %>% group_by(!!rlang::sym(var)) %>% summarize(mean = mean(pred)),
              aes(y = mean), color = 'red', size = 2) +
    scale_y_log10()
})
```



Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction



Revenue Actual vs Predicted: Log Scale
Successful Prediction



Revenue Actual vs Predicted: Log Scale Successful Prediction

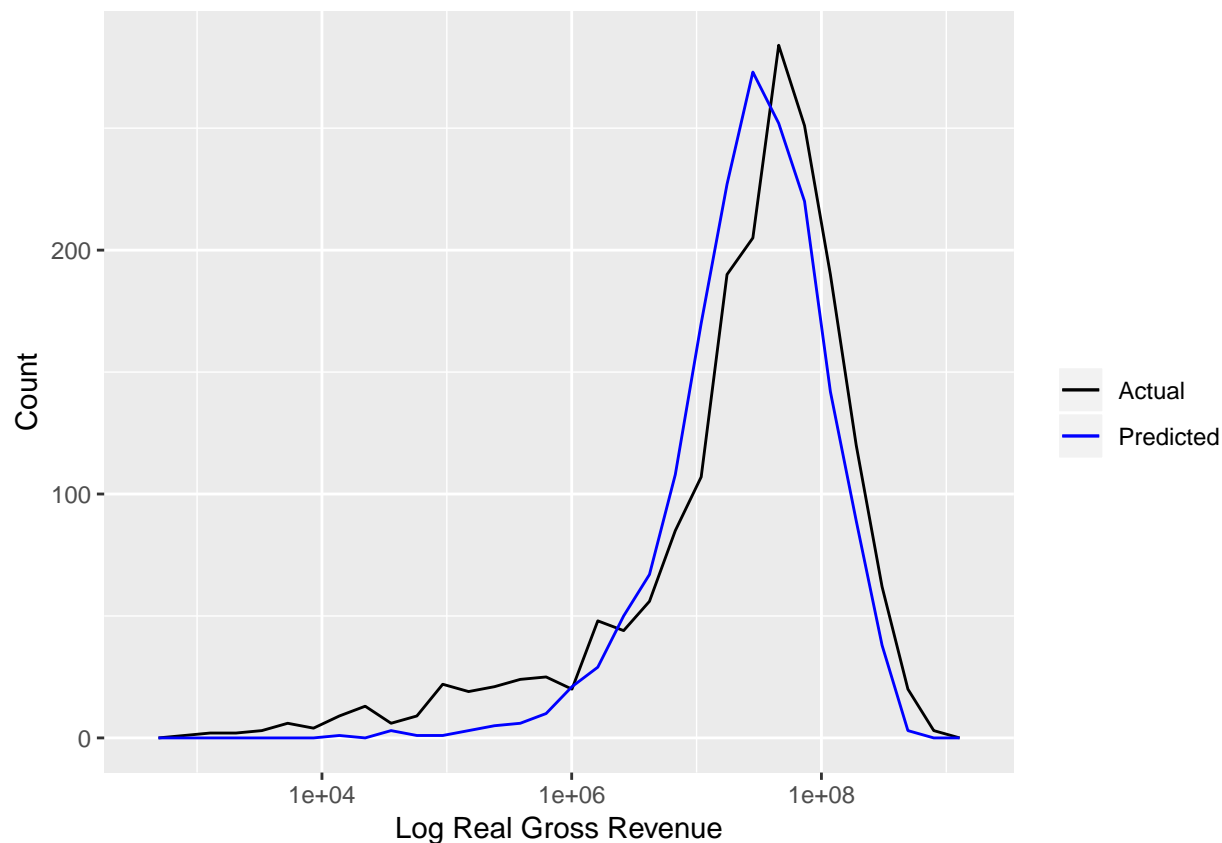


```
train %>%
  add_predictions(mod_all, 'lpred') %>%
  mutate(pred = 10^lpred) %>%
  ggplot() +
    geom_freqpoly(aes(x = real_gross, color = 'Actual')) +
    geom_freqpoly(aes(x = pred, color = 'Predicted')) +
    scale_x_log10() +
    labs(x = 'Log Real Gross Revenue', y = 'Count') +
    scale_color_manual(name = '', values = c(Actual = 'black', Predicted = 'blue'))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 132 rows containing non-finite values (stat_bin).
```



Fit Model From Scratch

Log budget and log revenue have the strongest and most linear relationship based on EDA. Also it is very logical that they would be related. Start with budget.

- Plot other variables against the residuals from the budget model.
- Non-random: year, IMDB score, content rating. All of these are somewhat subtle. Budget does a pretty good job of estimating this relationship by itself. However, high IMDB scores do have positive residuals, implying that actual > predicted revenue. Some years have skewed residuals. G movies also have positive residuals.
- + All genres have random relationships with residual. Do not need to include genre.
- Include the variables with non-random residuals (year, IMDB score, content rating) in a step-wise selection process. All included based off of step-wise.
- Look at new residuals of included and excluded relationships. All fairly random.
- Look at predictions for each individual variable. Good job even for variables not included in the model.
 - Year isn't great again.
- Look at predictions overall. Looks good.

Not shown here, but I also fit a step wise from absolute scratch. Result was the same four variables plus Comedy and Mystery.

Also did step wise where all of the genre variables had to be included together. Got the same four variables and no genre.

```
mod_simple <- lm(real_gross_log ~ real_budget_log, data = train)
summary(mod_simple)
```

```
##
## Call:
## lm(formula = real_gross_log ~ real_budget_log, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4222 -0.2401  0.0963  0.3667  3.5856
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.63498    0.18025   3.523 0.000438 ***
## real_budget_log 0.91363    0.02437  37.486 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6755 on 1747 degrees of freedom
## (102 observations deleted due to missingness)
## Multiple R-squared:  0.4458, Adjusted R-squared:  0.4455
## F-statistic: 1405 on 1 and 1747 DF,  p-value: < 2.2e-16

rmse(mod_simple, data = valid)

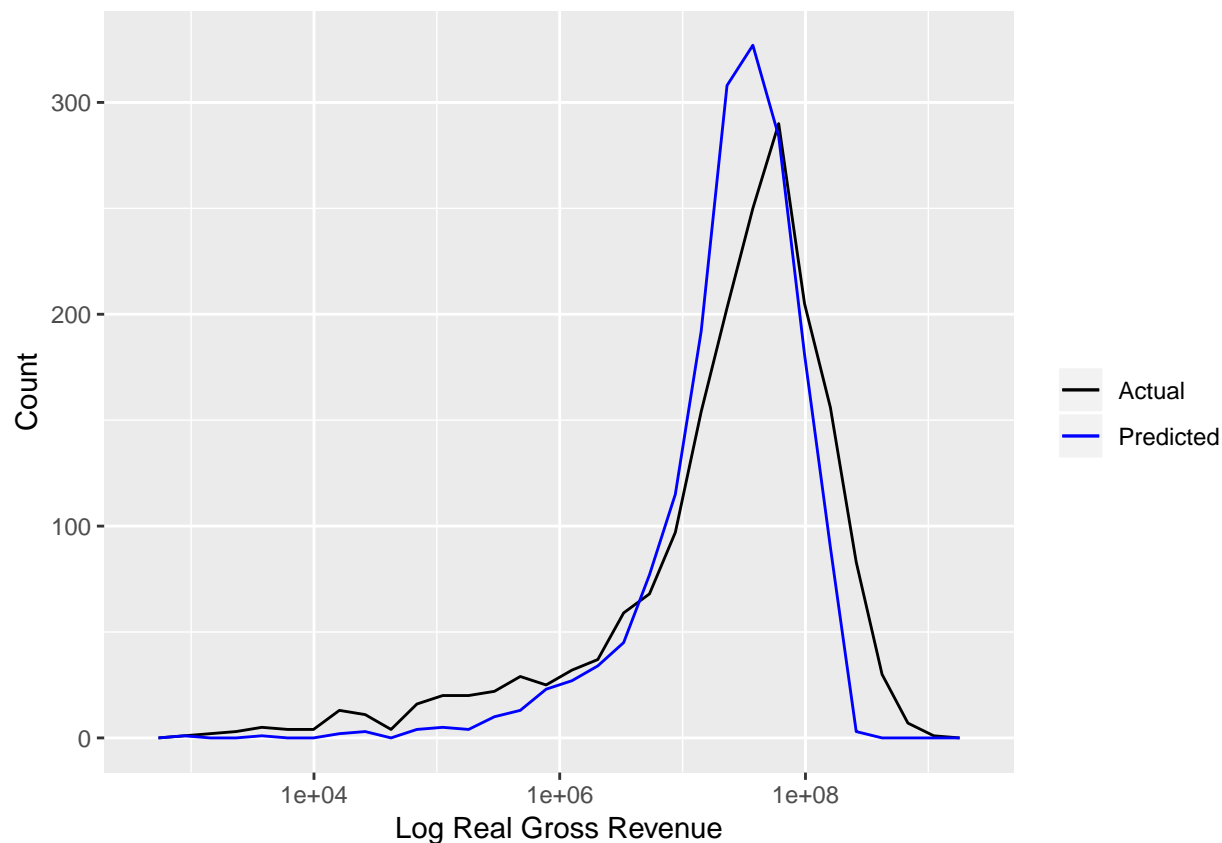
## [1] 0.6497192

gr_resid(mod_simple)
```

Take a quick look at the predictions from this very simple model with just budget. Decent, but was better with more variables.

```
train %>%
  add_predictions(mod_simple, 'lpred') %>%
  mutate(pred = 10^lpred) %>%
  ggplot() +
  geom_freqpoly(aes(x = real_gross, color = 'Actual')) +
  geom_freqpoly(aes(x = pred, color = 'Predicted')) +
  scale_x_log10() +
  labs(x = 'Log Real Gross Revenue', y = 'Count') +
  scale_color_manual(name = '', values = c(Actual = 'black', Predicted = 'blue'))

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 102 rows containing non-finite values (stat_bin).
```

```

train_resid_simple <- train %>%
  add_residuals(mod_simple, 'lresid')

gr_con <- function(v, title, xtitle) {
  lapply(v, function(var) {
    train_resid_simple %>%
      ggplot() +
      geom_point(aes_string(str_c(var, '_log'), y = 'lresid')) +
      geom_hline(aes(yintercept = 0)) +
      labs(y = 'Log Residual', title = title, x = xtitle)
  })
}

gr_cat <- function(v, title, xtitle, flip = F) {
  lapply(v, function(var) {
    gr <- train_resid_simple %>%
      filter(!is.na(!rlang::sym(var))) %>%
      ggplot() +
      geom_jitter(aes_string(var, y = 'lresid'), alpha = .3) +
      geom_hline(aes(yintercept = 0)) +
      labs(y = 'Log Residual', title = title, x = xtitle)
    if (flip) {
      gr + coord_flip()
    } else {
      gr
    }
  })
}

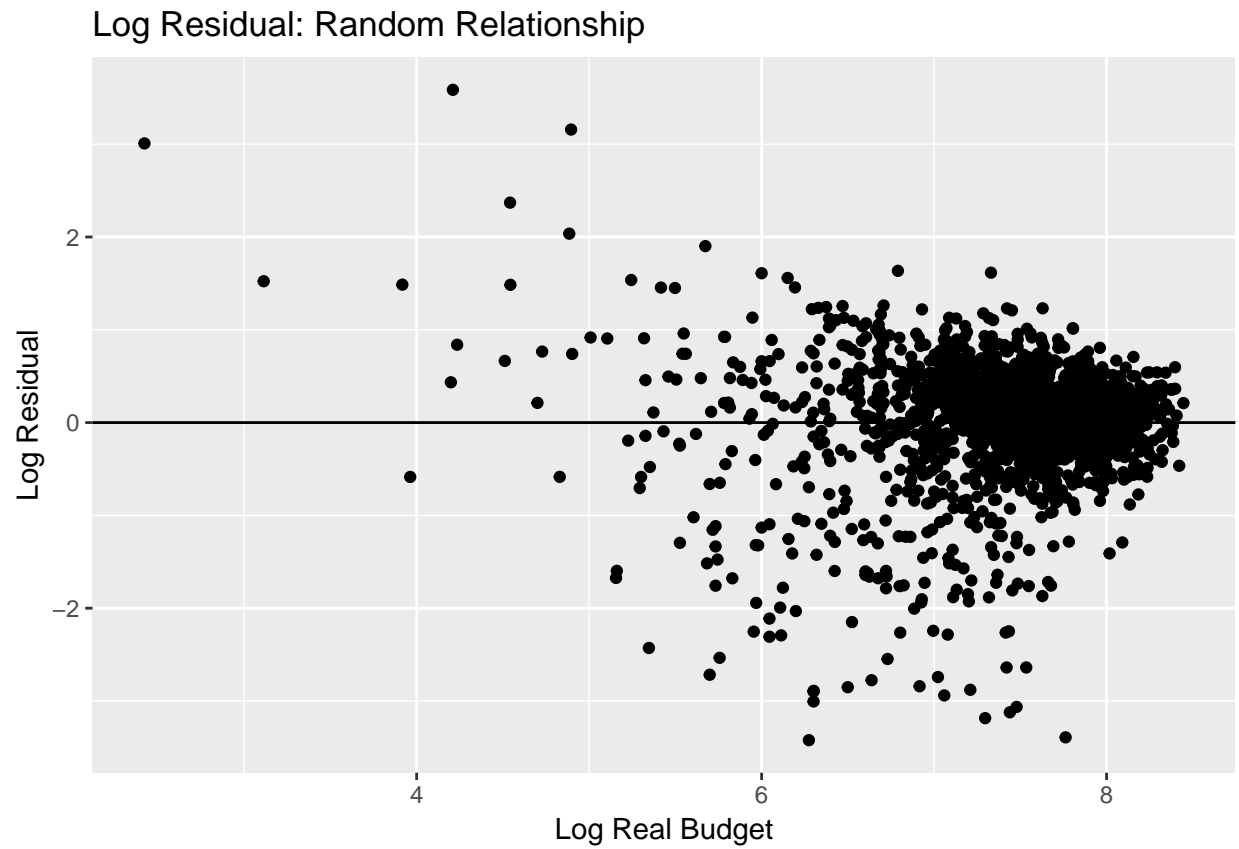
```

```

})
}

gr_con('real_budget', 'Log Residual: Random Relationship', 'Log Real Budget')

```




```

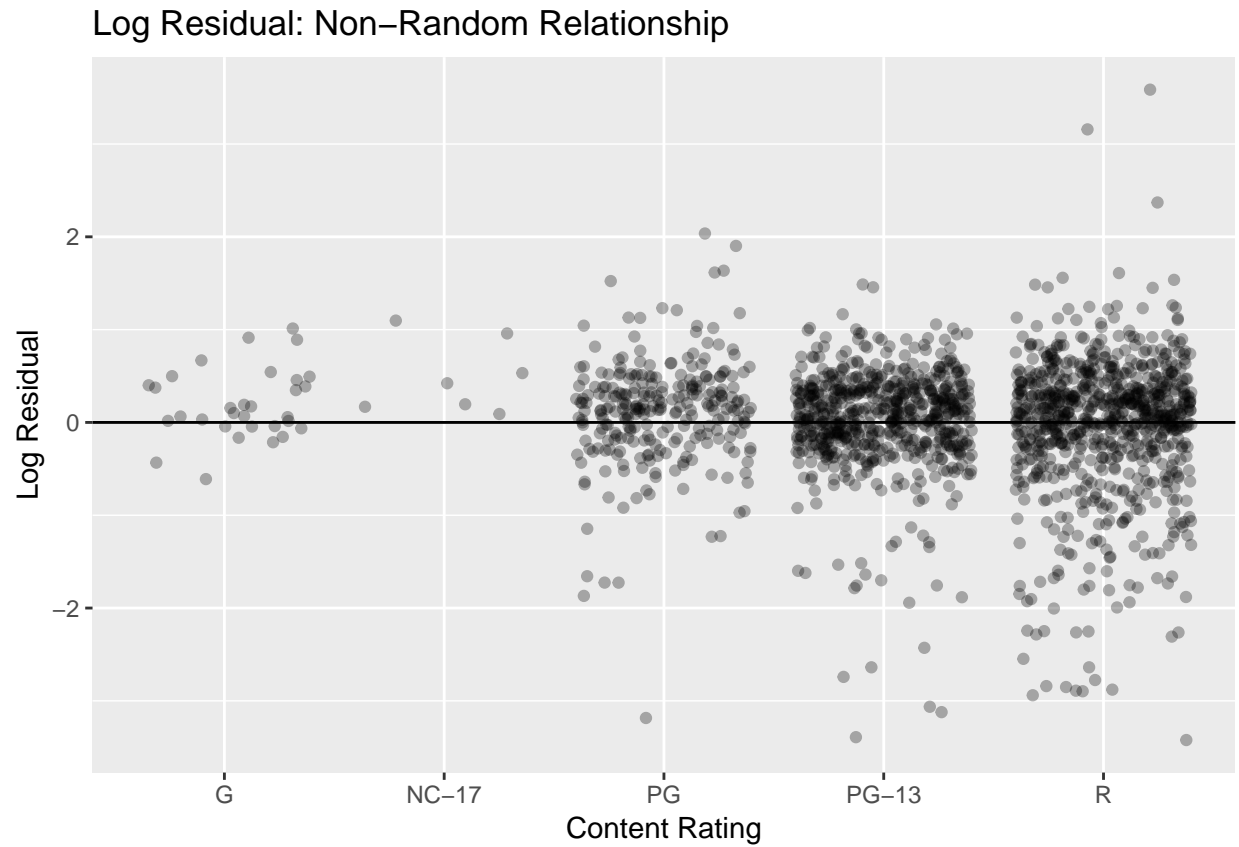
gr_con('imdb_score', 'Log Residual: Non-Random Relationship', 'Log IMDB Score')

```

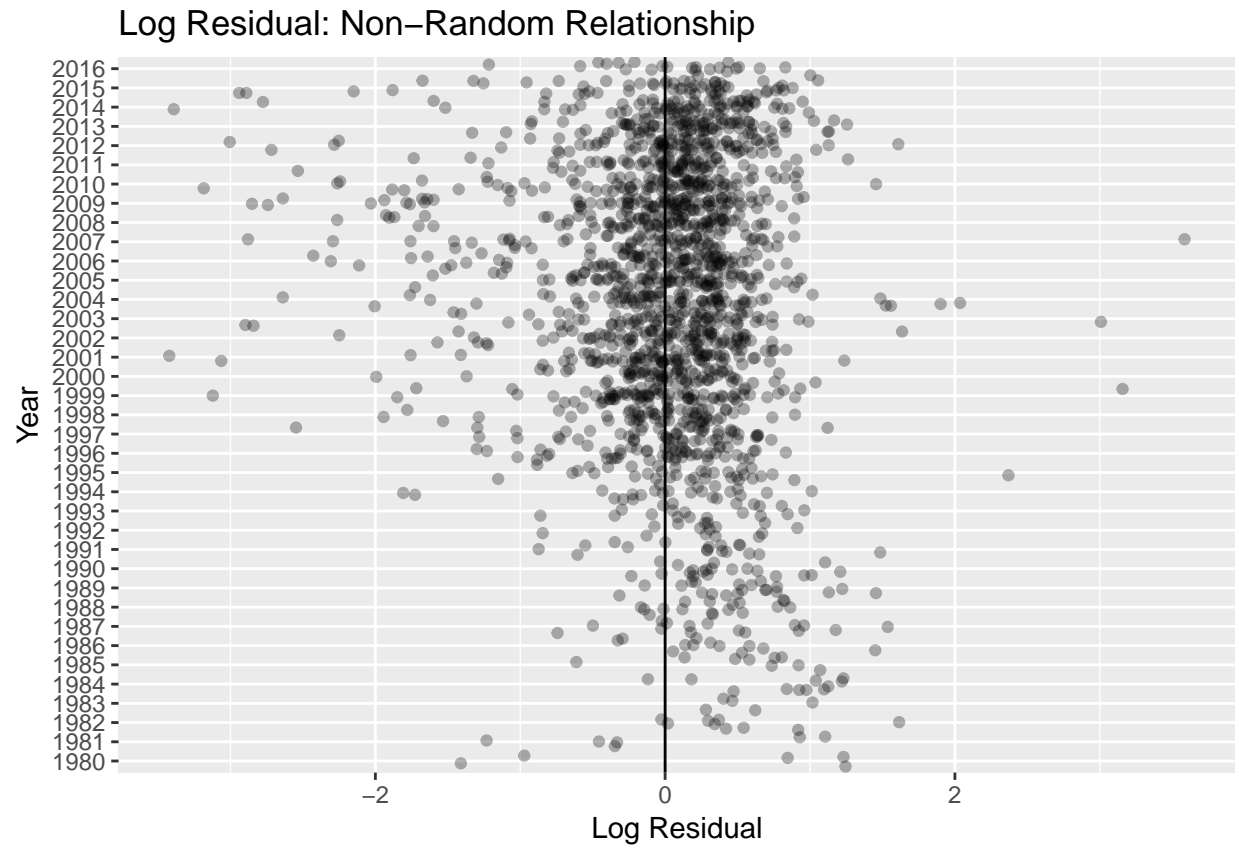
Log Residual: Non-Random Relationship



```
gr_cat('content_rating', 'Log Residual: Non-Random Relationship', 'Content Rating')
```




```
gr_cat('year', 'Log Residual: Non-Random Relationship', 'Year', flip = T)
```




```
gr_cat('Adventure', 'Log Residual: Random Relationship', 'Adventure')
```

Log Residual: Random Relationship



```
gr_cat('Action', 'Log Residual: Random Relationship', 'Action')
```

Log Residual: Random Relationship




```
# stepwise  
# ALL potentially relevant variables  
rmse_lst <- step_wise_loop(df = train %>% select(real_budget_log, imdb_score_log, year,  
                                                content_rating))
```

```
## real_budget_log  
##      0.6497192  
## [1] 1  
## imdb_score_log  
##      0.6382915  
## [1] 2  
##      year  
## 0.6281035  
## [1] 3  
## content_rating  
##      0.623201
```

```
# graph RMSE vs number of variables  
fit_rmse <- tibble(nvar = 1:length(rmse_lst),  
                  rmse = rmse_lst)  
ggplot(fit_rmse, aes(x = nvar, y = rmse)) + geom_line() +  
  scale_x_continuous(breaks = seq(1, length(rmse_lst), by = 1)) +  
  labs(x = 'Number of Variables', y = 'RMSE',  
       title = 'RMSE vs Number of Variables: Include All 4')
```

RMSE vs Number of Variables: Include All 4



after var 4, decreases too small or increase

```
mod_simple2 <- lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating,
                  data = train)
```

```
summary(mod_simple2)
```

```
##
## Call:
## lm(formula = real_gross_log ~ real_budget_log + imdb_score_log +
##     year + content_rating, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4983 -0.2230  0.0973  0.3455  3.4758
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.444778   0.391868   1.135  0.25653
## real_budget_log  0.821725   0.026058  31.534 < 2e-16 ***
## imdb_score_log  1.626517   0.199024   8.172 5.89e-16 ***
## year1981      -0.238907   0.376557  -0.634  0.52587
## year1982       0.297881   0.346750   0.859  0.39043
## year1983       0.235272   0.393319   0.598  0.54981
## year1984       0.504909   0.336175   1.502  0.13331
## year1985       0.335770   0.346892   0.968  0.33322
```



```

## year1986      0.111891  0.331663  0.337  0.73589
## year1987      0.212045  0.322869  0.657  0.51143
## year1988      0.226528  0.316488  0.716  0.47424
## year1989      0.285472  0.311721  0.916  0.35991
## year1990      0.178997  0.312725  0.572  0.56714
## year1991      0.014859  0.316338  0.047  0.96254
## year1992      0.045736  0.321585  0.142  0.88692
## year1993     -0.003696  0.312675 -0.012  0.99057
## year1994     -0.184559  0.306846 -0.601  0.54761
## year1995     -0.042381  0.298579 -0.142  0.88714
## year1996     -0.208917  0.291490 -0.717  0.47365
## year1997     -0.131576  0.291460 -0.451  0.65173
## year1998     -0.246108  0.290619 -0.847  0.39721
## year1999     -0.232583  0.287331 -0.809  0.41837
## year2000     -0.141456  0.288697 -0.490  0.62421
## year2001     -0.245089  0.286678 -0.855  0.39271
## year2002     -0.222859  0.286430 -0.778  0.43665
## year2003     -0.176337  0.287847 -0.613  0.54022
## year2004     -0.193313  0.287370 -0.673  0.50123
## year2005     -0.210828  0.286412 -0.736  0.46177
## year2006     -0.308781  0.286785 -1.077  0.28177
## year2007     -0.267002  0.288303 -0.926  0.35452
## year2008     -0.323941  0.286439 -1.131  0.25825
## year2009     -0.325844  0.285759 -1.140  0.25434
## year2010     -0.338852  0.286066 -1.185  0.23637
## year2011     -0.217375  0.288402 -0.754  0.45112
## year2012     -0.110811  0.287004 -0.386  0.69948
## year2013     -0.037216  0.286569 -0.130  0.89669
## year2014     -0.103257  0.288515 -0.358  0.72047
## year2015     -0.216276  0.290823 -0.744  0.45718
## year2016     -0.069711  0.307146 -0.227  0.82048
## content_ratingNC-17 -0.197326  0.272730 -0.724  0.46946
## content_ratingPG   -0.123677  0.119665 -1.034  0.30150
## content_ratingPG-13 -0.175106  0.115920 -1.511  0.13109
## content_ratingR    -0.346131  0.115941 -2.985  0.00287 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6214 on 1676 degrees of freedom
## (132 observations deleted due to missingness)
## Multiple R-squared:  0.4685, Adjusted R-squared:  0.4552
## F-statistic: 35.18 on 42 and 1676 DF, p-value: < 2.2e-16

rmse(mod_simple2, data = valid)

## [1] 0.623201

# when consider factors as one variable, they are significant
anova(mod_simple2)

## Analysis of Variance Table
##
## Response: real_gross_log
##          Df Sum Sq Mean Sq    F value    Pr(>F)
## real_budget_log    1 491.67   491.67 1273.3487 < 2.2e-16 ***

```

```
## imdb_score_log      1  28.18   28.18   72.9937 < 2.2e-16 ***
## year                36  36.51    1.01    2.6269 6.334e-07 ***
## content_rating       4  14.13    3.53    9.1519 2.589e-07 ***
## Residuals          1676 647.14    0.39
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# number of observations
nobs(mod_simple2)
```

```
## [1] 1719
```

```
# compare with anova to more complex models
anova(mod_all, mod_simple2)
```

```
## Analysis of Variance Table
##
## Model 1: real_gross_log ~ Adventure + Action + Family + Mystery + Documentary +
##      Drama + History + Romance + real_budget_log + imdb_score_log +
##      year + content_rating
## Model 2: real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      1668 614.85
## 2      1676 647.14 -8     -32.291 10.95 3.744e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# to compare to mod_simple, need some number of observations: filter out missing content rating and the
mod_simple <- lm(real_gross_log ~ real_budget_log, data = train %>% filter(!is.na(content_rating)))
anova(mod_simple, mod_simple2)
```

```
## Analysis of Variance Table
##
## Model 1: real_gross_log ~ real_budget_log
## Model 2: real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      1717 725.98
## 2      1676 647.14 41      78.834 4.9797 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
gr_resid(mod_simple2)
```

```
train %>%
  add_residuals(mod_simple2, 'lresid') %>%
  ggplot(aes(sample = lresid)) +
    geom_qq() +
    geom_qq_line() +
    labs(title = 'Residual QQQPlot: Deviation at Tails',
         x = 'Theoretical Quantiles', y = 'Sample Quantiles')
```

```
## Warning: Removed 132 rows containing non-finite values (stat_qq).
```

```
## Warning: Removed 132 rows containing non-finite values (stat_qq_line).
```

Residual QQPlot: Deviation at Tails




```
train_pred_simple <- train %>%
  add_predictions(mod_simple2, 'lpred') %>%
  mutate(pred = 10^lpred)


# point because many factors. won't just be a single coefficient determining
# SEE BOOK where did prediction based on a bunch of variables
lapply(c('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes', 'imdb_score'), function(var) {
  train_pred_simple %>%
    ggplot(aes_string(x = var)) +
    geom_point(aes(y = real_gross, color = 'Actual')) +
    geom_point(aes(y = pred, color = 'Predicted')) +
    scale_y_log10() + scale_x_log10()
})

# predictions against other genres
lapply(c('content_rating', 'year', 'total_oscars_actor', 'total_oscars_director', all_genre_vars), function(var) {
  train_pred_simple %>%
    ggplot(aes_string(x = var)) +
    geom_boxplot(aes(y = real_gross)) +
    # include mean
    stat_summary(aes(y = real_gross), fun.y = mean, geom = 'point', size = 4) +
    geom_point(data = train_pred %>% group_by(!!rlang::sym(var)) %>% summarize(mean = mean(pred)),
              aes(y = mean), color = 'red', size = 2) +
    scale_y_log10()
})
```


Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction



Revenue Actual vs Predicted: Log Scale
Successful Prediction



Revenue Actual vs Predicted: Log Scale Successful Prediction




```
train %>%
  add_predictions(mod_simple2, 'lpred') %>%
  mutate(pred = 10^lpred) %>%
  ggplot() +
    geom_freqpoly(aes(x = real_gross, color = 'Actual')) +
    geom_freqpoly(aes(x = pred, color = 'Predicted')) +
    scale_x_log10() +
    labs(x = 'Log Real Gross Revenue', y = 'Count') +
    scale_color_manual(name = '', values = c(Actual = 'black', Predicted = 'blue'))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 132 rows containing non-finite values (stat_bin).
```



Compare Models

- Somewhat of a toss-up between model without genre and model with some genre variables. Predictions similar, random relationships with residuals.
- Histogram of actual vs predicted revenue is, if anything, more similar for the no genre model.
- Very similar RMSE (.621 vs .623) and R-squared (.48 vs .46)
- Anova shows that the more complex genre model is significantly different from the simpler model without genre.
- However, given that the predictions etc. are so similar, it is better to go with the simpler model. Worry about over-fitting with the genre model. Especially because some of the genres don't have that many observations.
- Main take away is that budget has the strongest, most significant relationship with revenue. The model with just budget had a RMSE that was only slightly higher (.65 vs .62) and an R-squared that was only slightly lower (.45 vs .46) than the models with more variables. Also, the other variables that we add to the model had fairly close to random relationships with the residuals from the model with just budget.
- However, predicted vs actual histogram of real gross shows that it is not perfect. Model with the three extra variables does a better job predicting.

Fit Final Model on the Test Set

- Similar results. Predictions are occasionally a bit off, but by in large on track. Variables have random relationship with residuals.
- Year becomes insignificant (or significant at a low level) in anova
- Only two NC-17 movies in test set

```
mod_simple2 <- lm(real_gross_log ~ real_budget_log + imdb_score_log + year + content_rating,
  data = test)
```

```
summary(mod_simple2)
```

```
##
## Call:
## lm(formula = real_gross_log ~ real_budget_log + imdb_score_log +
##     year + content_rating, data = test)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-2.84764	-0.22899	0.06723	0.32426	1.59717

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.10393	0.60847	1.814	0.0702 .
real_budget_log	0.79996	0.04546	17.597	< 2e-16 ***
imdb_score_log	1.22228	0.27850	4.389	1.37e-05 ***
year1981	-0.60094	0.60421	-0.995	0.3204
year1982	-0.09994	0.74187	-0.135	0.8929
year1983	-0.76350	0.60738	-1.257	0.2093
year1984	-0.44646	0.55057	-0.811	0.4178
year1986	-0.29362	0.60302	-0.487	0.6265
year1987	-0.29062	0.52190	-0.557	0.5779
year1988	-0.18739	0.49264	-0.380	0.7038
year1989	-0.33162	0.50487	-0.657	0.5116
year1990	-0.05923	0.55134	-0.107	0.9145
year1991	-0.30015	0.55008	-0.546	0.5855
year1992	-0.20184	0.48420	-0.417	0.6769
year1993	-0.44070	0.47808	-0.922	0.3570
year1994	-0.61541	0.47137	-1.306	0.1923
year1995	-0.37526	0.45398	-0.827	0.4088
year1996	-0.38177	0.45414	-0.841	0.4009
year1997	-0.64346	0.45007	-1.430	0.1534
year1998	-0.42600	0.45331	-0.940	0.3478
year1999	-0.64173	0.43923	-1.461	0.1446
year2000	-0.71488	0.44239	-1.616	0.1067
year2001	-0.59091	0.44307	-1.334	0.1829
year2002	-0.45233	0.44085	-1.026	0.3053
year2003	-0.49016	0.44627	-1.098	0.2725
year2004	-0.74089	0.44106	-1.680	0.0936 .
year2005	-0.77633	0.44707	-1.736	0.0830 .
year2006	-0.75604	0.44412	-1.702	0.0893 .
year2007	-0.96033	0.44295	-2.168	0.0306 *
year2008	-0.49162	0.44472	-1.105	0.2694
year2009	-0.55675	0.44373	-1.255	0.2101
year2010	-0.82217	0.44436	-1.850	0.0648 .
year2011	-0.77645	0.44055	-1.762	0.0786 .
year2012	-0.57905	0.44518	-1.301	0.1939
year2013	-0.55307	0.44998	-1.229	0.2196
year2014	-0.48477	0.44118	-1.099	0.2723
year2015	-0.73253	0.45148	-1.623	0.1053
year2016	-0.47767	0.45780	-1.043	0.2972

```

## content_ratingNC-17 -0.49842    0.46461   -1.073    0.2839
## content_ratingPG     0.23849    0.15955    1.495    0.1356
## content_ratingPG-13  0.13385    0.15007    0.892    0.3728
## content_ratingR     -0.12547    0.15169   -0.827    0.4085
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6021 on 545 degrees of freedom
## (32 observations deleted due to missingness)
## Multiple R-squared:  0.4971, Adjusted R-squared:  0.4592
## F-statistic: 13.14 on 41 and 545 DF,  p-value: < 2.2e-16
anova(mod_simple2)

## Analysis of Variance Table
##
## Response: real_gross_log
##           Df Sum Sq Mean Sq F value    Pr(>F)
## real_budget_log    1 160.558  160.558 442.9189 < 2.2e-16 ***
## imdb_score_log     1   6.931    6.931  19.1201 1.471e-05 ***
## year              35  17.158    0.490   1.3523 0.08858 .
## content_rating      4  10.600    2.650   7.3105 9.709e-06 ***
## Residuals         545 197.563    0.363
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
rmse(mod_simple2, data = test)


## [1] 0.5801406
gr_resid(mod_simple2, df = test)

test %>%
  add_residuals(mod_simple2, 'lresid') %>%
  ggplot(aes(sample = lresid)) +
    geom_qq() +
    geom_qq_line() +
    labs(title = 'Residual QQPLOT: Deviation at Tails',
         x = 'Theoretical Quantiles', y = 'Sample Quantiles')

## Warning: Removed 32 rows containing non-finite values (stat_qq).
## Warning: Removed 32 rows containing non-finite values (stat_qq_line).

```

Residual QQPlot: Deviation at Tails




```
test_pred <- test %>%
  add_predictions(mod_simple2, 'lpred') %>%
  mutate(pred = 10^lpred)


# point because many factors. won't just be a single coefficient determining
# SEE BOOK where did prediction based on a bunch of variables
lapply(c('real_budget', 'director_facebook_likes', 'cast_total_facebook_likes', 'imdb_score'), function(var) {
  test_pred %>%
    ggplot(aes_string(x = var)) +
    geom_point(aes(y = real_gross, color = 'Actual')) +
    geom_point(aes(y = pred, color = 'Predicted')) +
    scale_y_log10() + scale_x_log10()
})

# predictions against other genres
lapply(c('content_rating', 'year', 'total_oscars_actor', 'total_oscars_director', all_genre_vars), function(var) {
  test_pred %>%
    ggplot(aes_string(x = var)) +
    geom_boxplot(aes(y = real_gross)) +
    # include mean
    stat_summary(aes(y = real_gross), fun.y = mean, geom = 'point', size = 4) +
    geom_point(data = train_pred %>% group_by(!rlang::sym(var)) %>% summarize(mean = mean(pred)),
              aes(y = mean), color = 'red', size = 2) +
    scale_y_log10()
})
```


Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction




Revenue Actual vs Predicted: Log Scale
Successful Prediction



Revenue Actual vs Predicted: Log Scale
Successful Prediction



Revenue Actual vs Predicted: Log Scale Successful Prediction




```
test %>%
  add_predictions(mod_simple2, 'lpred') %>%
  mutate(pred = 10^lpred) %>%
  ggplot() +
  geom_freqpoly(aes(x = real_gross, color = 'Actual')) +
  geom_freqpoly(aes(x = pred, color = 'Predicted')) +
  scale_x_log10() +
  labs(x = 'Log Real Gross Revenue', y = 'Count') +
  scale_color_manual(name = '', values = c(Actual = 'black', Predicted = 'blue'))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 32 rows containing non-finite values (stat_bin).
```



TO DO: * At the end, create those two labeled graphs with the two major outliers in residuals. * Put model results into a table (LaTeX?) * Add manual legends to boxplot prediction: black vs red dot